




# Prof. Florian Künzner

Technical University of Applied Sciences Rosenheim, Computer Science

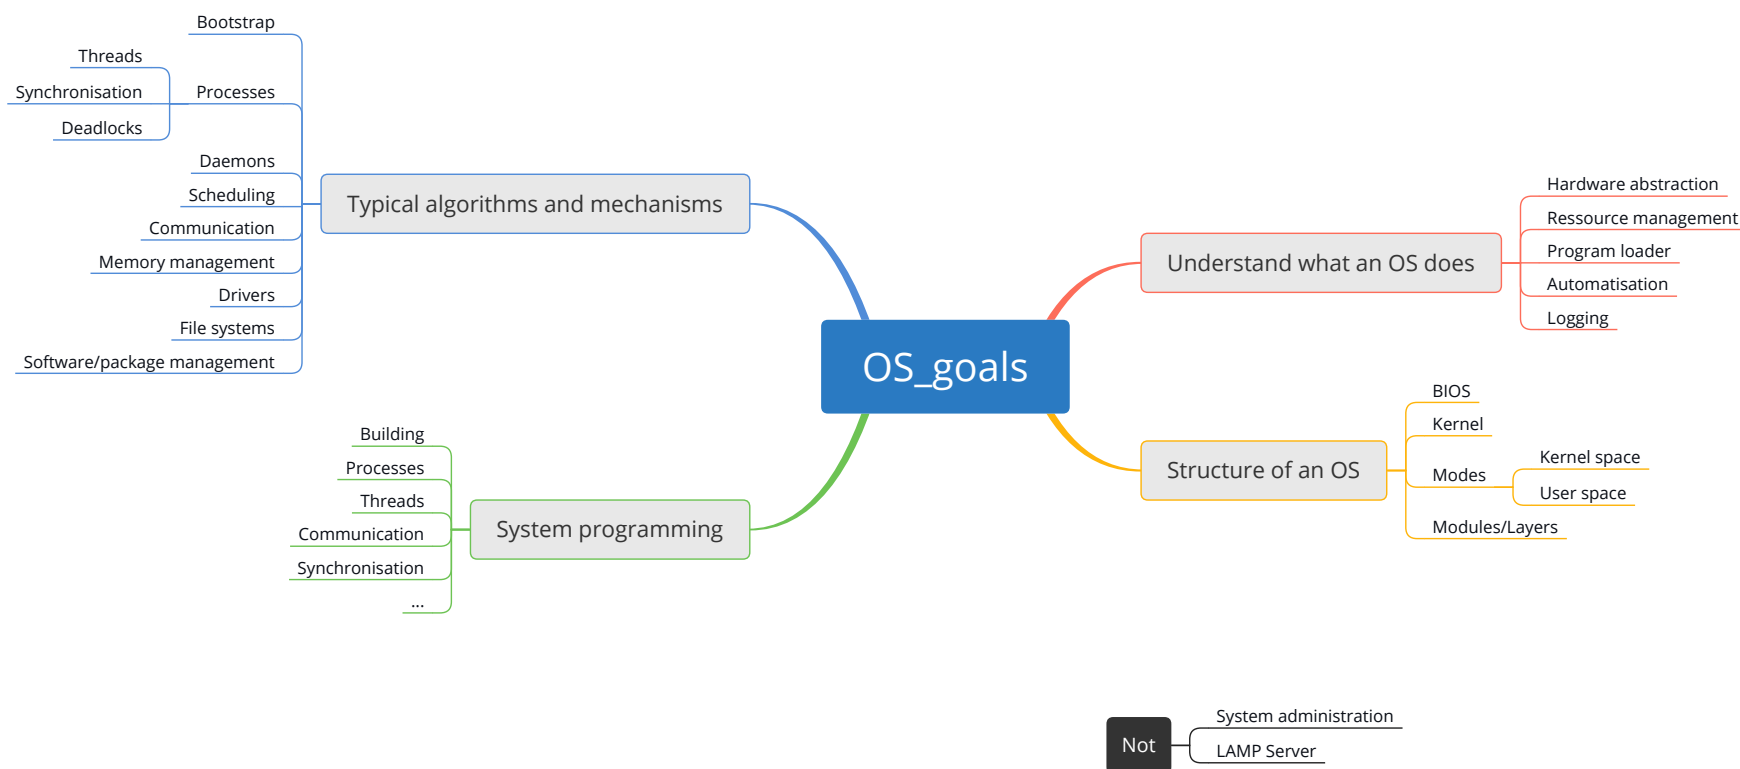
## OS 4 – System boot; OS architecture



source: [iconspng.com](https://www.iconspng.com)

The lecture is based on the work and the documents of Prof. Dr. Ludwig Frank

# Goal



# Goal

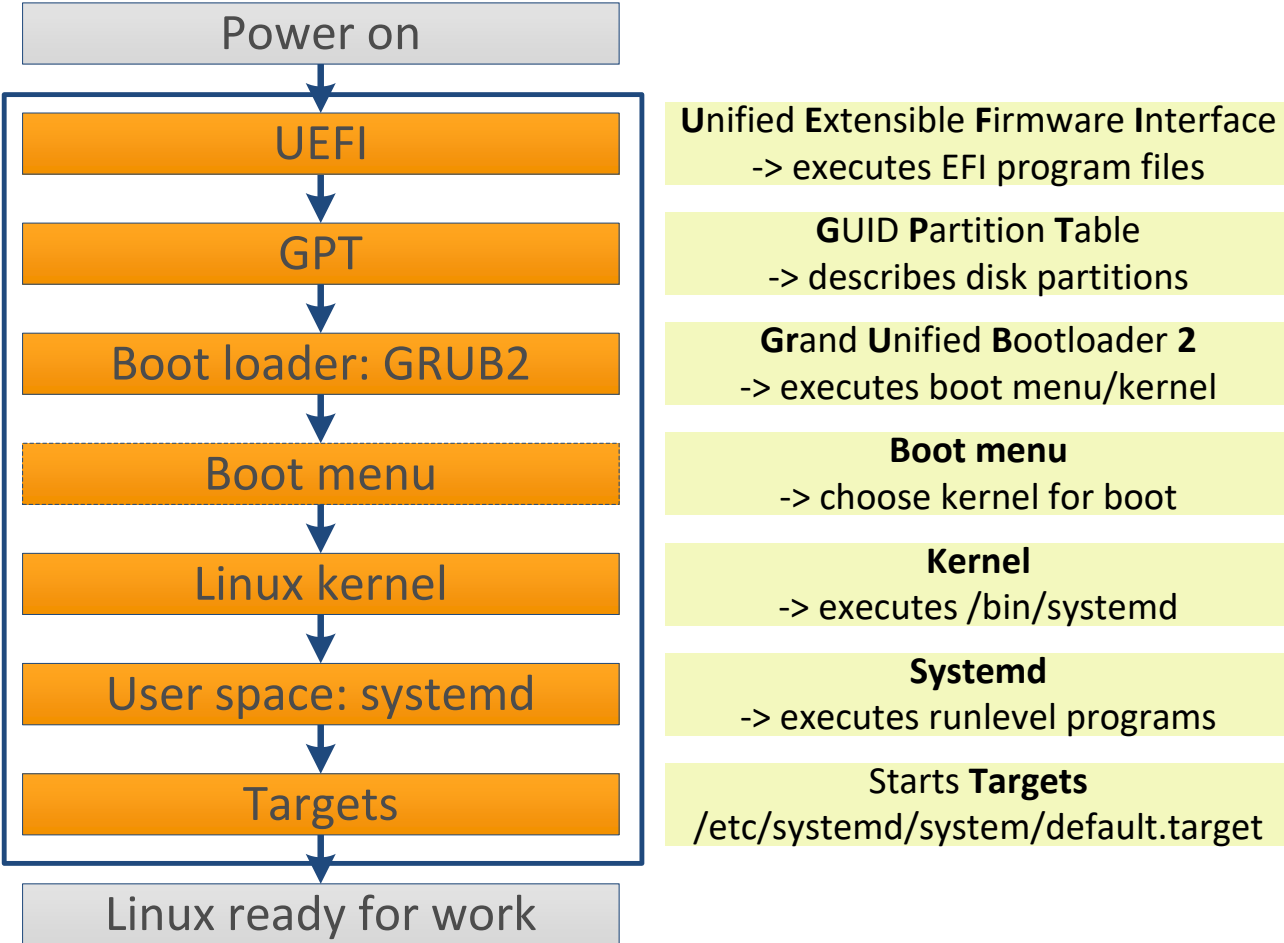
## OS::System boot; OS architecture

- Boot procedure
- BIOS/UEFI
- MBR/GPT
- Boot loader: Grub2
- init/systemd
- OS architecture

# Boot

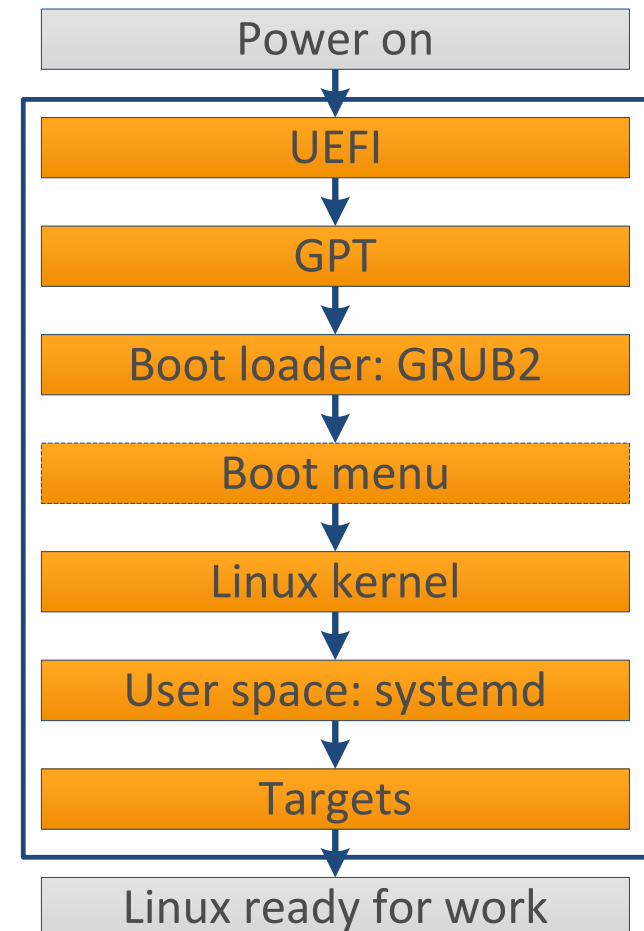
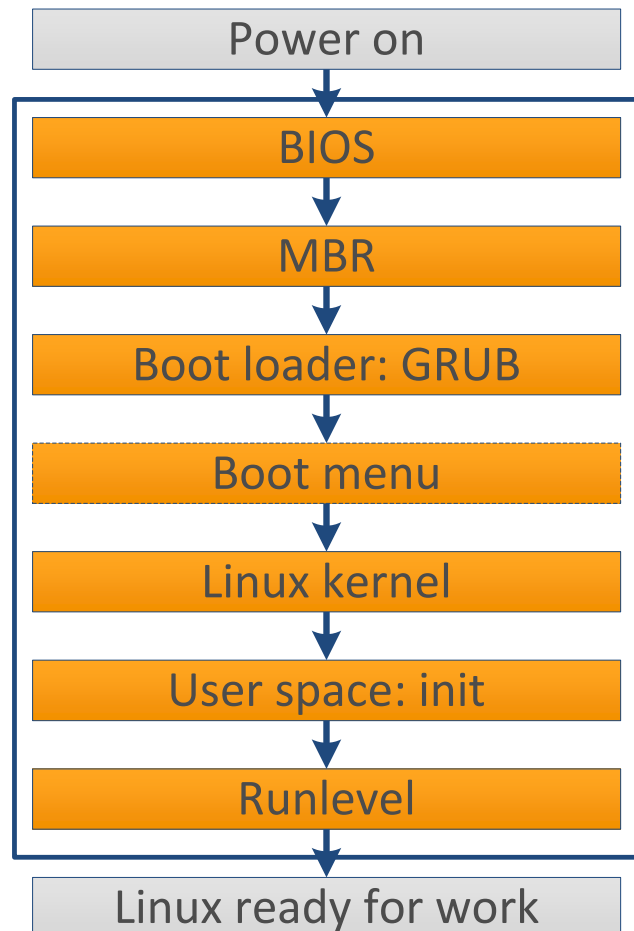
How does a PC boot into the operating system?

# Boot procedure: UEFI (typically)





# Boot procedure: BIOS vs. UEFI



# UEFI - Unified ext. firmware interface

- Since 2010 in consumer products
- Graphical user interface (with mouse)
- Fast boot: cache + hibernation (Windows only)
- Secure boot:
  - Protection against malware
  - Prevents against execution of unsigned code
- Network boot
- Modular interface for applications and devices (EFI drivers)
- Supported modes:
  - UEFI mode: Requires an EFI partition on boot device
  - BIOS mode: Old way of booting

More details: <https://www.marksei.com/bios-uefi-explained>



# MBR - Master Boot Record

## MBR (512 bytes)

Master boot code  
(446 bytes)

### Partition table (64 bytes)

1st entry	2nd entry	3rd entry	4th entry
--------------	--------------	--------------	--------------

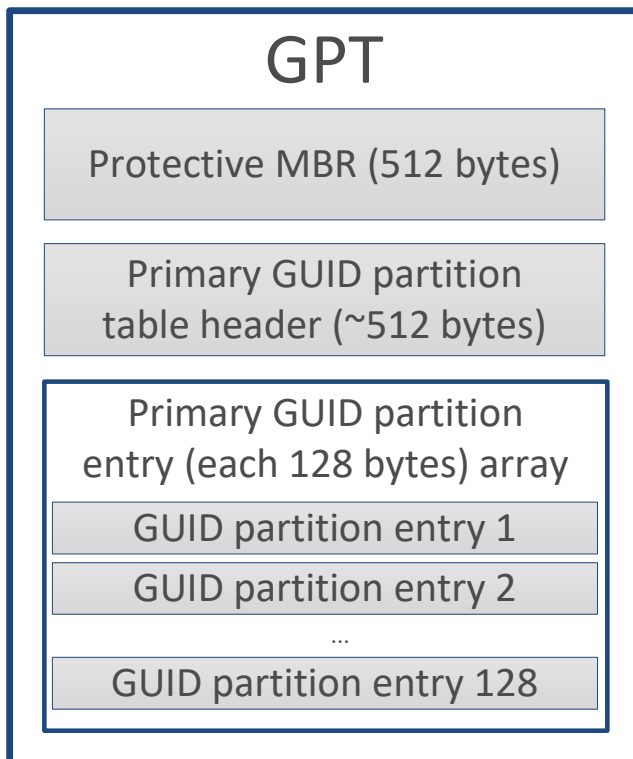
Magic Number: 0x55, 0xAA

- Supports up to **2 TiB disks**
- Supports up to **2 TiB partitions**
- No safety (**no checksum**)
- Supports **4 primary partitions**
- Supports **one extended partition** (not bootable)

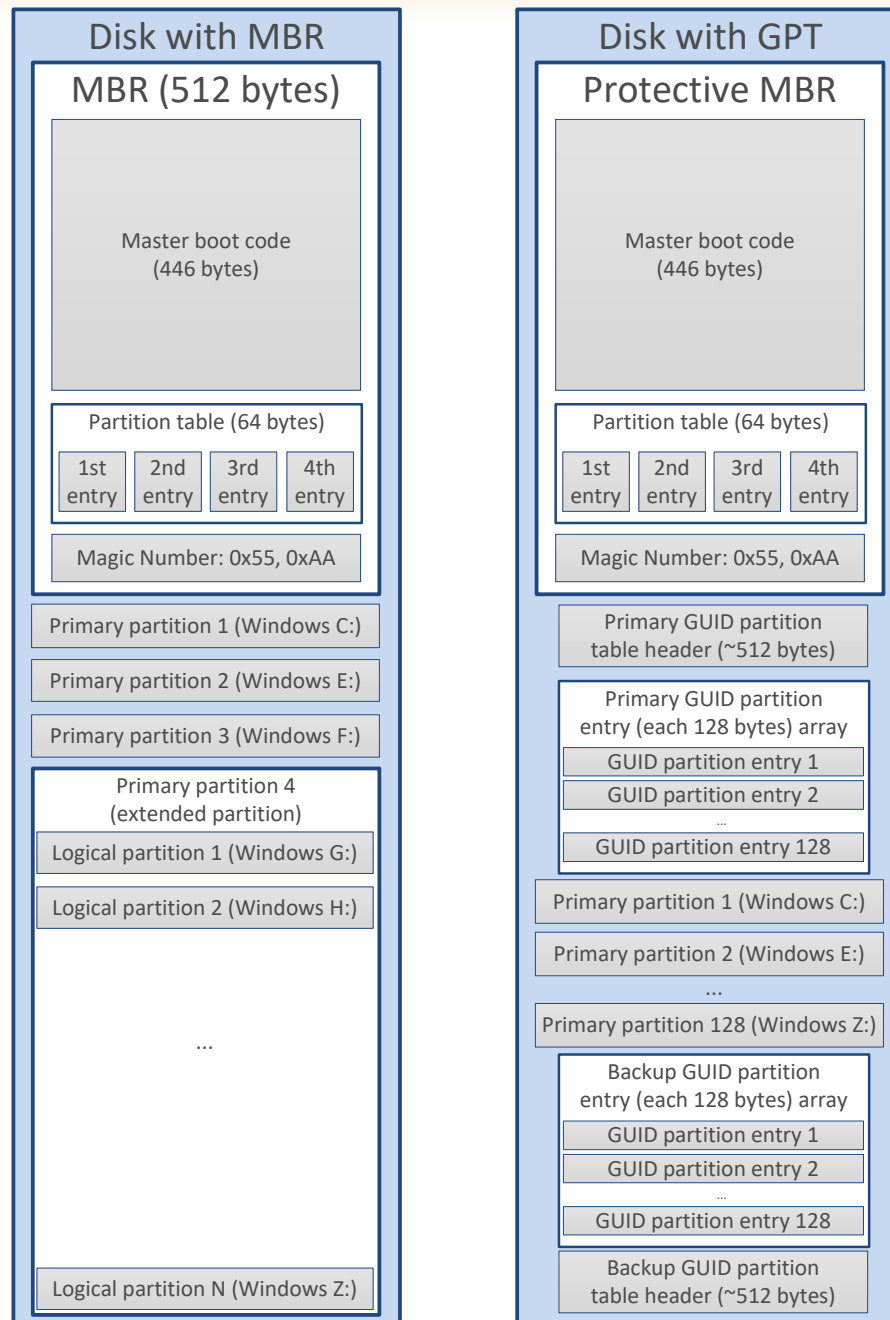




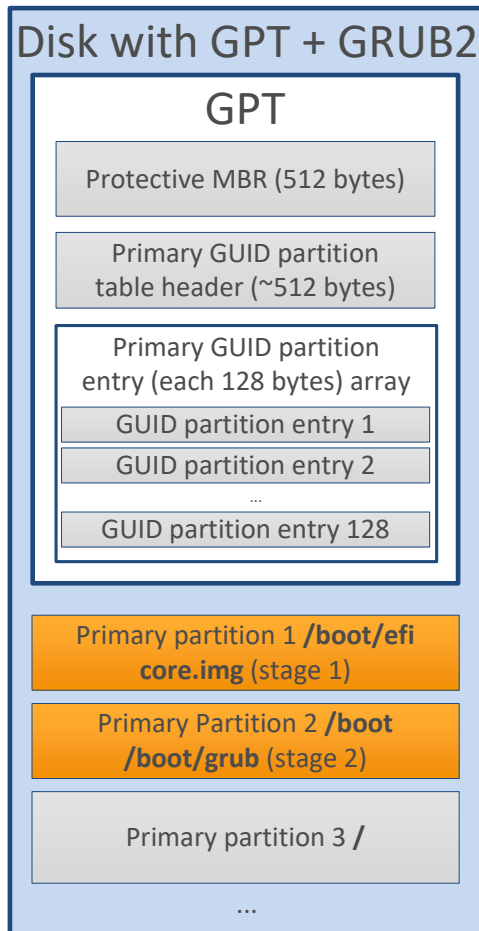
# GPT - GUID Partition Table



- Supports up to **18 EiB disks**
- Supports up to **18 EiB partitions**
- Safety (**checksum, backup**)
- Version number
- Supports **128 primary partitions**
- Does not have a boot code
- UEFI **boots** from an **EFI partition**



# GRUB2 - GRand Unified Bootloader 2



- Boots in stages:
  - **stage1** (core.img): Loads directly stage 2 (/boot/grub), usually from /boot partition. Contains file system drivers.
  - **stage2** (/boot/grub): Loads the default configuration file and any other modules needed.
- Themes, **graphical menus**, scripting support
- Uses **UUIDs** to identify disks
- Supports additionally: PowerPC
- Supports LLVM and RAID
- Boots live CD images from hard drive
- Automated search for other OS (like Windows)

# User space: systemd

- **systemd** is the **first process** started by the kernel
- The mother of all subprocesses
- Has always **PID 1**
- Looks in `/etc/systemd/system/default.target` for default target
- **Executes** the default **target** (unit) scripts
- **Starts** the user space **processes** on boot:
  - Daemons (crond, httpd, sshd, inetd, syslogd, ...)
  - Terminals
  - Graphical desktop
  - ...
- Speed-up the boot: starts processes in **parallel** (where possible)



# User space: systemd (commands)

- Get default target: `systemctl get-default`
- Change runlevel: `systemctl isolate poweroff.target`
- Control a daemon:  
`/etc/systemd/system/<daemon>.service`
  - `service daemon start`
  - `service daemon stop`
  - `service daemon reload`
  - `service daemon restart`
  - `service daemon status`
  - ...
- Enable a daemon: `systemctl enable daemon.service`
- Disable a daemon: `systemctl disable daemon.service`



# User space: systemd (example)

```
1 [Unit]
2 Description=Demo daemon that does something useful
3
4 [Service]
5 Type=forking
6 ExecStart=/usr/sbin/daemon
7 ExecStopPost=/bin/rm /var/run/daemon.pid
8
9 [Install]
10 WantedBy=multi-user.target
```

More details: <https://www.digitalocean.com/community/tutorials/understanding-systemd-units-and-unit-files>

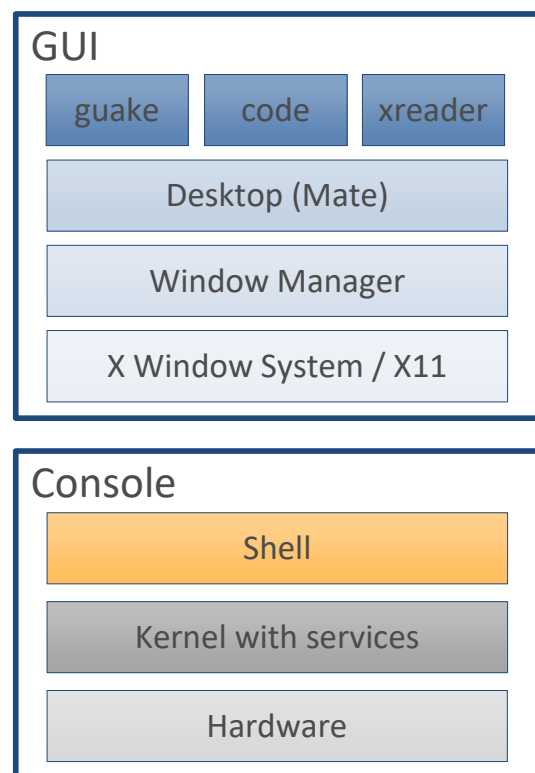


# Init runlevel vs. systemd targets

init runlevel	systemd target	systemd target aliases	Description
	halt.target		Halts the system without powering it down.
0	<b>poweroff.target</b>	runlevel0.target	Halts the system and turns the power off.
S	emergency.target		Single user mode. No services are running; filesystems are not mounted.
1	rescue.target	runlevel1.target	A base system including mounting the filesystems with only the most basic services running and a rescue shell on the main console.
2		runlevel2.target	Multiuser, without NFS but all other non-GUI services running.
3	multi-user.target	runlevel3.target	All services running but command line interface (CLI) only.
4		runlevel4.target	Unused.
5	<b>graphical.target</b>	runlevel5.target	Multiuser with a GUI.
6	reboot.target	runlevel6.target	Reboot.
	default.target		This target is always aliased with a symbolic link to either multi-user.target or graphical.target. systemd always uses the default.target to start the system.

More details: <https://opensource.com/article/17/2/linux-boot-and-startup>

# Linux high level overview (1)



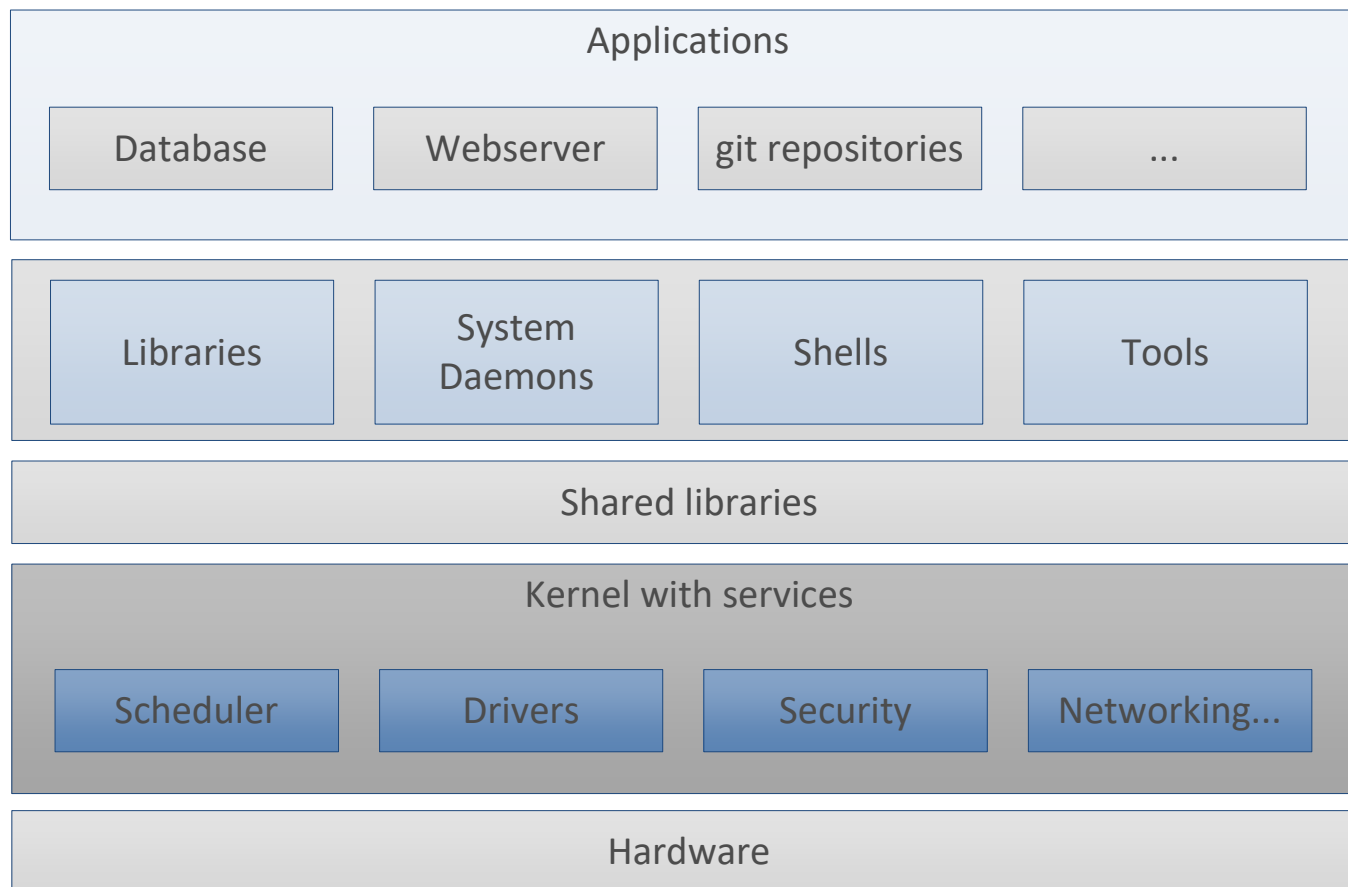
## OS tasks

- Execute (graphical) user applications
- Provide desktop environment
- Draw windows (graphical elements)
- Provide shells
- Manage resources
- Support, abstract and virtualise hardware



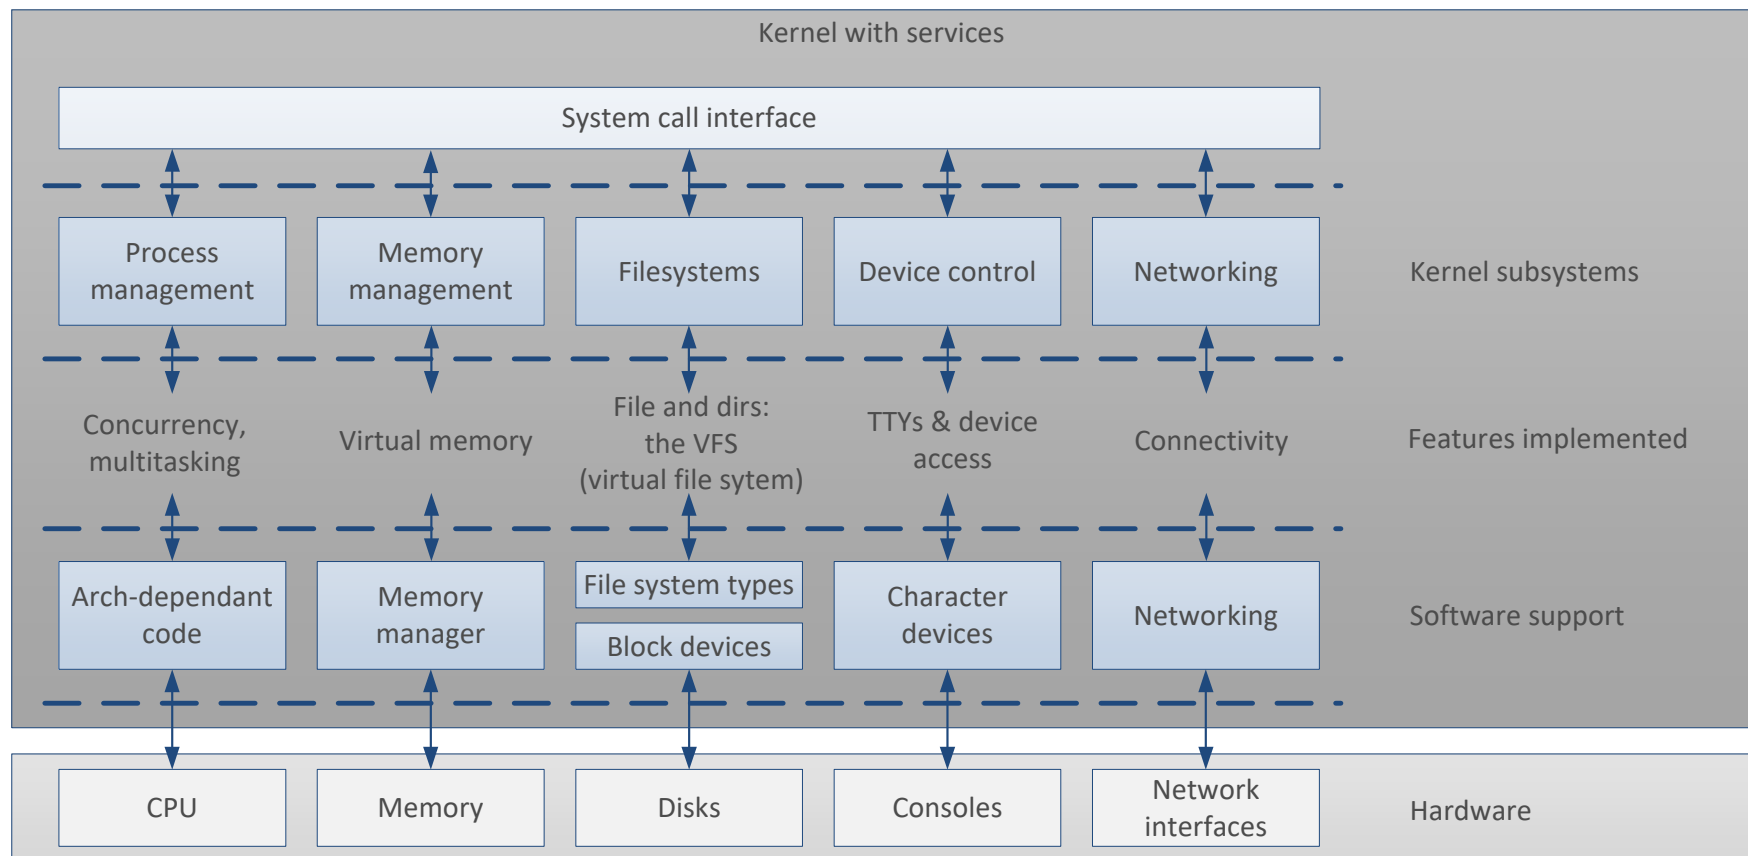


# Linux high level overview (2)

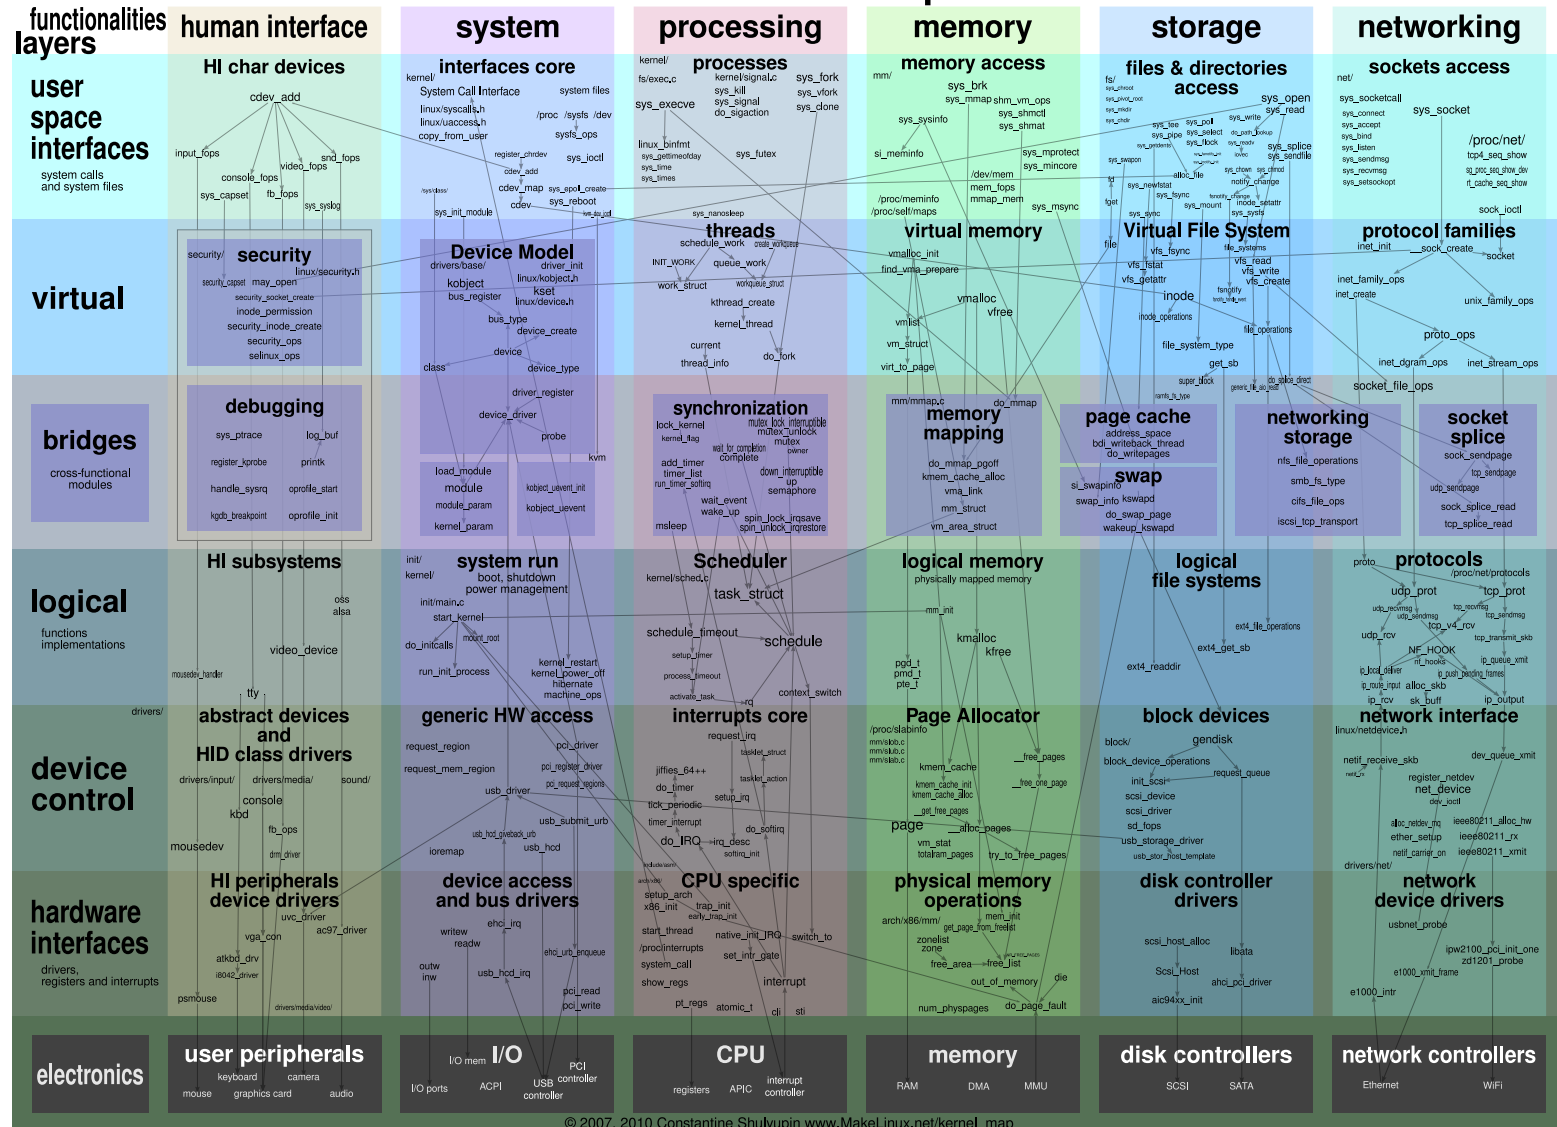




# OS overview: Linux kernel

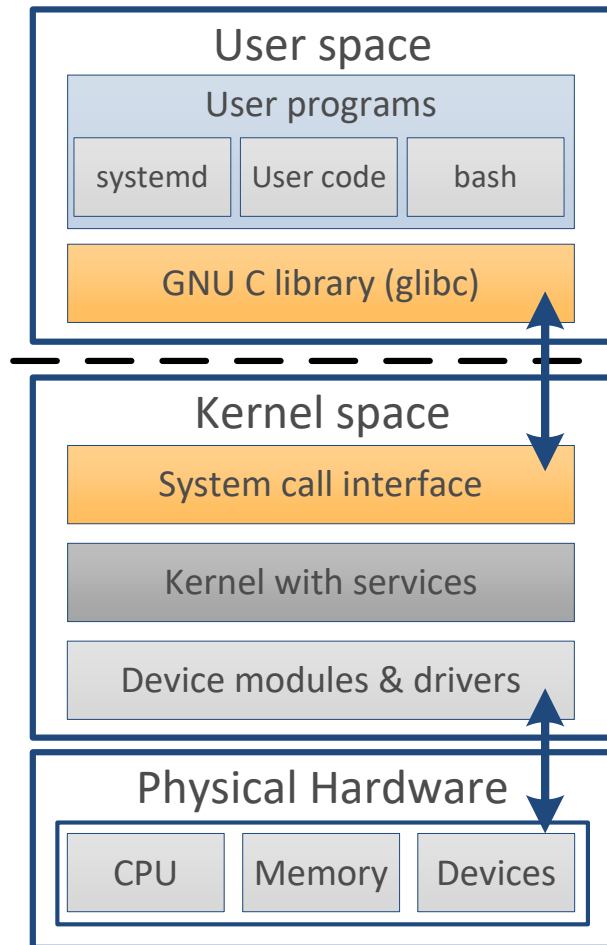


# Linux<sup>2.6.36</sup> kernel map





# Protection: user vs. kernel space



## User space

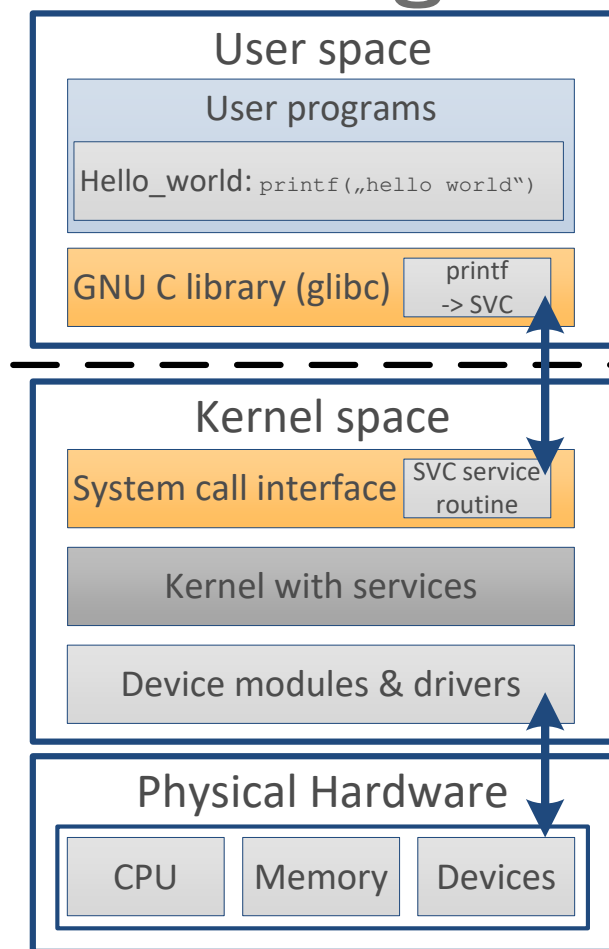
- All code outside the kernel
- Also called “userland”
- **Restricted** (encapsulated) access to the **hardware**
- Can **only** use a **subset of CPU instructions**
- Can **only** access the **assigned memory** addresses
- **Crashes** in a user process: **only stops these process.**

## Kernel space

- **Complete** and **unrestricted** access to the **hardware**
- Can execute **any CPU instruction**
- Can access **any memory address**
- **Crashes** in kernel are “catastrophic”: **system stop!**



# Accessing the kernel: supervisor call

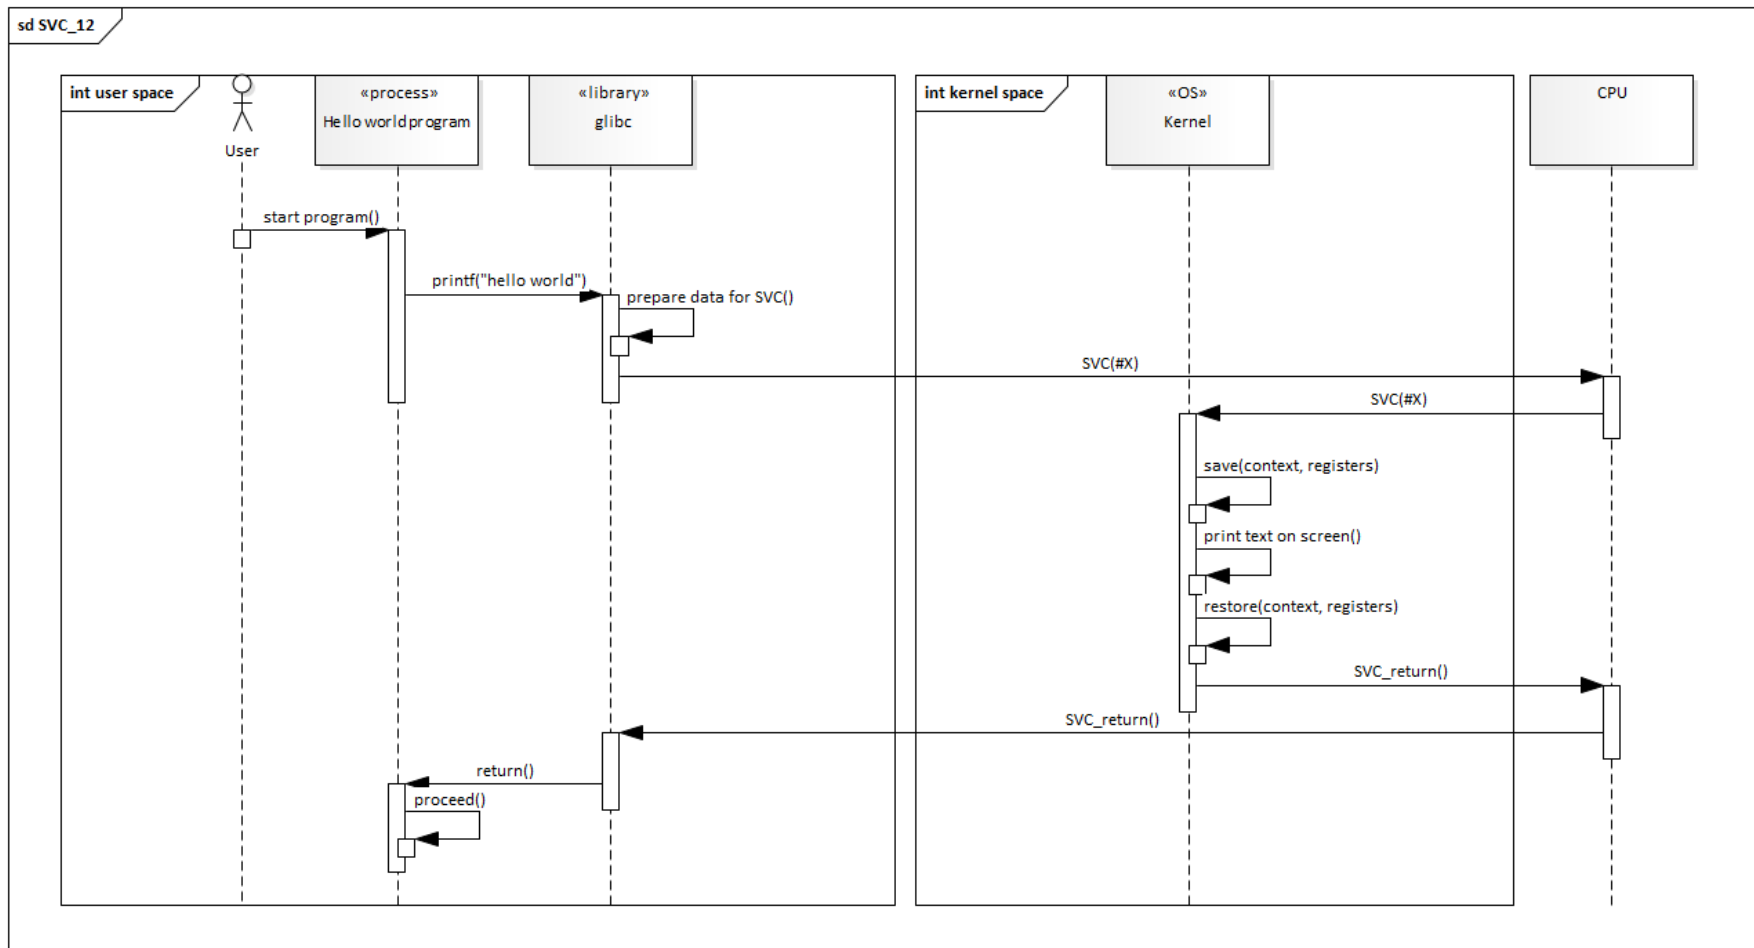


## Supervisor call (SVC)

- **CPU instruction** to give control to the OS/kernel
- **Requests for an OS service:**
  - Start process
  - Allocate memory
  - File open/close/read/write/...
  - Print something on terminal (TTY) or screen
  - Send data over network
  - ...
- SVCs are **numbered**
- The calling **process is interrupted** while the kernel executes the SVC.



# Supervisor call (SVC) example sequence



# Summary and outlook

## Summary

- Boot procedure
- BIOS/UEFI
- MBR/GPT
- Boot loader: Grub2
- init/systemd
- OS architecture

## Outlook

- Processes vs. Threads
- Parallelisation