

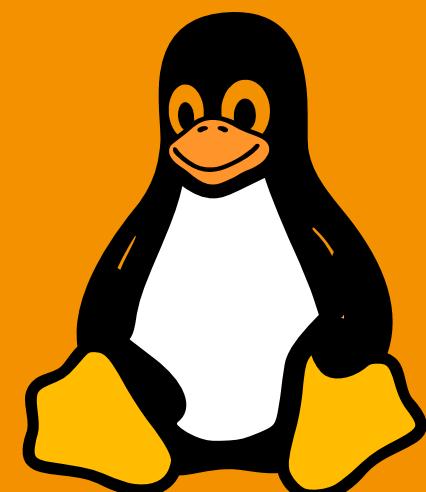
Start: 8:01



Prof. Dr. Florian Künzner

Technical University of Applied Sciences Rosenheim, Computer Science

OS 12 – Memory Management

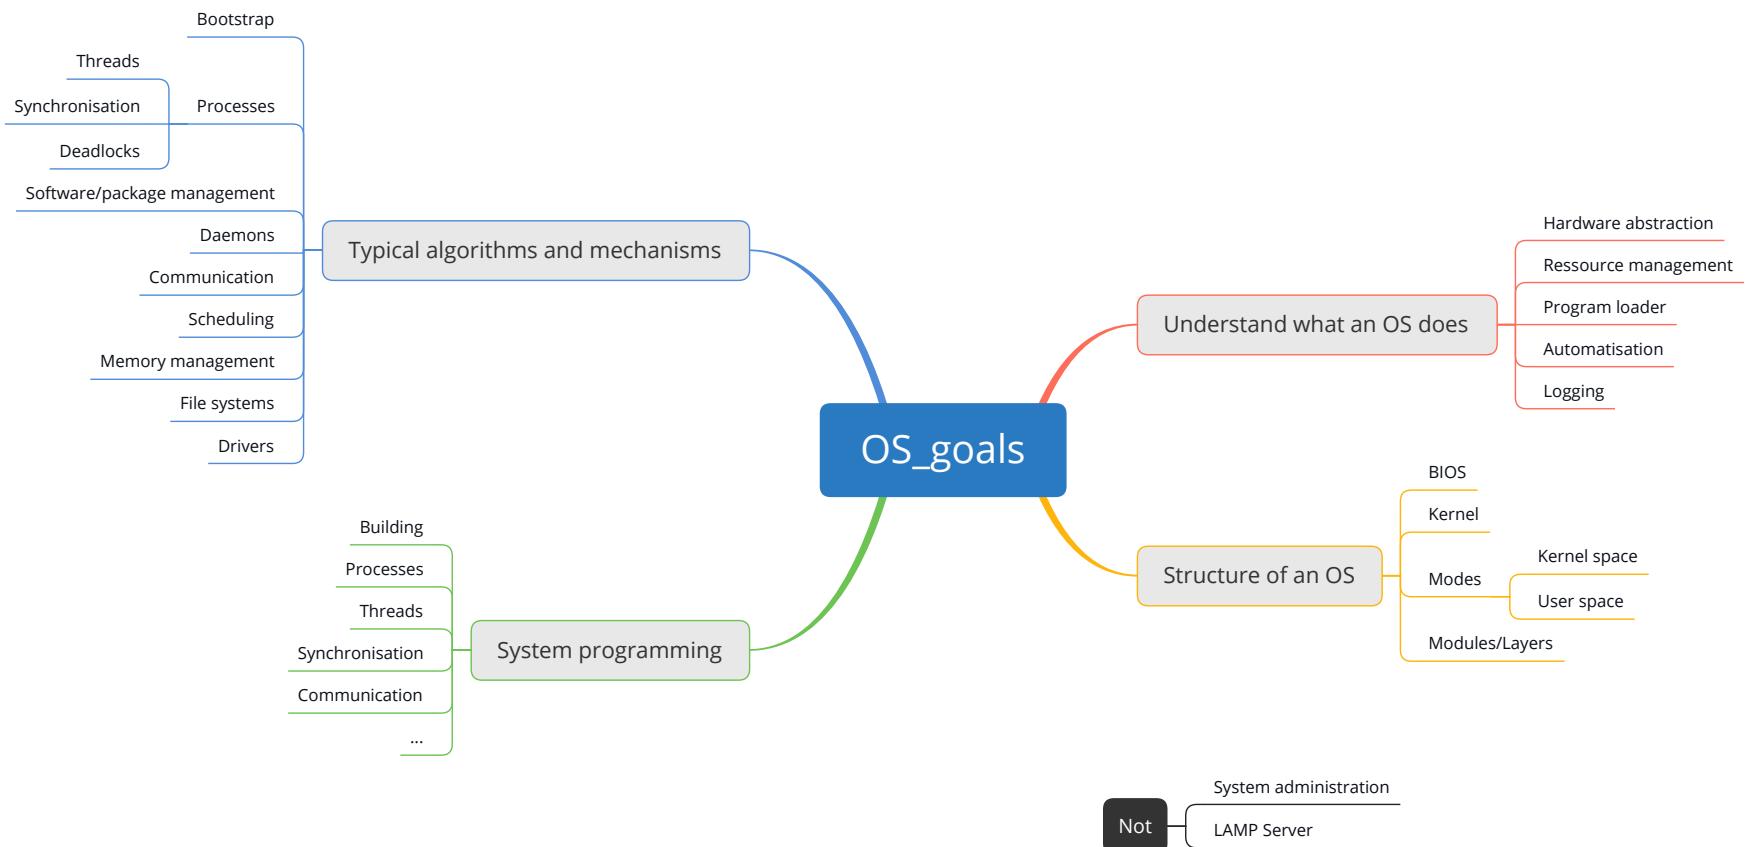


source: iconspng.com

The lecture is based on the work and the documents of Prof. Dr. Ludwig Frank



Goal





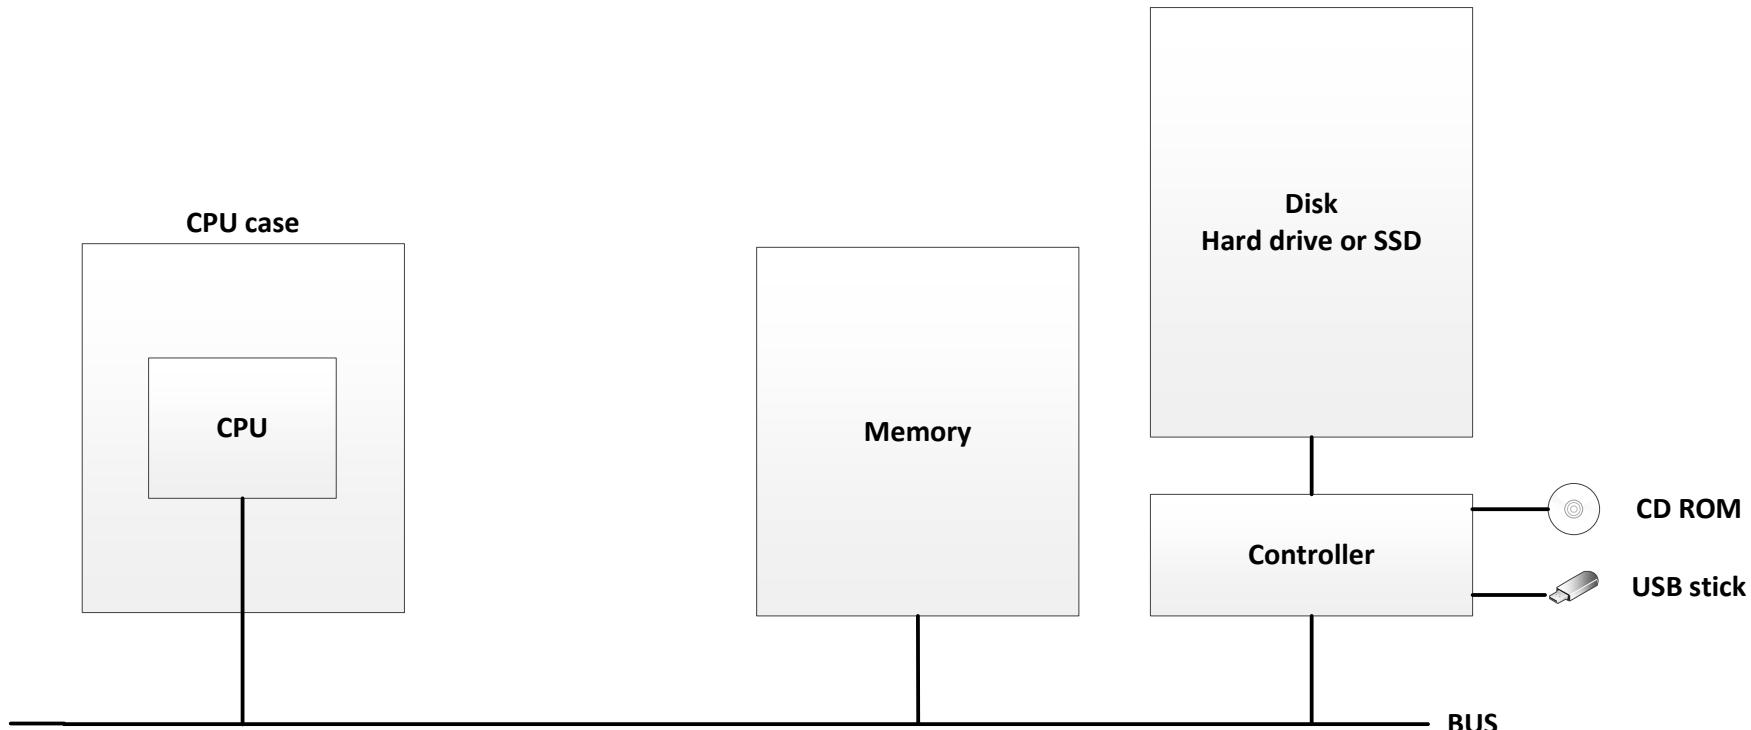
Goal

OS::Memory Management

- Caching
- Partitioning
- Fragmentation
- Allocation strategies
- Swapping
- Virtual memory



Memory overview





Memory management requirements

Requirement	Description
-------------	-------------



Memory management requirements

Requirement

Retrievability

Description

Who is the **owner** of the **memory area**?



Memory management requirements

Requirement

Retrievability

Protection

Description

Who is the **owner** of the **memory area**?

Protect the memory areas of the different processes against unwanted interference (read/write).



Memory management requirements

Requirement

Retrievability

Protection

Efficient usage

Description

Who is the **owner** of the **memory area**?

Protect the memory areas of the different processes against unwanted interference (read/write).

Memory is (always) too **small**.



Memory management requirements

Requirement

Retrievability

Protection

Efficient usage

Sharing

Description

Who is the **owner** of the **memory area**?

Protect the memory areas of the different processes against unwanted interference (read/write).

Memory is (always) too **small**.

Share memory between processes (**shared memory**).



Memory management requirements

Requirement

Retrievability

Protection

Efficient usage

Sharing

Logical organisation

Description

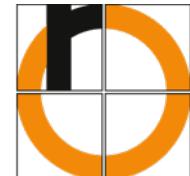
Who is the **owner** of the **memory area**?

Protect the memory areas of the different processes against unwanted interference (read/write).

Memory is (always) too **small**.

Share memory between processes (**shared memory**).

A process wants to see its memory as a **sequence of bytes**.



Memory management requirements

Requirement

Retrievability

Protection

Efficient usage

Sharing

Logical organisation

Transparency

Description

Who is the **owner** of the **memory area**?

Protect the memory areas of the different processes against unwanted interference (read/write).

Memory is (always) too **small**.

Share memory between processes (**shared memory**).

A process wants to see its memory as a **sequence of bytes**.

The **programmer doesn't know** at programming time **how much memory** on the target computer **exists**.
The memory management should transparently manage this.



Intro

Towards modern memory management...



Caching

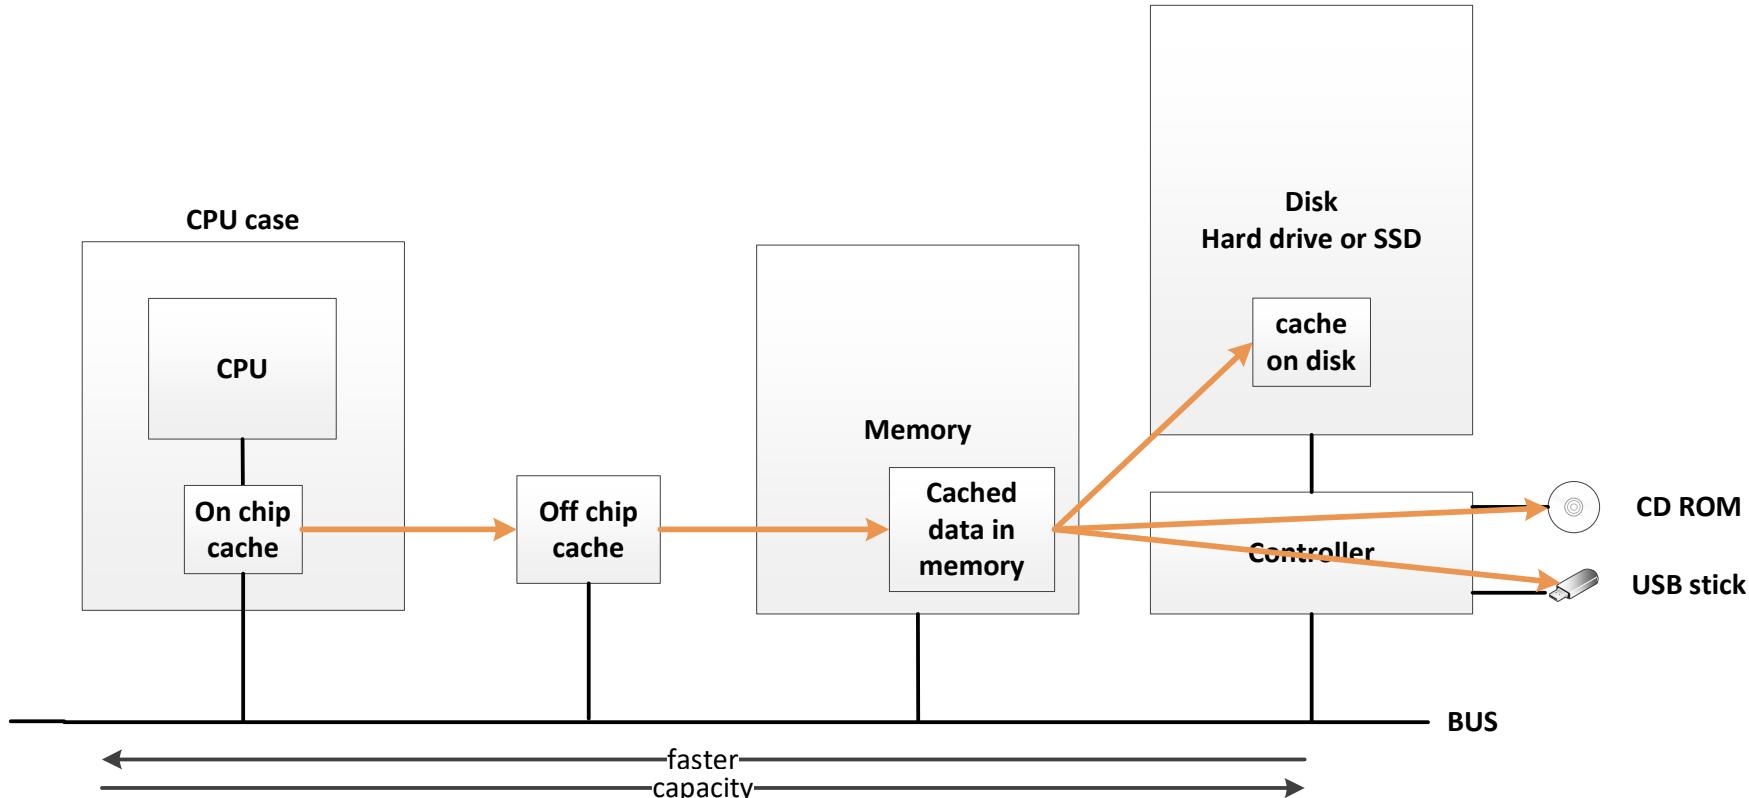
Problem

The **CPU** is much faster than the memory and all the other peripherals (e.g. network).



Caching

Idea Introduce caches (data buffers) on various places.





Caching terms

Hit

A hit occurs on repeated access of the same memory address.

Fault

Occurs on the first access. The cache is too small to buffer all data.

Locality

Caching terms

Hit

A hit occurs on repeated access of the same memory address.

Fault

Occurs on the first access. The cache is too small to buffer all data.

Locality

Caching terms

Hit

A hit occurs on repeated access of the same memory address.

Fault

Occurs on the first access. The cache is too small to buffer all data.

Locality

- **temporal:** If a memory address is accessed, there is a high probability that it will be **accessed again** very soon.
- **spatial:** If a memory address is accessed, it is very likely that **surrounding addresses will be accessed**, too.

Caching terms

Hit

A hit occurs on repeated access of the same memory address.

Fault

Occurs on the first access. The cache is too small to buffer all data.

Locality

- **temporal:** If a memory address is accessed, there is a high probability that it will be **accessed again** very soon.
- **spatial:** If a memory address is accessed, it is very likely that **surrounding addresses will be accessed**, too.



Caching terms

Hit

A hit occurs on repeated access of the same memory address.

Fault

Occurs on the first access. The cache is too small to buffer all data.

Locality

- **temporal**: If a memory address is accessed, there is a high probability that it will be **accessed again** very soon.
- **spatial**: If a memory address is accessed, it is very likely that **surrounding addresses will be accessed**, too.

Caching examples

On chip cache

A small cache inside the CPU: L1, L2, and L3. The average access time appears faster.

Off chip cache

A cache outside of the CPU (e.g. a special PCIe device).

Disk cache

A cache inside the disk or inside the OS. Speeds up the access time to the data on disk.

Internet cache

An internet proxy (e.g. squid) placed in the local network. If everyone accesses the same website or downloads the same file.



Caching examples

On chip cache

A small cache inside the CPU: L1, L2, and L3. The average access time appears faster.

Off chip cache

A cache outside of the CPU (e.g. a special PCIe device).

Disk cache

A cache inside the disk or inside the OS. Speeds up the access time to the data on disk.

Internet cache

An internet proxy (e.g. squid) placed in the local network. If everyone accesses the same website or downloads the same file.

Caching examples

On chip cache

A small cache inside the CPU: L1, L2, and L3. The average access time appears faster.

Off chip cache

A cache outside of the CPU (e.g. a special PCIe device).

Disk cache

A cache inside the disk or inside the OS. Speeds up the access time to the data on disk.

Internet cache

An internet proxy (e.g. squid) placed in the local network. If everyone accesses the same website or downloads the same file.

Caching examples

On chip cache

A small cache inside the CPU: L1, L2, and L3. The average access time appears faster.

Off chip cache

A cache outside of the CPU (e.g. a special PCIe device).

Disk cache

A cache inside the disk or inside the OS. Speeds up the access time to the data on disk.

Internet cache

An internet proxy (e.g. squid) placed in the local network. If everyone accesses the same website or downloads the same file.



Questions?

All right? \Rightarrow



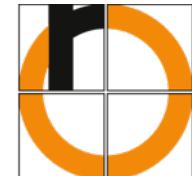
Question? \Rightarrow



and use **chat**

or

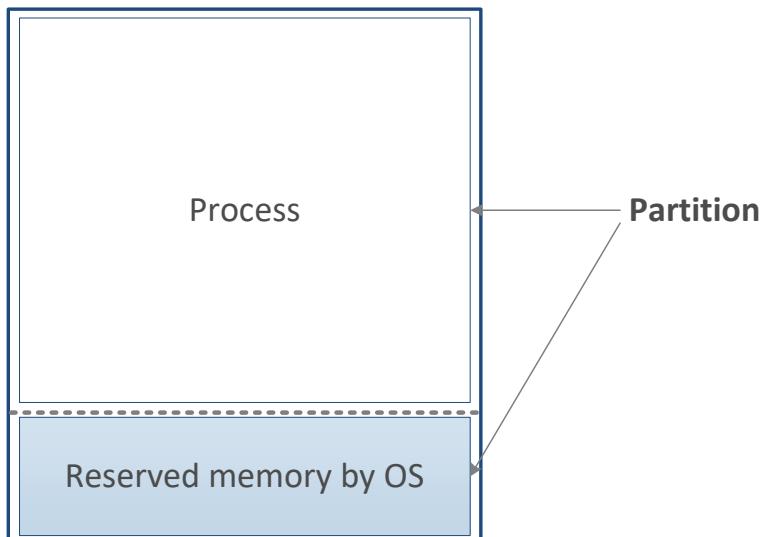
*speak after I
ask you to*



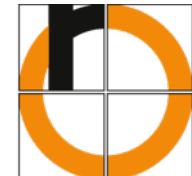
Partitioning

Single process mode

Main memory



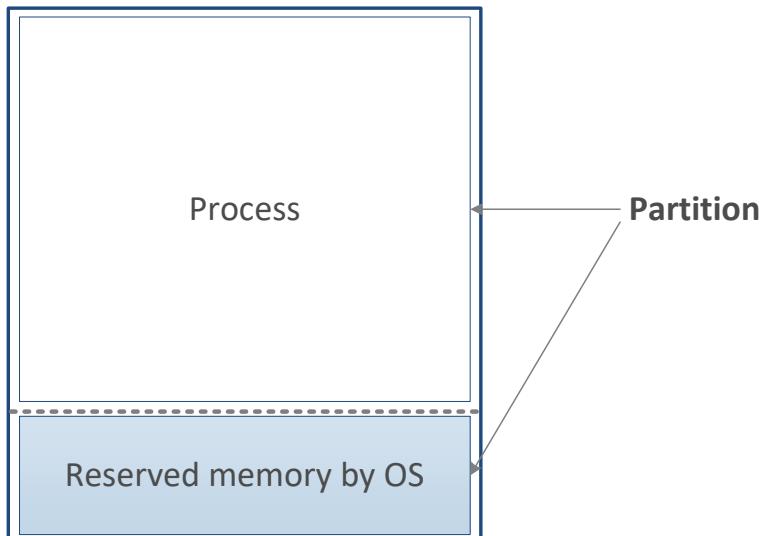
The main memory has to be divided into the OS part and the process part. If the process needs more memory than is physically available—it can't run.



Partitioning

Single process mode

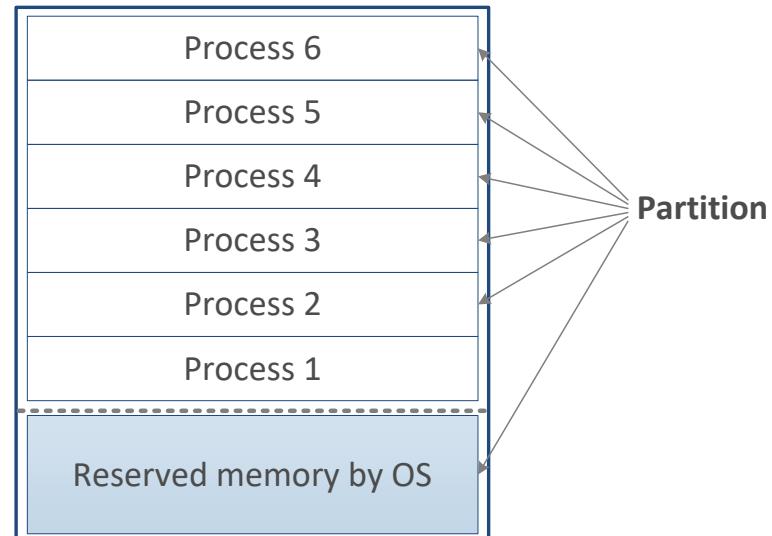
Main memory



The main memory has to be divided into the OS part and the process part. If the process needs more memory than is physically available—it can't run.

Multi process mode

Main memory



The main memory is divided into the OS part and the rest is partitioned between the processes.



Problem 1: memory protection

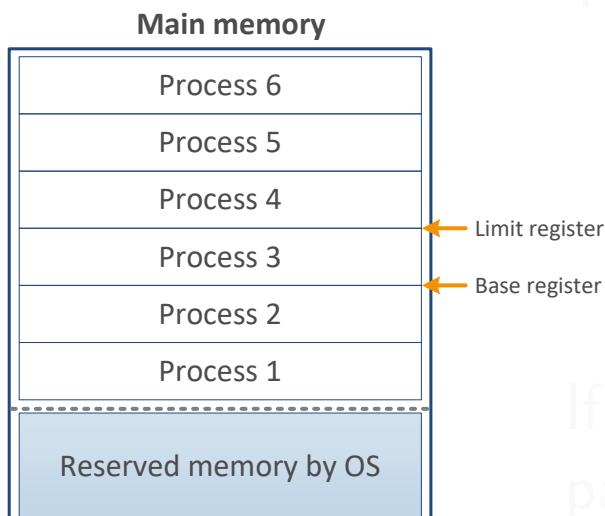
How to detect that a process accesses memory that it does not own?



Problem 1: memory protection

How to detect that a process accesses memory that it does not own?

The hardware offers two additional registers:



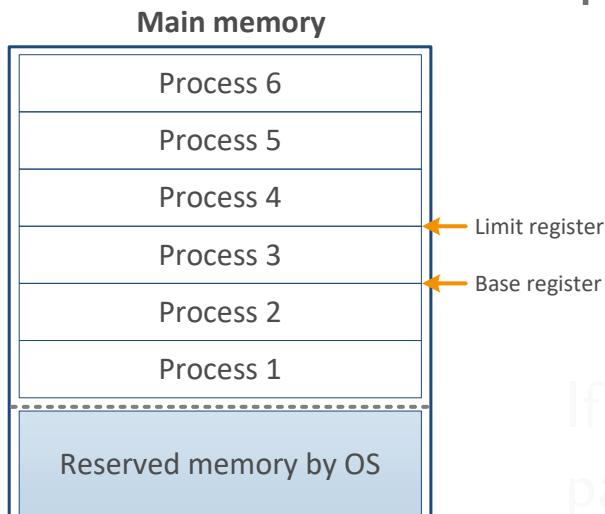
If the process access an address outside the partition, it is interrupted: the OS does than the error handling.



Problem 1: memory protection

How to detect that a process accesses memory that it does not own?

The hardware offers two additional registers:



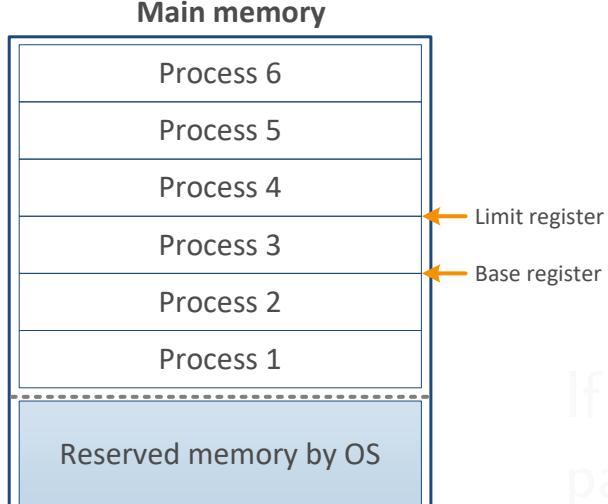
If the process access an address outside the partition, it is interrupted: the OS does than the error handling.



Problem 1: memory protection

How to detect that a process accesses memory that it does not own?

The hardware offers two additional registers:



■ **Base register:** start address of memory partition

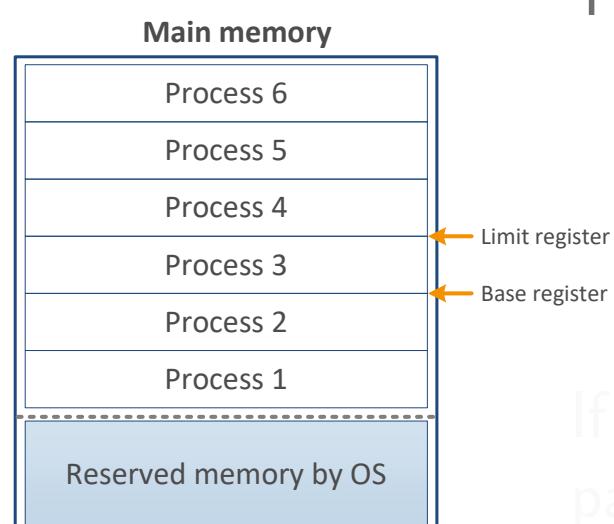
■ **Limit register:** end address of memory partition

If the process access an address outside the partition, it is interrupted: the OS does than the error handling.



Problem 1: memory protection

How to detect that a process accesses memory that it does not own?

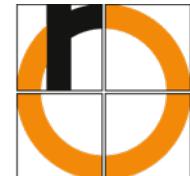


The hardware offers two additional registers:

- **Base register:** start address of memory partition

- **Limit register:** end address of memory partition

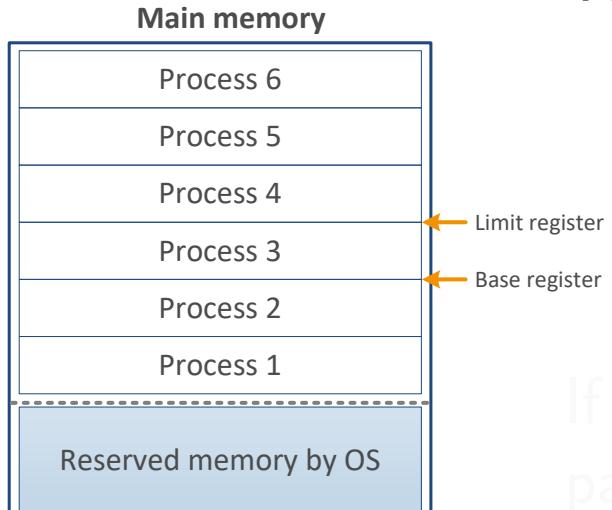
If the process access an address outside the partition, it is interrupted: the OS does than the error handling.



Problem 1: memory protection

How to detect that a process accesses memory that it does not own?

The hardware offers two additional registers:

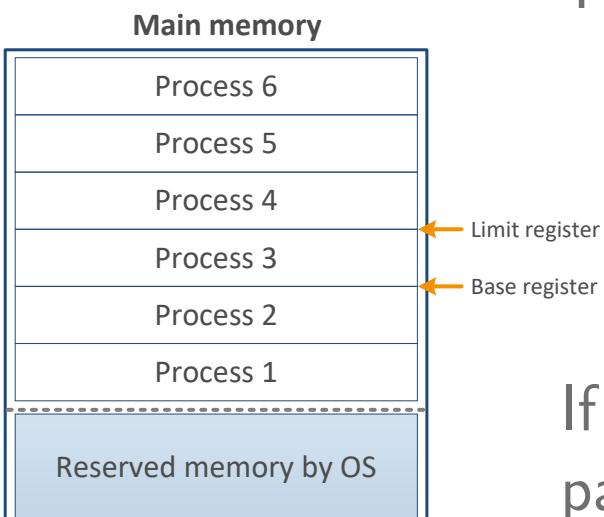


- **Base register:** start address of memory partition
- **Limit register:** end address of memory partition

If the process access an address outside the partition, it is interrupted: the OS does than the error handling.

Problem 1: memory protection

How to detect that a process accesses memory that it does not own?



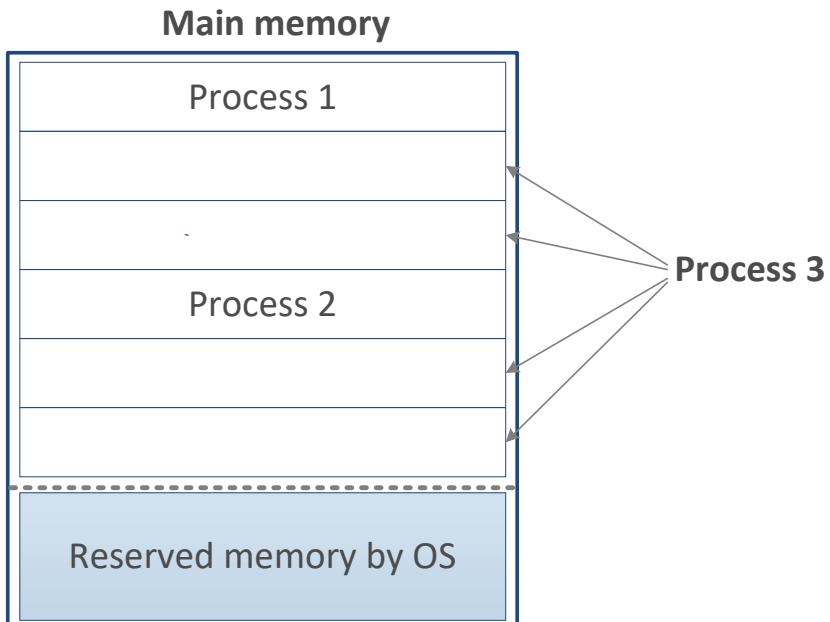
The hardware offers two additional registers:

- **Base register:** start address of memory partition
- **Limit register:** end address of memory partition

If the process access an address outside the partition, it is interrupted: the OS does than the error handling.

Problem 2: position dependent code

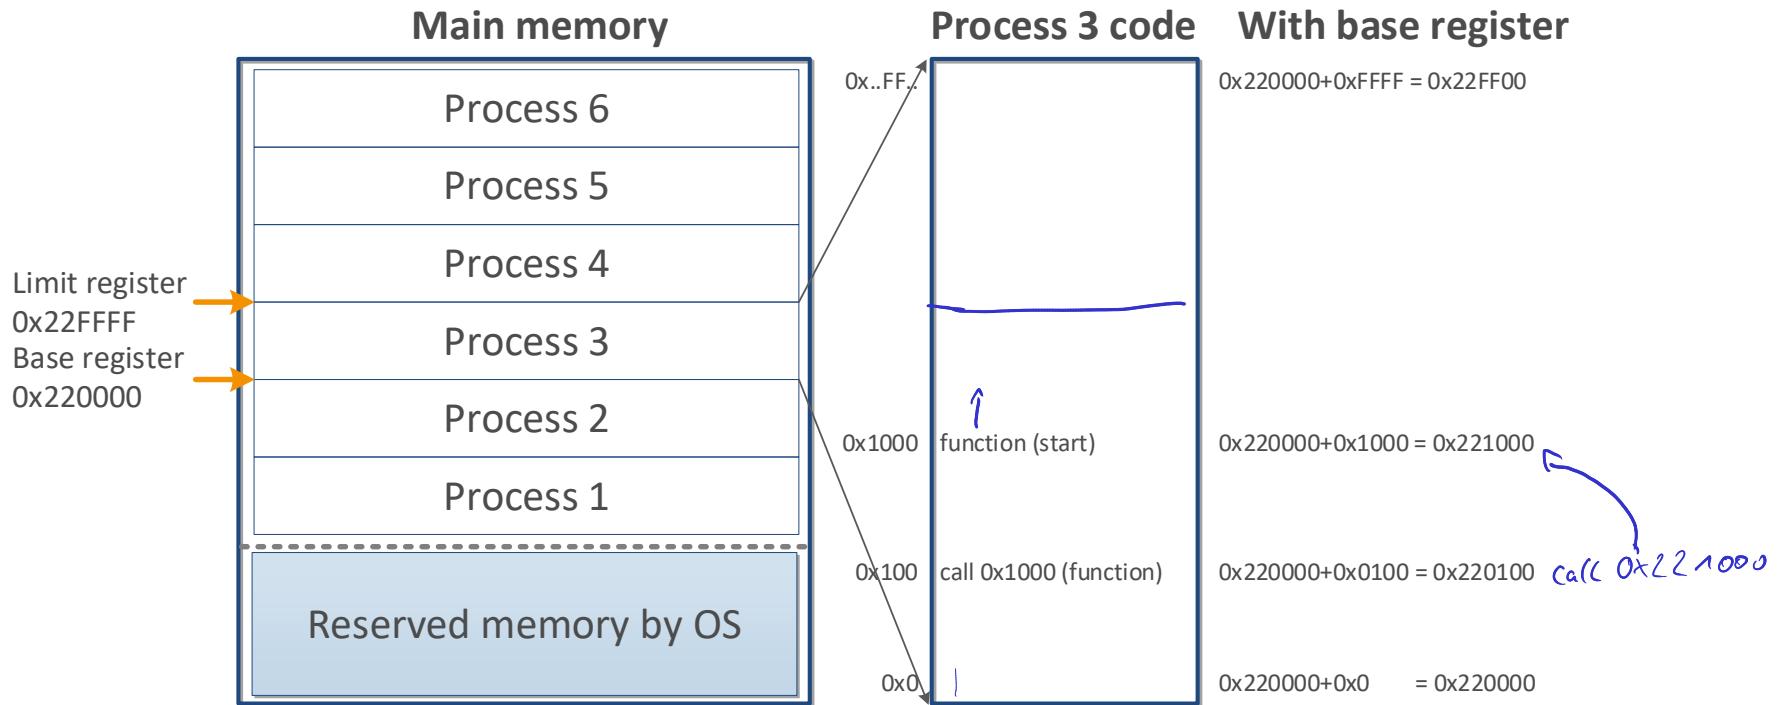
Where should the program be loaded? What happens with the addresses inside the process?





Problem 2: position independent code

Build the program as it would start at address 0. On every memory access, the base register is added to the address.





Questions?

All right? \Rightarrow



Question? \Rightarrow



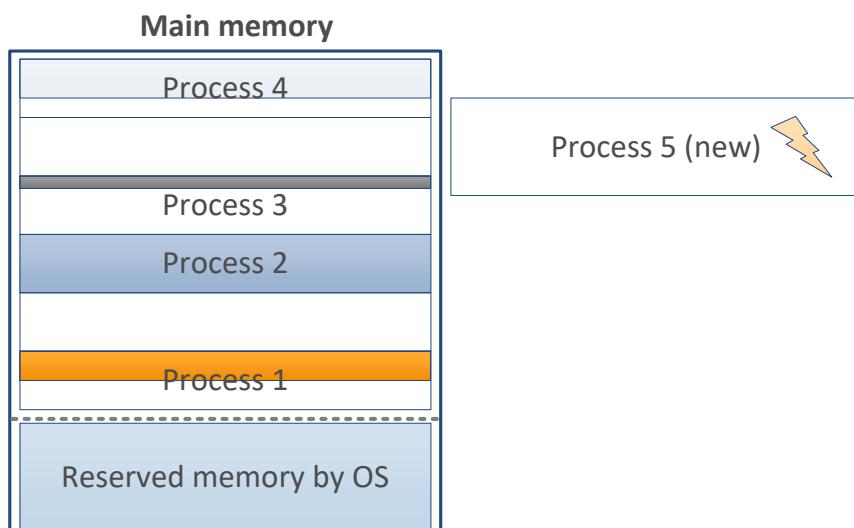
and use **chat**

or

*speak after I
ask you to*



Fragmentation



Internal fragmentation

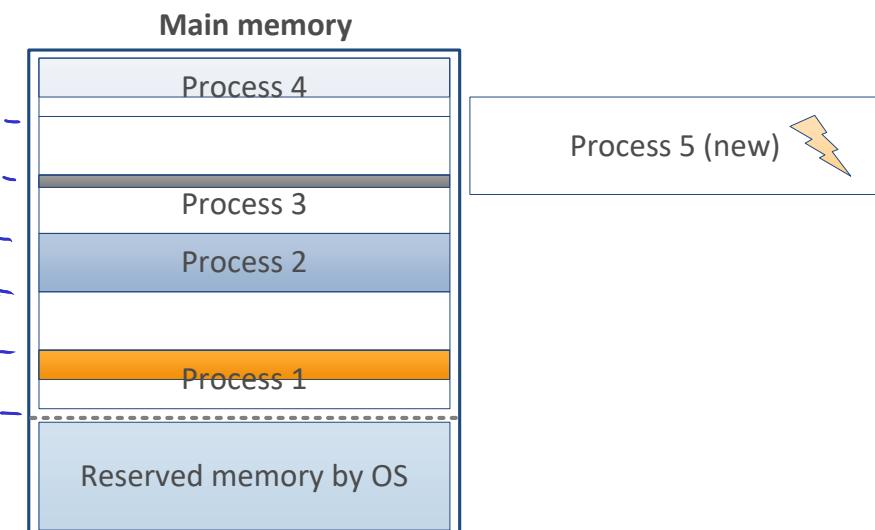
Partitions are allocated by processes, but not fully used. May be solved by variable partition size.

External fragmentation

A lot of free partitions can't be used, because they are **too small** for some processes. May be solved by dynamically move the partitions (base and limit register required). But this will take some time.



Fragmentation



Internal fragmentation

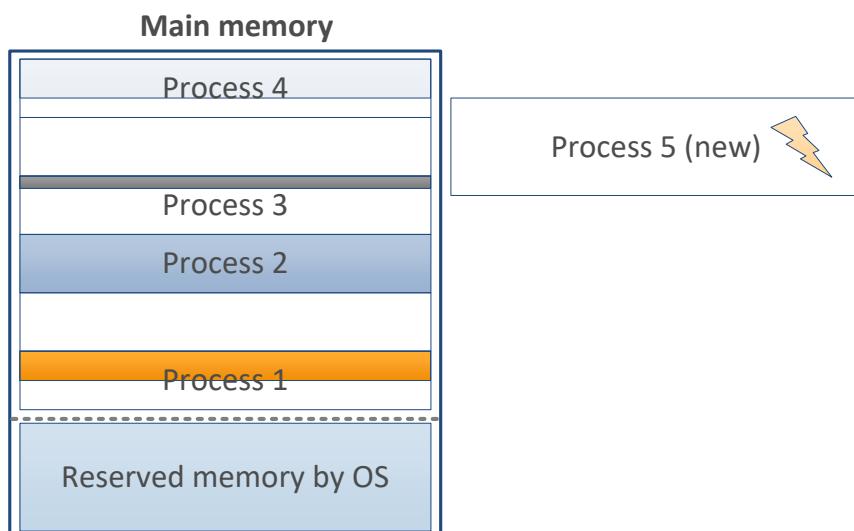
Partitions are allocated by processes, but **not fully used**. May be solved by variable partition size.

External fragmentation

A lot of free partitions can't be used, because they are **too small** for some processes. May be solved by dynamically move the partitions (base and limit register required). But this will take some time.



Fragmentation



Internal fragmentation

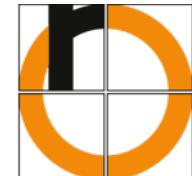
Partitions are allocated by processes, but **not fully used**. May be solved by variable partition size.

External fragmentation

A lot of free partitions can't be used, because they are **too small** for some processes. May be solved by dynamically move the partitions (base and limit register required). But this will take some time.

Allocation strategies

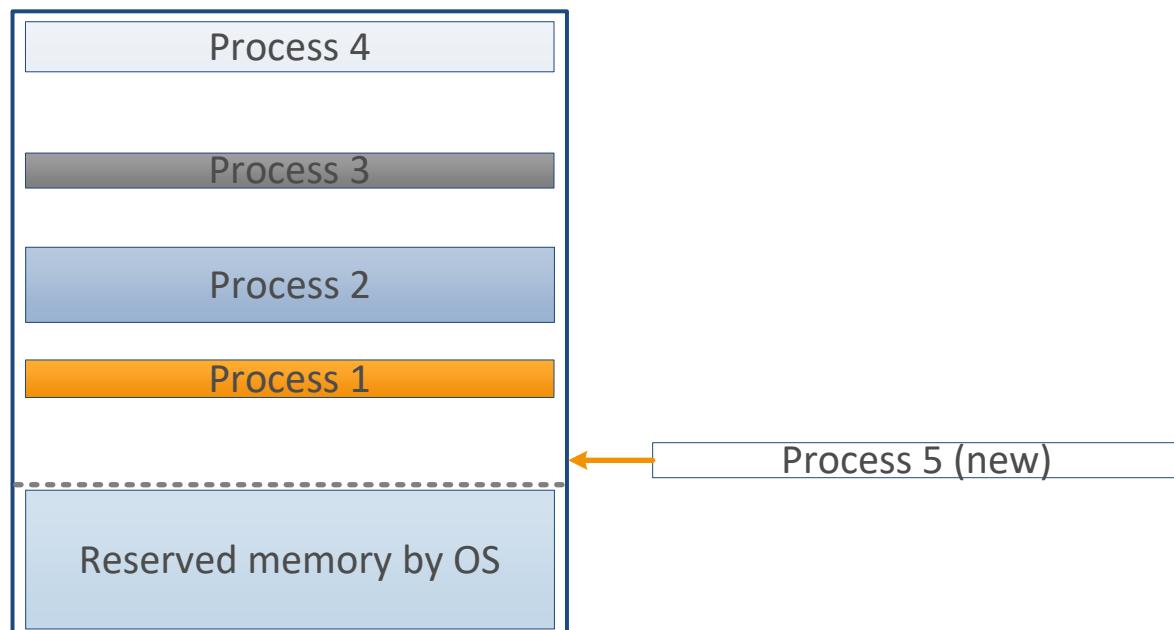
How to find the best place for the new partitions in the memory?

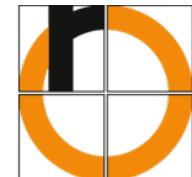


First fit

Starts at the beginning of the main memory and takes the first memory area that fits.

Main memory

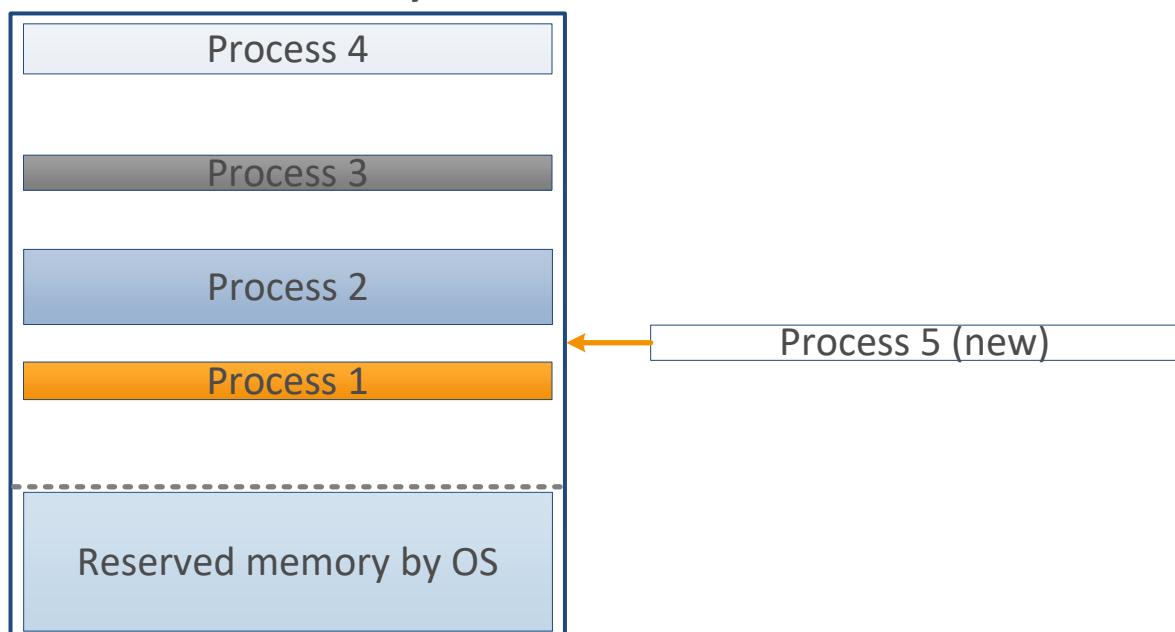


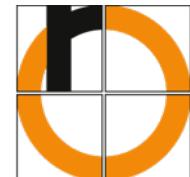


Best fit

Looks into all free memory areas and takes the memory area that fits best.

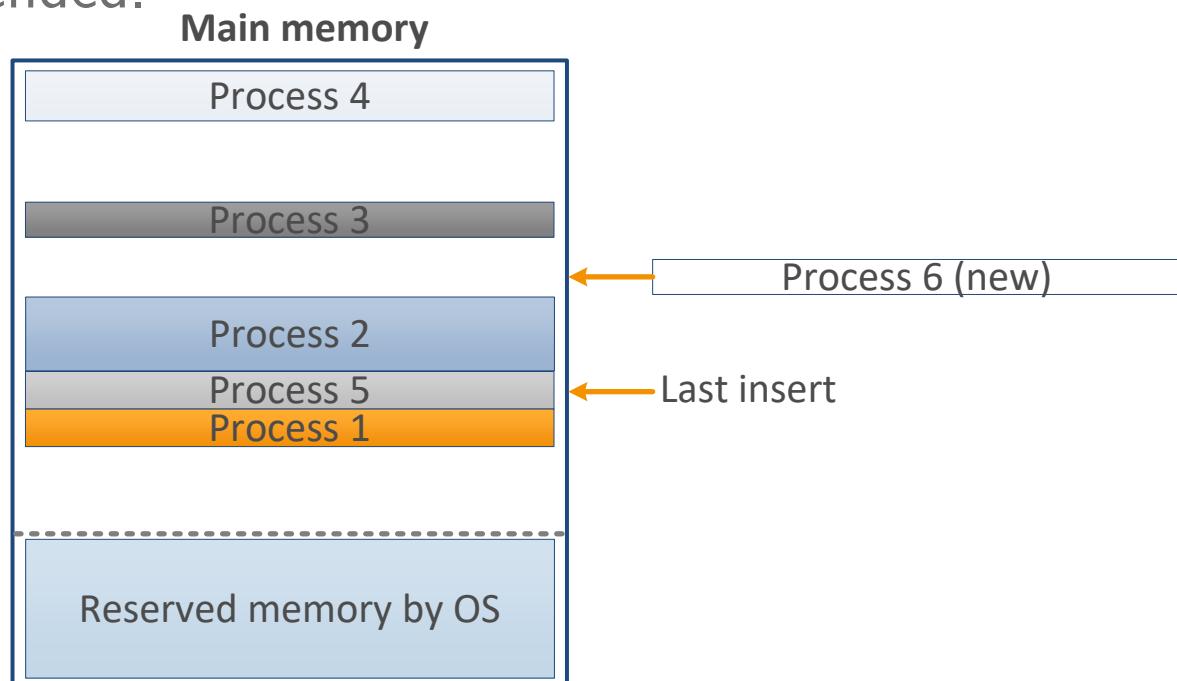
Main memory

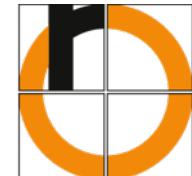




Next fit

The same as first fit, but it starts at the point, where the last search ended.

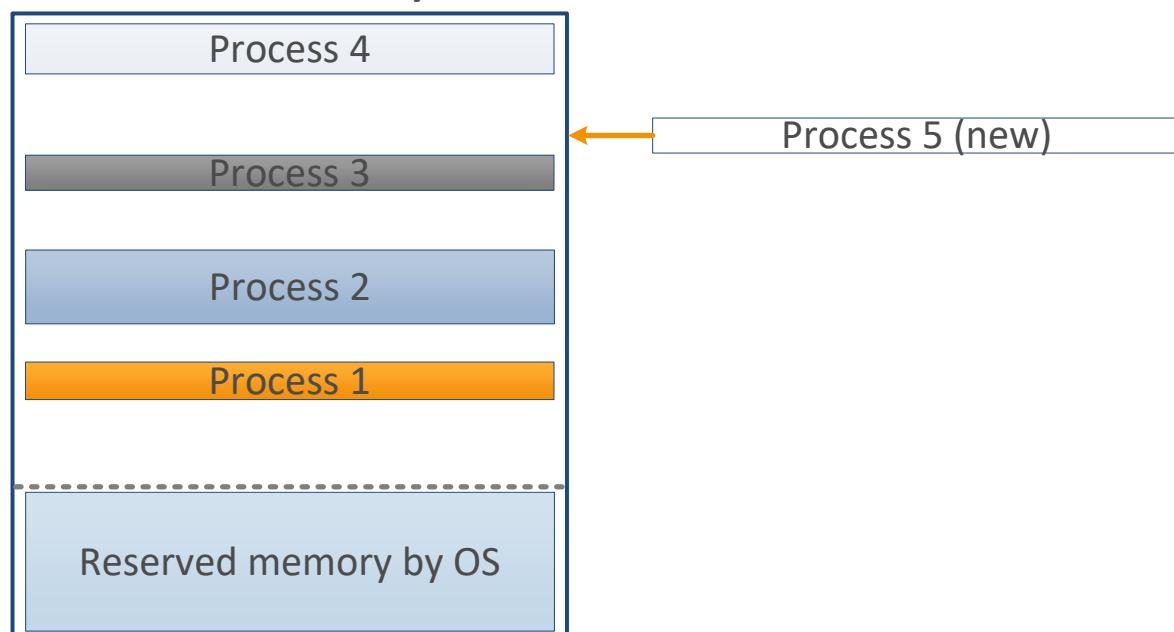




Worst fit

Looks into all free memory areas and takes the memory area that fits worst.

Main memory





Allocation strategies summary

Simulations show that generally best fit is worse than first and next fit.

Worst fit is the worst.

Allocation strategies summary

Simulations show that generally best fit is worse than first and next fit.

Worst fit is the worst.



Questions?

All right? \Rightarrow



Question? \Rightarrow



and use **chat**

or

*speak after I
ask you to*



Still existing problems

- The number of processes is limited by the number of partitions.
- The processes often don't use the memory allocated to them for a long time.
- A lot of processes are required to fully utilise the CPU.



Still existing problems

- The number of processes is limited by the number of partitions.
- The processes often don't use the memory allocated to them for a long time.
- A lot of processes are required to fully utilise the CPU.

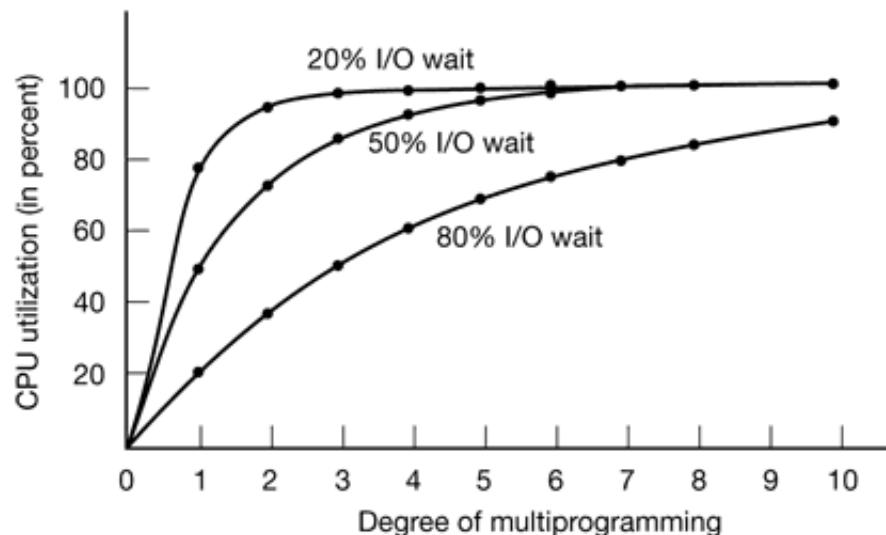


Still existing problems

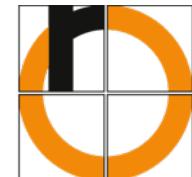
- The number of processes is limited by the number of partitions.
- The processes often don't use the memory allocated to them for a long time.
- A lot of processes are required to fully utilise the CPU.

Still existing problems

- The number of processes is limited by the number of partitions.
- The processes often don't use the memory allocated to them for a long time.
- A lot of processes are required to fully utilise the CPU.

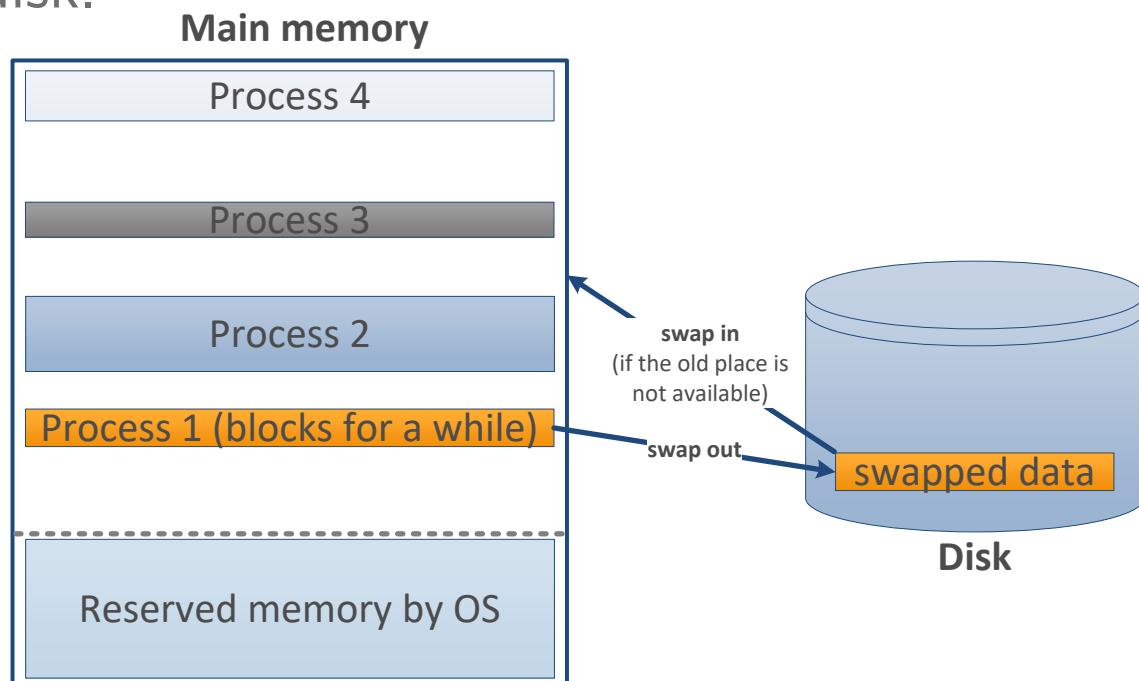


Source: Modern operating systems (Tanenbaum, second edition)

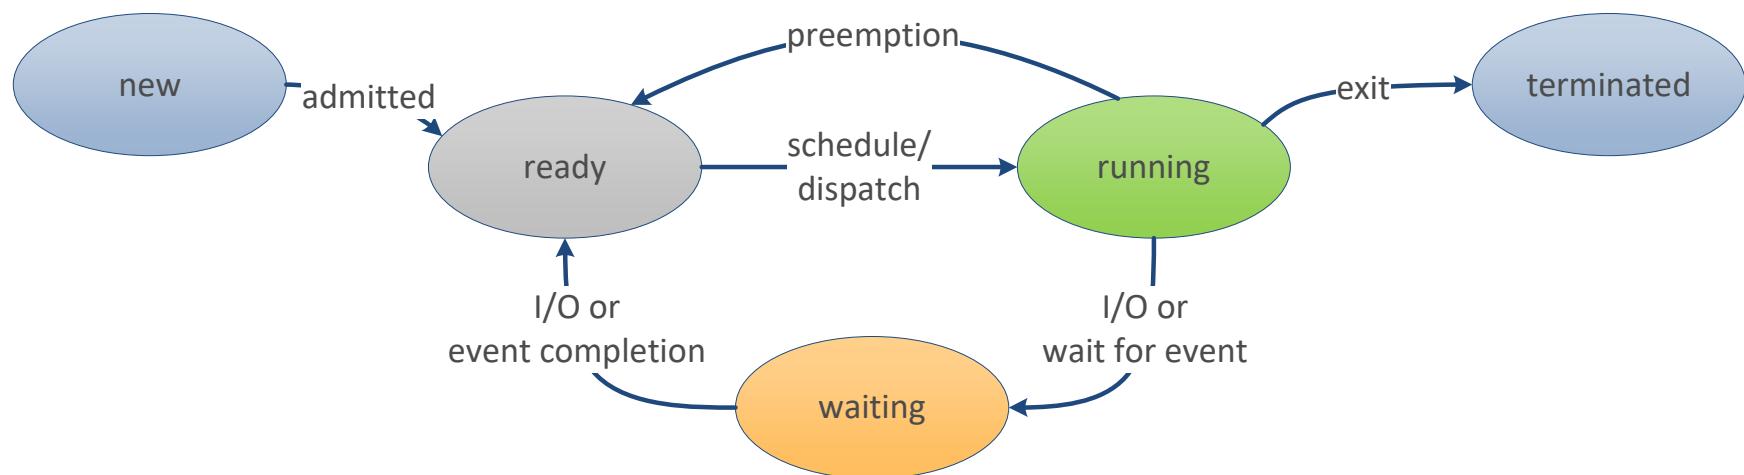


Swapping idea

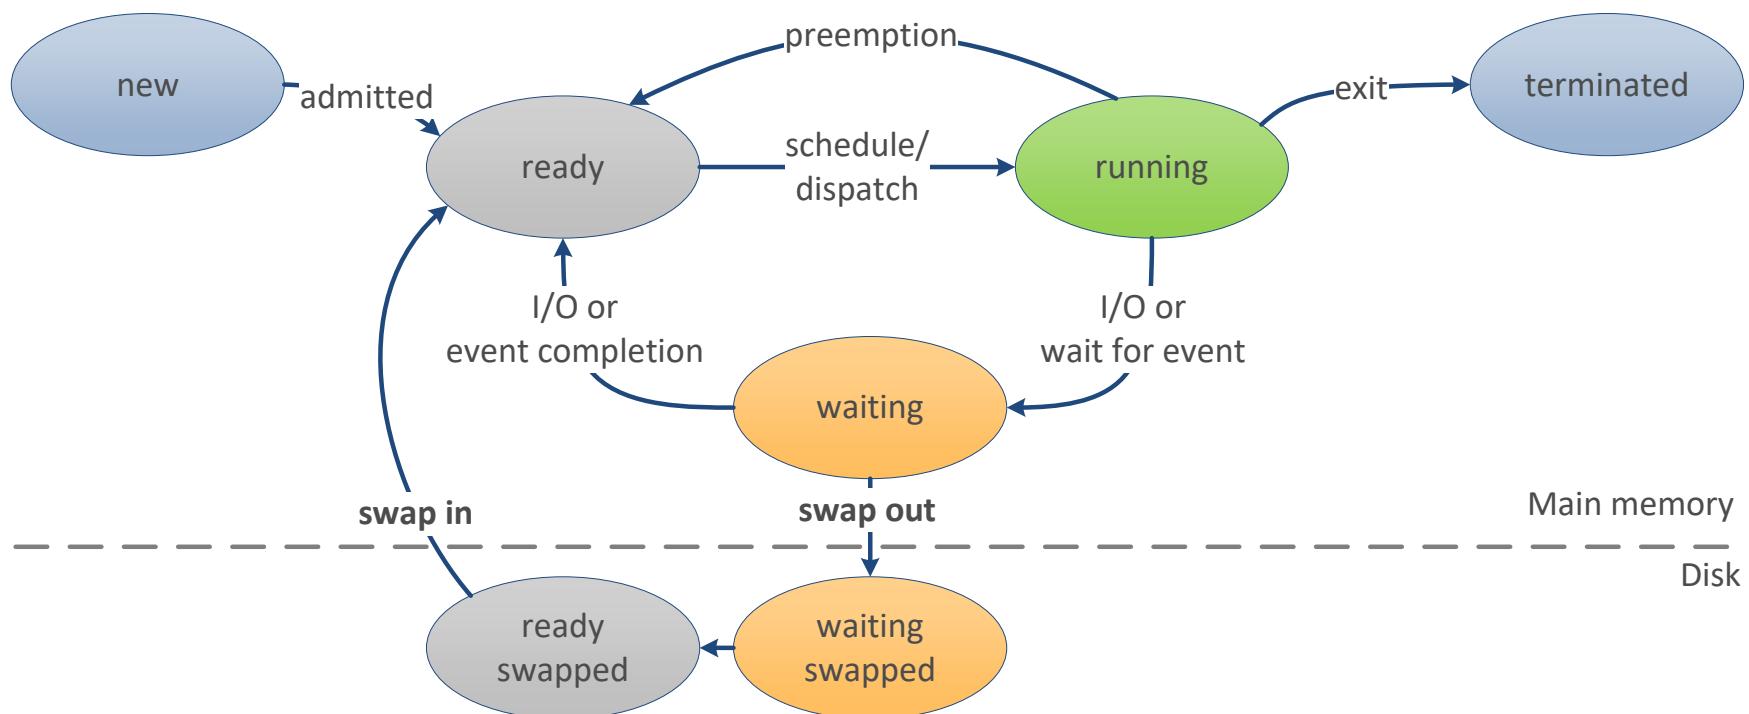
Swap the memory of a process that is not actively running to the disk.



Enhanced process states for swapping



Enhanced process states for swapping





Questions?

All right? \Rightarrow



Question? \Rightarrow



and use **chat**

or

*speak after I
ask you to*



Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.



Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.



Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.



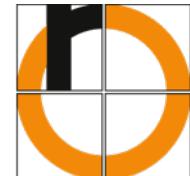
Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.



Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.

Where are we?

Achieved so far

The sum of the memory for all processes can be larger than the physically available memory.

Still existing problems

- Every process is limited by the physically available memory size.
- The processes often don't use the memory allocated to them for a long time.
- Problem with fragmentation still exists (dynamically move of partitions takes a long time).
- Swapping is time consuming. The advantage is often difficult to evaluate.



Virtual memory

Idea

Only load the required memory parts of a process into the memory.

Requirement

But neither the user nor the programmer should notice anything about it.

=> Virtual memory with virtual addresses!



Virtual memory

Idea

Only load the required memory parts of a process into the memory.

Requirement

But neither the user nor the programmer should notice anything about it.

=> Virtual memory with virtual addresses!



Virtual memory

Idea

Only load the required memory parts of a process into the memory.

Requirement

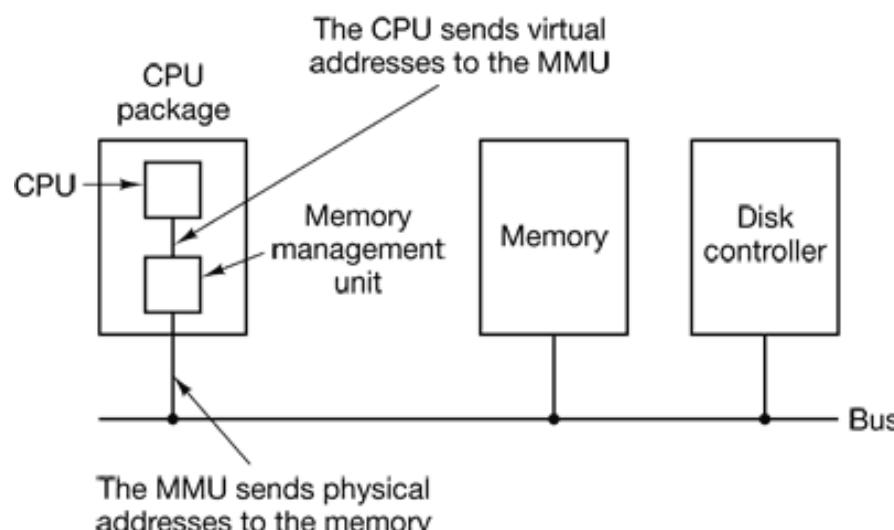
But neither the user nor the programmer should notice anything about it.

=> Virtual memory with virtual addresses!



Virtual address

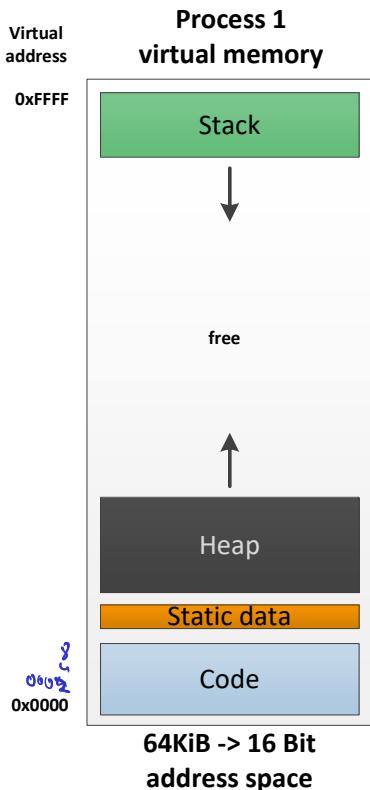
The MMU (a special hardware inside the CPU) supports the address calculation: virtual address -> real address.



Source: Modern operating systems (Tanenbaum, second edition)



Virtual address space

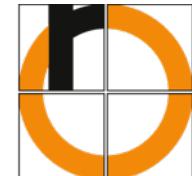


Virtual address space

Each process has its own virtual address space. The theoretically available memory.

Real address space

Physically available address space



Virtual address space



Virtual address space

Each process has its own virtual address space. The theoretically available memory.

Real address space

Physically available address space



Virtual address space

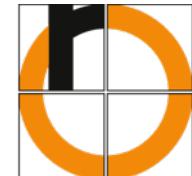


Virtual address space

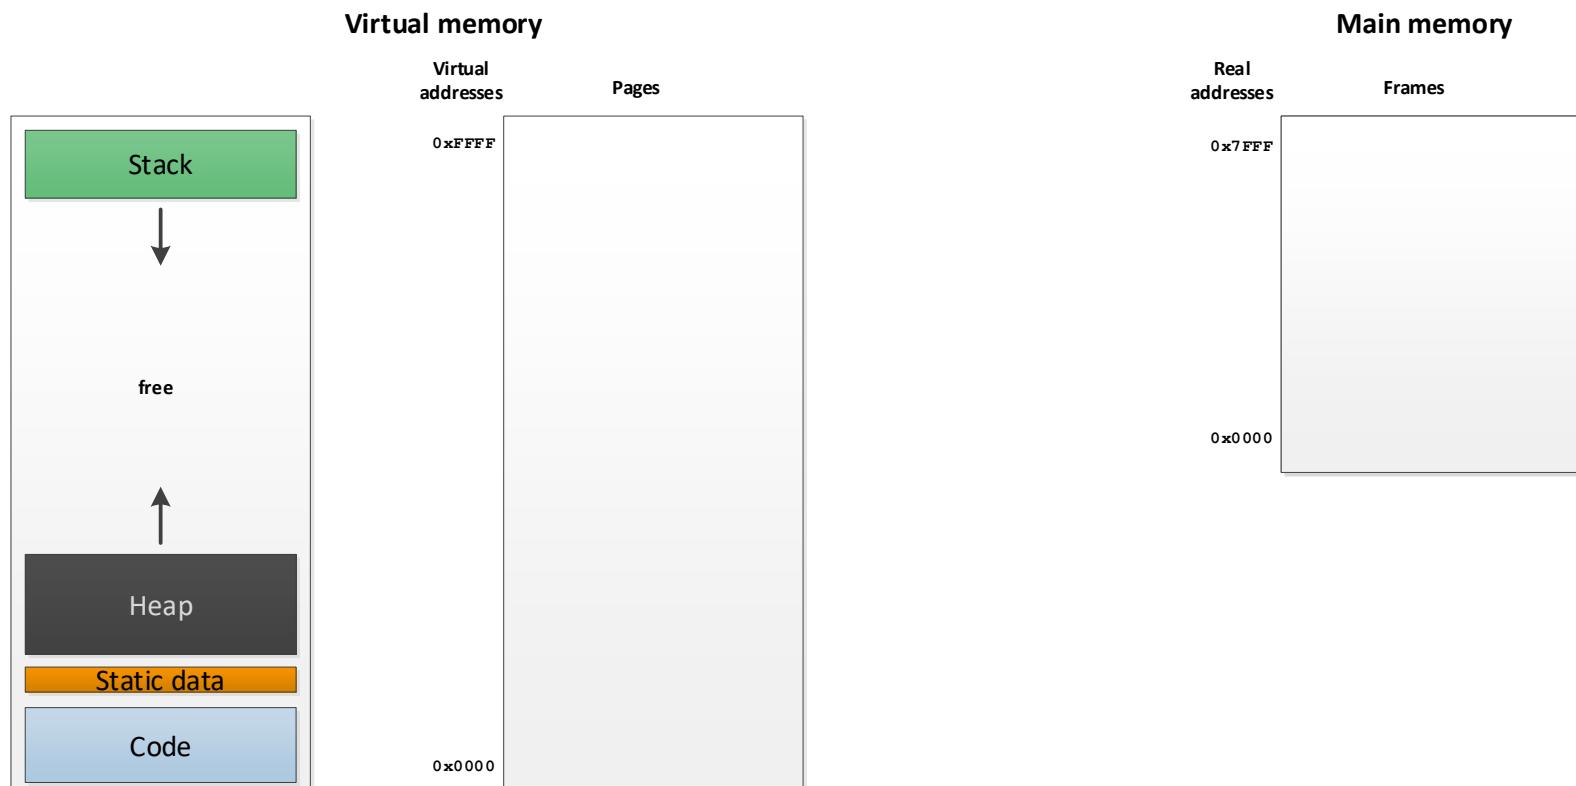
Each process has its own virtual address space. The theoretically available memory.

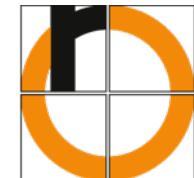
Real address space

Physically available address space

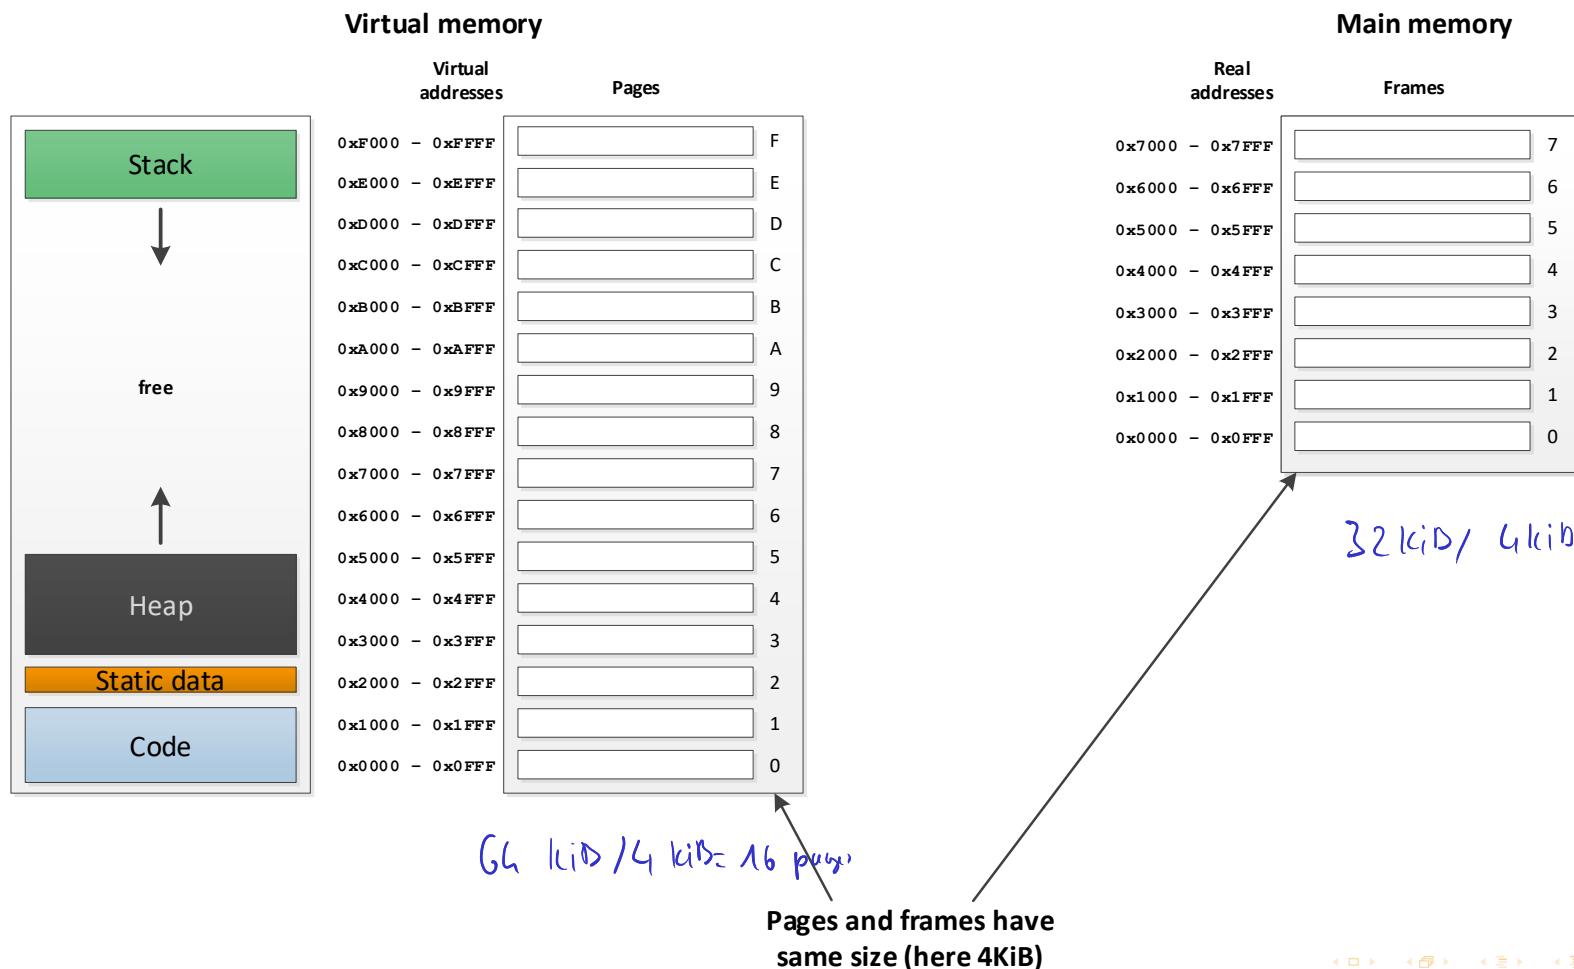


Pages and frames

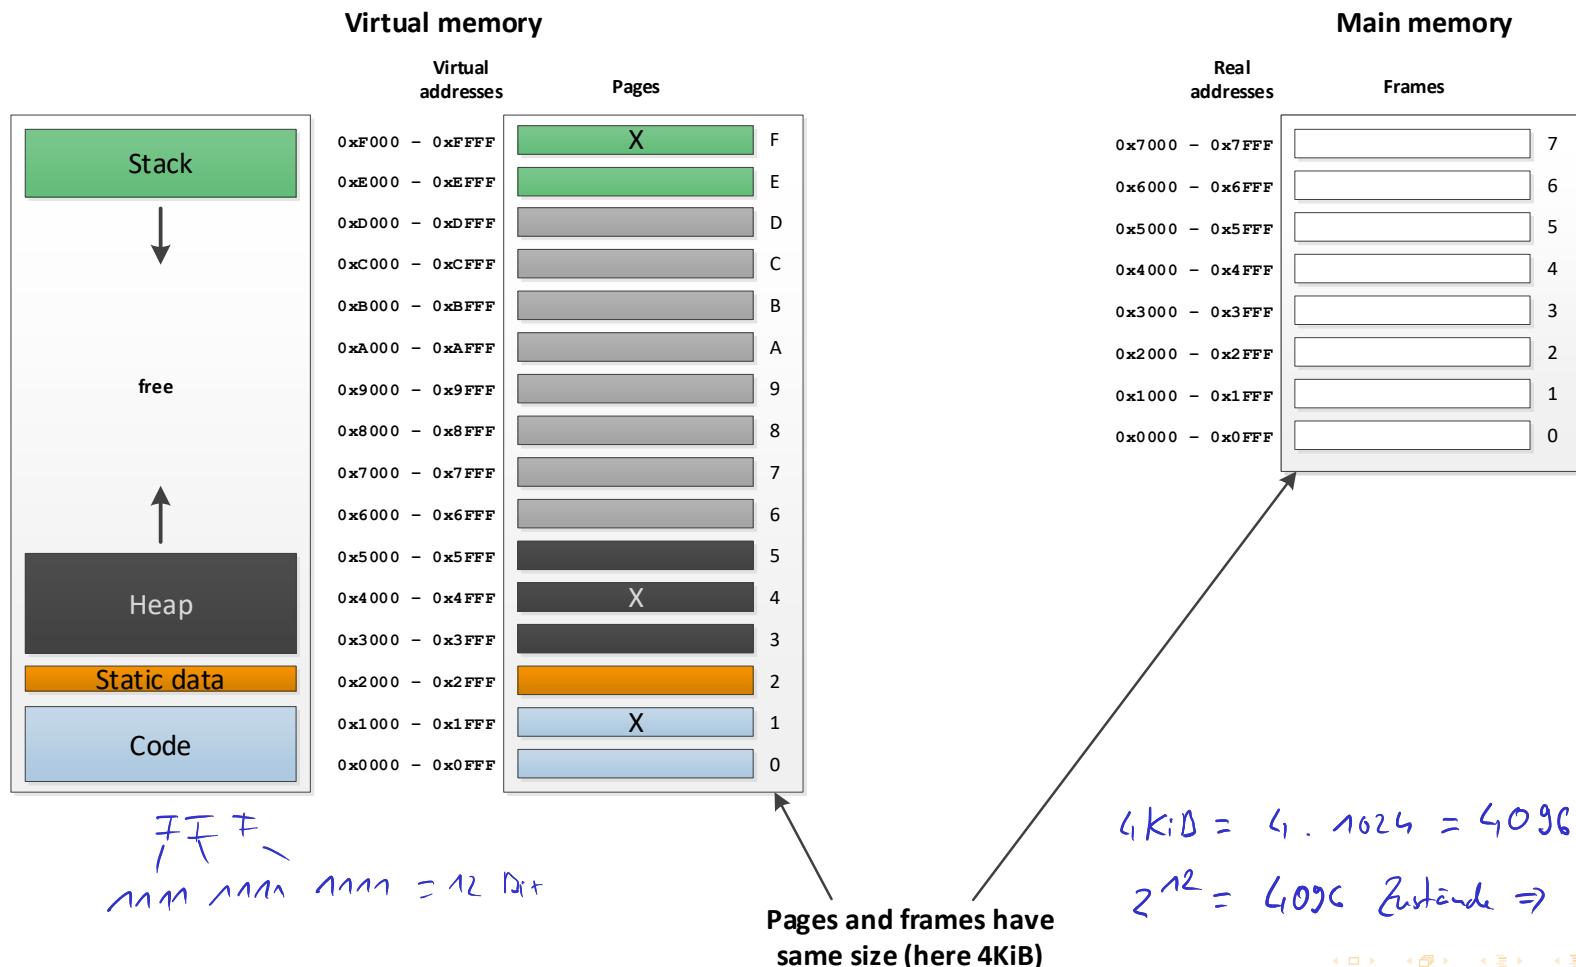




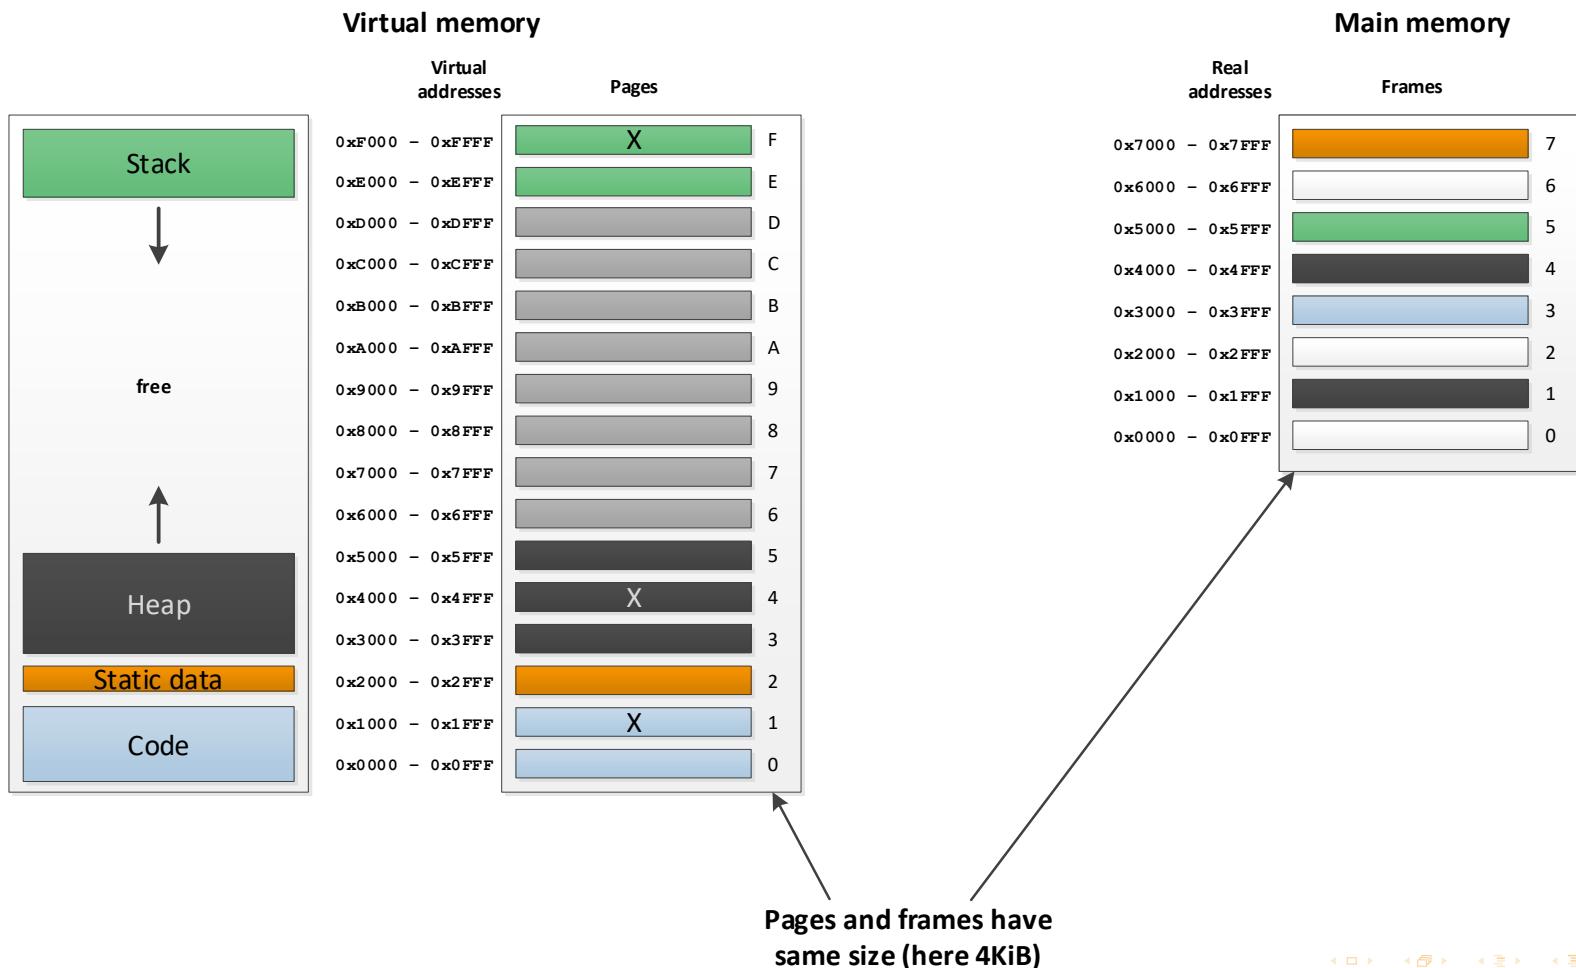
Pages and frames



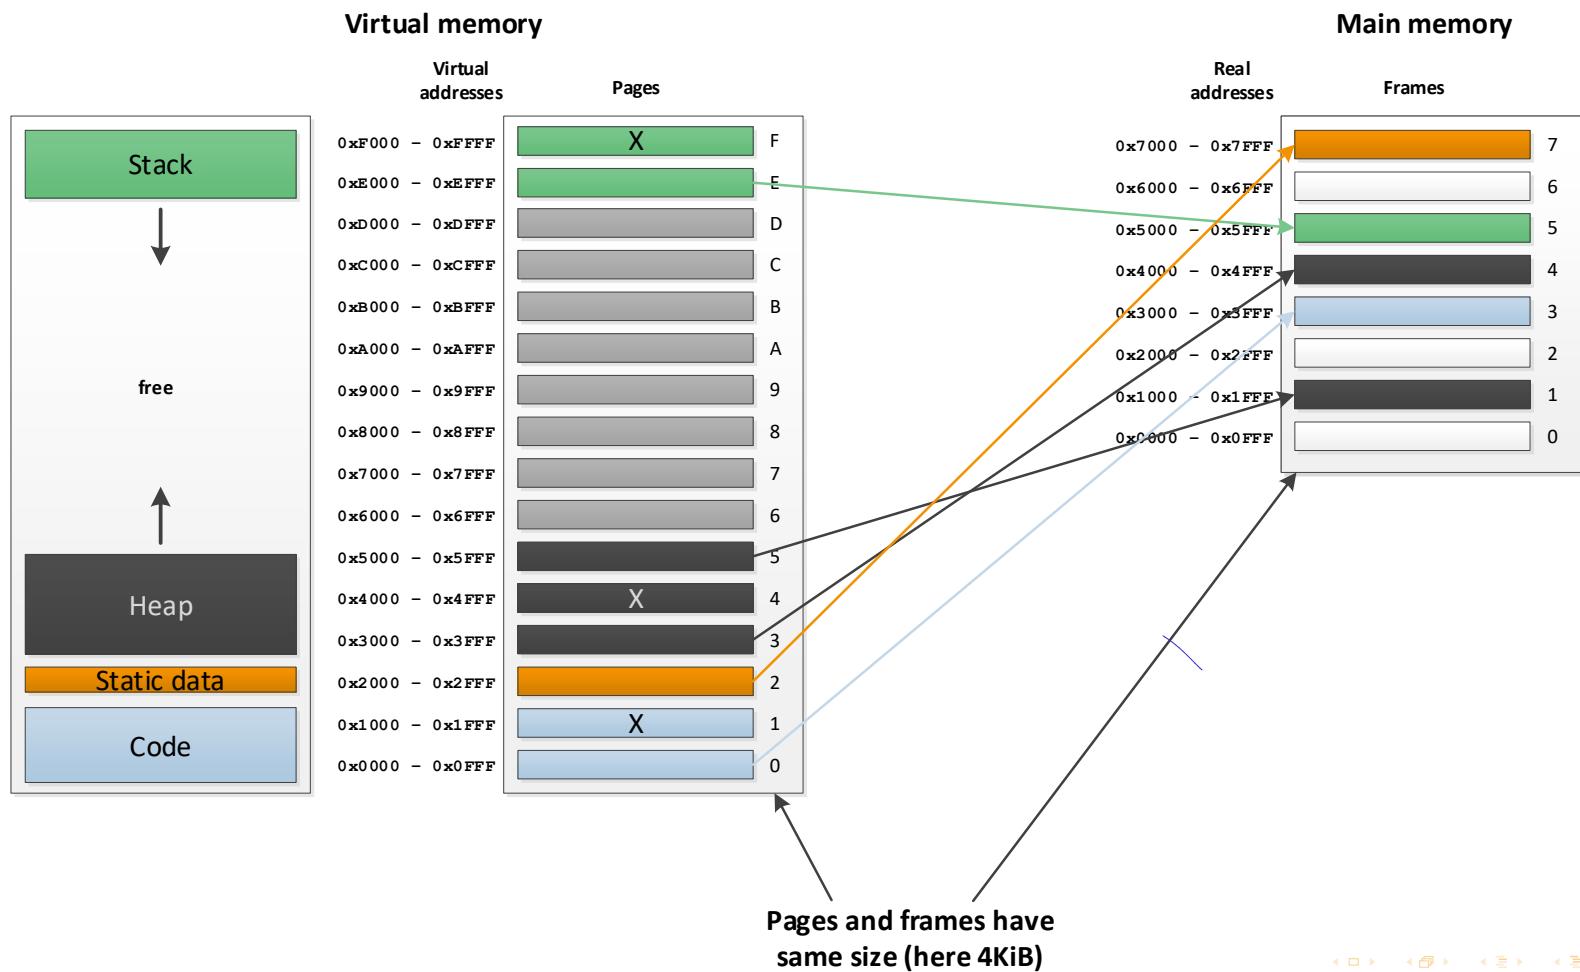
Pages and frames



Pages and frames



Pages and frames



Page table

Virtual memory

Virtual base addresses	Pages
0xF000	F
0xE000	E
0xD000	D
0xC000	C
0xB000	B
0xA000	A
0x9000	9
0x8000	8
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

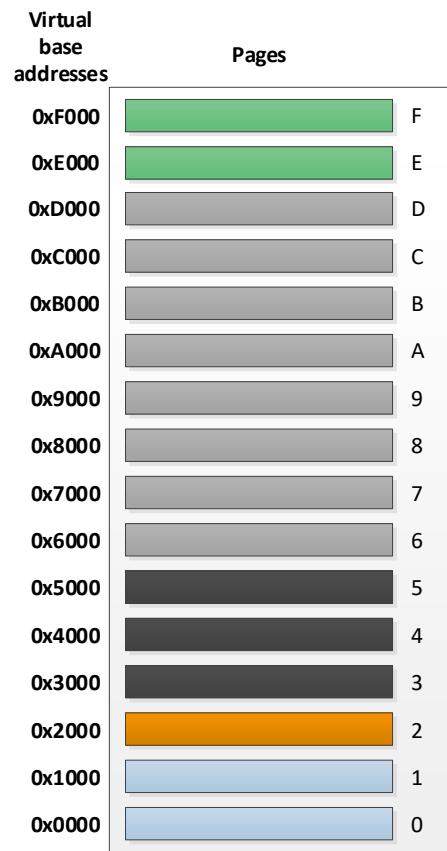
Each process has its own virtual memory

Main memory

Frame base addresses	Frames
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Page table

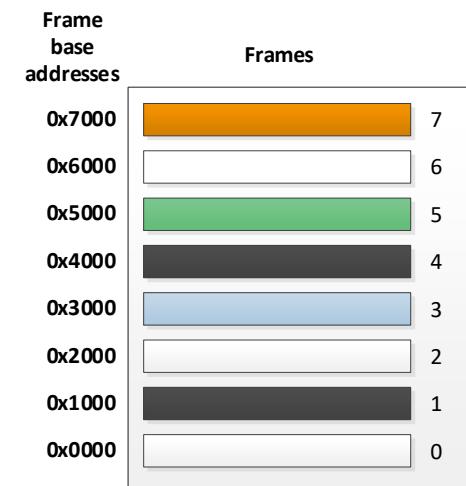
Virtual memory



Each process has its own virtual memory



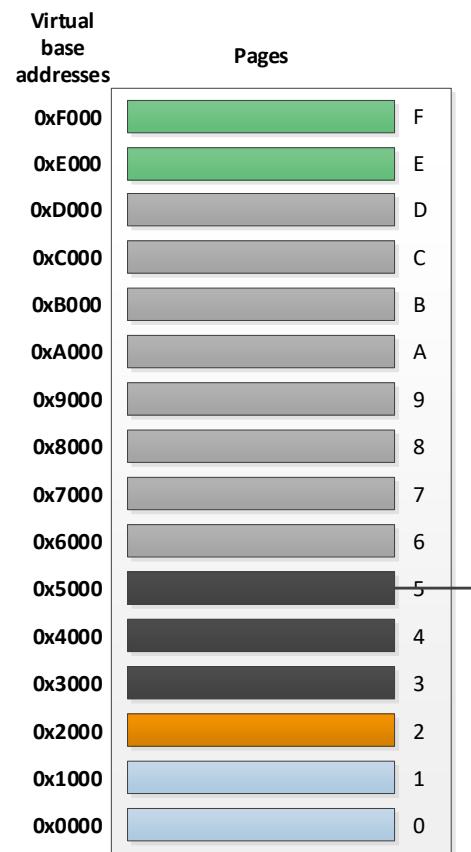
Each process has its own page table



0x5123

Page table

Virtual memory



Virtual addresses

0x 5 123

Real addresses

0x [] []

Page table

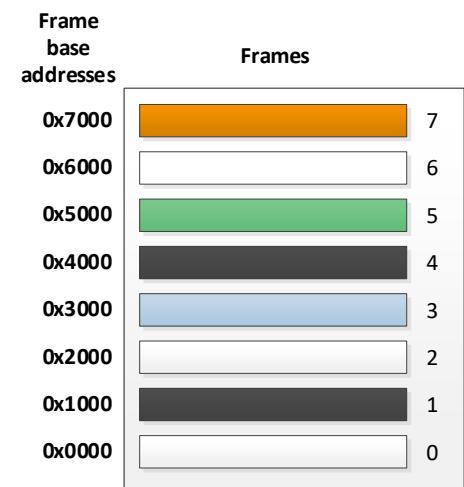
F	X (on disk)
E	
D	
C	
B	
A	
9	
8	
7	
6	
5	
4	X (on disk)
3	
2	X (on disk)
1	
0	

Each process has its own virtual memory



Technical University of Applied Sciences

Main memory

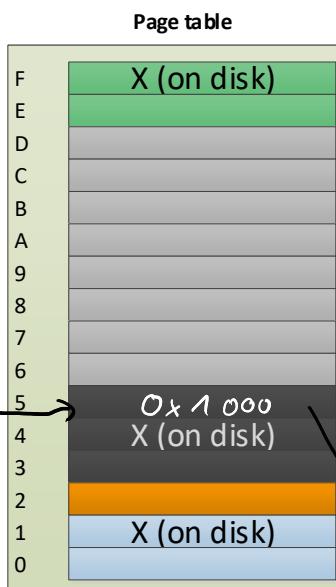
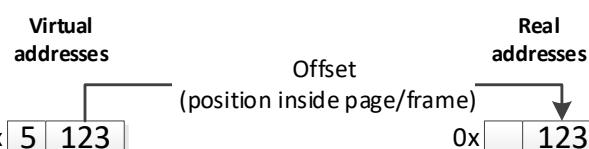


Page table

Virtual memory

Virtual base addresses	Pages
0xF000	F
0xE000	E
0xD000	D
0xC000	C
0xB000	B
0xA000	A
0x9000	9
0x8000	8
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Each process has its own virtual memory



Each process has its own page table

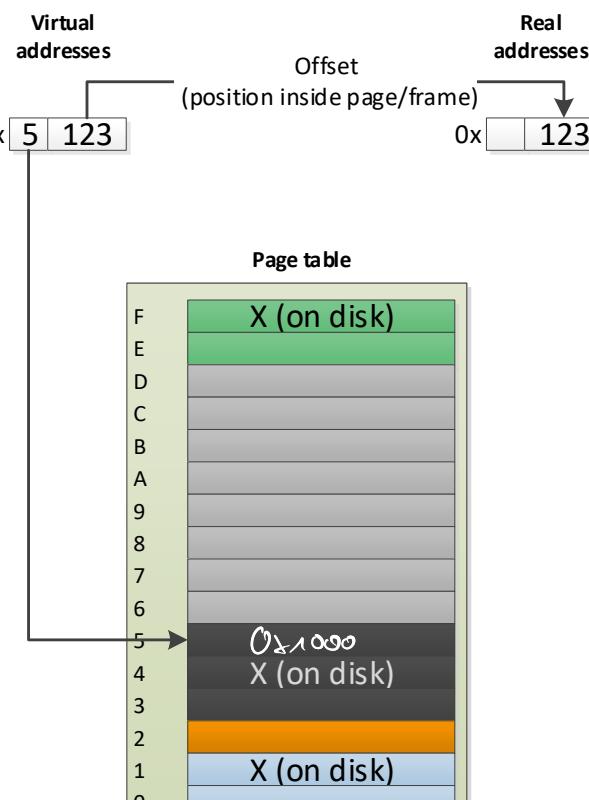
Frame base addresses	Frames
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Page table

Virtual memory

Virtual base addresses	Pages
0xF000	F
0xE000	E
0xD000	D
0xC000	C
0xB000	B
0xA000	A
0x9000	9
0x8000	8
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Each process has its own virtual memory



Each process has its own page table

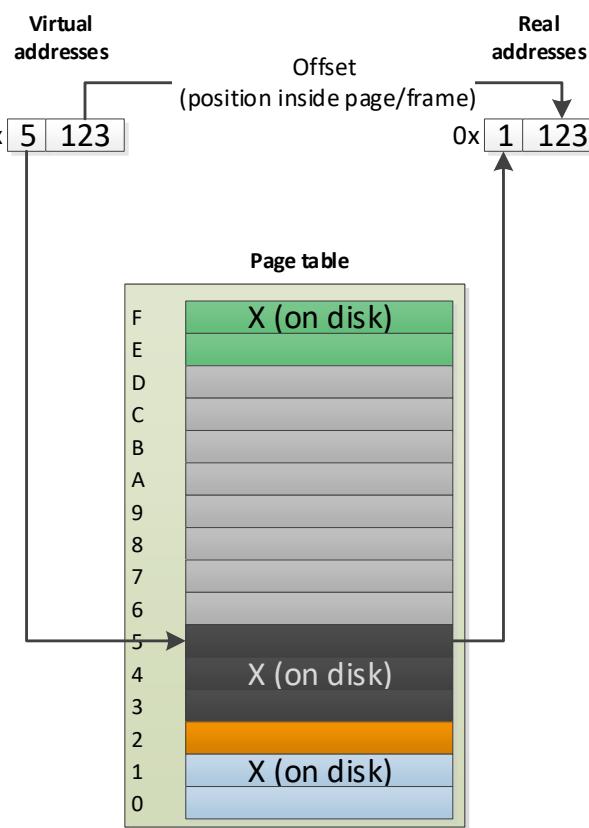
Frame base addresses	Frames
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Page table

Virtual memory

Virtual base addresses	Pages
0xF000	F
0xE000	E
0xD000	D
0xC000	C
0xB000	B
0xA000	A
0x9000	9
0x8000	8
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Each process has its own virtual memory

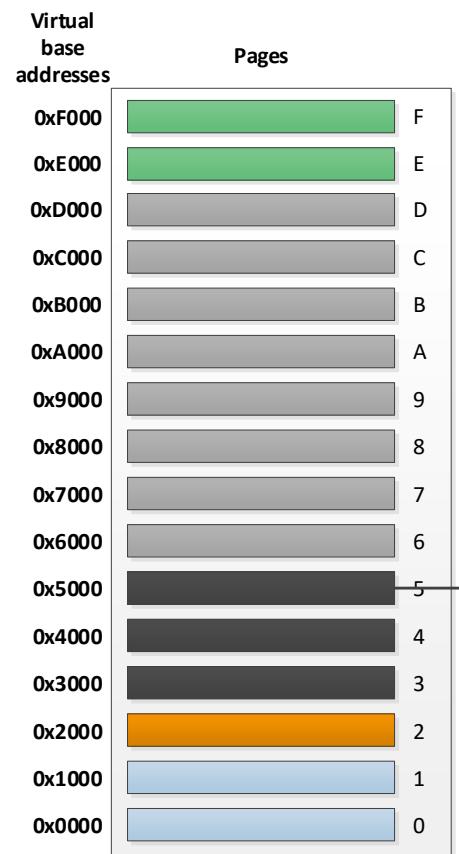


Each process has its own page table

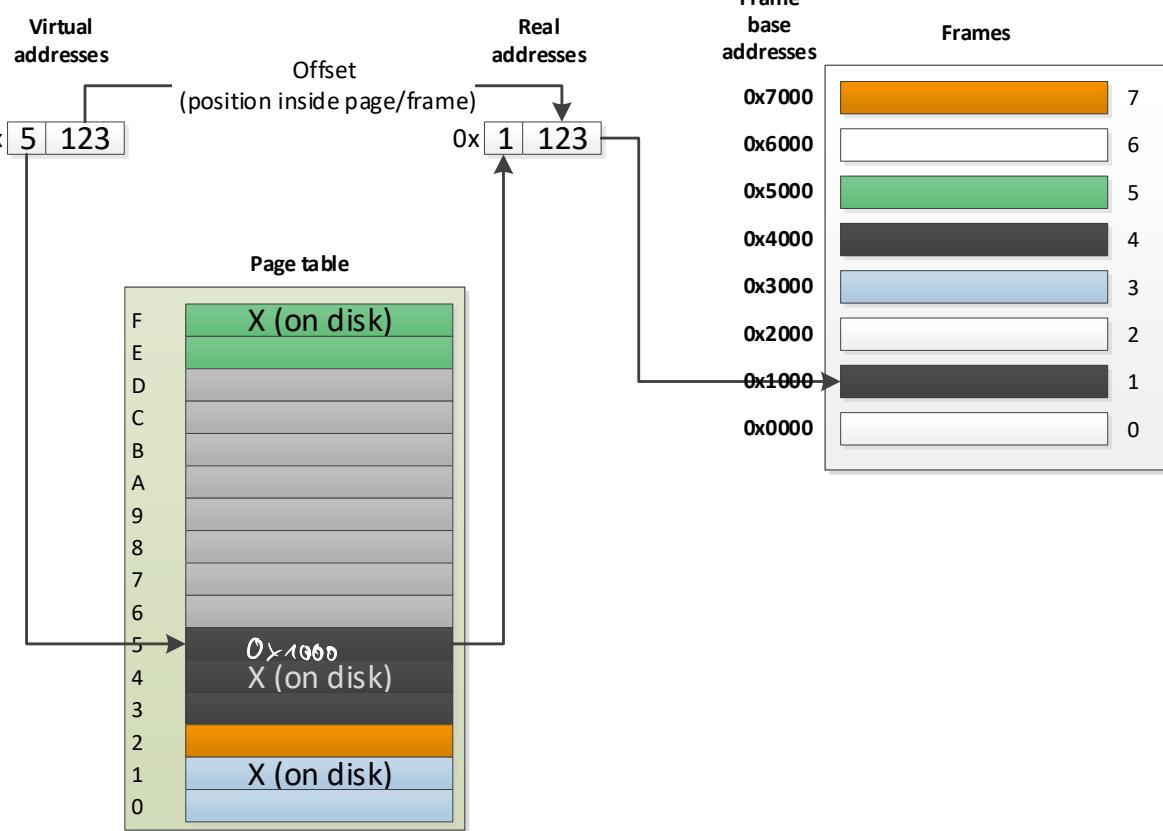
Frame base addresses	Frames
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Page table

Virtual memory



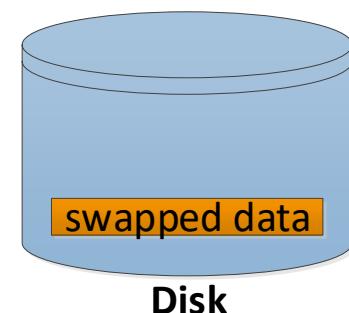
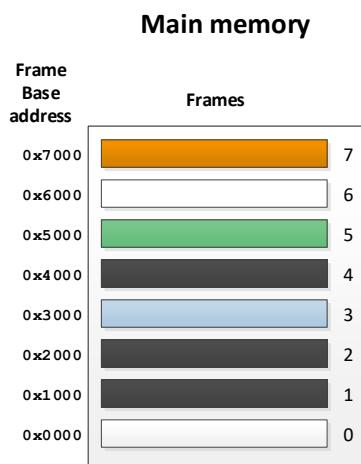
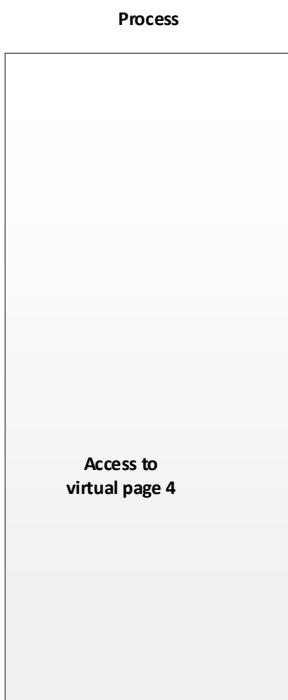
Each process has its own virtual memory



Each process has its own page table

Page fault

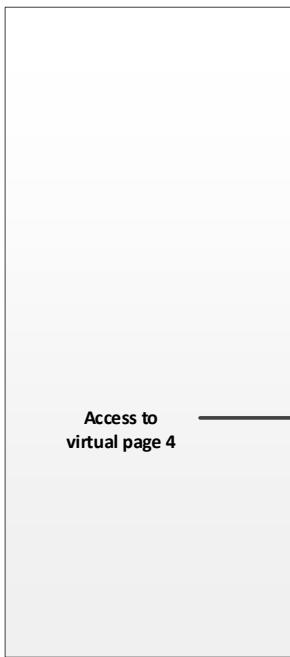
Virtual memory



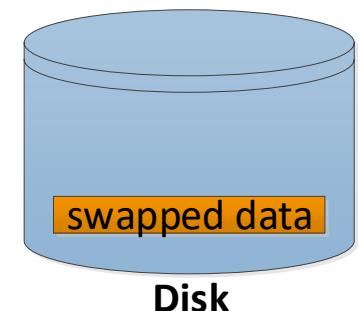
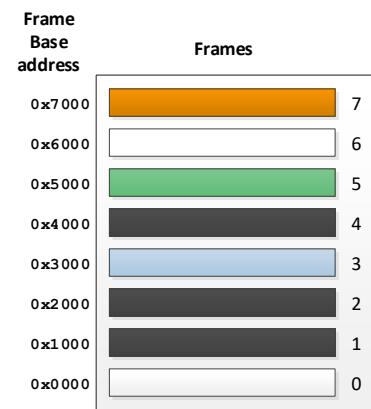
Page fault

Virtual memory

Process

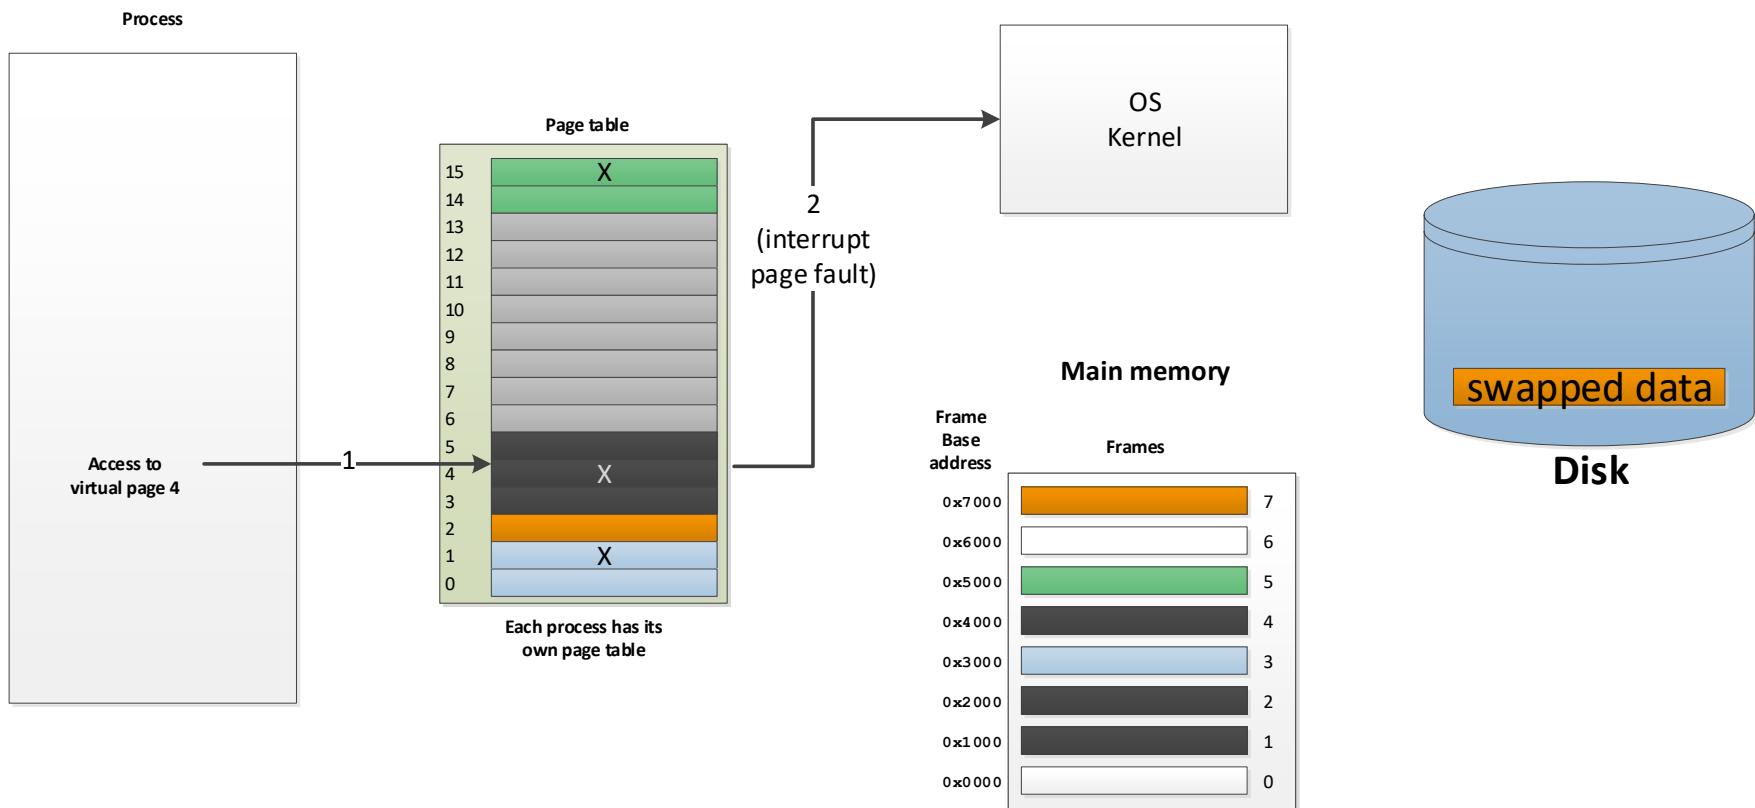


Main memory



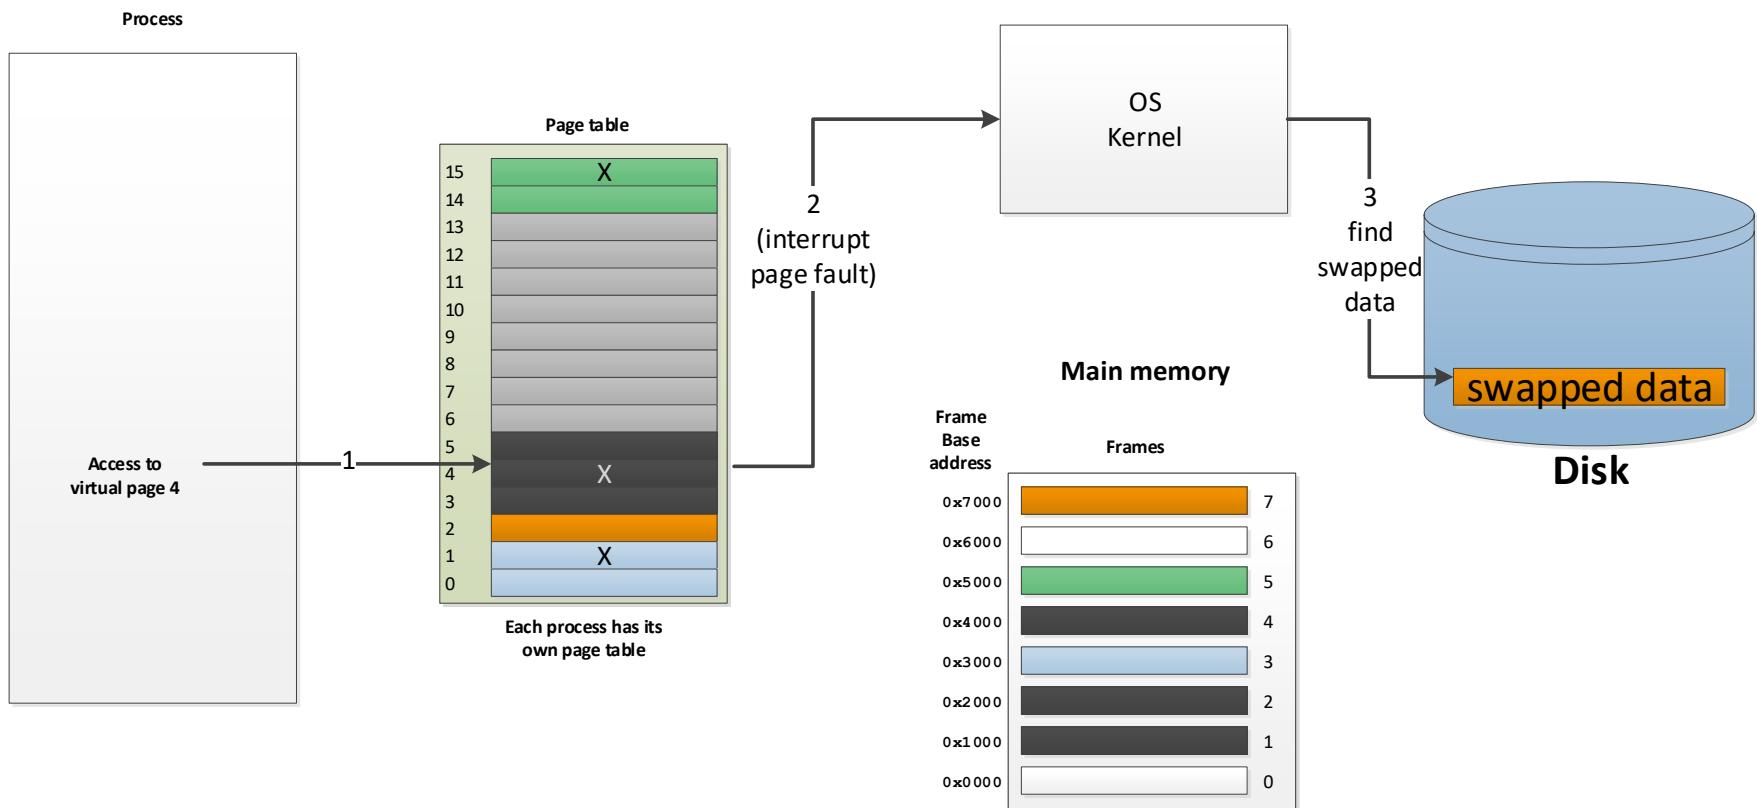
Page fault

Virtual memory



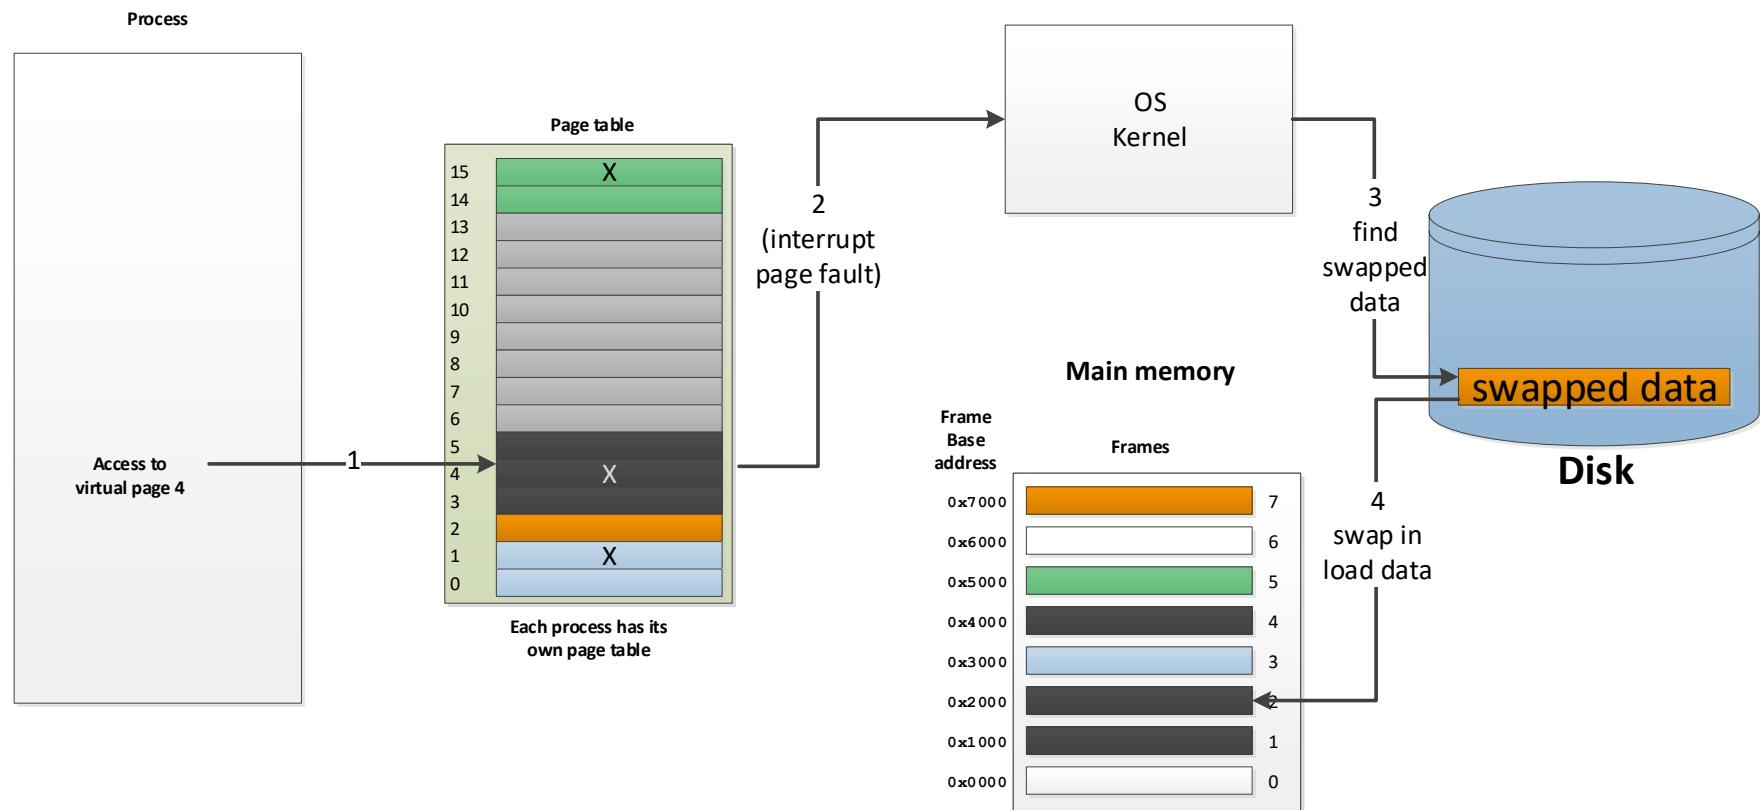
Page fault

Virtual memory



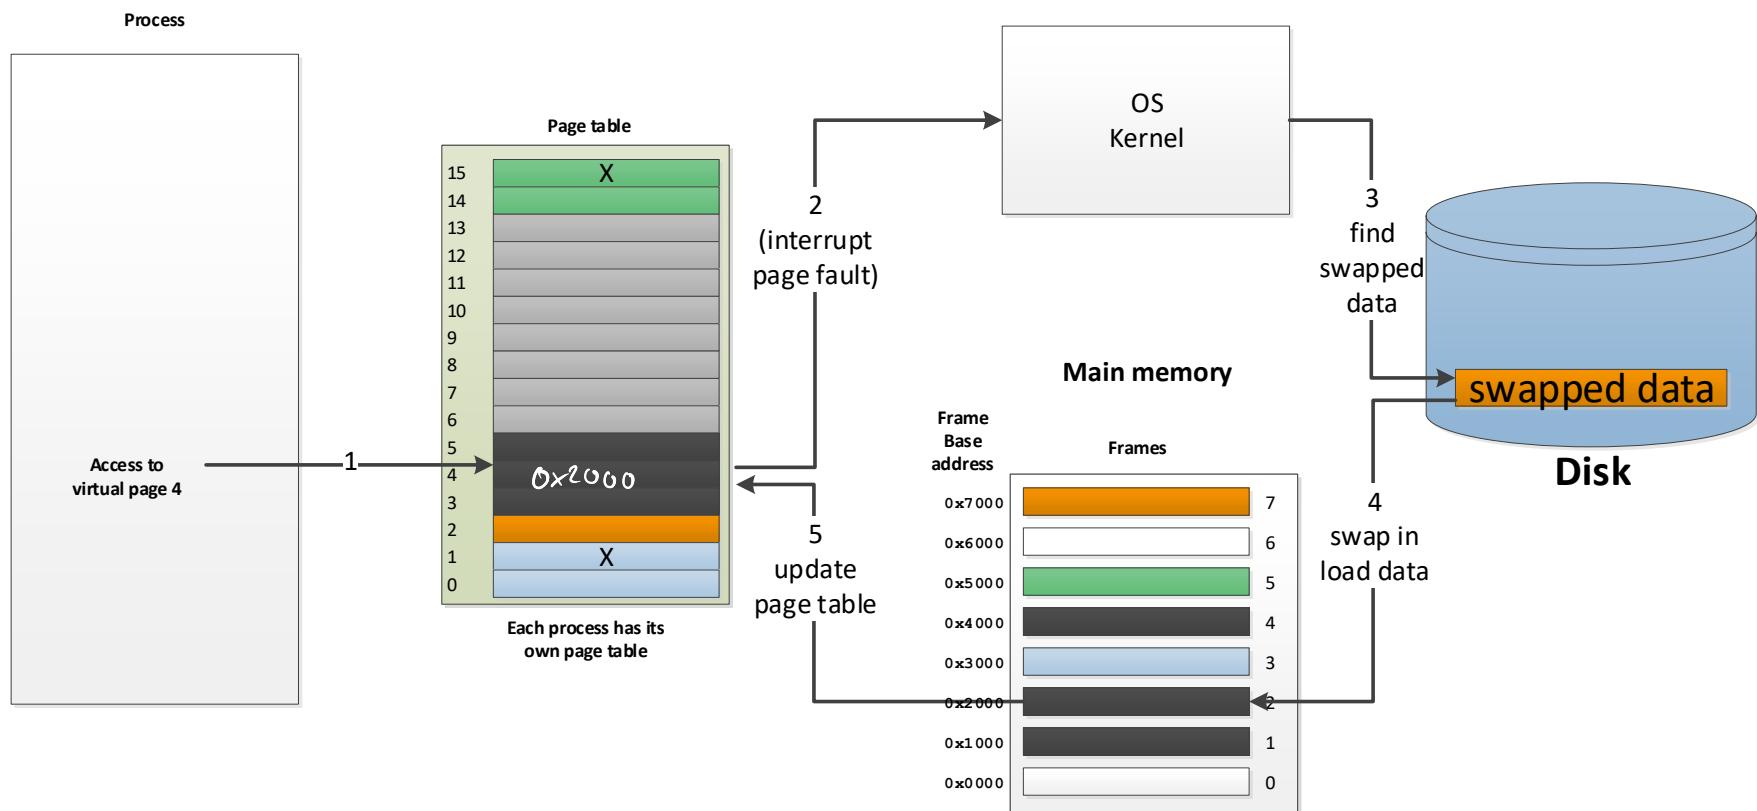
Page fault

Virtual memory



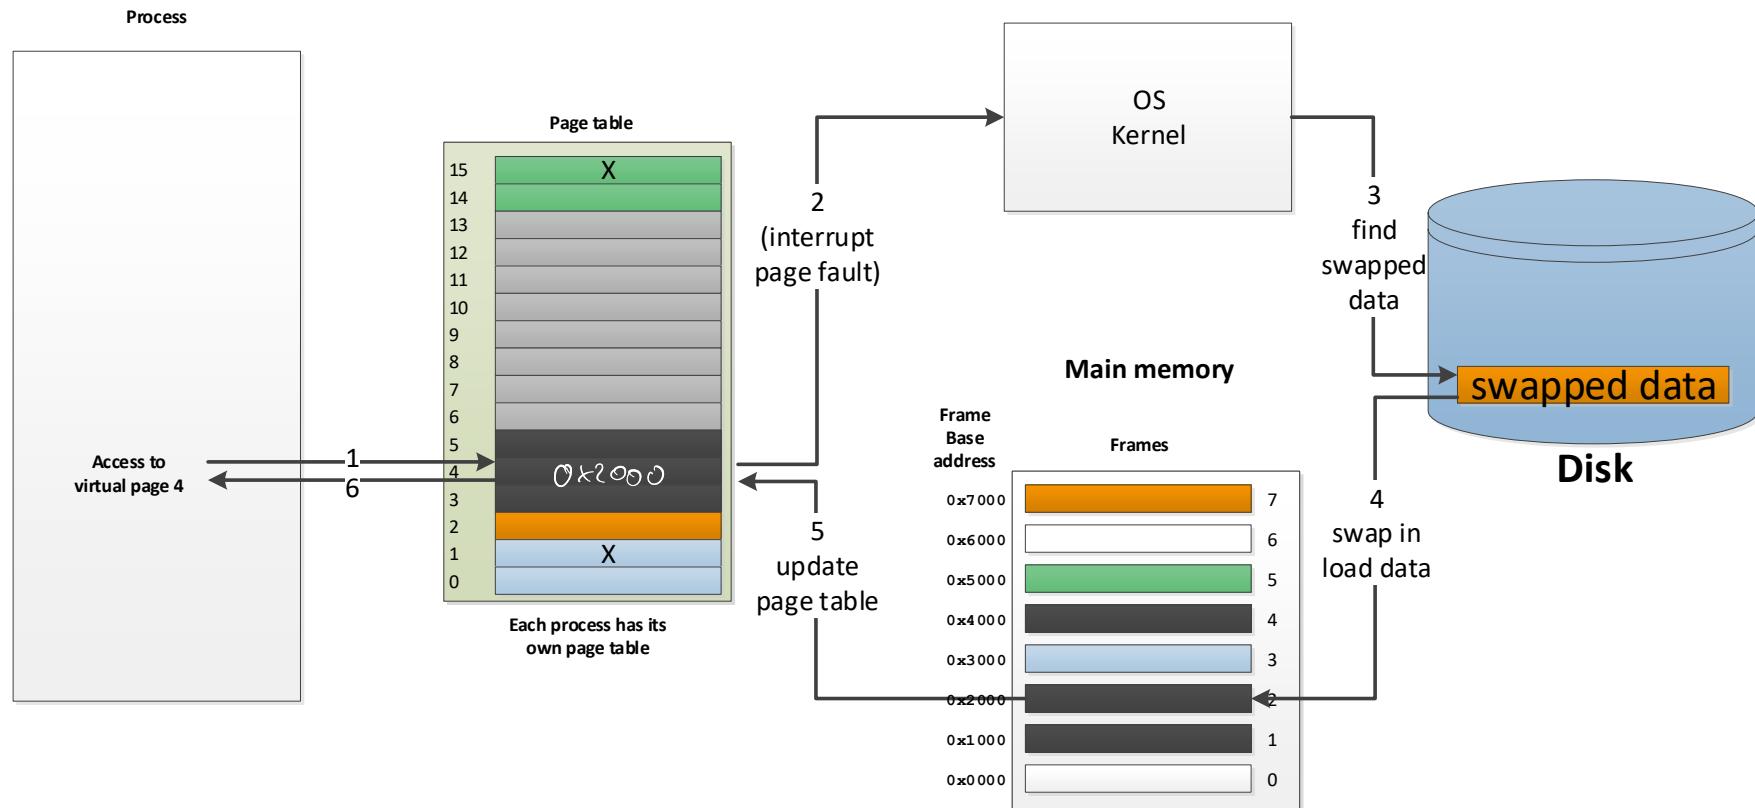
Page fault

Virtual memory



Page fault

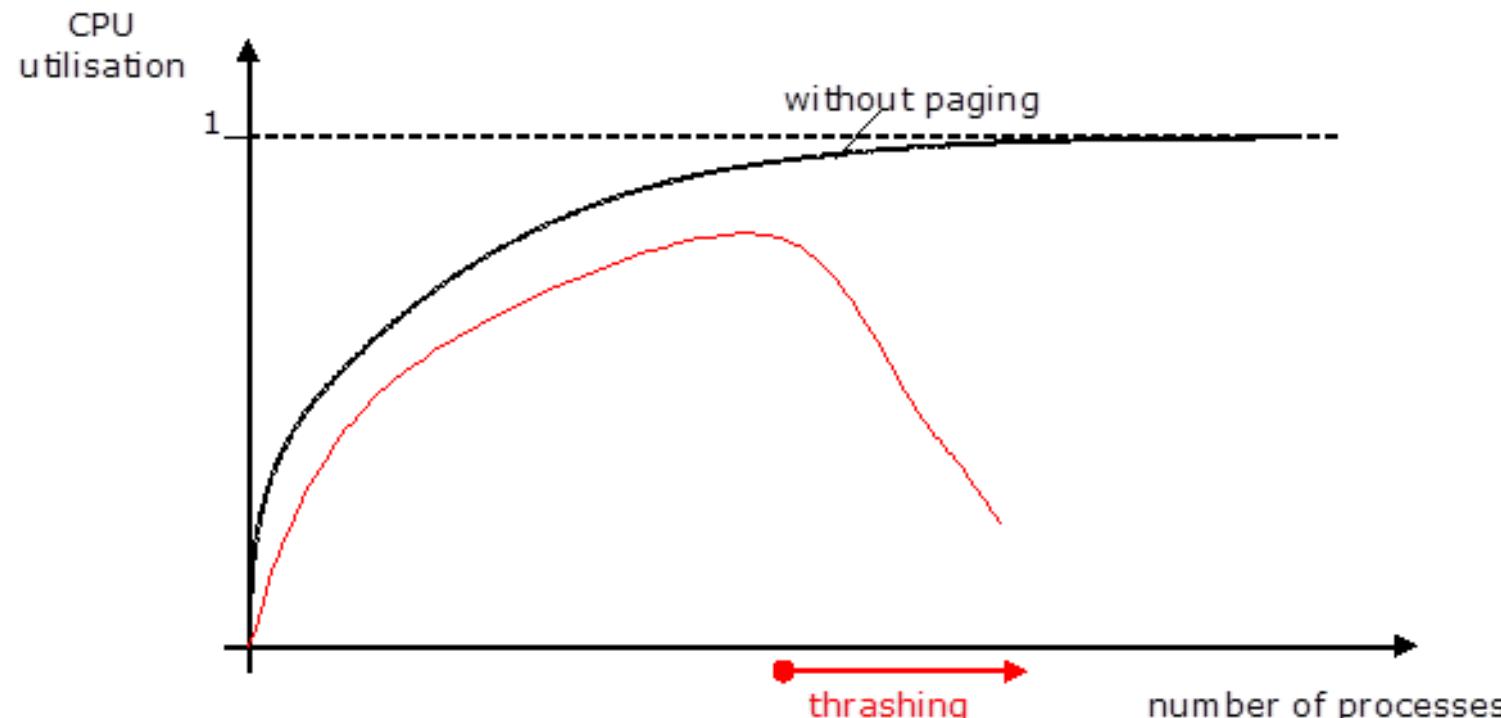
Virtual memory





Thrashing

Thrashing happens if processes with swapped pages are concurrently be activated and the OS needs swap in and swap out very often.





Questions?

All right? \Rightarrow



Question? \Rightarrow



and use **chat**

or

*speak after I
ask you to*



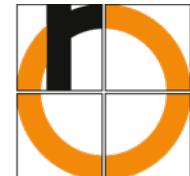
Summary and outlook

Summary

- Caching
- Partitioning
- Fragmentation
- Allocation strategies
- Swapping
- Virtual memory

Outlook

- Shared libraries
- User management
- File systems



Summary and outlook

Summary

- Caching
- Partitioning
- Fragmentation
- Allocation strategies
- Swapping
- Virtual memory

Outlook

- Shared libraries
- User management
- File systems