



Exercise sheet 12 – Memory management

Goals:

- Memory management

Exercise 12.1: Memory allocation strategies

Consider a main memory that contains the following free partitions: 10 KiB, 4 KiB, 20 KiB, 18 KiB, 7 KiB, 9 KiB, 12 KiB, and 15 KiB. Between the free partitions there are used partitions of an unknown size.

- (a) Visualise the situation: Draw a sketch of the memory view.

Now, the following subsequent requests for memory partitions occur: 12 KiB, 10 KiB, and 9 KiB.

- (b) Show the results within your memory sketch when *first fit* is used.
(c) Show the results within your memory sketch when *best fit* is used.
(d) Show the results within your memory sketch when *next fit* is used.
(e) Show the results within your memory sketch when *worst fit* is used.

Proposal for solution:

OS exercise: Memory allocation strategies					
	a) Memory	b) first fit	c) best fit	d) next fit	e) worst fit
15 KiB	free				9 KiB
	acquired				
12 KiB	free		12 KiB		
	acquired				
9 KiB	free		9 KiB	9 KiB	
	acquired				
7 KiB	free				
	acquired				
18 KiB	free	9 KiB		10 KiB	10 KiB
	acquired				
20 KiB	free	12 KiB		12 KiB	12 KiB
	acquired				
4 KiB	free				
	acquired				
10 KiB	free	10 KiB	10 KiB		
	acquired				

Exercise 12.2: Memory management programming and OS memory mechanism

(a) How and where can a process acquire (allocate) main memory in C (there are two possibilities)?

Proposal for solution:

- Stack memory: by declaring variables inside functions.
- Heap memory: by using `malloc()` or `calloc()`

(b) How can a process release memory (distinct two possibilities)?

Proposal for solution:

- Stack memory: when the function goes out of scope, all variables that are declared inside the function goes out of scope and are automatically released.
- Heap memory: by using `free()`.

(c) Write a small C program that shows how the main memory acquire (allocation) and release works (distinct two possibilities).

Proposal for solution:

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
4  void function_with_stack_data(int num_bytes)
5  {
6      char stack_array[num_bytes];
7
8      //when the function is left, stack_array is automatically deleted from stack.
9  }
10
11 void function_with_heap_data(int num_bytes)
12 {
13     // acquire (allocate) heap memory
14     char* heap_array = (char*)malloc(sizeof(char)*num_bytes);
15
16     //free the heap memory
17     free(heap_array);
18     heap_array = NULL;
19 }
20
21 int main(int argc, char** argv)
22 {
23     if(argc < 2){
24         printf("Error: no size is specified\n");
25         printf("Usage: %s <size>\n", argv[0]);
26         exit(EXIT_FAILURE);
27     }
28
29     int num_bytes = atoi(argv[1]);
30     if(num_bytes <= 0){
31         printf("Error: size has to be > 0\n");
32         exit(EXIT_FAILURE);
33     }
34 }
```

```
35     function_with_stack_data(num_bytes);  
36     function_with_heap_data(num_bytes);  
37  
38     printf("The very end.\n");  
39  
40  
41     return 0;  
42 }
```

- (d) Is the operating system involved when acquire (allocation) and release of main memory is done by a process (distinct two possibilities)?

Proposal for solution:

- Stack memory: No, this is done by manipulating the stack pointer: decrease means acquire memory, increase means release memory.
- Heap memory: Yes, `malloc()` and `free()` involves SVCs and are handled by the OS.

- (e) Is the operating system involved when the process writes data into the main memory (distinct two possibilities)?

Proposal for solution:

- Stack memory: No, the process (machine code) directly writes into the main memory using virtual addresses.
- Heap memory: No, the process (machine code) directly writes into the main memory using virtual addresses.

Exercise 12.3: Memory management

- (a) What is a cache?

Proposal for solution: A cache is an associative memory, which stores data of a request or computation. The next time this data is needed, it is read from the cache instead, which is much faster.

- (b) In the context of caching: What is a hit and what is a fault?

Proposal for solution: A hit occurs if the requested data is already in the cache. A fault occurs if the requested data is not in the cache (usually at the first access).

- (c) What is position independent code (PIC)?

Proposal for solution: It means, that the code can be loaded on any location on the main memory.

- (d) Can fragmentation problems be solved by variable partition sizes (each process can choose its own required partition size)?

Proposal for solution: No, the problem is somewhat alleviated but it still exists, because the number of processes is limited by the partitions and there will be gaps formed automatically.

- (e) Can the operating system protect the memory of a process against others without the help of the CPU?

Proposal for solution: No, not really. Because memory access is done directly with CPU instructions.

(f) What is swapping?

Proposal for solution: Swap out the memory or parts of the memory (with virtual memory) of a process that is not actively running to the disk.

(g) Does swapping improve the performance?

Proposal for solution: It depends! Because transferring data from memory to the disk and vice versa takes a lot of time. The overall performance may go down. But it helps to execute a lot of processes where some can be swapped. And usually, a lot of processes are required to fully utilise the CPU.

(h) What happens on a page fault and how is the operating system involved?

Proposal for solution: The operating system has to perform a swap in first. More details: See `os_12_memory_management_handout.pdf` slide 35.

(i) Does thrashing help to improve the systems performance?

Proposal for solution: No, thrashing happens when the operating system is almost fully busy with swap in/swap out.

(j) Is it right, that the virtual memory has to be smaller than the real memory, because the operating systems also needs some memory?

Proposal for solution: No, the virtual memory has the theoretically available size and is usually greater than the real memory.

(k) What is a virtual address space?

Proposal for solution: The virtual address space consists of all virtual addresses.

(l) What is a page table and how is a virtual address transformed into a real address?

Proposal for solution: A page table contains the information which page (virtual memory) is mapped on which frame (real memory).

(m) Consider a system with virtual memory, MMU, and swapping. Is it required that the code of the executables (ELFs) is build with position independent code?

Proposal for solution: No, because every process has its own virtual memory where it can always use the same virtual addresses.

(n) If you have a system without a MMU, is it possible to have threads?

Proposal for solution: Yes, it is possible to have threads without a MMU at Linux. Systems like Arduinos (microcontrollers) with an Atmega or Cortex M processors also can have threads—but these small microcontrollers does not have an MMU.

(o) If you have a system without a MMU, is it possible to have “real” processes with all that security?



Proposal for solution: No, without a MMU it is not really possible to have processes that are very fast and very secure, because no MMU means no virtual memory.