



Prof. Dr. Florian Künzner

Technical University of Applied Sciences Rosenheim, Computer Science

CA 4 – Processor 1

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier

Processor properties

- Processor architecture
- Exception and interrupt handling
- Instruction set architecture
- Memory model (part of memory lectures)
- Endianness
- Registers
- Addressing modes
- Advanced concepts
 - Instruction scheduling
 - Pipelines
 - Superscalarity
 - VLIW
 - Out-of-order memory access

Goal




Goal

CA::Processor 1

- Processor architecture
- Exception and interrupt handling
- Instruction set architecture

Von Neumann architecture



[schematic, simplified view]

Von Neumann architecture

Properties:

- **Instructions and data** are located in the **same memory** or address space
- Von Neumann – **bottleneck**:
Instruction execution time < Memory access time

Von Neumann architecture


Properties:

- Instructions and data are located in the **same memory** or address space
- Von Neumann – **bottleneck**:
Instruction execution time < **Memory access time**

Is the Von Neumann bottleneck still relevant?


In order to avoid and mitigate such problems, various strategies have been developed (e.g. cache memory)–but its still there!

Von Neumann architecture – bottleneck




[schematic, simplified view]

Von Neumann architecture – bottleneck



[schematic, simplified view]

Harvard architecture



[schematic, simplified view]

Harvard architecture

Properties:

- **Separate memory for data and instructions**
- Data memory is usually **read- and writeable**
- **Instruction** memory is usually **read-only**. Can't be modified through runtime

Harvard architecture

Properties:

- **Separate memory** for **data** and **instructions**
- **Data** memory is usually **read- and writeable**
- **Instruction** memory is usually **read-only**. Can't be modified through runtime

Processor architecture

Discussion

Von Neumann vs Harvard architecture:
Does it play a role nowadays?

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be manipulated (e.g. in supervisor mode by the OS kernel)

Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)

Exceptions

Examples for exceptions:

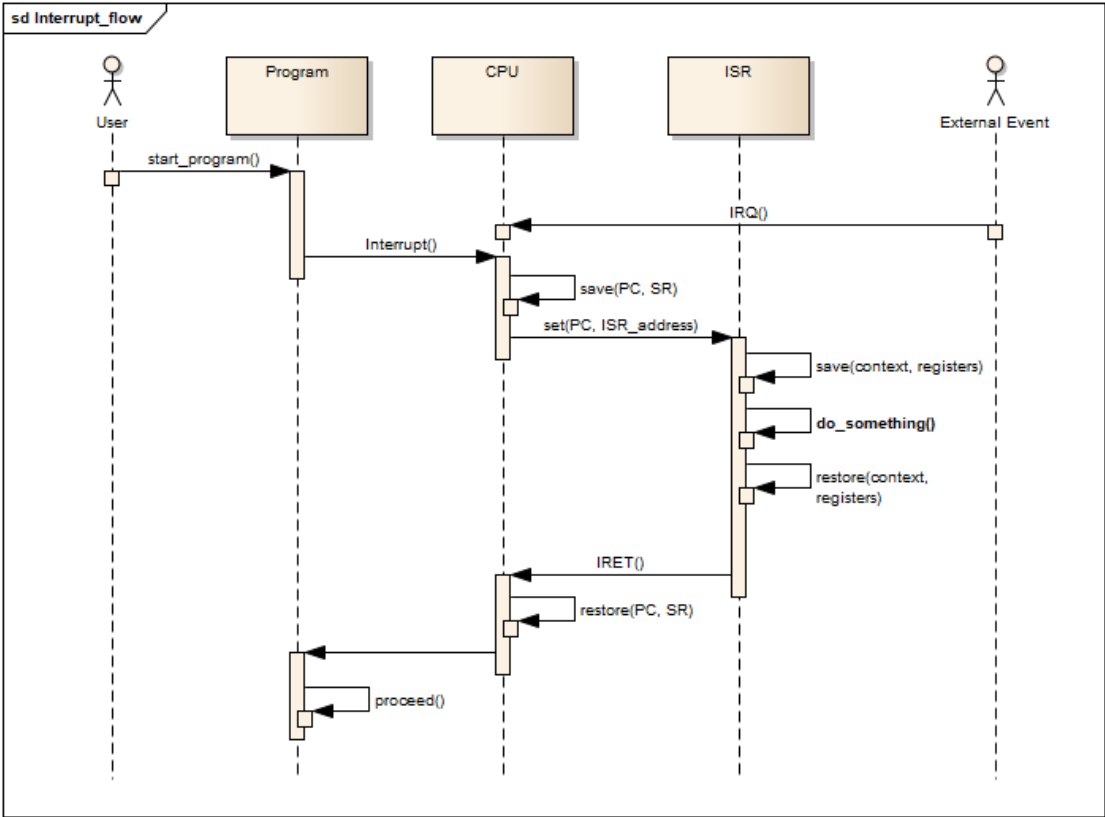
- Division by zero
- Illegal instruction code
- Load or store to an unaligned address
- Unauthorized memory access

Exceptions

Examples for exceptions:

- Division by zero
- Illegal instruction code
- Load or store to an unaligned address
- Unauthorized memory access


Interrupt flow



Sequence in the control unit

Save the old


Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
 - (e.g. on the stack)


Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
 - (e.g. on the stack)


Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
 (e.g. on the stack)

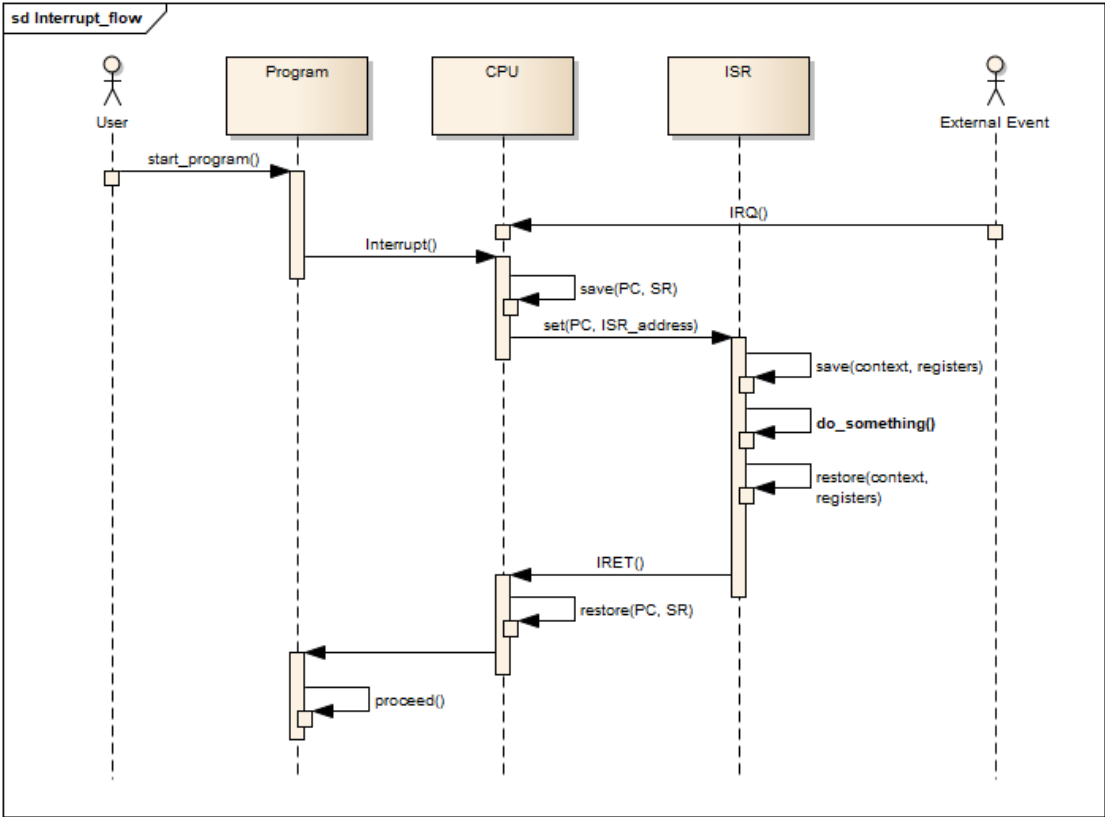
Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
 (e.g. on the stack)
- 2 Assign new values to


Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
(e.g. on the stack)
- 2 Assign new values to
 - PC (program counter) and
 - SR (status register)
from a fixed address ("interrupt vector")

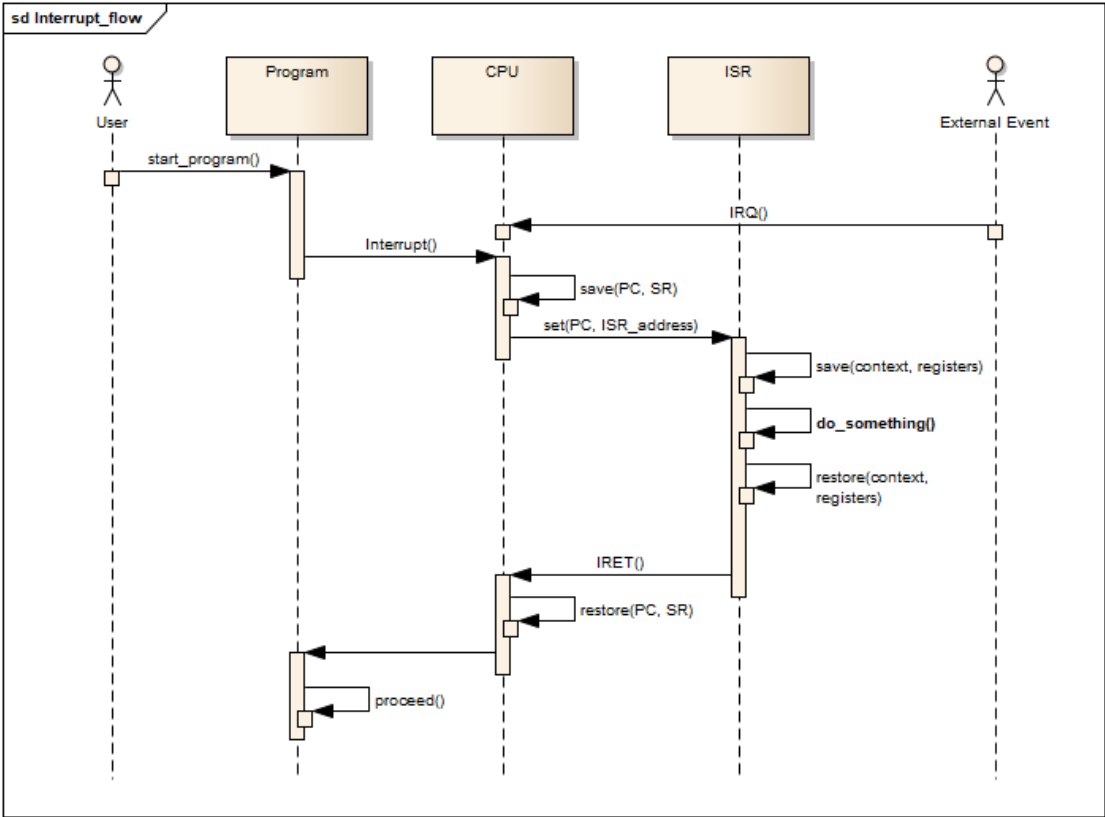
Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
(e.g. on the stack)
- 2 Assign new values to
 - PC (program counter) and
 - SR (status register)
from a fixed address ("interrupt vector")

Interrupt flow



Sequence in the control unit

- 1 Save the old
 - PC (program counter) and
 - SR (status register)
 (e.g. on the stack)
- 2 Assign new values to
 - PC (program counter) and
 - SR (status register)
 from a fixed address ("interrupt vector")

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interruption in the middle of an instruction?

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interruption in the middle of an instruction?

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interruption in the middle of an instruction?

- Usually at the end of a CPU instruction
- But: exceptions can interrupt a running CPU instruction

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interruption in the middle of an instruction?

- Usually at the end of a CPU instruction
- But: exceptions can interrupt a running CPU instruction

Interrupt details

Who saves the registers?

- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

Interruption in the middle of an instruction?

- Usually at the end of a CPU instruction
- But: exceptions can interrupt a running CPU instruction

ISA - instruction set architecture

Do you know some common ISAs?

ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

Degrees of freedom in instruction set architecture (ISA) design

Operations: How many? Which? How complex?

ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

Degrees of freedom in instruction set architecture (ISA) design

- Operations:** How many? Which? How complex?
- Data types:** Which data types can the operations handle?
- Instruction format:** Instruction length in bytes? Number of addresses? Size of the address fields?

ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

Degrees of freedom in instruction set architecture (ISA) design

- Operations:** How many? Which? How complex?
- Data types:** Which data types can the operations handle?
- Instruction format:** Instruction length in bytes? Number of addresses? Size of the address fields?
- Endianness:** Byte order within a long word: Big endian vs little endian?

ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

Degrees of freedom in instruction set architecture (ISA) design

- Operations:** How many? Which? How complex?
- Data types:** Which data types can the operations handle?
- Instruction format:** Instruction length in bytes? Number of addresses? Size of the address fields?
- Endianness:** Byte order within a long word: Big endian vs little endian?
- Register:** How many? Usable in which way?

ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

Degrees of freedom in instruction set architecture (ISA) design

- Operations:** How many? Which? How complex?
- Data types:** Which data types can the operations handle?
- Instruction format:** Instruction length in bytes? Number of addresses? Size of the address fields?
- Endianness:** Byte order within a long word: Big endian vs little endian?
- Register:** How many? Usable in which way?
- Addressing:** Addressing types for the operands? Can be combined with the operations arbitrary ("orthogonal") or restricted?

ISA - instruction set architecture

Instruction formats:
Instruction address Example Operands

ISA - instruction set architecture

Instruction formats:

Instruction address	Example
Zero-address	ADD

Operands

Operands and result on stack!

ISA - instruction set architecture

Instruction formats:

Instruction address	Example	Operands
Zero-address	ADD	Operands and result on stack!
One-address	ADD X	$A = A + X$ ($A = \text{"accu"}$)

ISA - instruction set architecture

Instruction formats:

Instruction address	Example	Operands
Zero-address	ADD	Operands and result on stack!
One-address	ADD X	$A = A + X$ ($A = \text{"accu"}$)
Two-addresses	ADD X, Y	$X = X + Y$ or $Y = X + Y$

ISA - instruction set architecture

Instruction formats:

Instruction address	Example	Operands
Zero-address	ADD	Operands and result on stack!
One-address	ADD X	$A = A + X$ ($A = \text{"accu"}$)
Two-addresses	ADD X, Y	$X = X + Y$ or $Y = X + Y$
Three-addresses	ADD X, Y, Z	$X = Y + Z$ or ...

ISA - special instructions

Synchronisation instructions:

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
 - Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
 - In the hardware single, atomic (i.e., non-interruptible) operation
 - Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area	Protection through
Critical area in a process	P and V operation
P-Operation in the operating system	TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
- Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area

Critical area in a process

P-Operation in the operating system

Protection through

P and V operation

TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
- Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area

Critical area in a process

P-Operation in the operating system

Protection through

P and V operation

TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
- Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area

Critical area in a process

P-Operation in the operating system

Protection through

P and V operation

TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
- Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area

Critical area in a process

P-Operation in the operating system

Protection through

P and V operation

TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions

Synchronisation instructions:

- TAS (test and set) or similar
- Tests a memory cell ($? = 0$) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

Synchronisation area:

Area

Critical area in a process

P-Operation in the operating system

Protection through

P and V operation

TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare_and_Swap, keyword "Lock-free Programming".

ISA - special instructions: TAS

TAS example for Freescale ColdFire architecture

```
1 byte LOCK = 0; // =0 -> lock is free; !=0 -> locked
2 //...
3 __asm { ;Inline assembly block with assembler instructions
4   GetLock:
5       TAS.B LOCK    ;Sets the N- or Z-Bit depending on LOCK and
6                       ;always sets LOCK = 0x80
7       BNE GetLock   ;If LOCK was != 0 then try it again (loop)
8                       ;(BNE, branch not equal)
9   }
10 // Now, we have the LOCK and we are inside the critical section
11 // ...
12 __asm { ;Inline assembly block with assembler instructions
13   ReleaseLock:
14       CLR.B LOCK    // Sets LOCK = 0
15   }
```


ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

- Less, simple instructions
- One instruction should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

- Less, simple instructions
- One instruction should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

- Less, simple instructions
- One instruction should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

- Less, simple instructions
- One instruction should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

ISA - CISC vs RISC

CISC - complex instruction set computer

- Instructions should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for an instruction
- Support for a lot of addressing modes
- Allows compact encoding of programs (one instruction does a lot)

RISC - reduced instruction set computer

- Less, simple instructions
- One instruction should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

