



**Prof. Dr. Florian Künzner**

Technical University of Applied Sciences Rosenheim, Computer Science

## CA 5 – Processor 1

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier

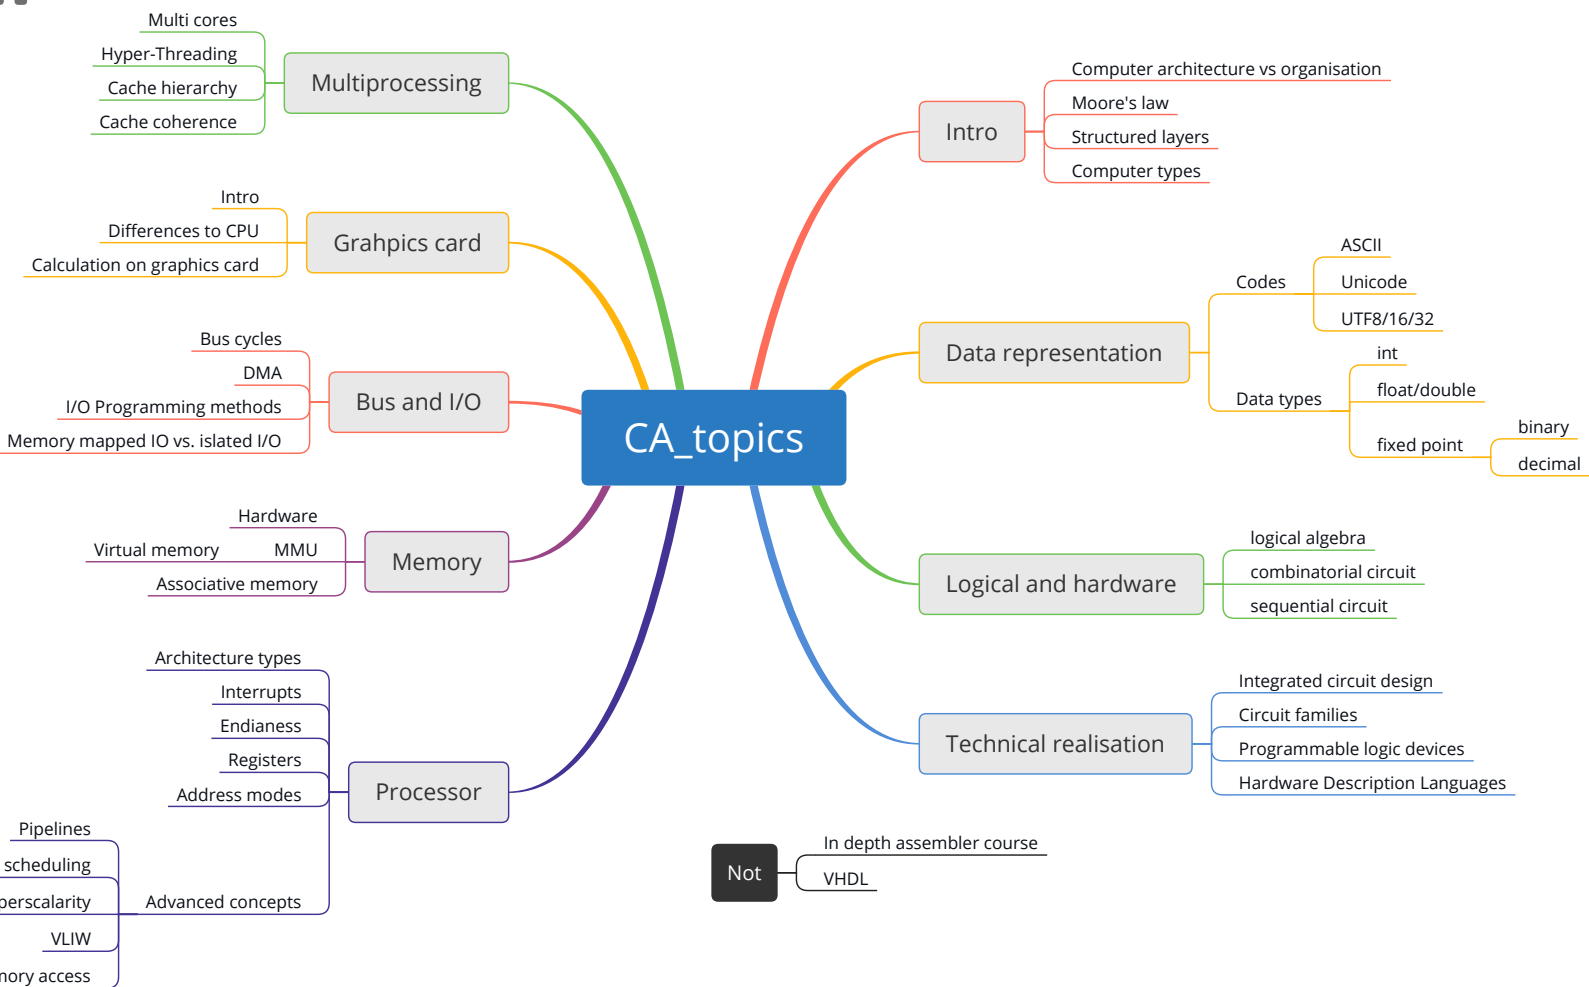
# Overview

What are the properties of a processor?

# Processor properties

- Processor architecture
- Exception and interrupt handling
- Instruction set architecture
- Memory model (part of memory lectures)
- Endianness
- Registers
- Addressing modes
- Advanced concepts
  - Instruction scheduling
  - Pipelines
  - Superscalarity
  - VLIW
  - Out-of-order memory access

# Goal

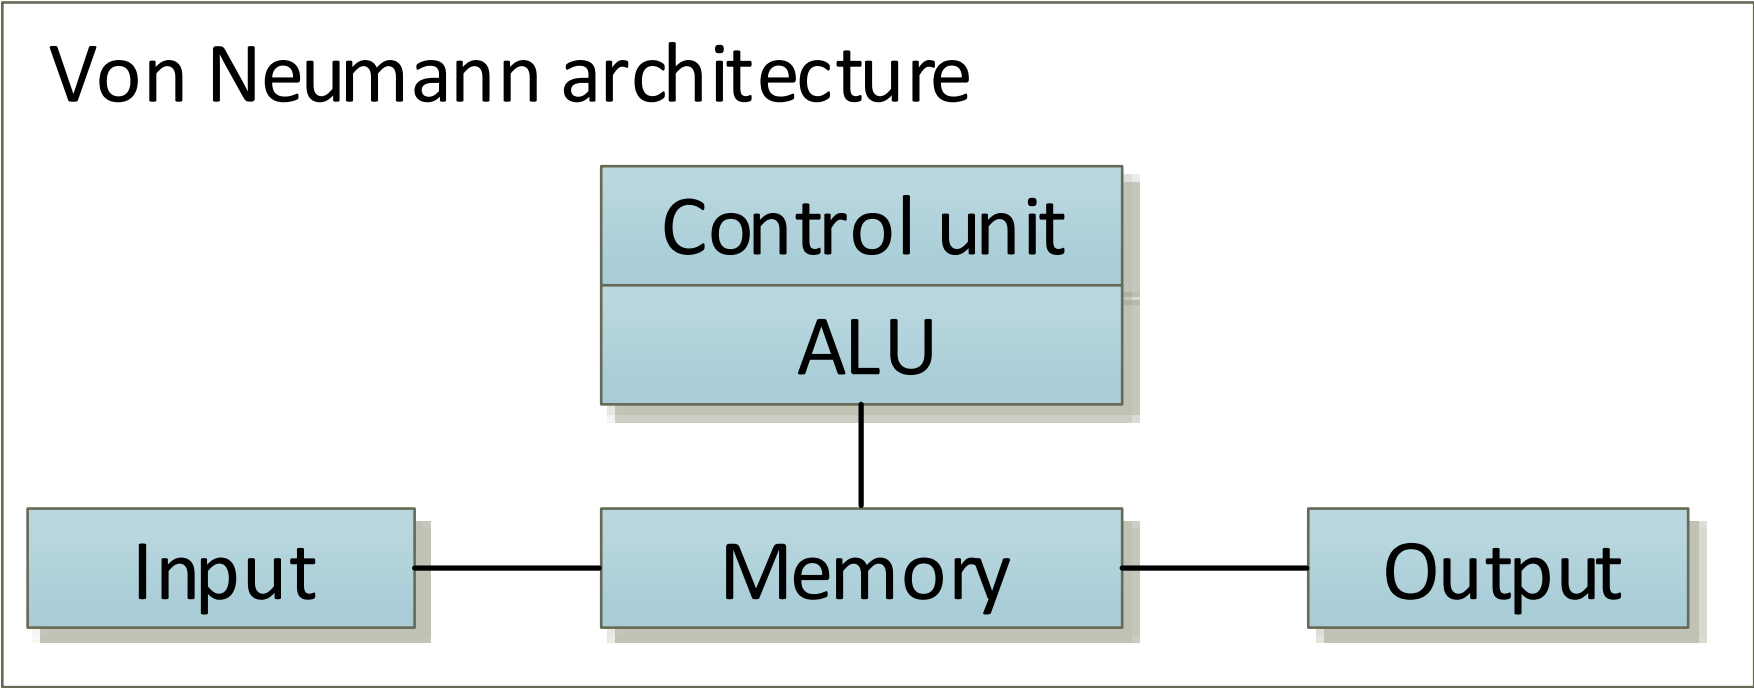


# Goal

## CA::Processor 1

- Processor architecture
- Exception and interrupt handling
- Instruction set architecture
- Endianness

# Von Neumann architecture



[schematic, simplified view]

# Von Neumann architecture

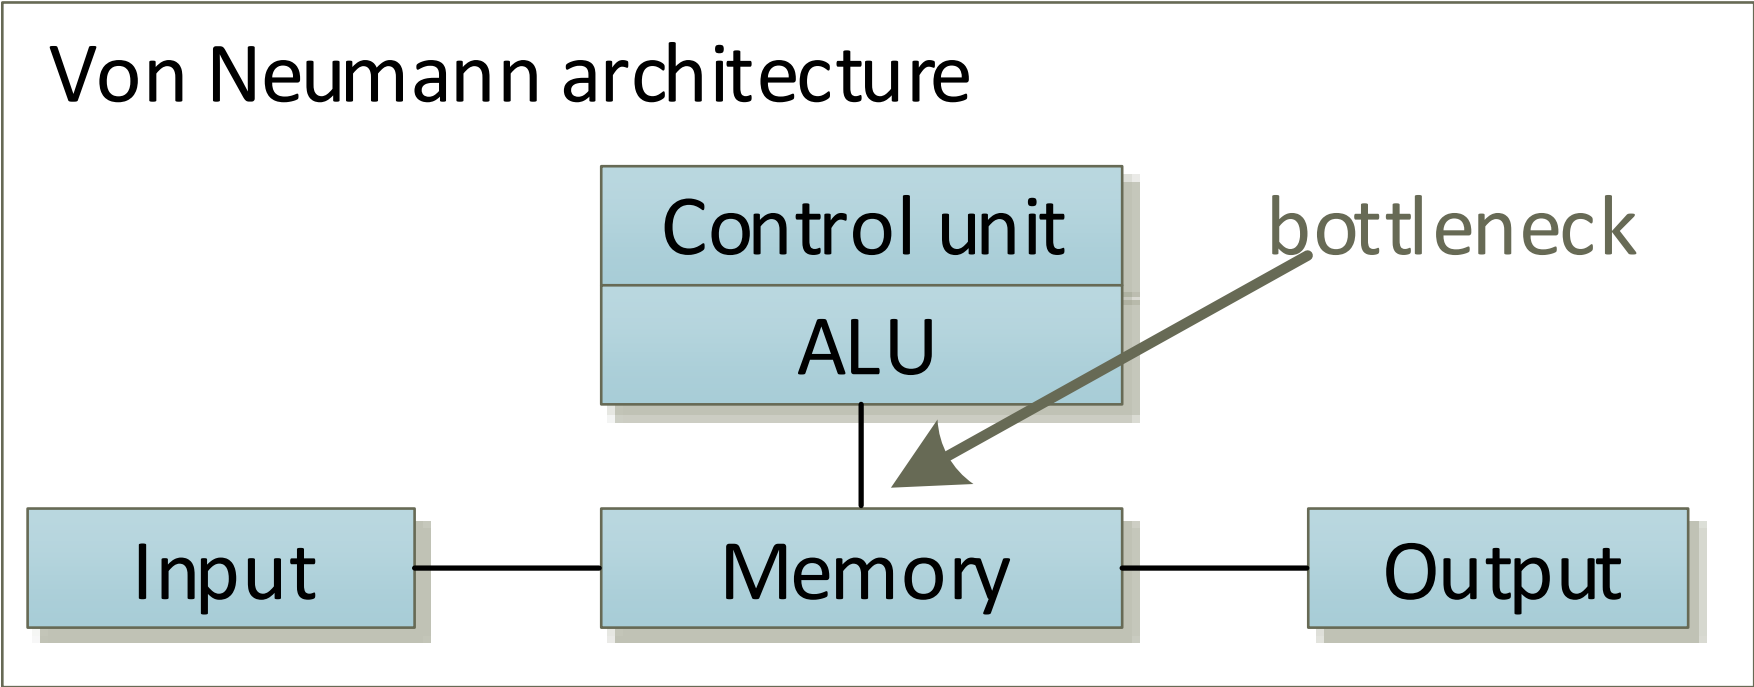
## Properties:

- Instructions and data are located in the **same memory** or address space
- Von Neumann – **bottleneck**:  
**Command execution time < Memory access time**

## Is the Von Neumann bottleneck still relevant?

In order to avoid and mitigate such problems, various strategies have been developed (e.g. cache memory)–but its still there!

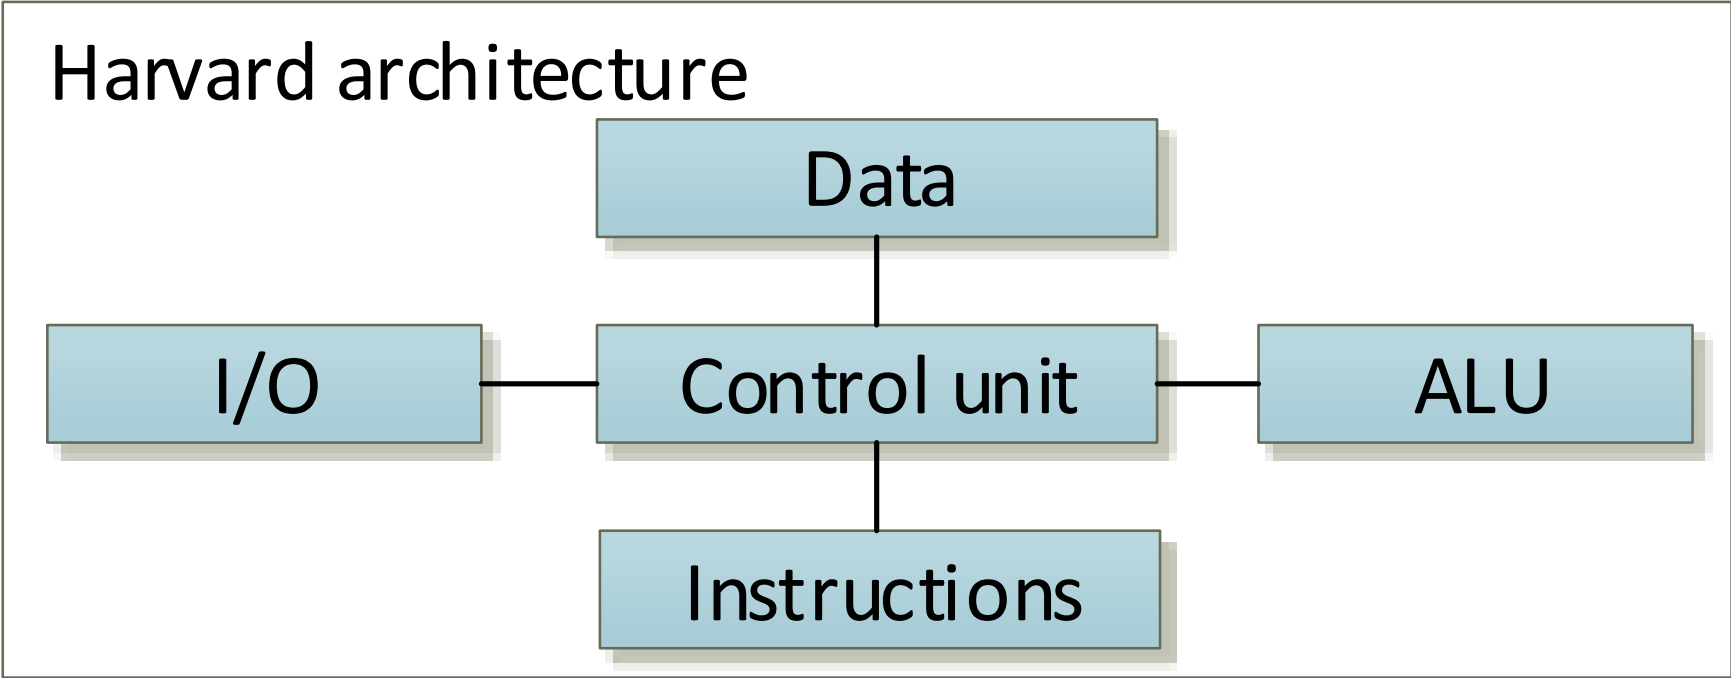
# Von Neumann architecture – bottleneck



[schematic, simplified view]



# Harvard architecture



[schematic, simplified view]

# Harvard architecture

## Properties:


- **Separate memory** for **data** and **instructions**
- **Data** memory is usually **read- and writeable**
- **Instruction** memory is usually **read-only**. Can't be modified through runtime


# Processor architecture

## Discussion

Von Neumann vs. Harvard architecture:  
Does it play a role nowadays?

# Questions?

All right? ⇒ 

Question? ⇒  and use **chat**

or

**speak** *after* I ask you to


# Control unit


Pseudo C code of **control unit** inside the CPU:

```
1 while(true){  
2     fetch_next_command();  
3     decode_command();  
4     execute_command();  
5     if(interrupt_is_requested()) {  
6         save_PC_and_SR();  
7         load_new_PC();  
8     }  
9 }
```

**Instruction cycle:** in principle, it's an **endless loop**

# Questions?

All right? ⇒ 

Question? ⇒  and use **chat**

or

**speak** *after* I ask you to

# Interrupts

What is an interrupt?

# Interrupt handling

An **interrupt request (IRQ)** causes the processor to **stop the current workflow** and to **process with a predefined interrupt service routine (ISR)**.

- **IRQ** -> Interrupt request
- **ISR** -> Interrupt service routine
- An ISR is a **function** that is called when an interrupt occurs
- It's just a **memory address** where the function (ISR) starts
- There exist **different types** of interrupts
- The ISR addresses are managed within the CPU with the **interrupt vector table**
- The **interrupt vector table** can be **manipulated** (e.g. in supervisor mode by the OS kernel)



# Interrupt types

Name	Usual name	Reason, cause	Arrival	Comment
Reset	Reset	external	asynchronous	reset potentially at any point in an instruction
Interrupt	Real interrupt, IRQ	external (e.g. I/O)	asynchronous	is usually handled at the end of a command
Exception	Internal interrupt	internal (e.g. command error)	synchronous	is usually cancelled and may be repeated later
System call	SVC, supervisor call, Trap, Software interrupt	internal	synchronous	SVC n (n is the number of the SVC)
Timer	Timer, SysTick, Watchdog	external (clock)	asynchronous	is usually handled at the end of a command

...

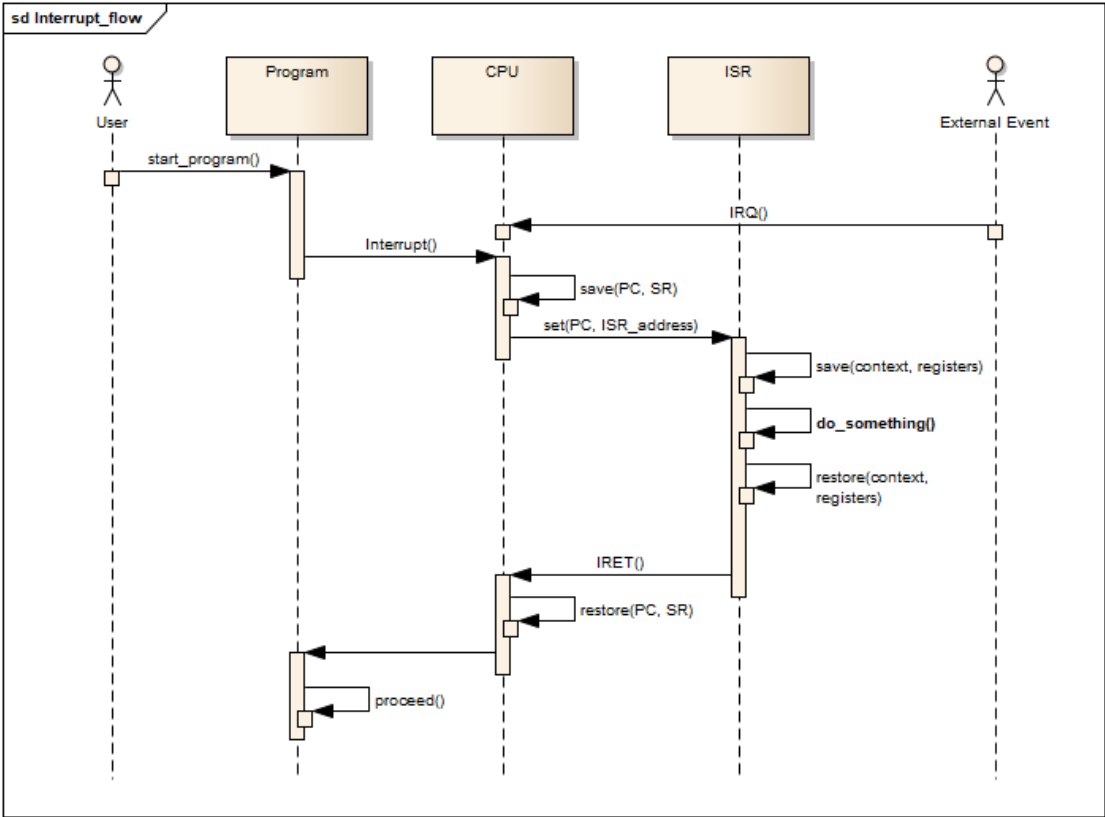
[for example: ARM Cortex-M Exception handlers]

# Exceptions

## Examples for exceptions:

- Division by zero
- Illegal command code
- Load or store to an unaligned address.
- Unauthorized memory access


# Interrupt flow




## Sequence in the control unit

- 1 Save the old
  - PC (program counter) and
  - SR (status register)(e.g. on the stack)
- 2 Assign new values to
  - PC (program counter) and
  - SR (status register)from a fixed address ("interrupt vector")

# Questions?

All right?  $\Rightarrow$  

Question?  $\Rightarrow$   and use **chat**

or

**speak** *after* I ask you to

# Interrupt vector table

Exception number	IRQ number	Vector	Offset
16+n	n	IRQn	0x40+4n
⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮
18	2	IRQ2	0x48
17	1	IRQ1	0x44
16	0	IRQ0	0x40
15	-1	SysTick, if implemented	0x3C
14	-2	PendSV	0x38
13		Reserved	
12			
11	-5	SVCall	0x2C
10			
9			
8			
7			
6			
5			
4			
3	-13	HardFault	0x10
2	-14	NMI	0x0C
			0x08
		Reset	0x04
1		Initial SP value	0x00

- Example: **Cortex M0**
- Interrupt request (IRQ)
- For each IRQ one entry
- Each entry contains the address of an ISR

[source: Cortex-M0 Generic User Guide (2009), p. 2-22]

# Interrupt details


## Who saves the registers?


- The CPU saves and restores PC and SR
- The ISR has to save the other registers at the beginning
- The ISR has to restore the other registers at the end
- Usually, the operating system does this (if you have one) in advance

## Interruption in the middle of a command?

- Usually at the end of a CPU command
- But: exceptions can interrupt a running CPU command

# Questions?

All right? ⇒ 

Question? ⇒  and use **chat**

or

**speak** *after* I ask you to

# ISA - instruction set architecture

Do you know some common ISAs?



# ISA - instruction set architecture

The interface for low-level programming, very close to the hardware.

## Degrees of freedom in instruction set architecture (ISA) design

- Operations:** How many? Which? How complex?
- Data types:** Which data types can the operations handle?
- Command format:** Command length in bits? Number of addresses? Size of the address fields?
- Register:** How many? Usable in which way?
- Addressing:** Addressing types for the operands? Can be combined with the operations arbitrary ("orthogonal") or restricted?

# ISA - instruction set architecture

## Command formats:

Command address	Example	Operands
Zero-address	ADD	Operands and result on stack!
One-address	ADD X	$A = A + X$ ( $A = \text{"accu"}$ )
Two-addresses	ADD X, Y	$X = X + Y$ or $Y = X + Y$
Three-addresses	ADD X, Y, Z	$X = Y + Z$ or ...

# ISA - special commands

## Synchronization commands:

- TAS (test and set) or similar
- Tests a memory cell ( $? = 0$ ) and sets it to one, if the test was successful.
- In the hardware single, atomic (i.e., non-interruptible) operation
- Uses a read-modify-write memory cycle that completes the operation without interruption.

## Synchronisation area:

Area	Protection through
Critical area in a process	P and V operation
P-Operation in the operating system	TAS instruction (or similar)

In other architectures, partly different, more efficient, much more complicated solutions. Similar: Compare\_and\_Swap, keyword "Lock-free Programming".

# ISA - special commands: TAS

## TAS example for Freescale ColdFire architecture

```

1 byte LOCK = 0; // =0 -> lock is free; !=0 -> locked
2
3 __asm { ;Inline assembly block with assembler instructions
4   GetLock:
5       TAS.B LOCK    ;Sets the N- or Z-Bit depending on LOCK and
6                       ;always sets LOCK = 0x80
7       BNE GetLock   ;If was LOCK != 0 then try it again (loop)
8                       ;(BNE, branch not equal)
9
10      ;Now we have the LOCK and we are inside the critical section
11      ;...
12
13   ReleaseLock:
14       CLR.B LOCK    // Sets LOCK = 0
15 }
```

# ISA - CISC vs. RISC


## CISC - complex instruction set computer


- Commands should be as close as possible to the high-level languages
- Can take a lot of internal CPU cycles for a command
- Support for a lot of addressing modes
- Allows compact encoding of programs (one command does a lot)

## RISC - reduced instruction set computer

- Less, simple commands
- One commands should only take a few internal CPU cycles
- Supports only simple addressing modes (load/store necessary)
- Usually fixed length of opcodes

# Questions?

All right? ⇒ 

Question? ⇒  and use **chat**

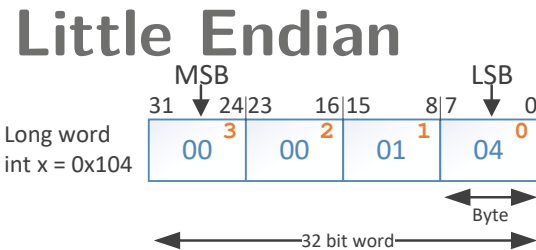
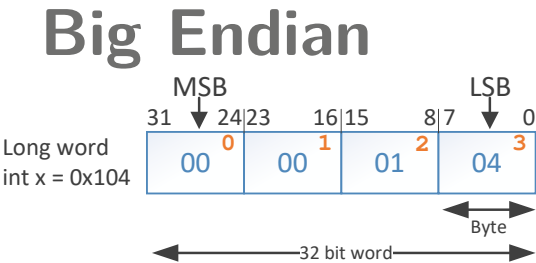
or

**speak** *after* I ask you to

# Endianness

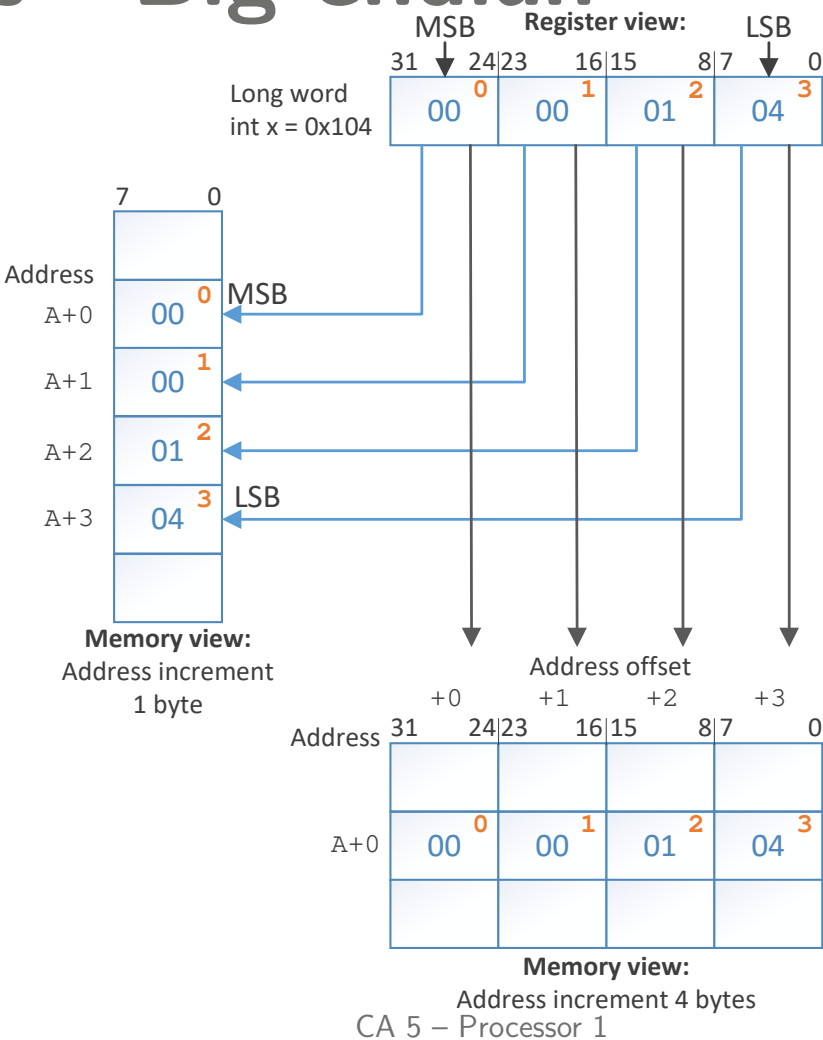
Endianness: The definition of the **byte order** within a **long word**.

Register view of a 32bit architecture:



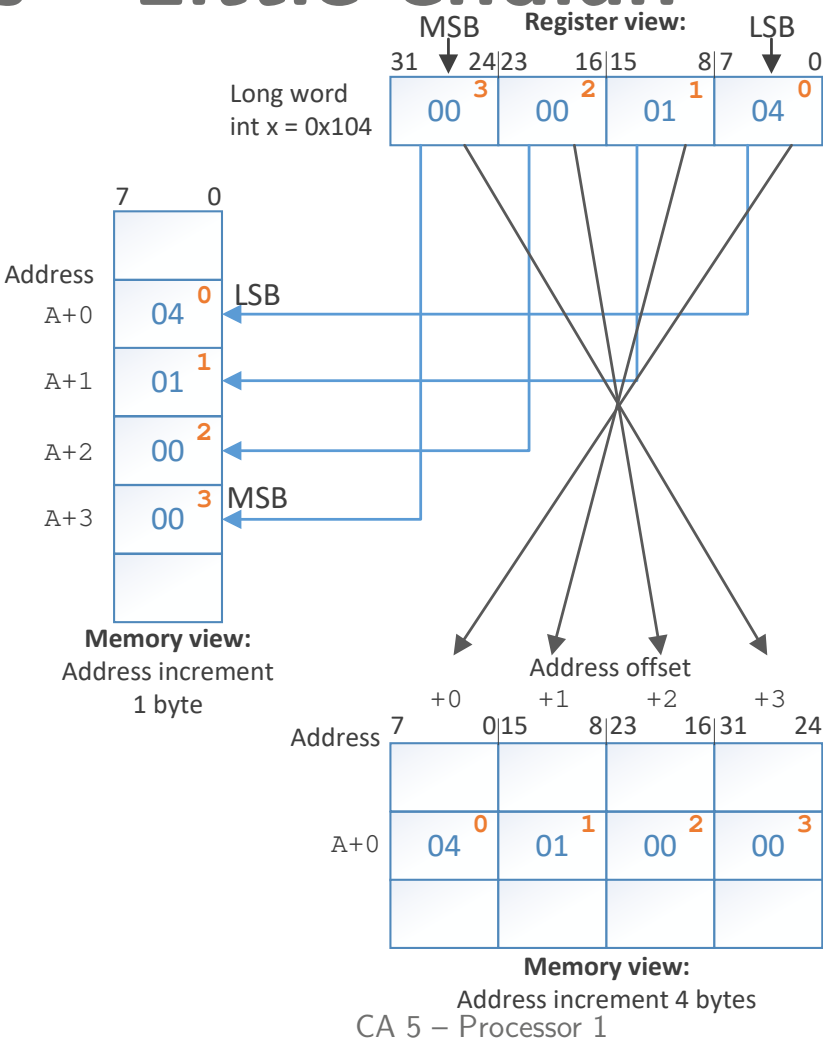
- MSB - Most significant byte
- LSB - Least significant byte

# Endianness - Big endian



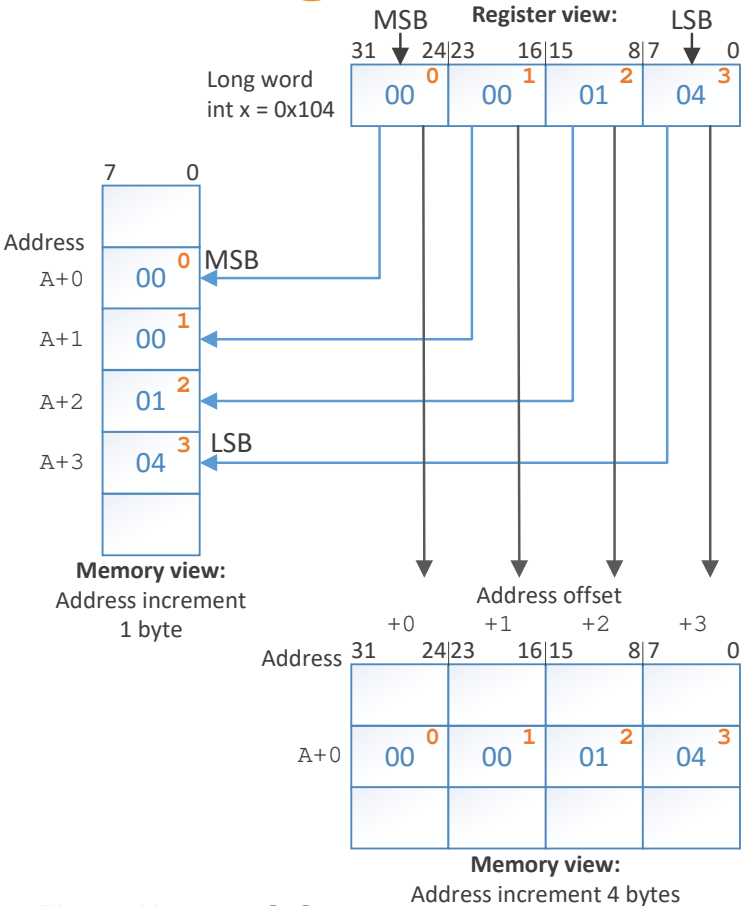


# Endianness - Little endian

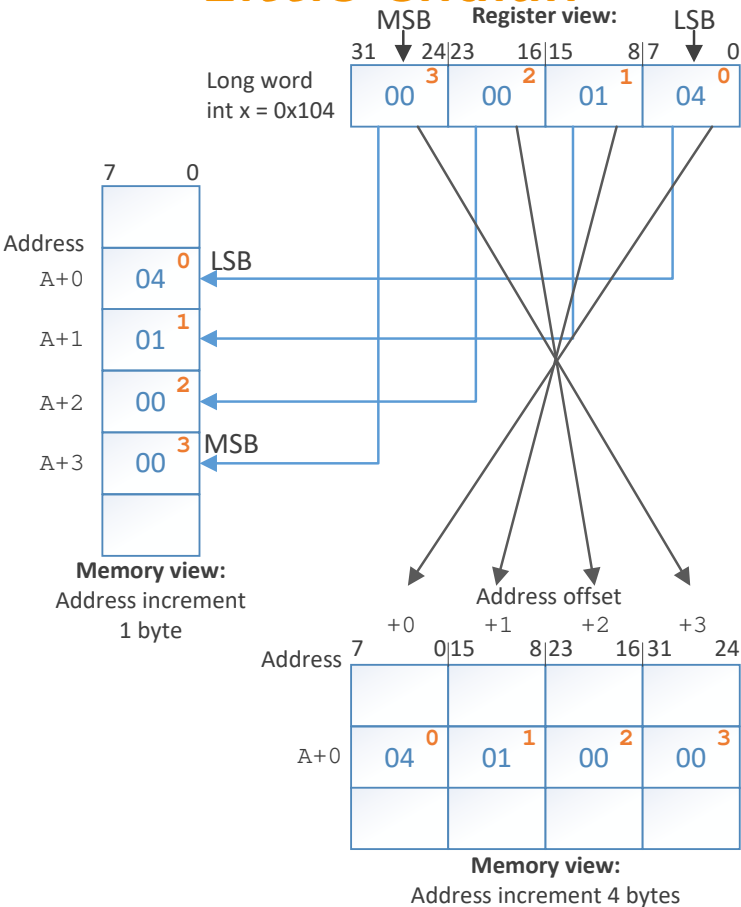


# Endianness - BE/LE

## Big endian



## Little endian



# Endianness - example BE

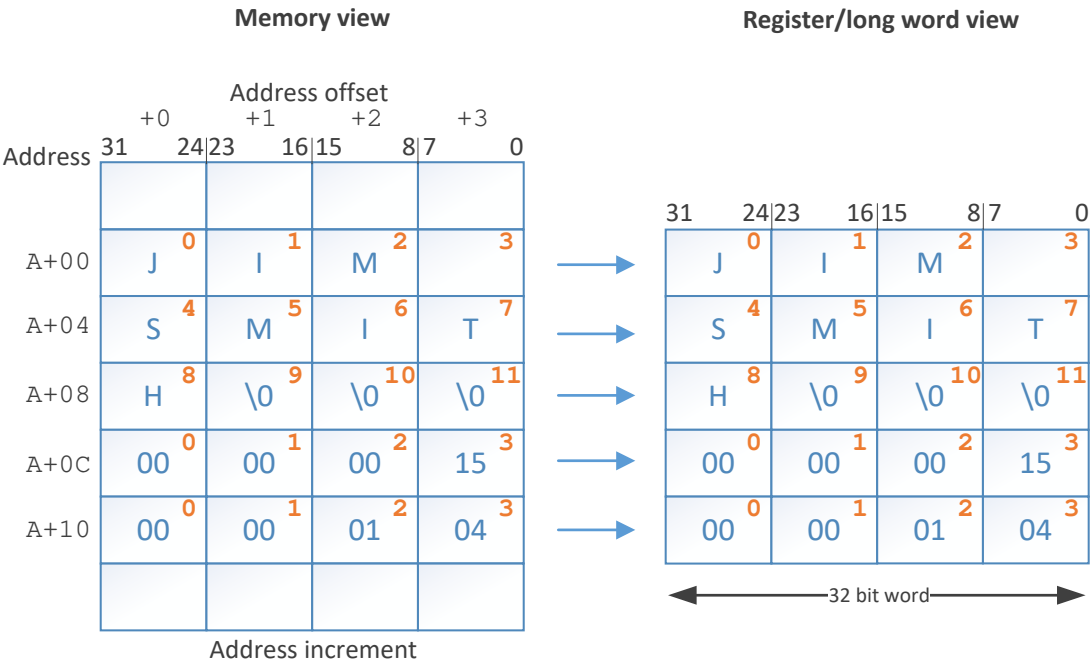
## Big endian memory -> Register

```

1  #include <stdlib.h>
2  #include <stdint.h>
3
4  int main()
5  {
6      struct employee {
7          char    name[12];
8          uint32_t age;
9          uint32_t dept_nr;
10     };
11
12     struct employee smith = {
13         .name    = "JIM SMITH",
14         .age     = 21,        //0x15
15         .dept_nr = 0x104     //260
16     };
17
18     return EXIT_SUCCESS;
19 }

```

[cmp: [1, p. 95-96]]

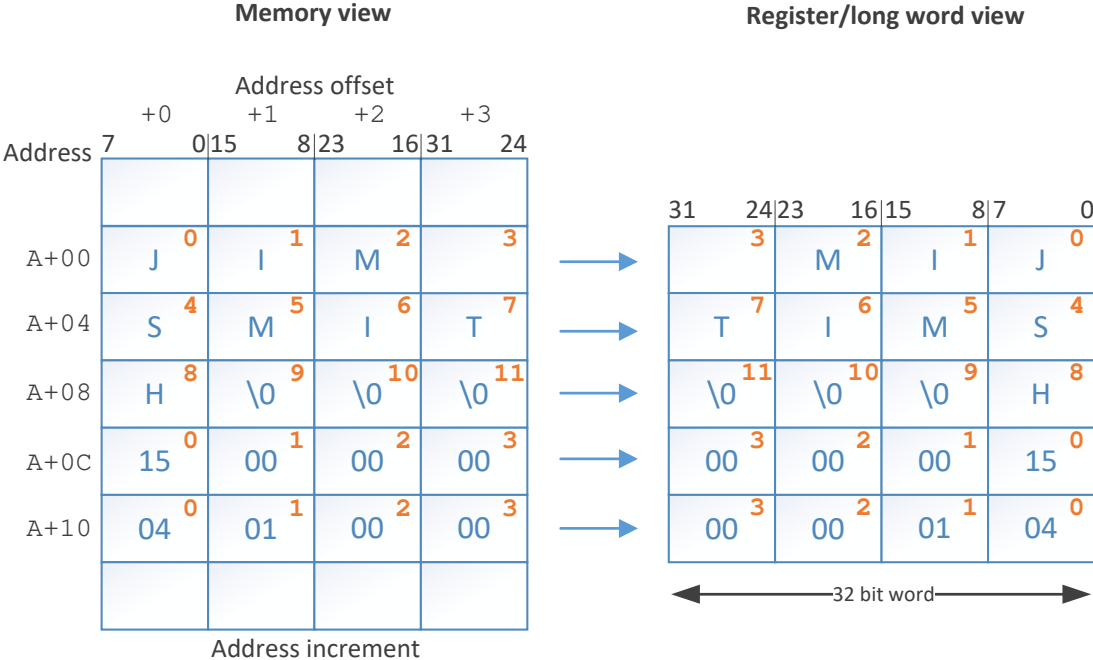


# Endianness - example LE

## Little endian memory

```
1 #include <stdlib.h>
2 #include <stdint.h>
3
4 int main()
5 {
6     struct employee {
7         char    name[12];
8         uint32_t age;
9         uint32_t dept_nr;
10    };
11
12    struct employee smith = {
13        .name    = "JIM SMITH",
14        .age     = 21,    //0x15
15        .dept_nr = 0x104 //260
16    };
17
18    return EXIT_SUCCESS;
19 }
```

[cmp: [1, p. 95-96]]



# Endianness - example BE/LE

## Big endian memory

		Address offset			
		+0	+1	+2	+3
Address	31	24 23	16 15	8 7	0
A+00	J <sup>0</sup>	I <sup>1</sup>	M <sup>2</sup>	<sup>3</sup>	
A+04	S <sup>4</sup>	M <sup>5</sup>	I <sup>6</sup>	T <sup>7</sup>	
A+08	H <sup>8</sup>	\0 <sup>9</sup>	\0 <sup>10</sup>	\0 <sup>11</sup>	
A+0C	00 <sup>0</sup>	00 <sup>1</sup>	00 <sup>2</sup>	15 <sup>3</sup>	
A+10	00 <sup>0</sup>	00 <sup>1</sup>	01 <sup>2</sup>	04 <sup>3</sup>	
		Address increment 4 bytes			

## Little endian memory

		Address offset						
		+0	+1	+2	+3			
Address	7	0 15	8 23	16 31	24			
A+00	J	0	I	1	M	2	3	
A+04	S	4	M	5	I	6	T	7
A+08	H	8	\0	9	\0	10	\0	11
A+0C	15	0	00	1	00	2	00	3
A+10	04	0	01	1	00	2	00	3
		Address increment						
		4 bytes						

[cmp: [1, p. 95-96]]]

# Endianness - usage


## Big endian


- IBM Mainframe
- Freescale ColdFire
- Atmel AVR/AVR32
- ARM Thumb and ARM64 (also Apple M1)

## Little endian

- Intel x86
- x86-64 (AMD64, Intel 64)
- RISC-V
- Qualcomm Hexagon

# Questions?

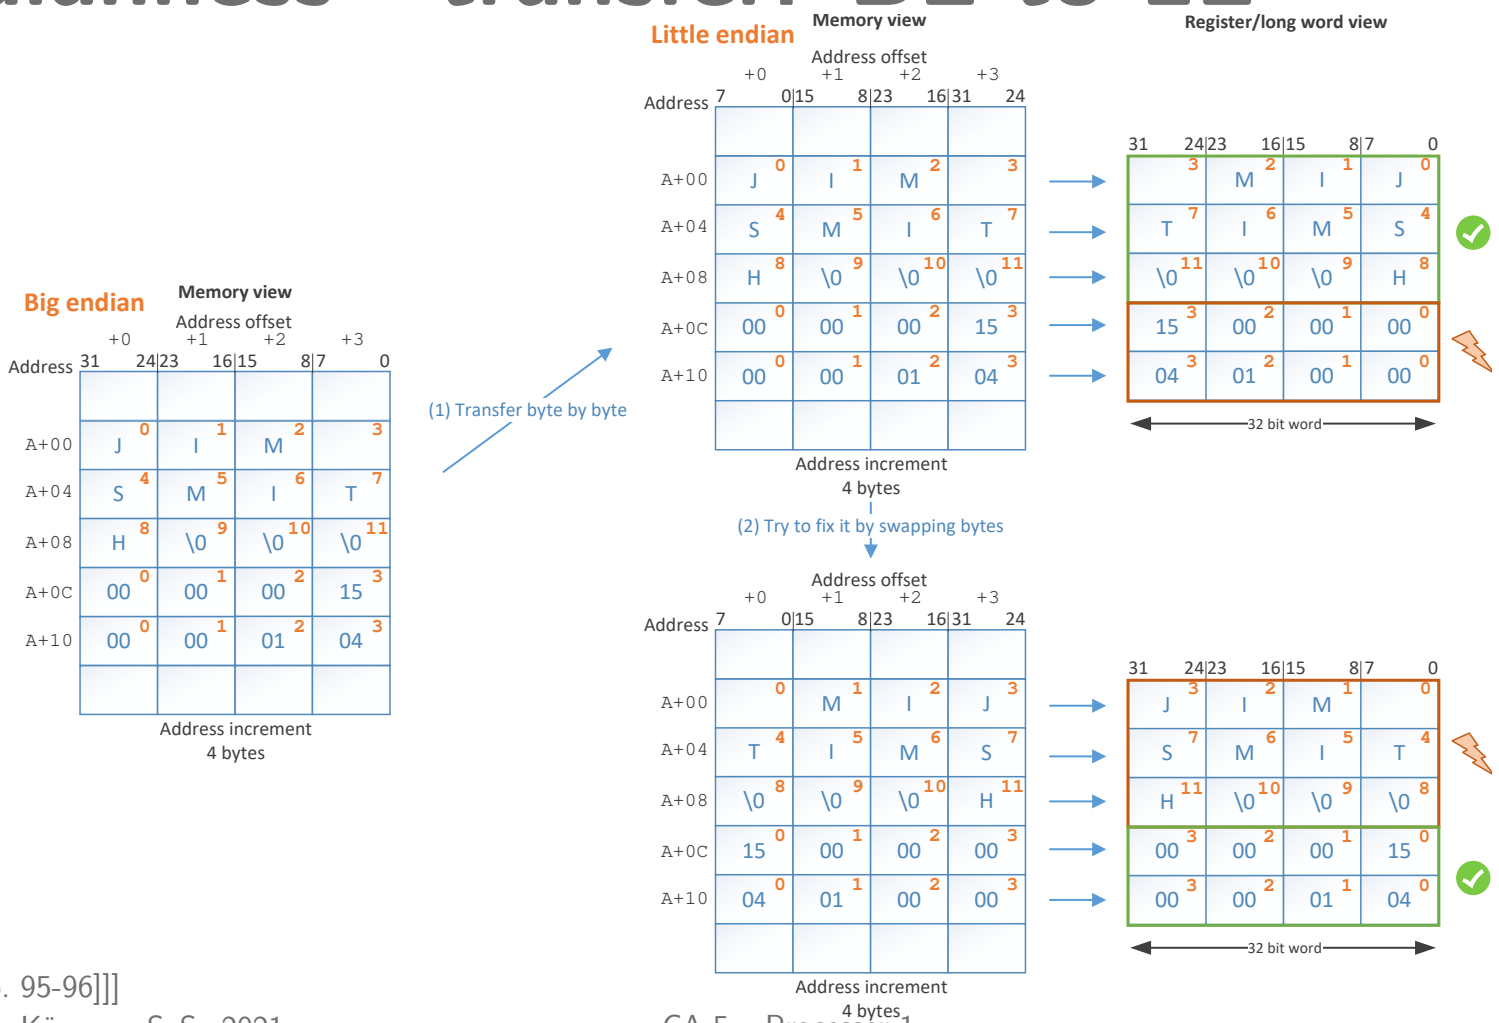
All right? ⇒ 

Question? ⇒  and use **chat**

or

**speak** *after* I ask you to

# Endianness - transfer: BE to LE





# Endianness - problem

## Problem can occur if

- Different data types are mixed: numbers, strings, or other data types
- Data type consists of more than one byte (multi byte,  $\geq 2$ )
- Data are transferred between BE/LE systems

## No problem occurs if

- Single byte data is transferred byte by byte (e.g. ASCII)
- Data is transferred within same endianness (LE  $\rightarrow$  LE, BE  $\rightarrow$  BE)

# Endianness - conclusion

Without the knowledge about the data types and the alignment, a transfer between BE/LE systems is not feasible.

Tanenbaum: „*There is no easy solution to this*“ [1, p. 96]

# Endianness - possible solutions


## Possible solution


- **Know** the endianness (e.g. **meta data!**)
- **Transfer** byte by byte (no problem for single byte data)
- If endianness is different and a long word is transferred: additionally **swap** the bytes

## Some examples:

- Network order: always BE
- Java: always BE; for transfer with others, `ByteOrder` can be set
- Unicode UTF-16/32: uses a BOM (byte order mark)
- TIF files: BE/LE identifier in header
- RPC (remote procedure call): marshalling (data as byte stream) solves the problem by using meta data

# Questions?

All right? ⇒ 

Question? ⇒  and use **chat**

or

**speak** *after* I ask you to

# Summary and outlook

## Summary

- Processor architecture
- Exception and interrupt handling
- Instruction set architecture
- Endianness

## Outlook

- Processor registers
- Processor examples
- Addressing modes