



# Prof. Dr. Florian Künzner

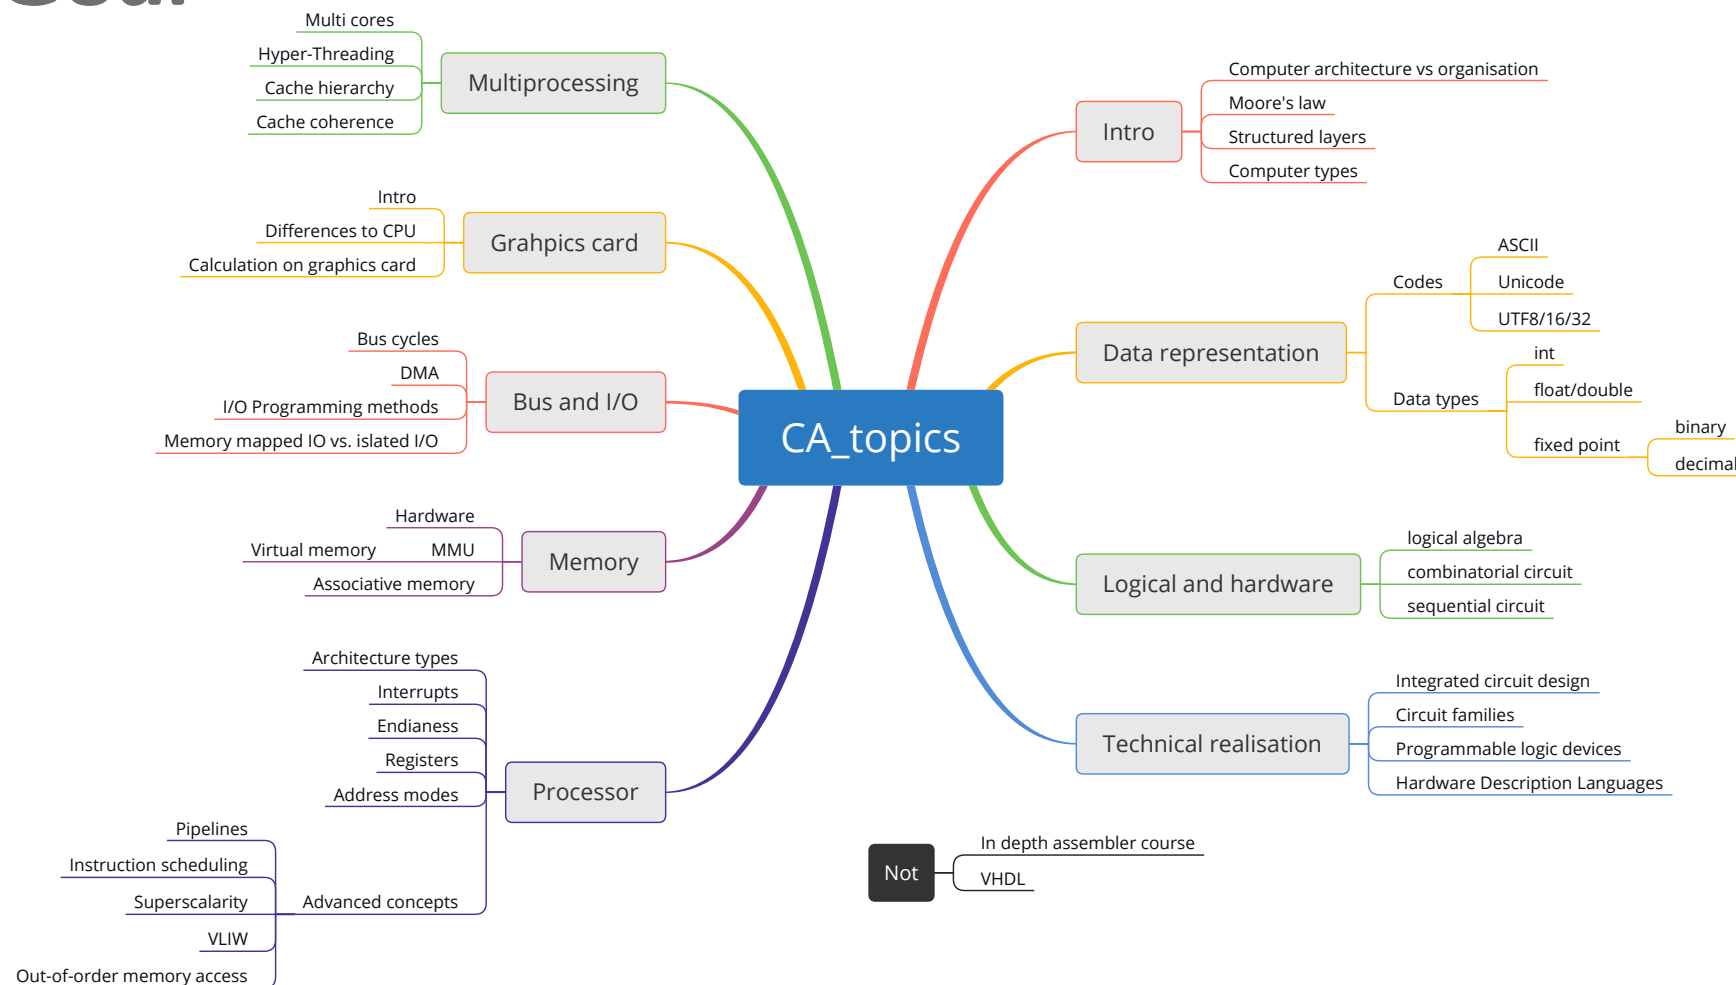
Technical University of Applied Sciences Rosenheim, Computer Science

## CA 7 – Processor 3

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier



# Goal



# Modern CPU techniques


Why are modern (single core)  
processors **so fast** nowadays?

# Goal

## CA::Processor 3

- Pipelining
- Instruction scheduling
- Superscalar architecture
- VLIW
- Out-of-order memory access

# Pipelining



## Idea of pipelining:

- Splitting an instruction processing cycle into individual phases
- Processing of these phases in subsequent stations
- Instructions run through the individual steps
- Each step does a partial processing of an instruction
- Time overlap of instruction processing!

**Pipelining is done in hardware!**

# Pipelining - example

This is a very simplified (classical RISC) pipeline, since modern processors contain many more pipeline phases.

1. Fetch instruction	2. Load operands	3. Execute instruction	4. Memory access	5. Save result into register	
INST1					$t_1$
INST2	INST1				$t_2$
INST3	INST2	INST1			$t_3$
INST4	INST3	INST2	INST1		$t_4$
INST5	INST4	INST3	INST2	INST1	$t_5$
INST6	INST5	INST4	INST3	INST2	$t_6$

time ↓

## Processing

- The processing of an instruction takes 5 clock cycles.
- On each clock cycle one instruction is started and one ends
- Without pipelining: Every 5 clock cycles, an instruction would be started and ended.

# Pipelining - Problems

The execution of instructions in a pipeline is disturbed by:

- Jump instructions
- Interrupts
- Data dependencies (the next instruction requires data from the previous)

**Special hardware** is often available for this purpose, in order to **minimize the disruption of the pipeline** as much as possible:

- On conditional jumps: instructions for both alternatives are fetched and evaluated
- A kind of cache remembers which alternative from the past seems more likely (**branch prediction**)
- Try to avoid conditional jumps—instead use **conditional instructions** (e.g. CMOV, MOVCC)

# Questions?

All right?



Question?



and use **chat**

or

**speak** *after* I  
ask you to



# Instruction scheduling

Idea: **The Compiler changes the order of the instructions**

- Change the processing order of the instructions to **maximize the pipeline utilisation**
- Important: The **semantics** of the high-level language program **remains intact** (as-if rule).

# Instruction scheduling - Example

Unscheduled:

Nr	Opcode	Operand	Comment	C code
1	LOAD	R7, (R30)	R7 = (R30)	x++;
2	ADD	R7, R7, 1	R7 = R7+1	
3	STORE	R7, (R30)	(R30) = R7	
4	LOAD	R8, (R31)	R8 = (R31)	y += 4;
5	ADD	R8, R8, 4	R8 = R8+4	
6	STORE	R8, (R31)	(R31) = R8	
7	CMP	R10, R6, R2	R10 = R6 < R2	if(R6 < R2) GOTO M if R10 < 0
8	BLT	R10, M	GOTO M if R10 < 0	

[source (with changes): Alsup: Motorola's 88000 Family Architecture, IEEE Micro, June 1990]

Pipeline (with 4 phases):

Fetch	Load	Execute	Save	
1				t <sub>1</sub>
2	1			t <sub>2</sub>
-	-	1		t <sub>3</sub>
-	-	-	1	t <sub>4</sub>
3	2	-	-	t <sub>5</sub>
-	-	2	-	t <sub>6</sub>
-	-	-	2	t <sub>7</sub>
4	3	-	-	t <sub>8</sub>
5	4	3	-	t <sub>9</sub>
-	-	4	3	t <sub>10</sub>
-	-	-	4	t <sub>11</sub>
6	5	-	-	t <sub>12</sub>
-	-	5	-	t <sub>13</sub>
-	-	-	5	t <sub>14</sub>
7	6	-	-	t <sub>15</sub>
8	7	6	-	t <sub>16</sub>
-	-	7	6	t <sub>17</sub>
-	-	-	7	t <sub>18</sub>
-	8	-	-	t <sub>19</sub>
-	-	8	-	t <sub>20</sub>
-	-	-	8	t <sub>21</sub>
M				t <sub>22</sub>

Total clock cycles: 21

# Instruction scheduling - Example

Scheduled:

Nr	Opcode	Operand	Old order
1	LOAD	R7, (R30)	1
2	CMP	R10, R6, R2	7
3	LOAD	R8, (R31)	4
4	ADD	R7, R7, 1	2
5	ADD	R8, R8, 4	5
6	STORE	R7, (R30)	3
7	BLT	R10, M	8
8	STORE	R8, (R31)	6

[source (with changes): Alsup: Motorola's 88000 Family Architecture, IEEE Micro, June 1990]


Pipeline (with 4 phases):

Fetch	Load	Execute	Save	
1				$t_1$
2	1			$t_2$
3	2	1		$t_3$
4	3	2	1	$t_4$
5	4	3	2	$t_5$
-	-	4	3	$t_6$
6	5	-	4	$t_7$
7	6	5	-	$t_8$
8	7	6	5	$t_9$
-	8	7	6	$t_{10}$
-	-	8	7	$t_{11}$
-	-	-	8	$t_{12}$
M				$t_{13}$

Total clock cycles: 12



# Example: AMD Bobcat pipelining



[Source: AMD Bobcat & Bulldozer Hot Chips Presentations 2011, available on youtube, last access 6.3.2015;  
Presentation: AMD Bobcat]

# Questions?

All right?



Question?




and use **chat**

or

**speak** *after* I  
ask you to

# Superscalar architecture

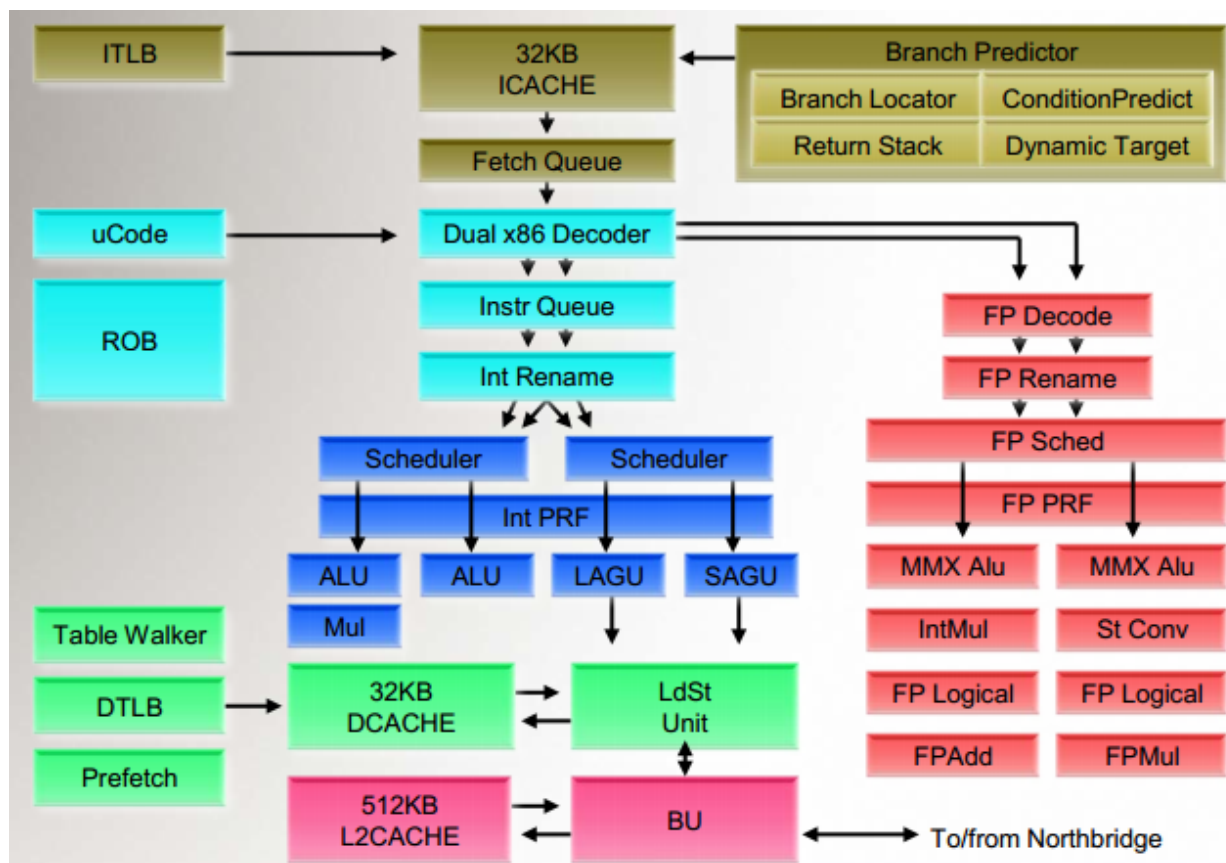


## Superscalar architecture:

- The CPU provides **several pipelines**
- The pipelines may be specialised for integer or floating point operations
- The CPU loads **multiple instructions** at each clock cycle
- Everything is done automatically inside the CPU: **programming model stays unchanged.**



# Example: AMD Bobcat superscalarity



[Source: AMD Bobcat & Bulldozer Hot Chips Presentations 2011, available on youtube, last access 6.3.2015]

Notice: AGU = address generation units, LAGU = AGU for loading operations, SAGU = AGU for storage operations

# Questions?

All right?



Question?



and use **chat**


or

**speak** *after* I  
ask you to





# VLIW - Very Long Instruction Word



## VLIW architecture:

- Parallel **operations are packed** into a very long instruction
- Packing is done by the **compiler at compile time**: **programming model changed!**
- The partial instructions of a VLIW must correspond to the functional units of the hardware.
- Its used by the **Intel Itanium** architecture (EPIC - Explicitly parallel instruction computer)

# Questions?

All right?



Question?



and use **chat**

or

**speak** *after* I  
ask you to

# Out-of-order memory access

**Idea:** The **hardware (CPU)** changes the order of the instructions: with the goal to better utilise the hardware

## Details:

- Memory access (instruction) is **not** done in the specified order
- Some accesses are “**too early**” (e.g. speculative load)
- Some accesses are “**too late**” (e.g. posted writes)
- Attention: The **semantics must be guaranteed!**

# Out-of-order memory access

## In-order execution

- 1 **Fetch** instruction
- 2 **Load** operands (may take some cycles)
- 3 **Execute** instruction by the appropriate functional unit
- 4 **Save** the results into the register

## Out-of-order execution

- 1 **Fetch** instruction
- 2 **Queue instruction** to an instruction queue (instruction buffer)
- 3 **Load** operands: **Instructions wait** in the queue until its input operands are available (a instruction can leave the queue before an older instructions)
- 4 **Execute** instruction by the appropriate functional unit
- 5 **Queue results**
- 6 **Save** the results into the register—**only after all older instructions have their results saved** to the register

# Out-of-order memory access - Example

## Core 1: (producer)

```
1 byte data[1000];  
2 byte shared_data[1000];  
3 bool flag = false;  
4  
5 // shared_data = data;  
6 memcpy(shared_data, data, 1000);  
7  
8 flag = true; //data ready
```

## Core 2: (consumer)

```
1 byte data_copy[1000];  
2 byte* shared_data; //attach  
3 bool* flag; //attach  
4  
5  
6  
7  
8  
9 while (flag != true) {} //busy wait  
10  
11 // data_copy = shared_data;  
12 memcpy(data_copy, shared_data, 1000);
```

Any problems with that code?

# Out-of-order memory access - Problem

**Problem:** Through **asynchronous execution** units (threads, processes, interrupts, ...), the out-of-order memory access can cause **unpredictable behaviour** in (concurrent programs, device drivers).

## Prevent **hardware** from **reordering**

- Use of a **memory barrier**: a special hardware instruction (`memory_barrier`, `FENCE`)
- Forces the completion of all outstanding (instructions) memory accesses

## Prevent **compiler** from **reordering** (instruction scheduling)

- Use of the **volatile** keyword in C
- Forces the compiler to
  - not omit reads from and writes to volatile memory locations,
  - nor may it reorder read/writes relative to other such actions for the same volatile location
- Notice: The **volatile** keyword does not prevent the hardware to reorder instructions

# Out-of-order memory access - Example

## Core 1: (producer)

```
1 byte data[1000];  
2 volatile byte shared_data[1000];  
3 volatile bool flag = false;  
4  
5 // shared_data = data;  
6 memcpy(shared_data, data, 1000);  
7 __asm { memory_barrier }  
8 flag = true; //data ready
```

## Core 2: (consumer)

```
1 byte data_copy[1000];  
2 volatile byte* shared_data; //attach  
3 volatile bool* flag; //attach  
4  
5  
6  
7  
8  
9 while (flag != true) {} //busy wait  
10 __asm { memory_barrier }  
11 // data_copy = shared_data;  
12 memcpy(data_copy, shared_data, 1000);
```

- The use of `__asm { memory_barrier }` and `volatile` should fix the problem.
- **Best solution:** use of a **semaphore** or a monitor, then all this is already solved for you (by OS software).

# Questions?

All right?



Question?



and use **chat**

or

**speak** *after* I  
ask you to



# Overview

## Technique

Pipelining

Instruction scheduling

Superscalarity

VLIW

Out-of-order memory access

## How

parallelisation of parts of an instruction

reordering of instructions

parallelisation on instruction level

parallelisation on instruction level

reordering of instructions

## Unit

hardware

compiler

hardware

compiler

hardware

## Time

runtime

compile time

runtime

compile time

runtime

# Questions?

All right?



Question?



and use **chat**

or

**speak** *after* I  
ask you to

# Summary and outlook

## Summary

- Pipelining
- Instruction scheduling
- Superscalar architecture
- VLIW
- Out-of-order memory access

## Outlook

- Memory