



Prof. Dr. Florian Künzner

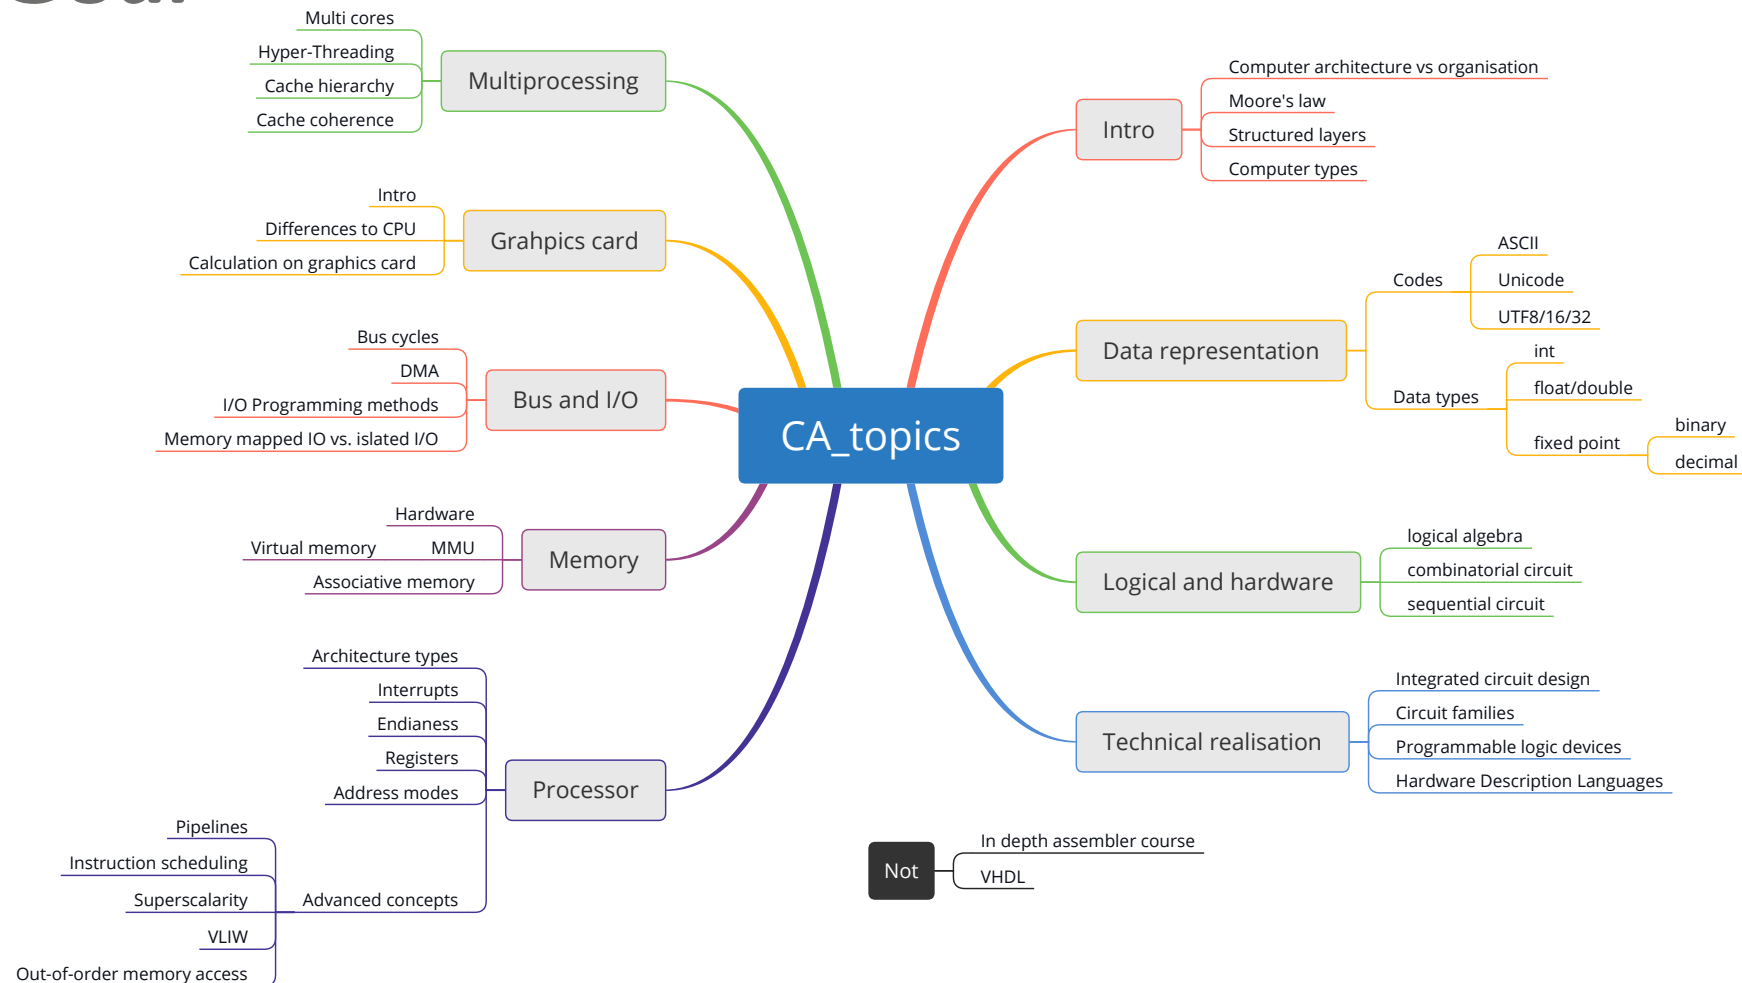
Technical University of Applied Sciences Rosenheim, Computer Science

CA 9 – MMU

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier



Goal




Goal

CA::MMU

- MMU
- Virtual addresses
- 1 level page table
- 2 level page table
- N level page table

Overview

MMU - Memory management unit



- Special hardware inside the CPU case
- Dynamic address conversion at runtime

Overview

Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
 - Only one type of addresses
 - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access



Situation without virtual addresses

We try to recap the situation without virtual addresses.

C example

```
1 int x;  
2 int y;  
3  
4 x = x + y;
```

Compiled assembler example

```
1 ;CODE of ADD  
2 ;real address of x  
3 ;real address of y  
4  
5 X: ;value of x  
6 Y: ;value of y
```

Situation without virtual addresses

Absolute addresses

Compiled assembler

```

1  ...
2  1000: ;code of ADD
3  1004: ;real address of x
4  1008: ;real address of y
5  ...
6  1024: ;value of x
7  1028: ;value of y

```

Facts

- An *.elf/*.exe with exactly this code and addresses (called image) is saved on the harddisk
- The *.elf/*.exe has to be loaded exactly at this memory addresses (starting with 1000)

Problems

- Program can't be started twice at the same time
- Different programs can't use the same address

Situation without virtual addresses

Position independent code (i.e. PC relative)

Compiled assembler

```

1  ...
2  1000: ;code of ADD
3  1004: 20 ;(offset of x to PC)
4  1008: 20 ;(offset of y to PC)
5  ...
6  1024: ;value of x
7  1028: ;value of y
  
```

Facts

- The operands are addressed with an relative offset to the PC
- The program can be loaded on any address in the memory
- It is also possible to relatively address to the SP

Problems

- Address translation required to obtain addresses of x and y at runtime

Situation without virtual addresses

Relocation

Compiled assembler

```

1  ...
2  2000: ;code of ADD
3  2004: 2024 ;real address of x
4  2008: 2028 ;real address of y
5  ...
6  2024: ;value of x
7  2028: ;value of y

```


Facts


- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

Problems

- Address translation required on startup

Questions?

All right? \Rightarrow 

Question? \Rightarrow  and use **chat**


or

speak *after* I
ask you to

Towards virtual addresses

Towards virtual addresses

Idea 1: Separation of virtual and real address space



Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

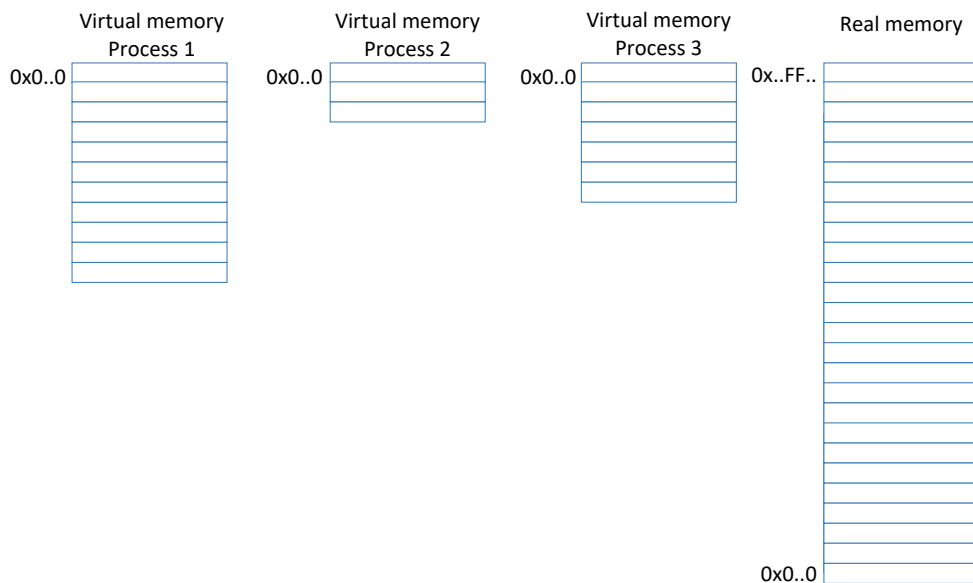
Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



Towards virtual addresses

Idea 2: Introduce pages and frames



Facts

- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space $>$ real address space

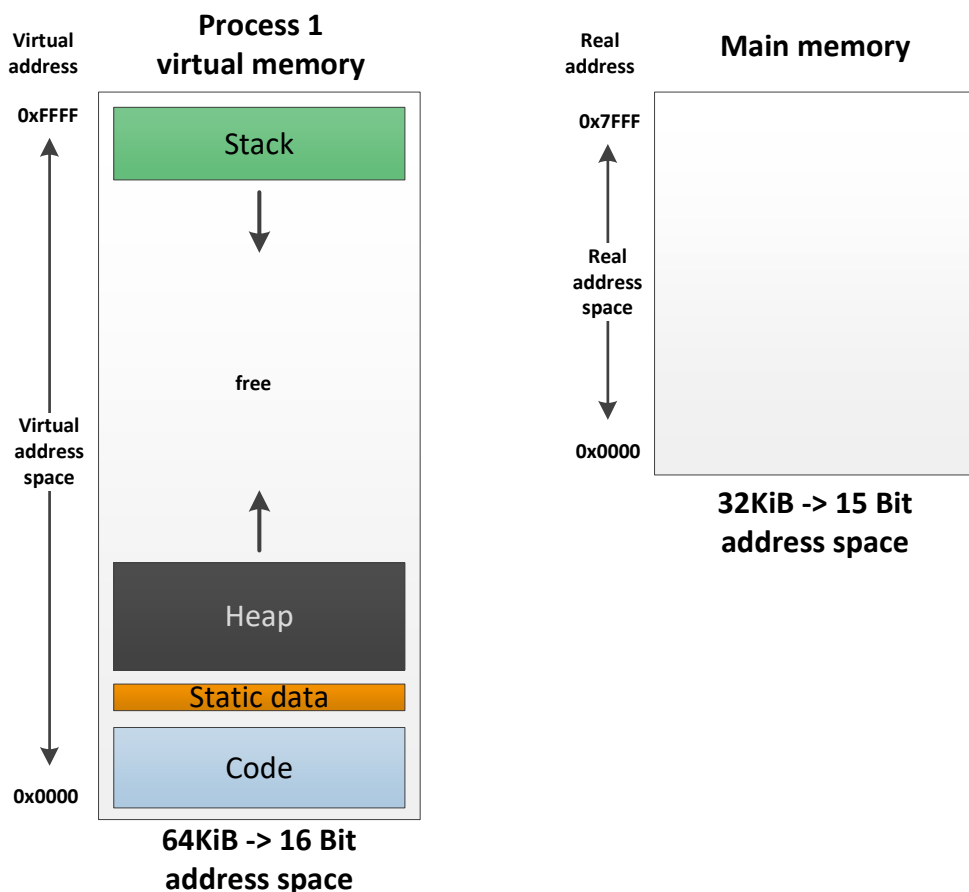
Goal

Understand in detail virtual addresses and its translation to real addresses

Procedure

- Introduction of terms
- 1 level page table: with a small memory 64 KiB (virtual) and 32 KiB (real)
- 2 level page table: Intel x86/32 bit architecture
- 4 level page table: Intel x86/64 architecture (AMD/Intel)

Virtual address space



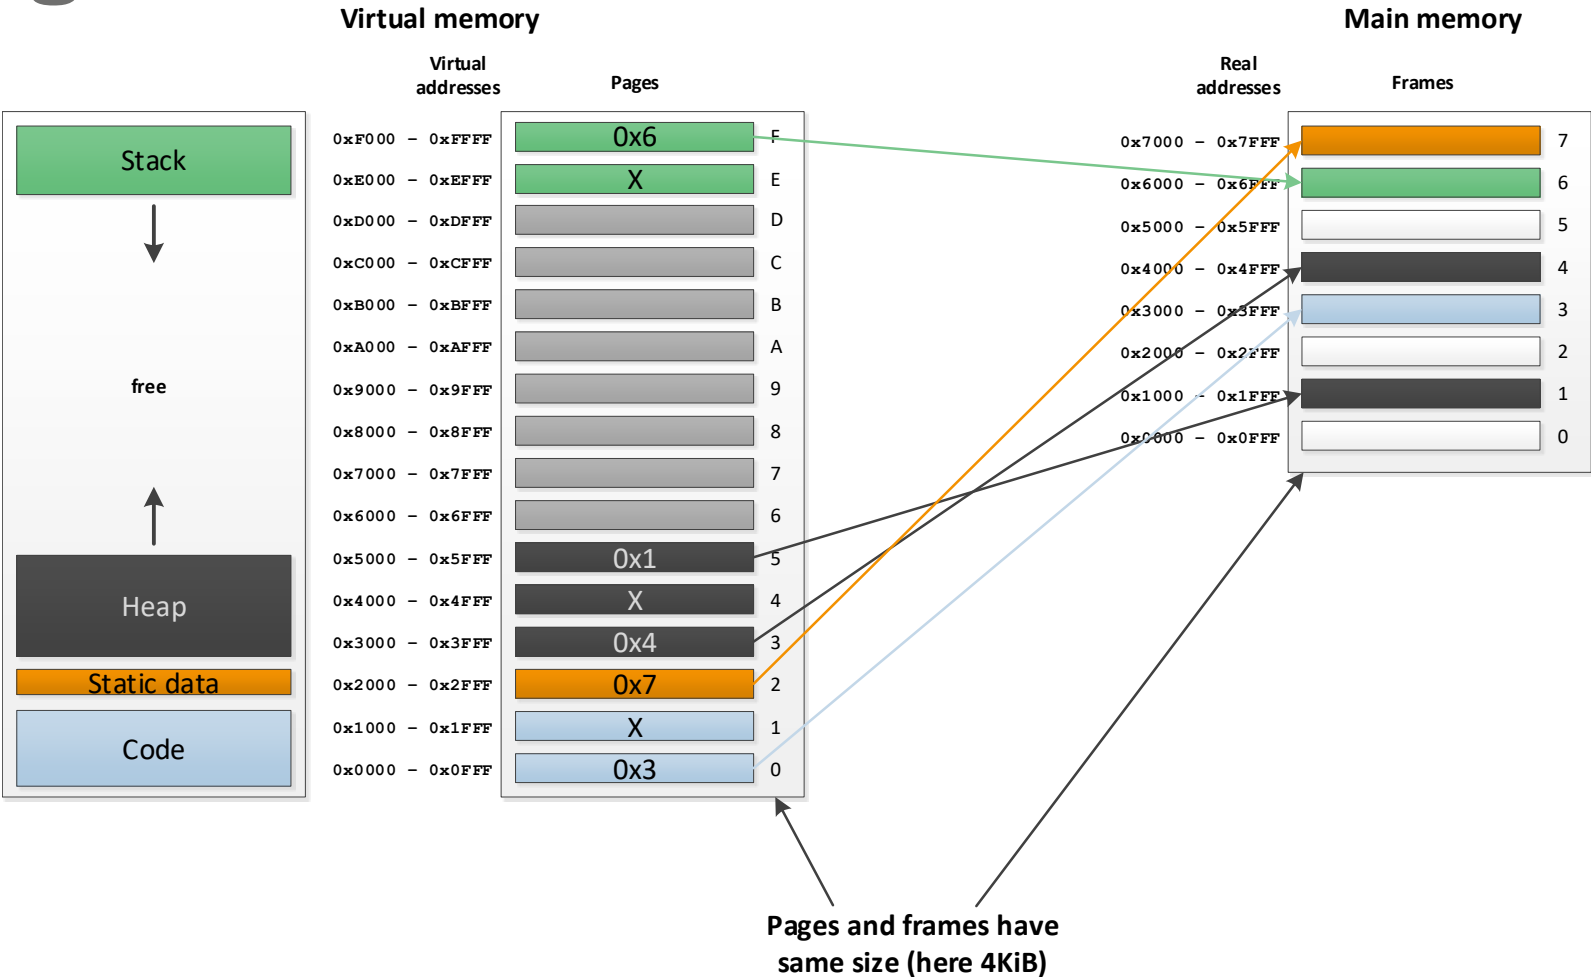
Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

Real address space

- All physically available addresses
- Size depends on available memory

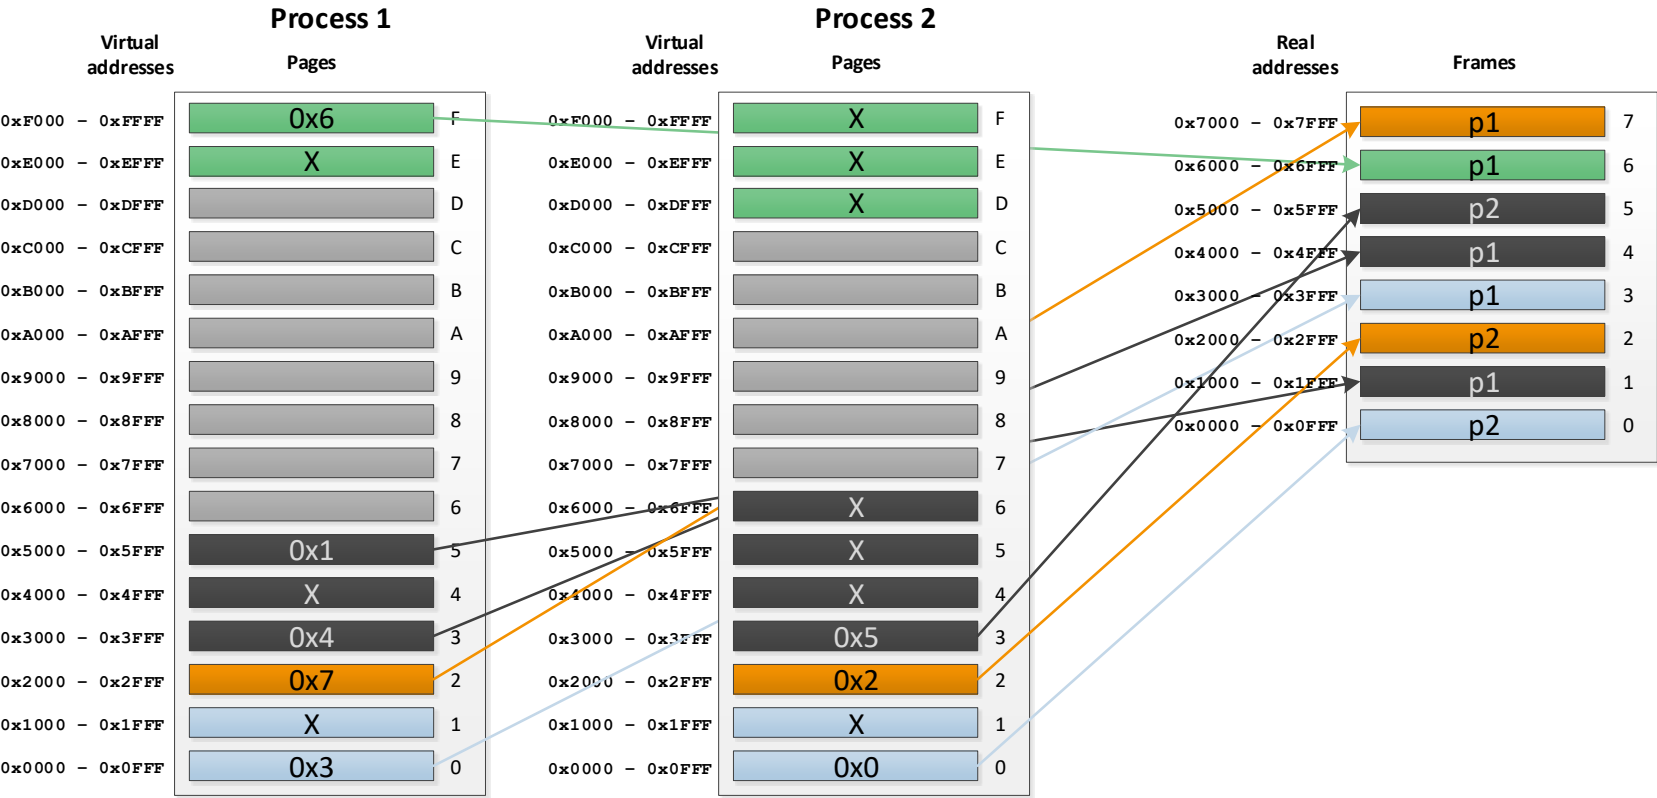
Pages and frames



Multiprogramming

Virtual memory

Main memory



Questions?

All right?



Question?




and use **chat**

or


speak *after* I
ask you to

Let's start with a simple
1 level page table.

Virtual address translation

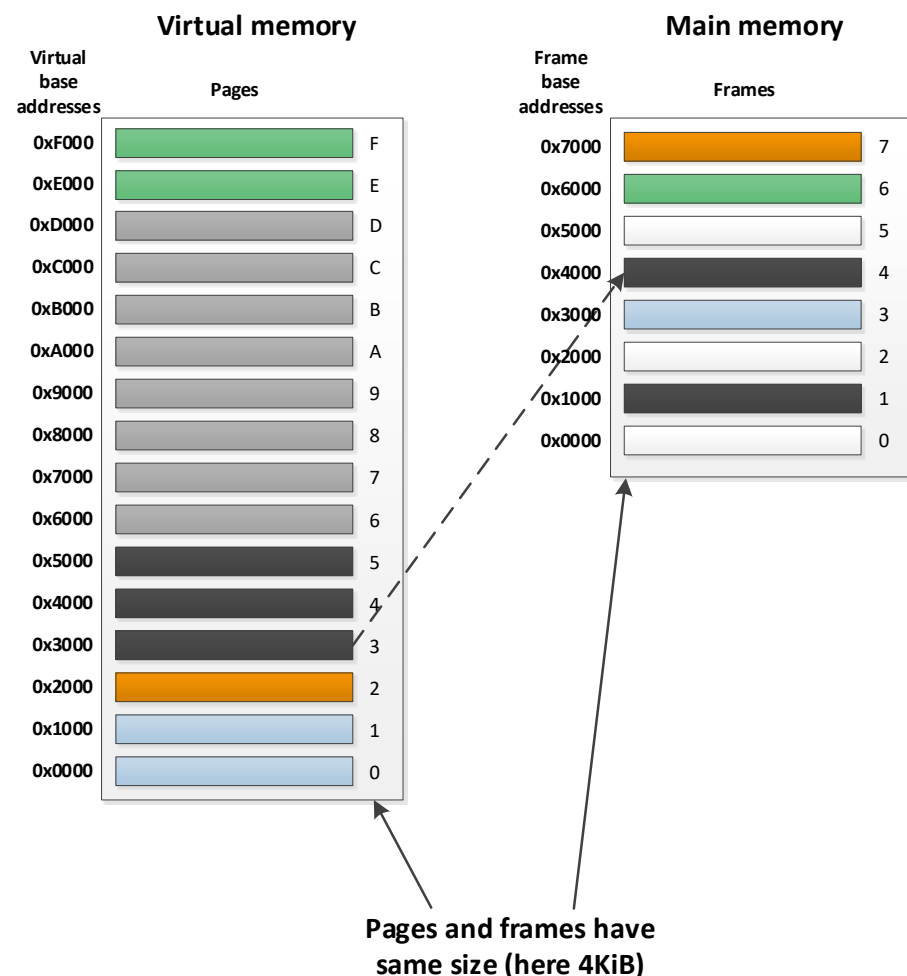


Page table





Virtual address translation: example



Given:

- Virtual address space: 64 KiB
- Real address space: 32 KiB
- Virtual address: 0x3D05

Questions:

- 1 Address width (virtual and real)?
- 2 Offset: number of bits?
- 3 Page/frame numbers: number of bits?
- 4 Number of pages and frames?
- 5 Real address?

Results:

- 1 virtual: 16 bit/real: 15 bit
- 2 offset: 12 bit
- 3 pages: 4 bit/frames: 3 bit
- 4 pages: $2^4 = 16$ /frames: $2^3 = 8$
- 5 real address: 0x4D05

1 level page table


Any problems with a 1 level page table?


- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

Possible solutions

- Increase page/frame size
 - i.e. 1 MiB, 2 MiB, ...
 - \Rightarrow less pages, less entries in page table, smaller page table
- Multilevel page table
 - 32 bit: 2 level page table with 4 KiB page size
 - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

Questions?

All right? \Rightarrow 


Question? \Rightarrow  and use **chat**

or

speak *after* I ask you to

Let's proceed with a
2 level page table.

2 level page table



Intel x86/32 bit

- Usually uses 2 level page tables
- Offset: 12 bit
- $2^{12} = 4 \text{ KiB}$ page size
- 1 page directory per process
- $2^{10} = 1024$ page tables per process

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

2 level page table entries

PDE - page directory entry

- Page-Table Base Address (20 bit)
= address of page table

PTE - page table entry

- Physical-Page Base Address (20 bit)
= address of frame




Figure 5-5. 4-Kbyte PDE—Non-PAE Paging Legacy-Mode




Figure 5-6. 4-Kbyte PTE—Non-PAE Paging Legacy-Mode


[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

Fields

- P = present (loaded into memory)
- R/W = read/write
- U/S = user/supervisor
- PWT = page level writethrough
- PCD = page level cache disable
- A = accessed
- D = dirty
- AVL = available to software (for OS)

Questions?

All right? \Rightarrow 


Question? \Rightarrow  and use **chat**

or

speak *after* I
ask you to

And finally with
N level page tables.

4 level page table



[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]




N level page table


More details:

AMD64 Architecture Programmer's Manual Volume 2:
System Programming (link)

Lets have a look on page 121, table 5-1.

Questions?

All right? \Rightarrow 

Question? \Rightarrow  and use **chat**

or

speak *after* I
ask you to

Summary and outlook

Summary

- MMU
- Virtual addresses
- 1/2/N level page table

Outlook

- Memory hierarchy
- Associative memory
- Transalation lookaside buffer
- Cache
- Memory protection