

Start: 8:01

Technische  
Hochschule  
**Rosenheim**

Technical University of Applied Sciences



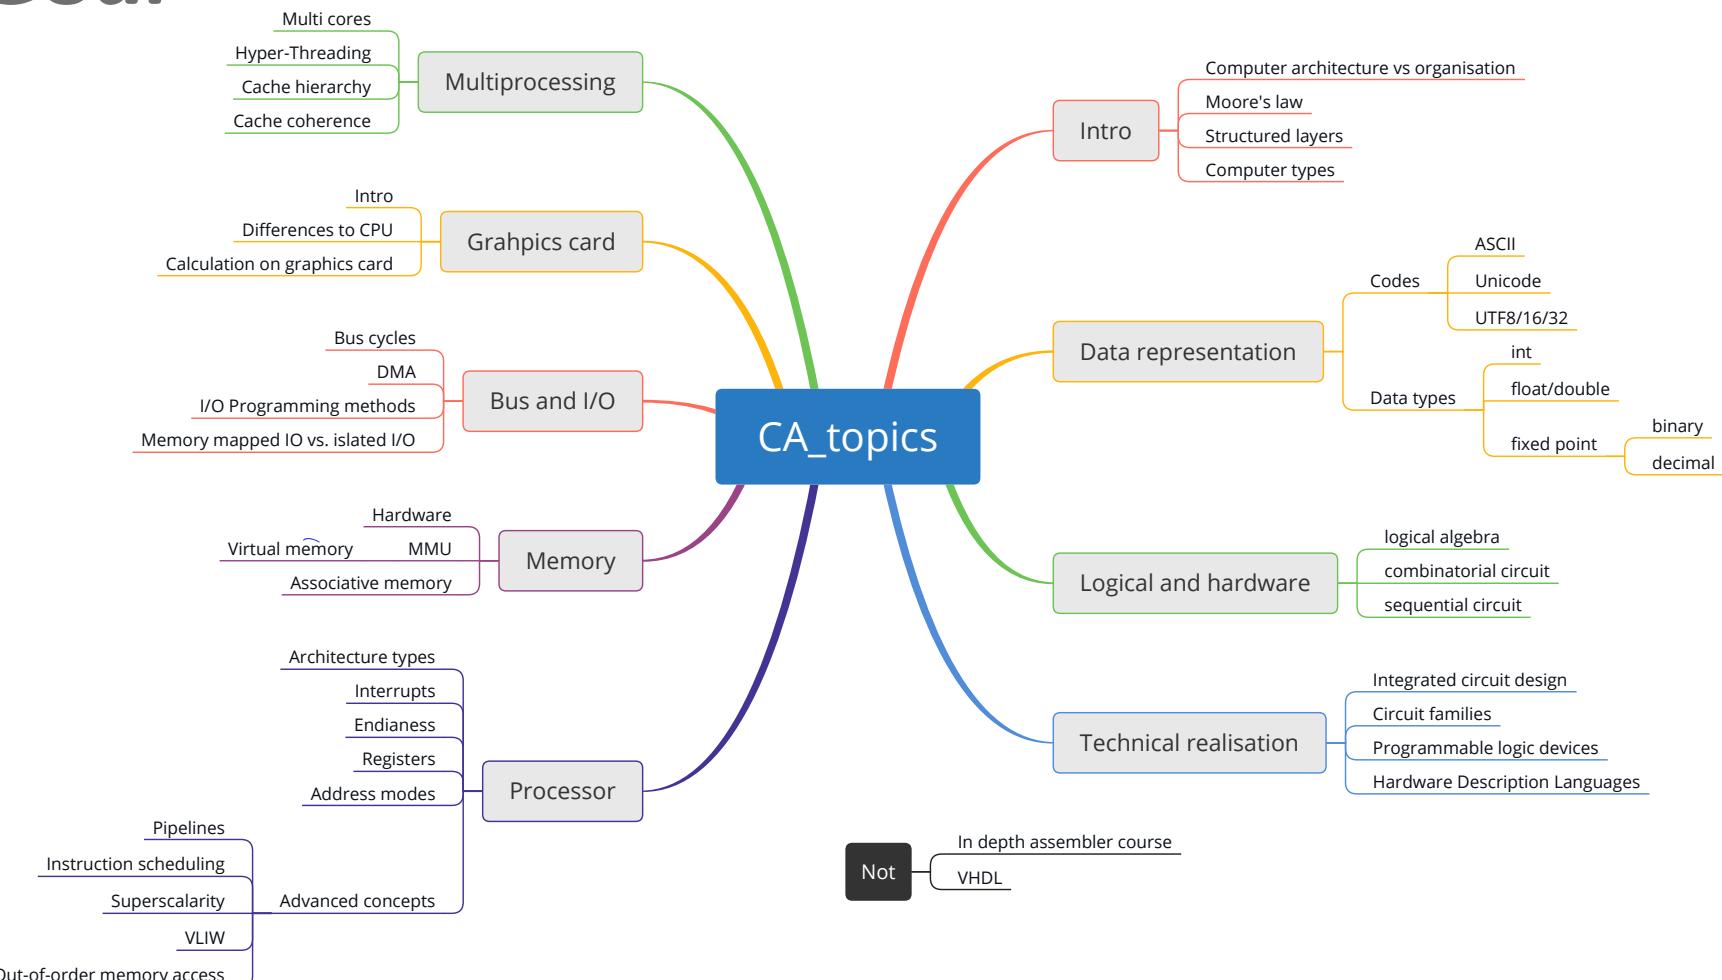
# Prof. Dr. Florian Künzner

Technical University of Applied Sciences Rosenheim, Computer Science

## CA 9 – MMU

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier

# Goal





# Goal

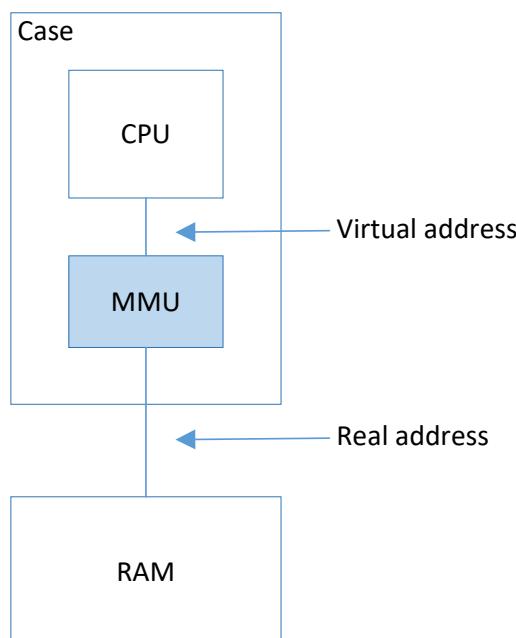
## CA::MMU

- MMU
- Virtual addresses
- 1 level page table
- 2 level page table
- N level page table



# Overview

## MMU - Memory management unit



- Special hardware inside the CPU case
- Dynamic address conversion at runtime



# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access



# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access



# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access

# Overview

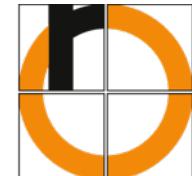
## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access

# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access



# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access

# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access

# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access

# Overview

## Goals of virtual memory management:

- Aggregation of different memory sources
- Suitable abstraction for developers
  - Only one type of addresses
  - Linearly accessible memory
- Use of full theoretical available memory
- Multiple processes: Multiprogramming
- Shared libraries (DLL, so)
- Memory protection
- Parallelism (multiple cores: parallel processes)
- High performance memory access



# Situation without virtual addresses

We try to recap the situation without virtual addresses.

## C example

```
1 int x;  
2 int y;  
3  
4 x = x + y;
```

## Compiled assembler example

```
1 ;CODE of ADD  
2 ;real address of x  
3 ;real address of y  
4  
5 X: ;value of x  
6 Y: ;value of y
```



# Situation without virtual addresses

We try to recap the situation without virtual addresses.

## C example

```
1 int x;  
2 int y;  
3  
4 x = x + y;
```

## Compiled assembler example

```
1 ;CODE of ADD  
2 ;real address of x  
3 ;real address of y  
4  
5 X: ;value of x  
6 Y: ;value of y
```



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x 1024  
4 1008: ;real address of y 1028  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\*.exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\*.exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\*.exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\*.exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\*.exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\*.exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\* .exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\* .exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\* .exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\* .exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\* .exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\* .exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address



# Situation without virtual addresses

## Absolute addresses

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: ;real address of x  
4 1008: ;real address of y  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- An \*.elf/\* .exe with exactly this code and addresses (called image) is saved on the harddisk
- The \*.elf/\* .exe has to be loaded exactly at this memory addresses (starting with 1000)

### Problems

- Program can't be started twice at the same time
- Different programs can't use the same address

# Situation without virtual addresses

## Position independent code (i.e. PC relative)

## Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

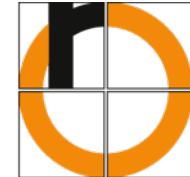
# Situation without virtual addresses

## Position independent code (i.e. PC relative)

## Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

## Facts



# Situation without virtual addresses

## Position independent code (i.e. PC relative)

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- The operands are addressed with an relative offset to the PC
- The program can be loaded on any address in the memory
- It is also possible to relatively address to the SP

### Problems

- Address translation required to obtain addresses of x and y at runtime

# Situation without virtual addresses

## Position independent code (i.e. PC relative)

# Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

## Facts

- The operands are addressed with an relative offset to the PC
  - The program can be loaded on any address in the memory
  - It is also possible to relatively address to the SP

# Situation without virtual addresses

## Position independent code (i.e. PC relative)

## Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

## Facts

- The operands are addressed with an relative offset to the PC
  - The program can be loaded on any address in the memory
  - It is also possible to relatively address to the SP

# Situation without virtual addresses

## Position independent code (i.e. PC relative)

## Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

## Facts

- The operands are addressed with an relative offset to the PC
  - The program can be loaded on any address in the memory
  - It is also possible to relatively address to the SP

## Problems



# Situation without virtual addresses

## Position independent code (i.e. PC relative)

### Compiled assembler

```
1 ...  
2 1000: ;code of ADD  
3 1004: 20 ;(offset of x to PC)  
4 1008: 20 ;(offset of y to PC)  
5 ...  
6 1024: ;value of x  
7 1028: ;value of y
```

### Facts

- The operands are addressed with an relative offset to the PC
- The program can be loaded on any address in the memory
- It is also possible to relatively address to the SP

### Problems

- Address translation required to obtain addresses of x and y at runtime



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Situation without virtual addresses

## Relocation

### Compiled assembler

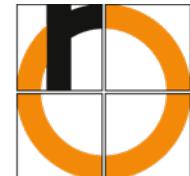
```
1 ...  
2 2000: ;code of ADD  
3 2004: 2024 ;real address of x  
4 2008: 2028 ;real address of y  
5 ...  
6 2024: ;value of x  
7 2028: ;value of y
```

### Facts

- The program is compiled with absolute addresses
- When the program is loaded into memory, the addresses are changed.
- The relocation table contains the information which addresses have to be changed

### Problems

- Address translation required on startup



# Questions?

All right?  $\Rightarrow$



Question?  $\Rightarrow$



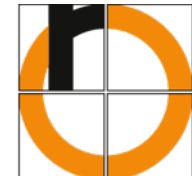
and use **chat**

or

*speak after I  
ask you to*

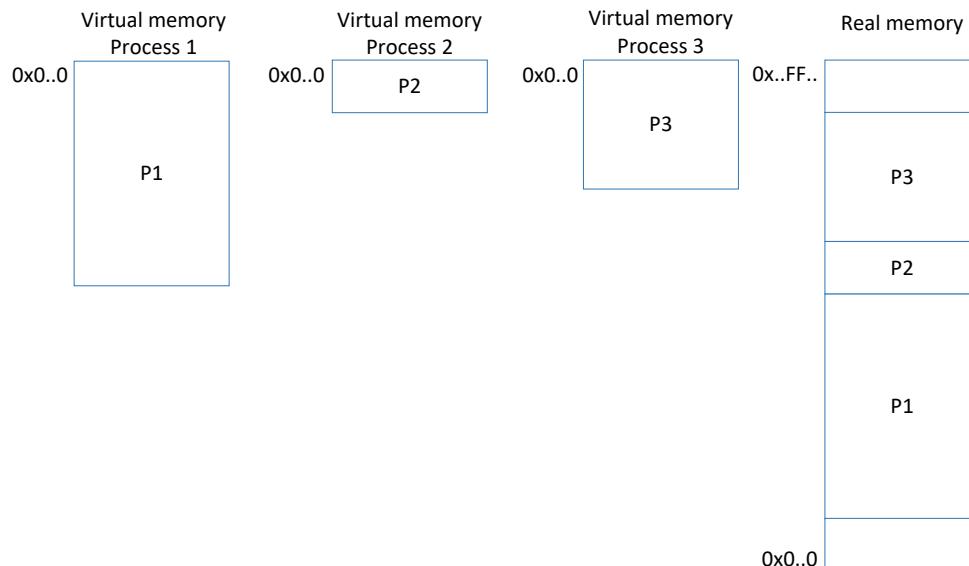


# Towards virtual addresses



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

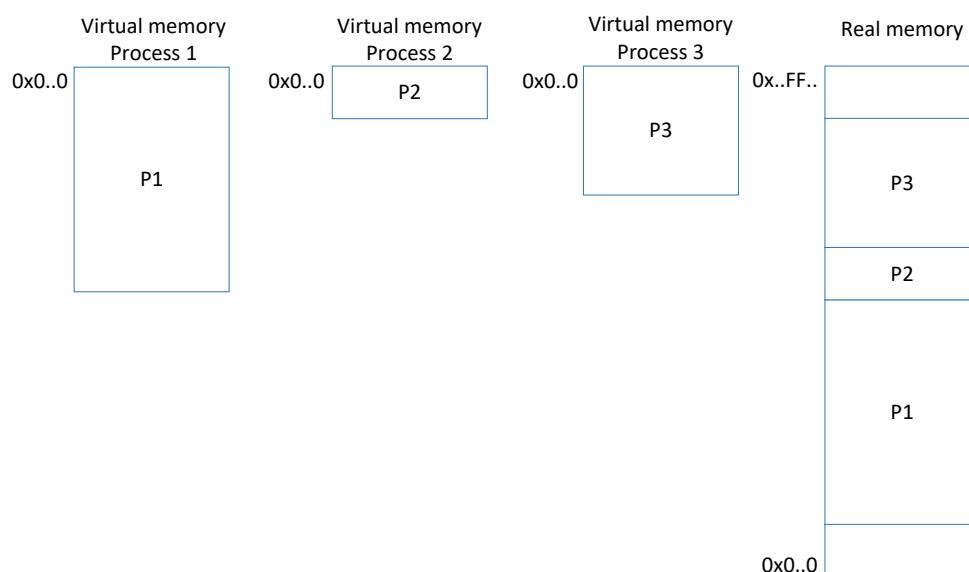
### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

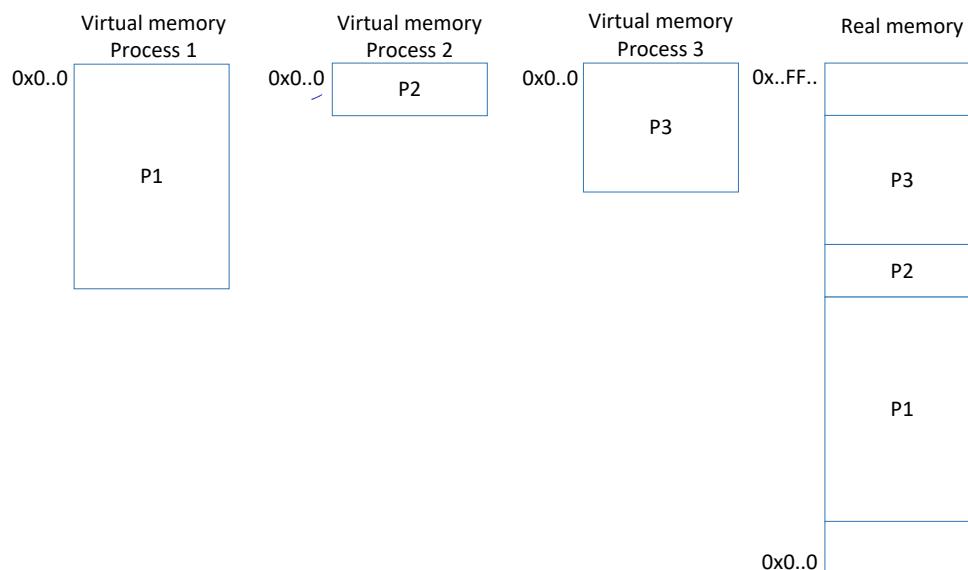
### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

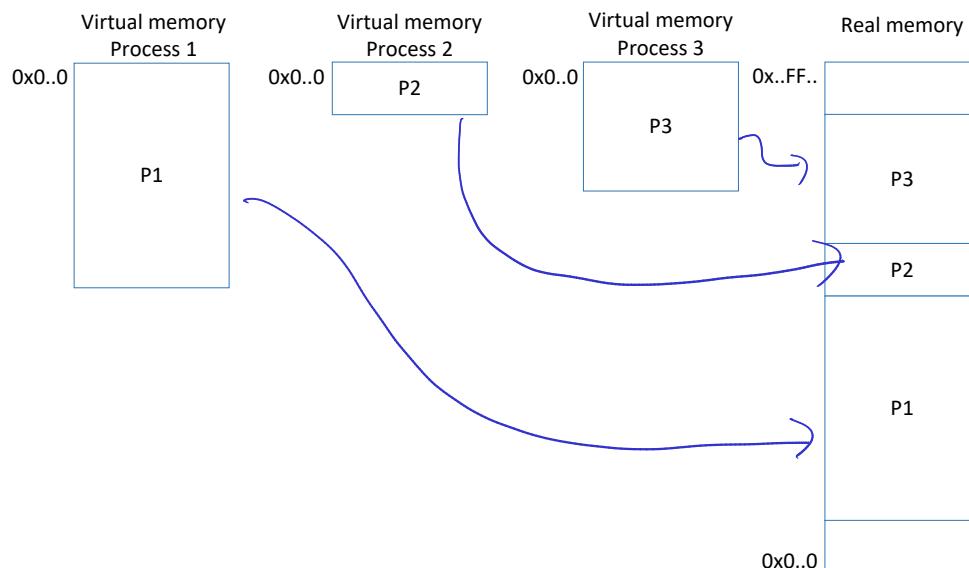
- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory

# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

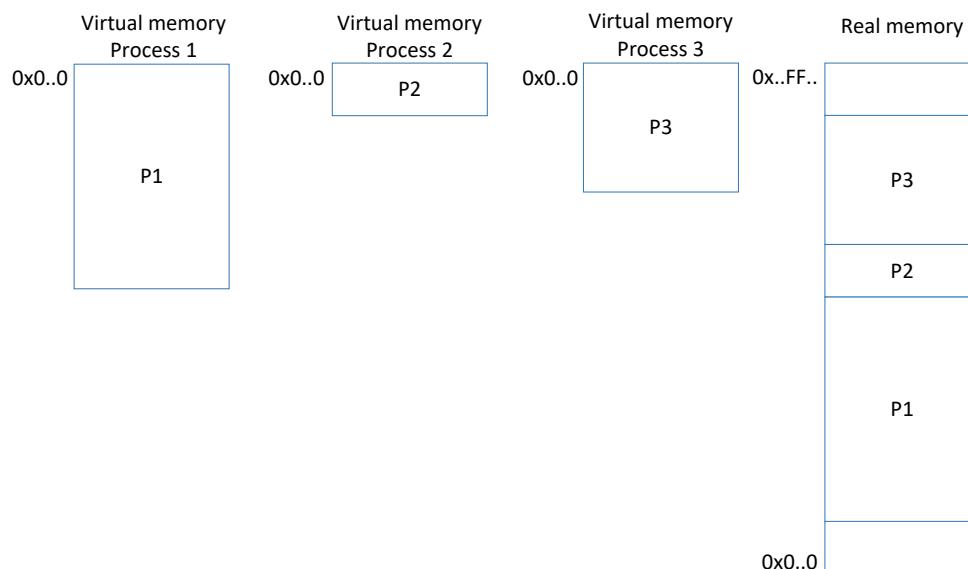
### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space

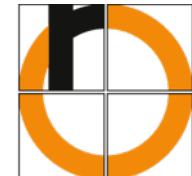


### Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

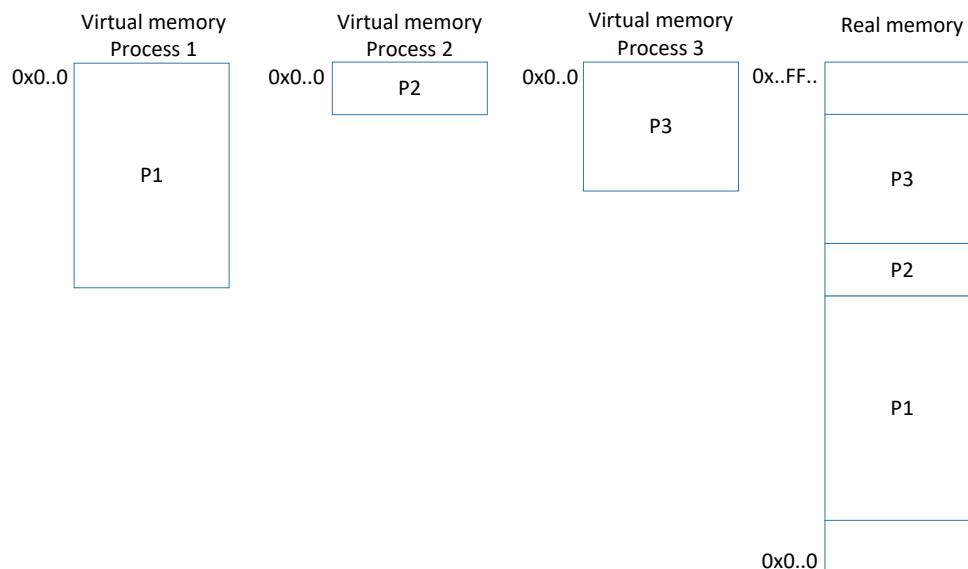
### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

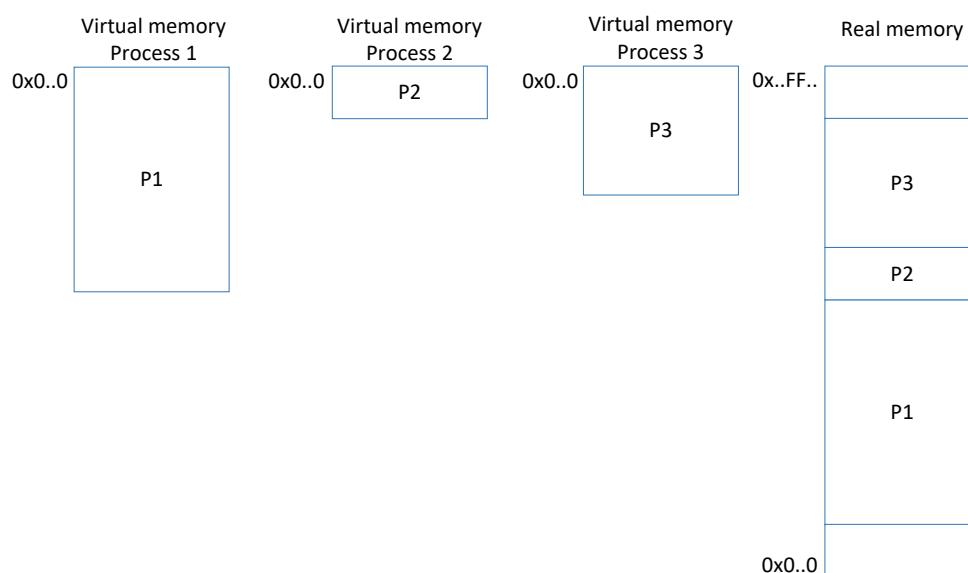
### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory



# Towards virtual addresses

## Idea 1: Separation of virtual and real address space



### Facts

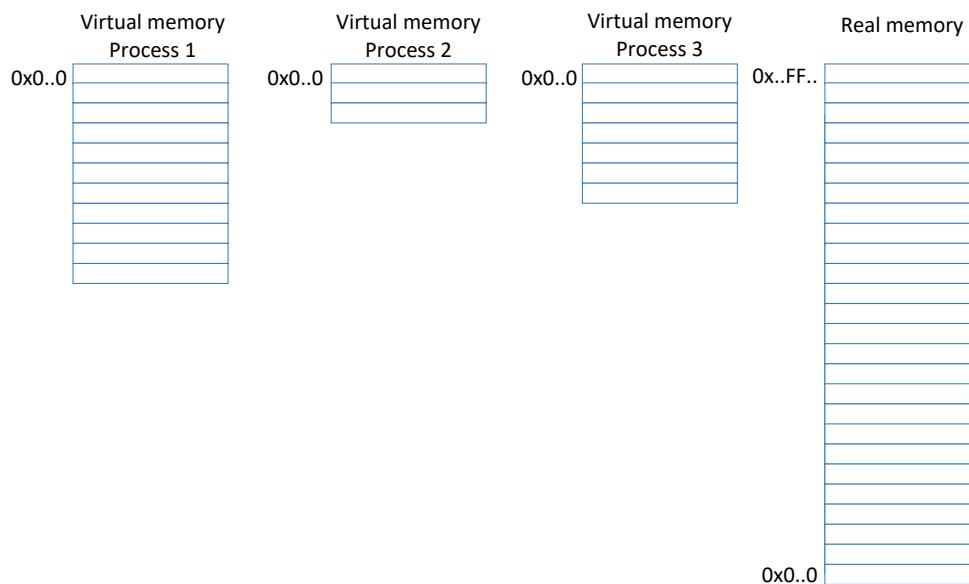
- Every process starts with address 0x0..0 until its upper limit
- The process needs to know its upper limit at compile time

### Problems

- Real memory fragmentation
- Maximum address space is limited by the real memory

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

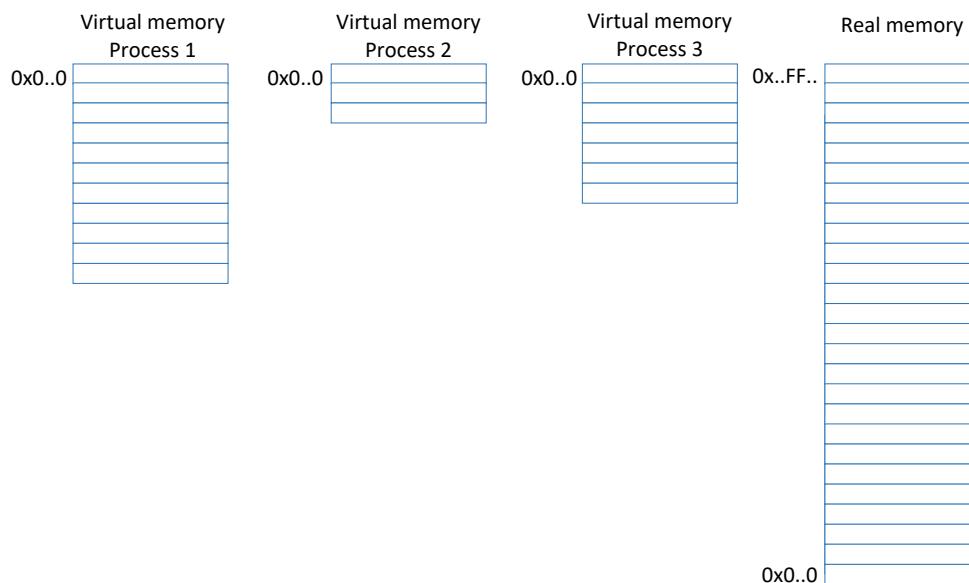
### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space



# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

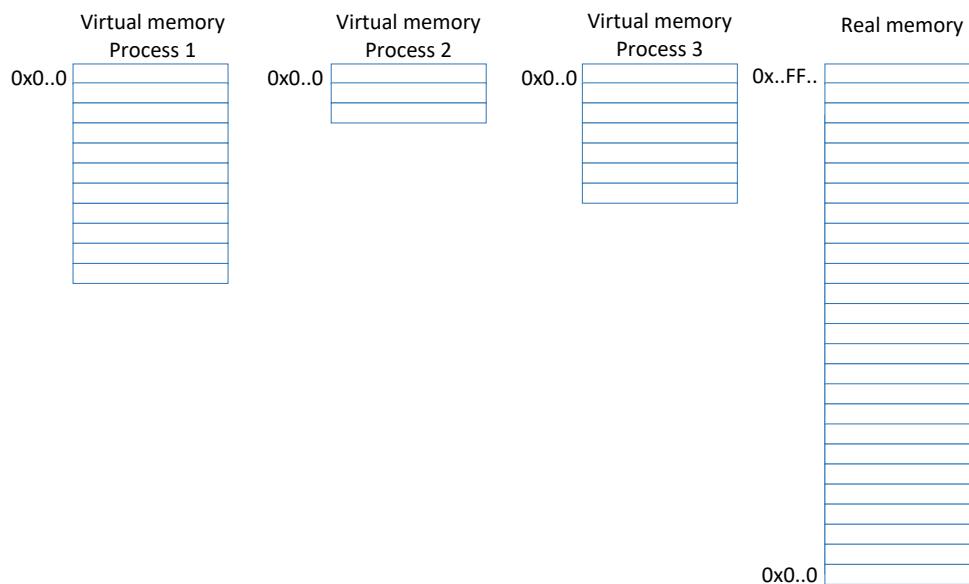
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

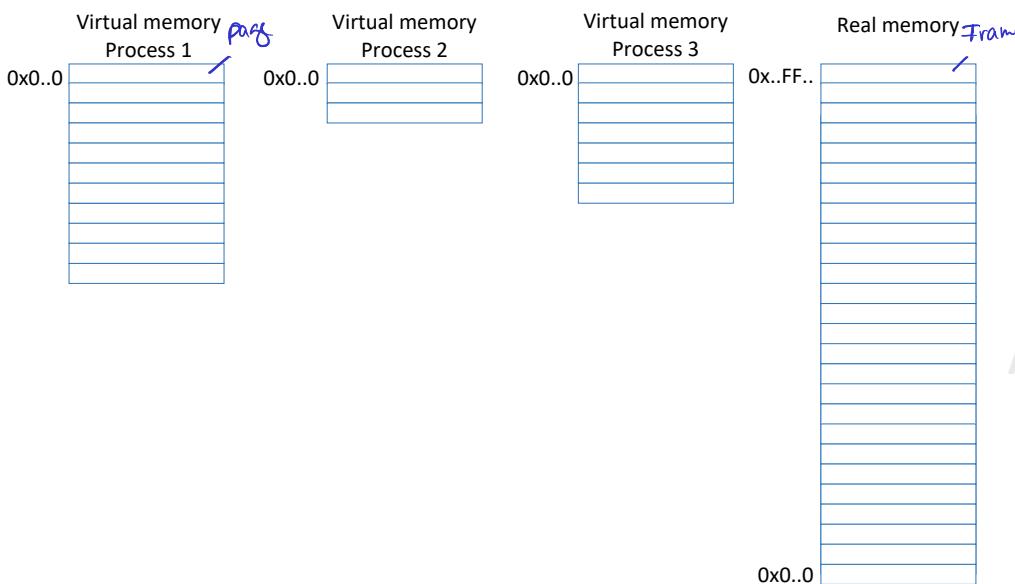
- The memory is divided into small pieces of **equal size** called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

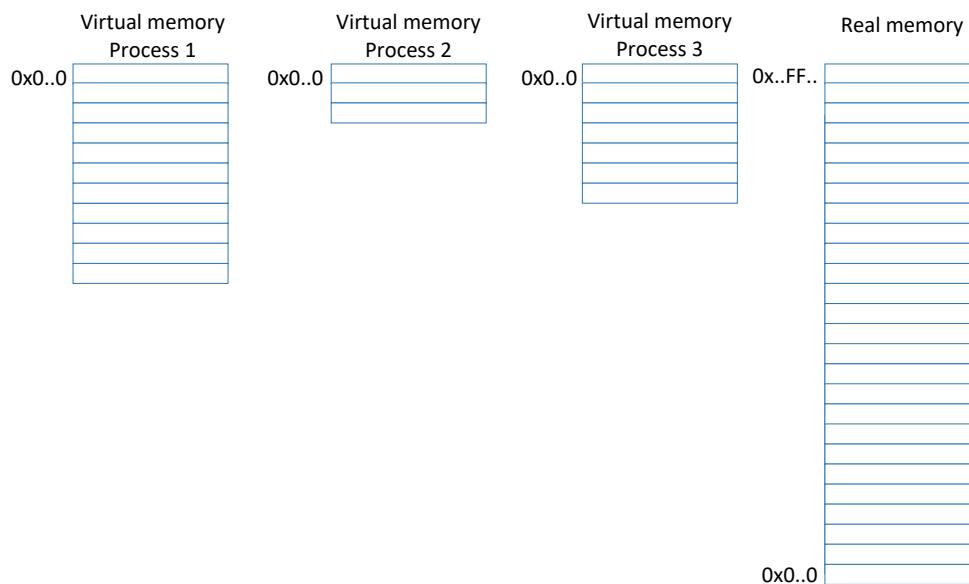
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

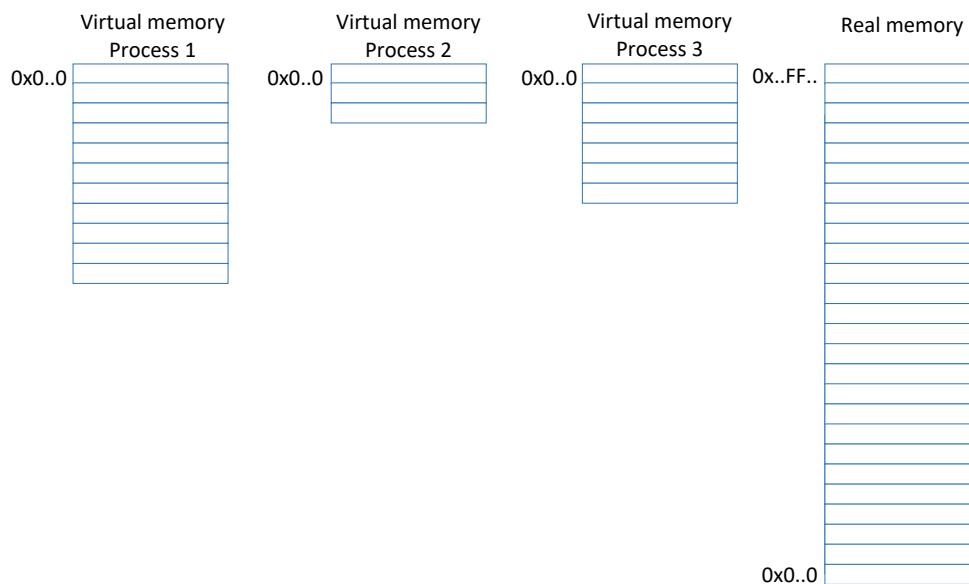
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

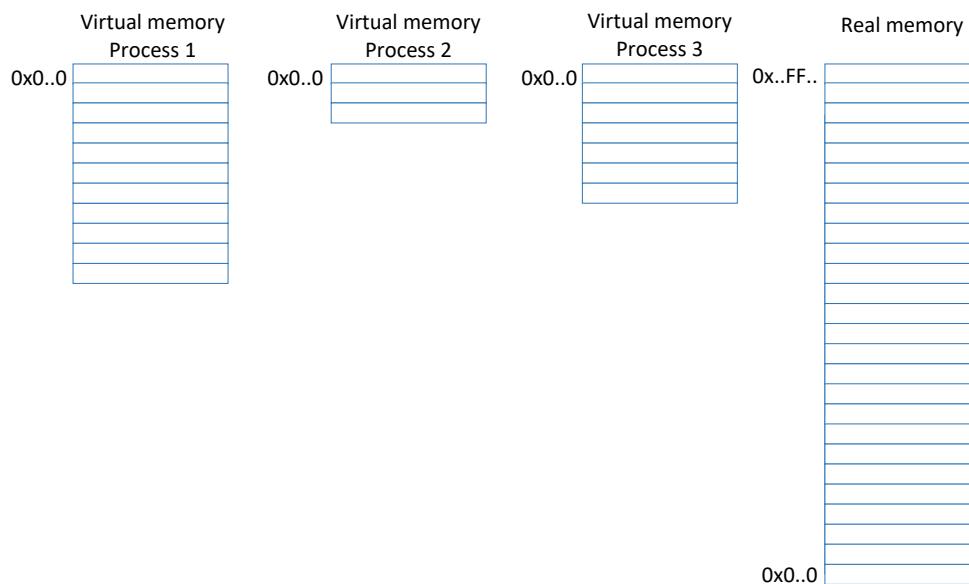
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

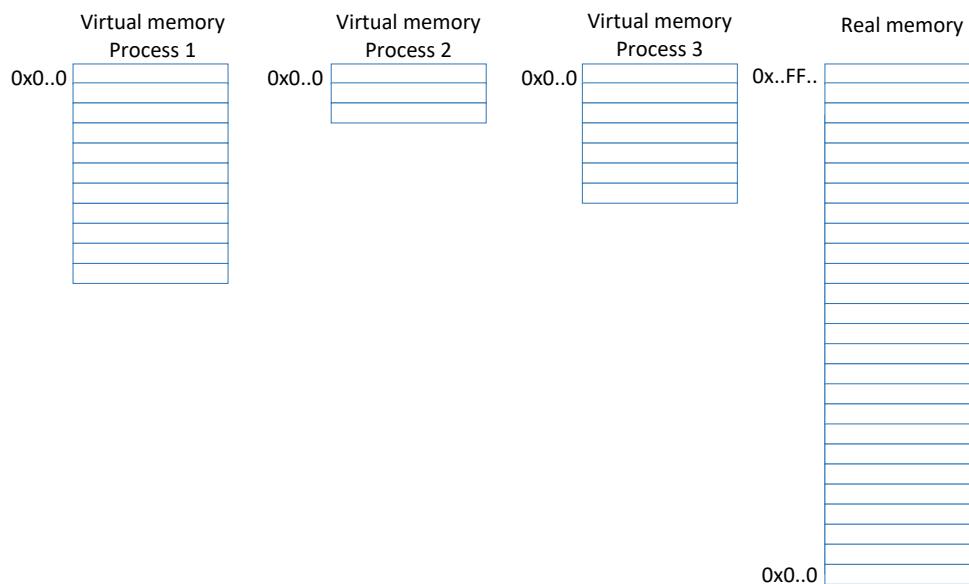
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



### Facts

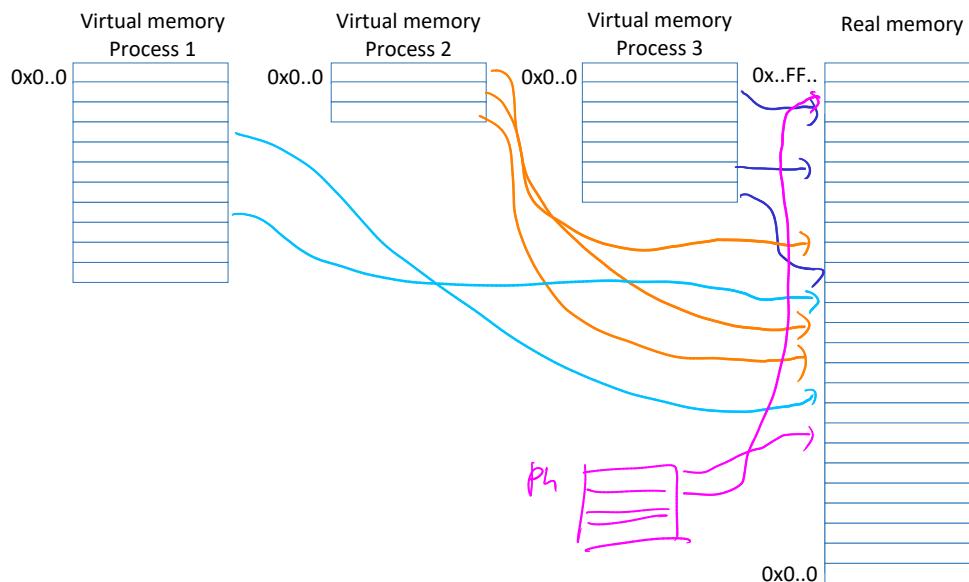
- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

### Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space

# Towards virtual addresses

## Idea 2: Introduce pages and frames



## Facts

- The memory is divided into small pieces of equal size called pages and frames
- Page: virtual address space
- Frame: real address space

## Advantages

- No (little) real memory fragmentation
- No upper memory limit required at compile time
- Virtual address space > real address space



# Goal

**Understand in detail virtual addresses and its translation to real addresses**

## Procedure

- Introduction of terms
- 1 level page table: with a small memory 64 KiB (virtual) and 32 KiB (real)
- 2 level page table: Intel x86/32 bit architecture
- 4 level page table: Intel x86/64 architecture (AMD/Intel)

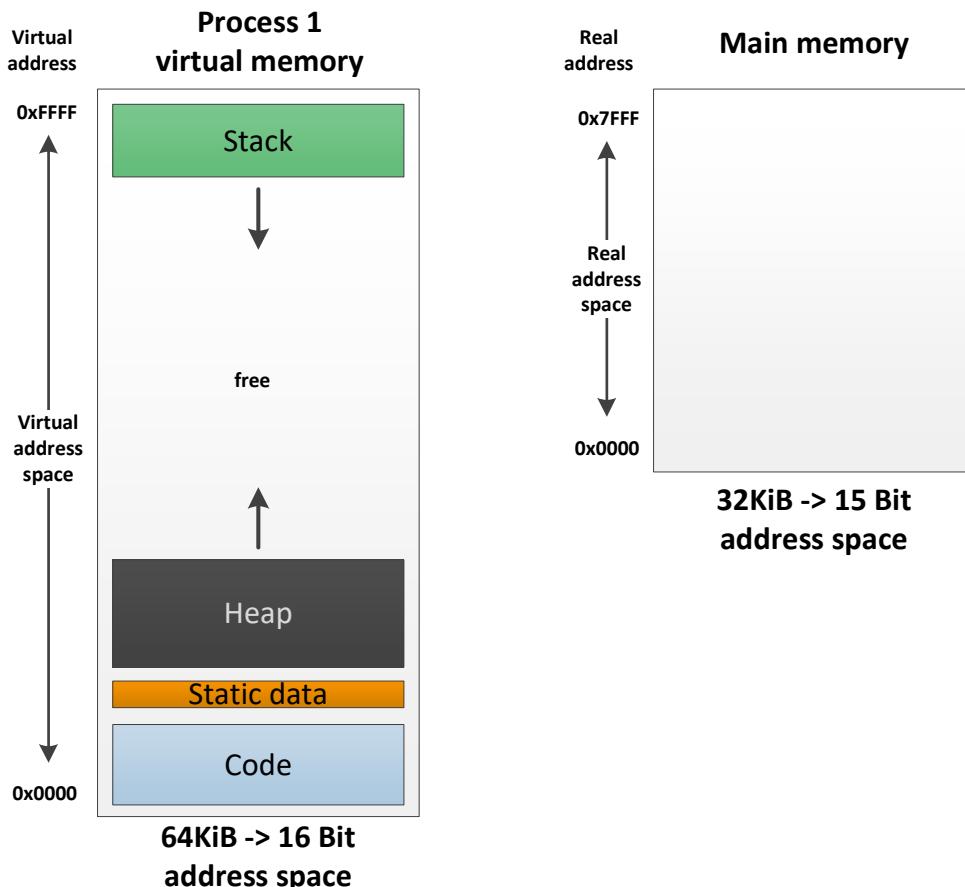
# Goal

**Understand in detail virtual addresses and its translation to real addresses**

## Procedure

- Introduction of terms
- 1 level page table: with a small memory 64 KiB (virtual) and 32 KiB (real)
- 2 level page table: Intel x86/32 bit architecture
- 4 level page table: Intel x86/64 architecture (AMD/Intel)

# Virtual address space



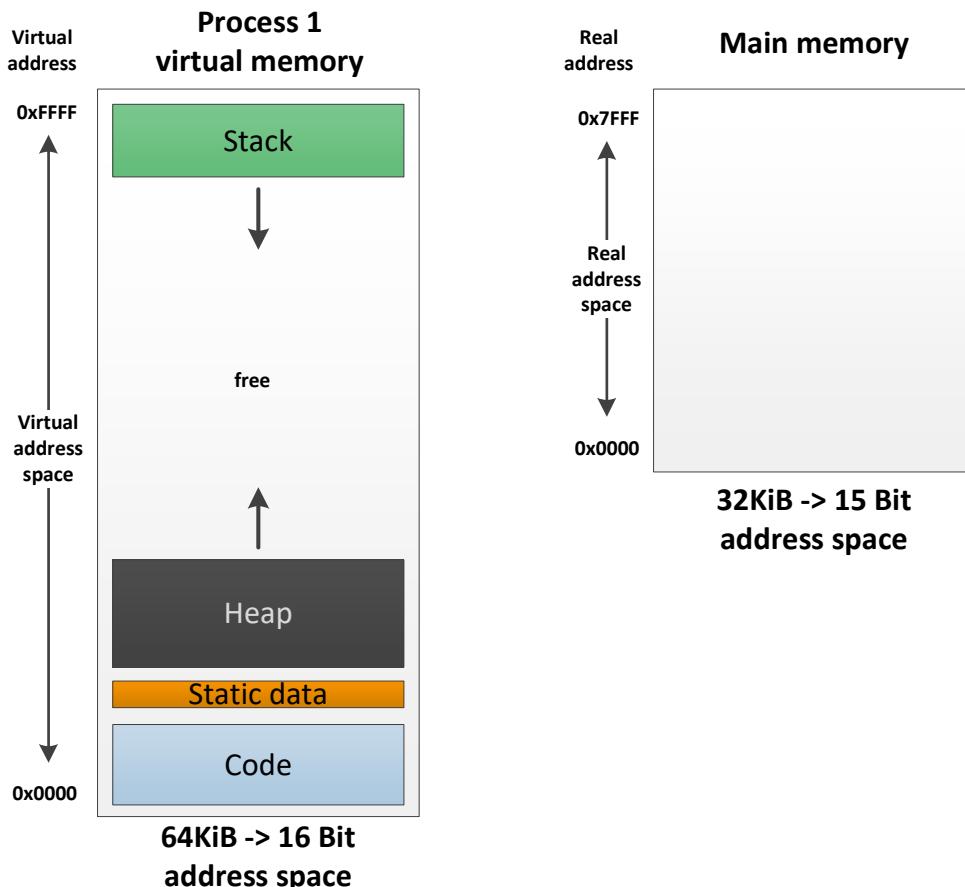
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



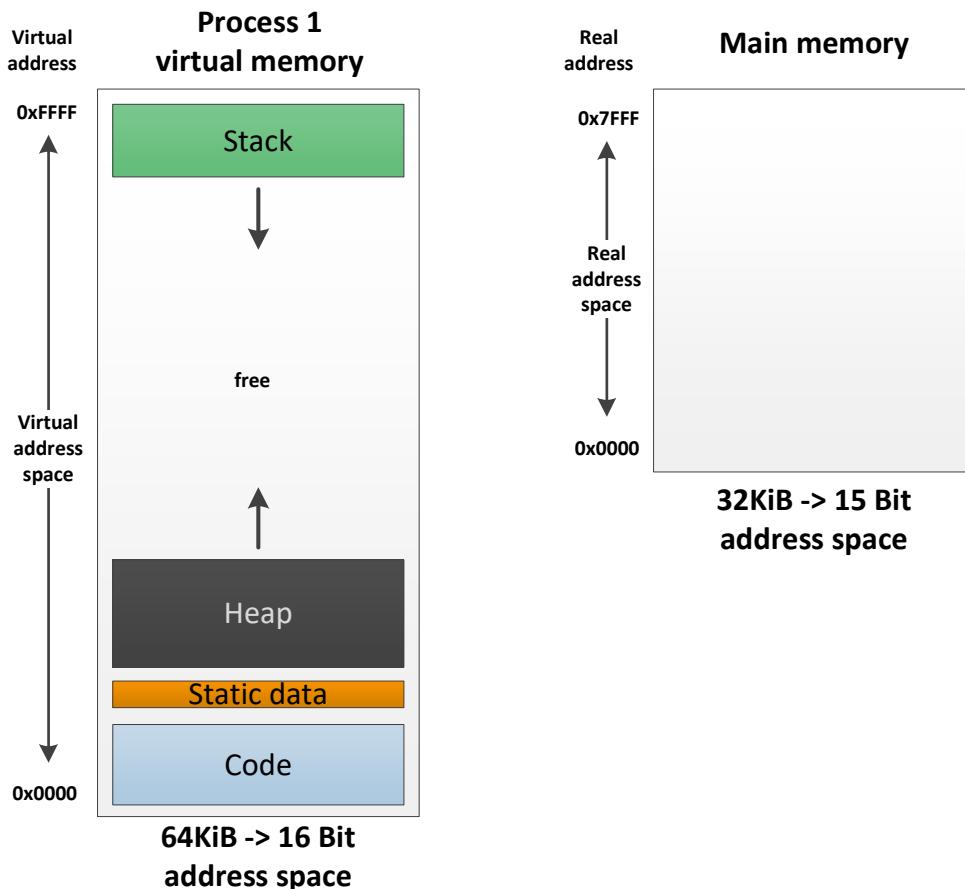
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



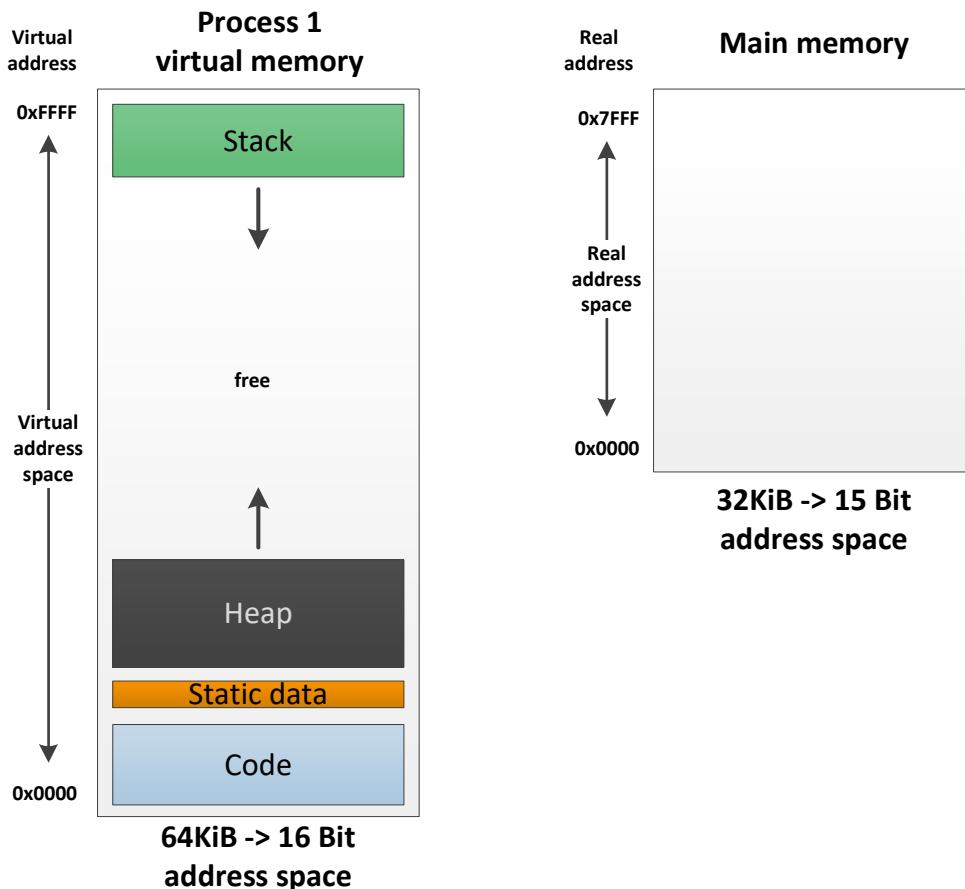
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



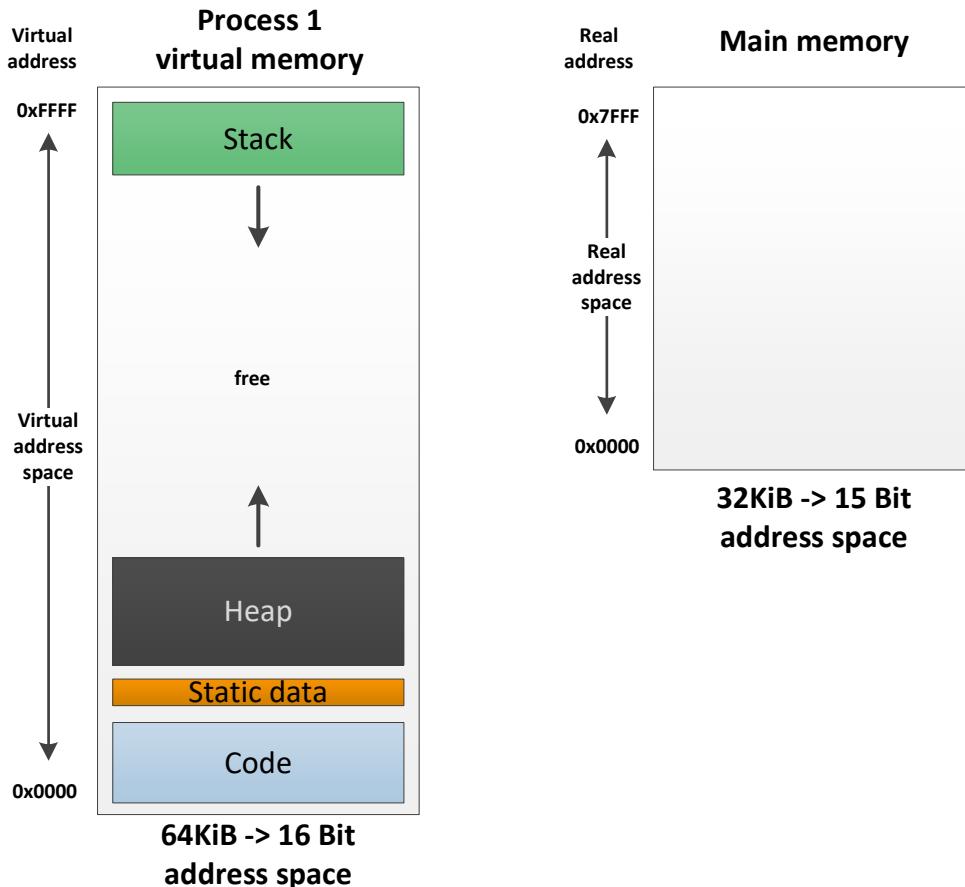
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



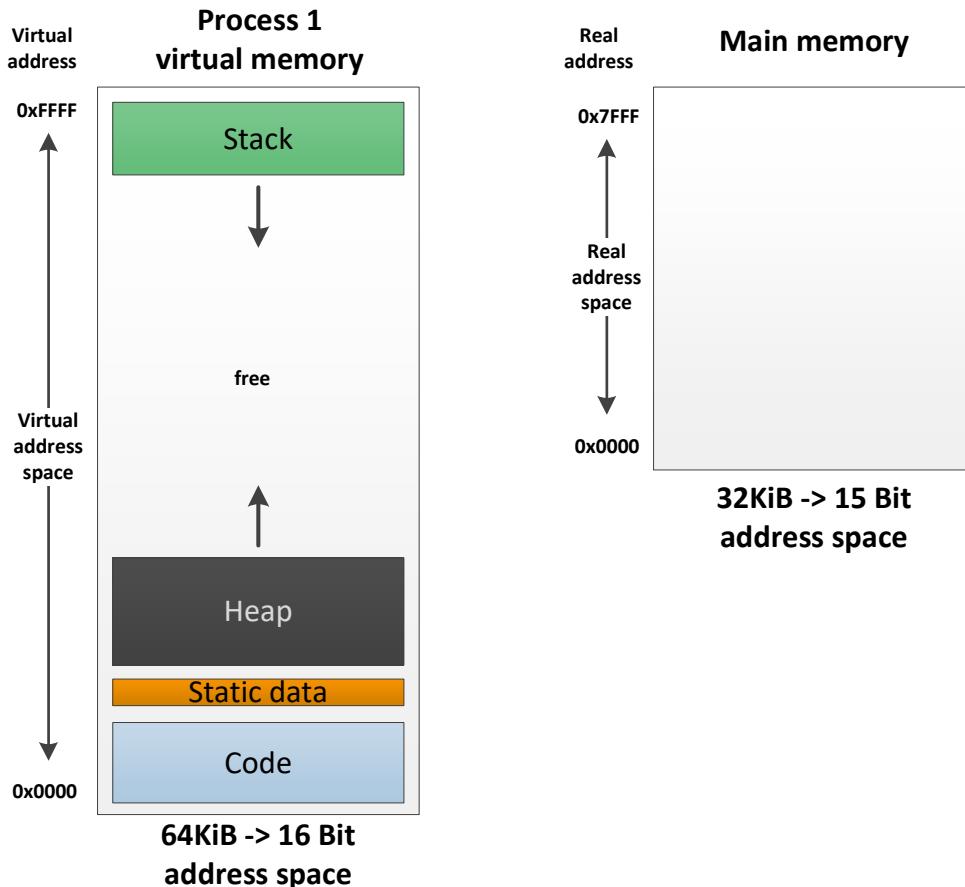
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



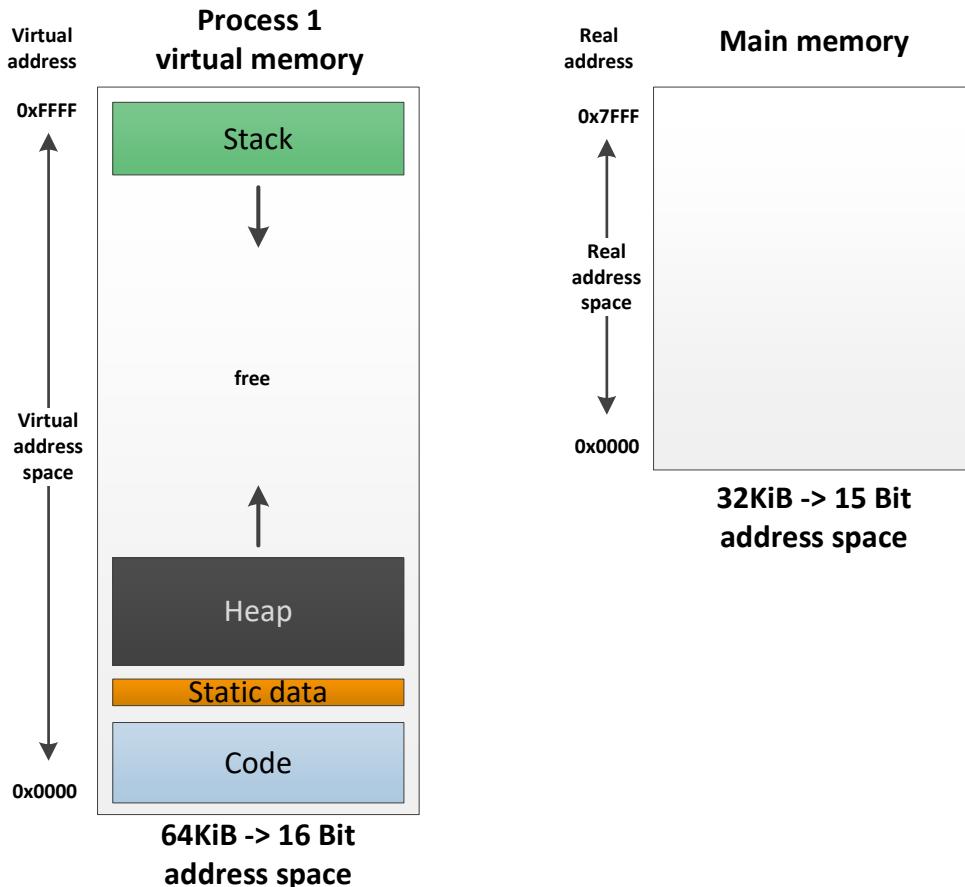
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



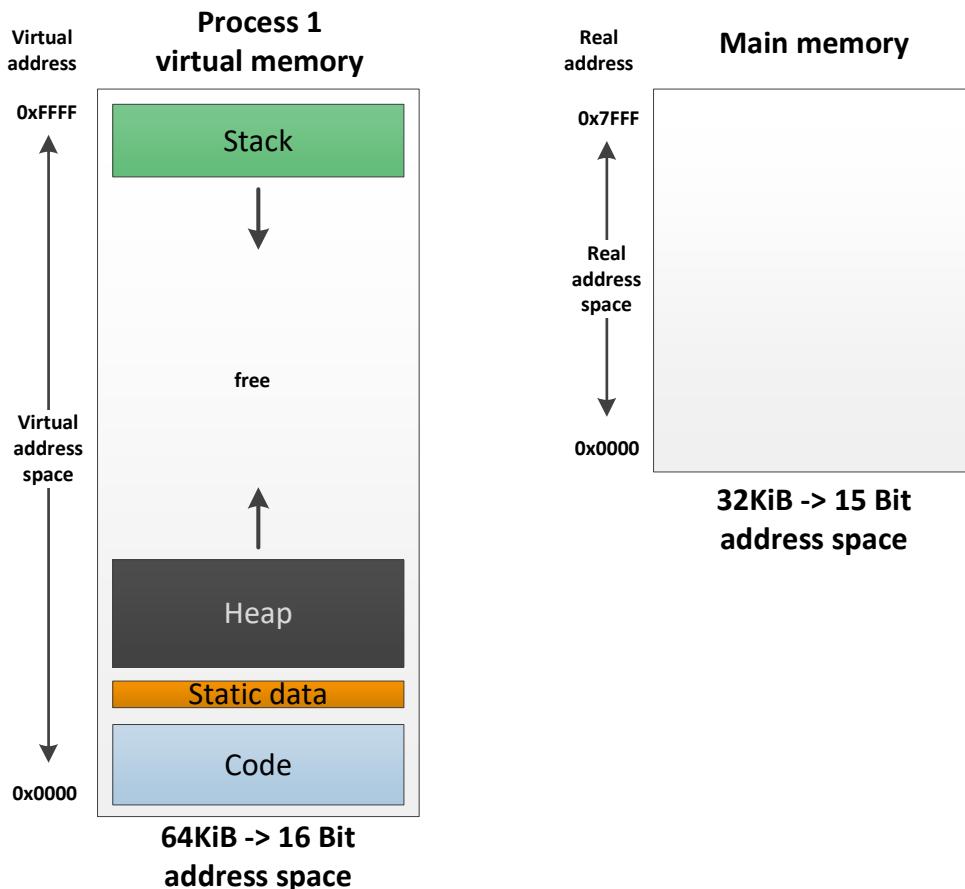
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



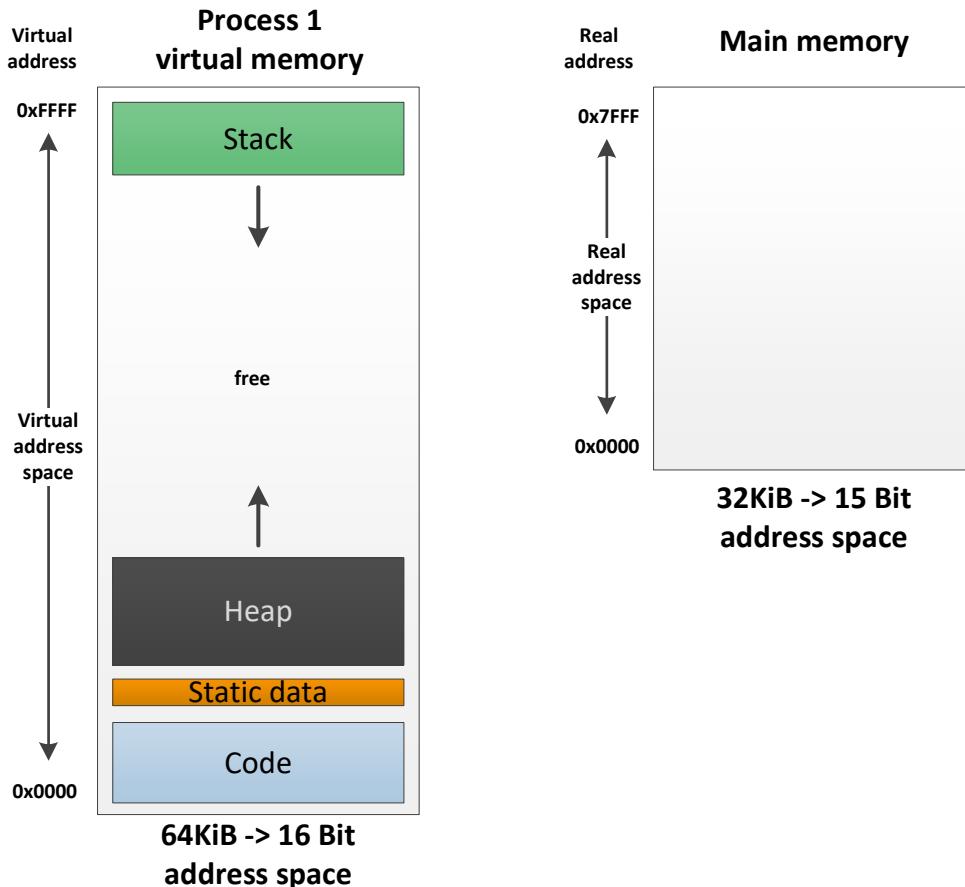
## Virtual address space

- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

- All physically available addresses
- Size depends on available memory

# Virtual address space



## Virtual address space

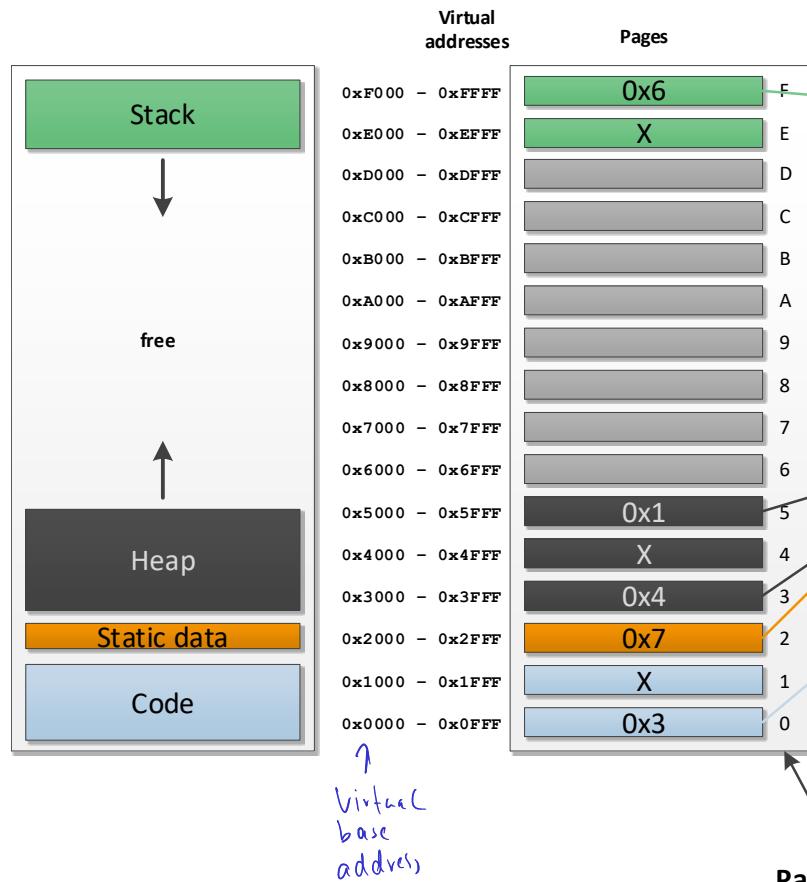
- Linear address space
- Starting from 0x0..0 to 0x..FF..
- All virtual addresses => virtual address space
- Every process has its own virtual address space

## Real address space

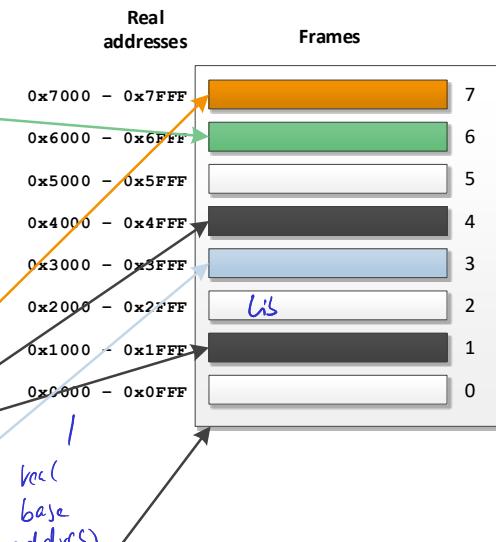
- All physically available addresses
- Size depends on available memory

# Pages and frames

Virtual memory



Main memory



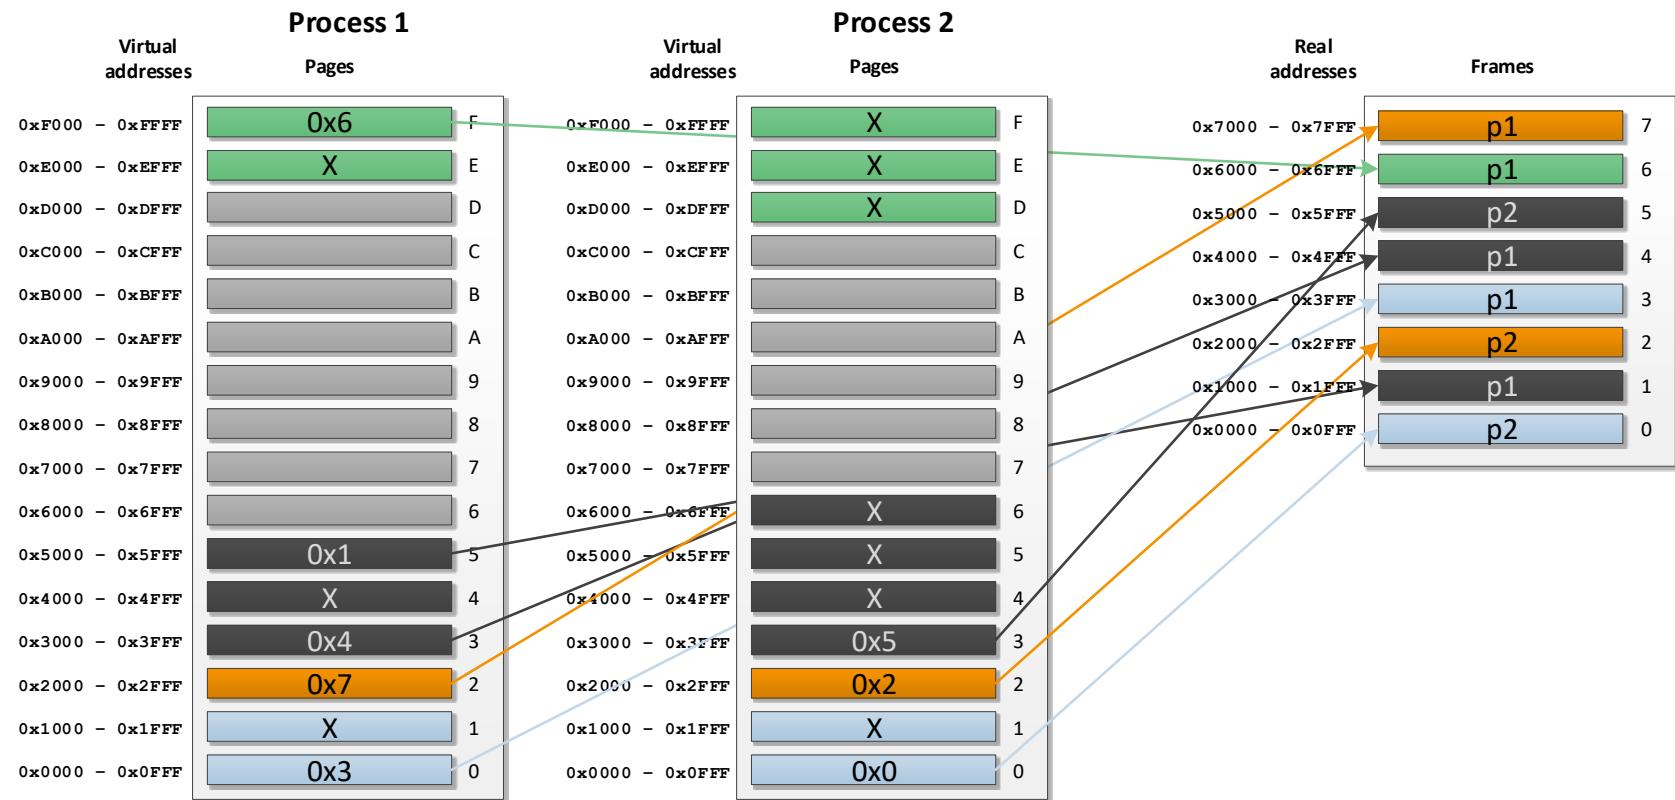
Pages and frames have same size (here 4KiB)

4096 Bytes

# Multiprogramming

Virtual memory

Main memory





# Questions?

All right?  $\Rightarrow$



Question?  $\Rightarrow$



and use **chat**

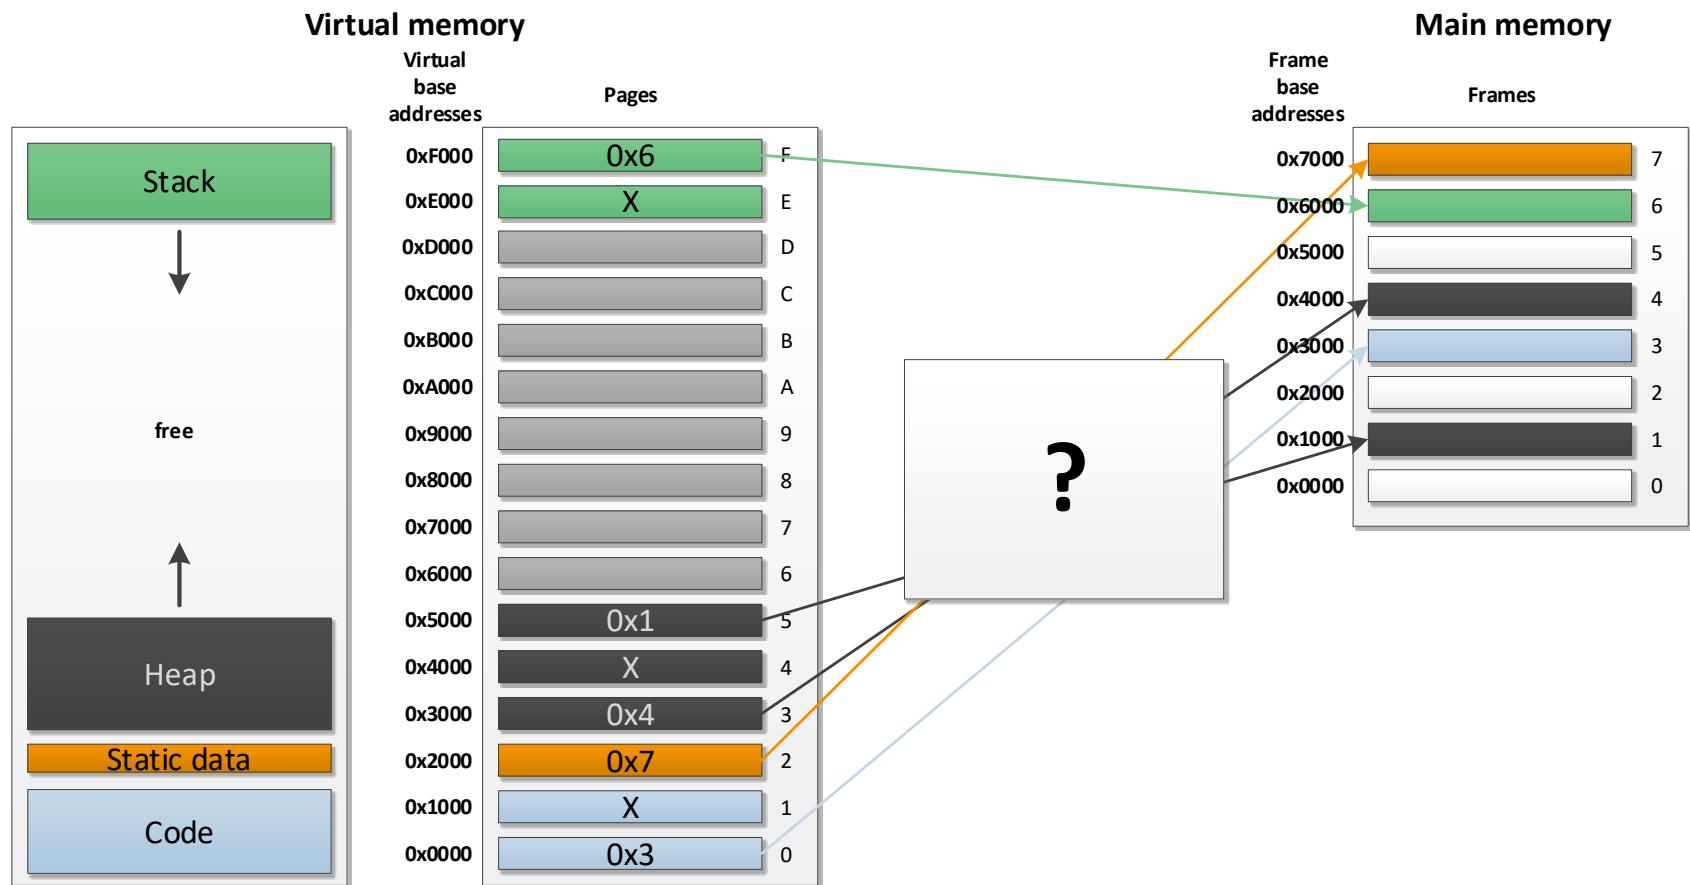
or

*speak after I  
ask you to*

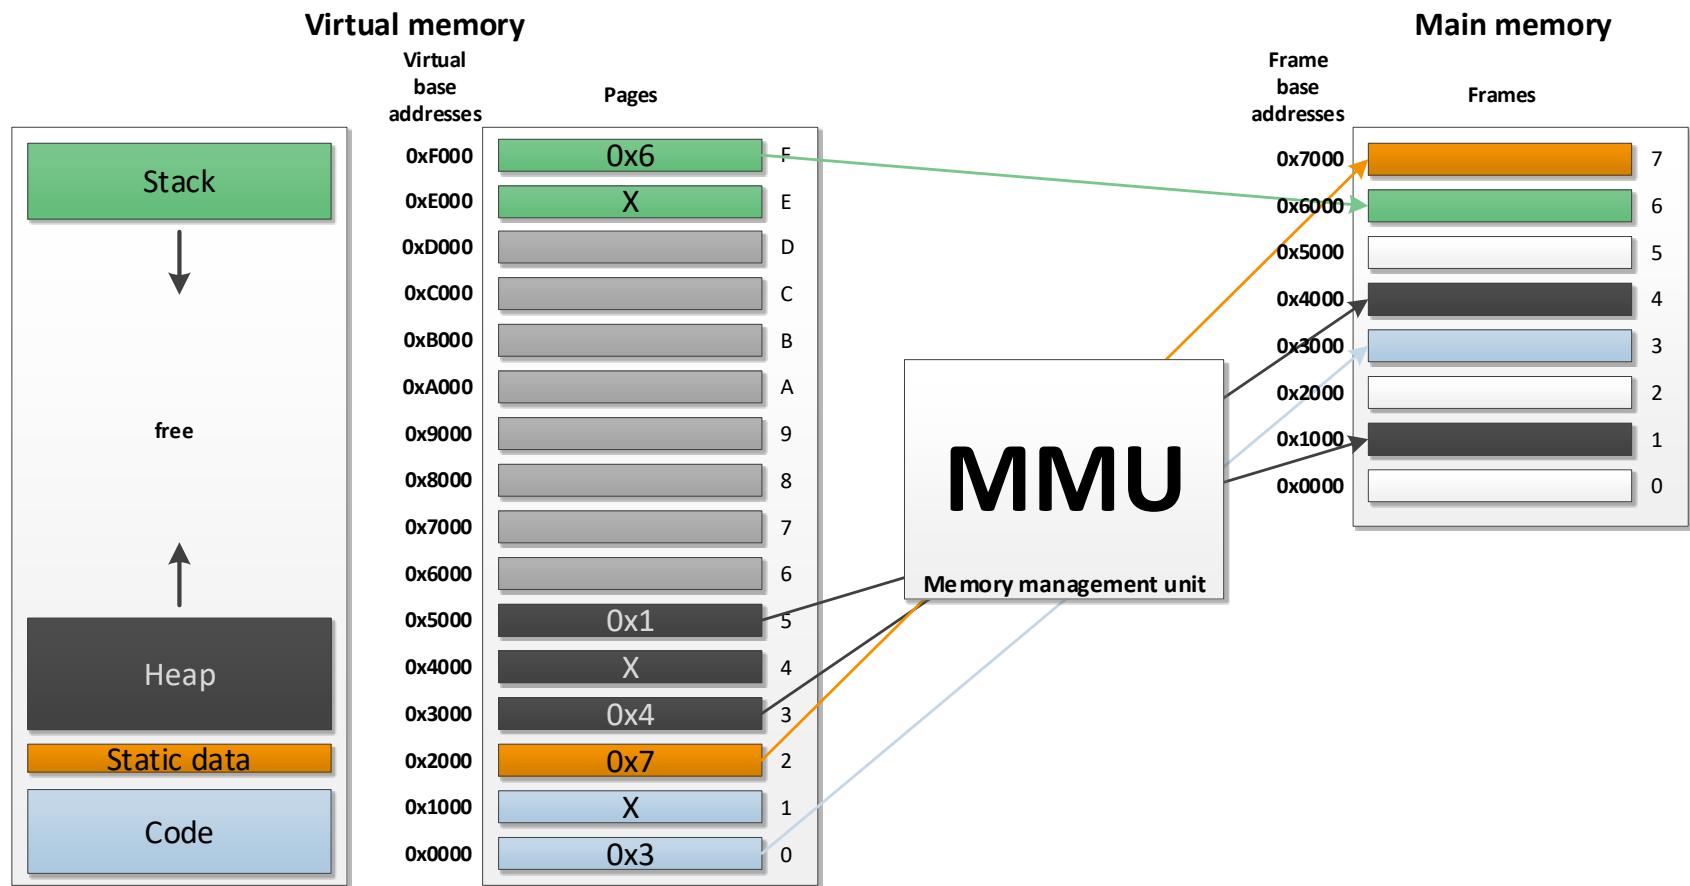


Let's start with a simple  
**1 level page table.**

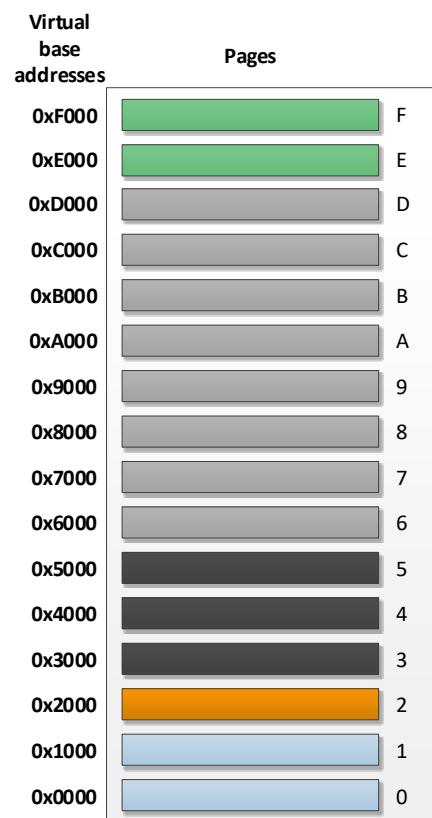
# Virtual address translation



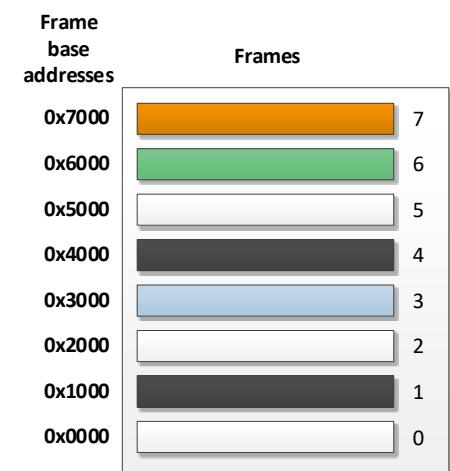
# Virtual address translation



# Page table

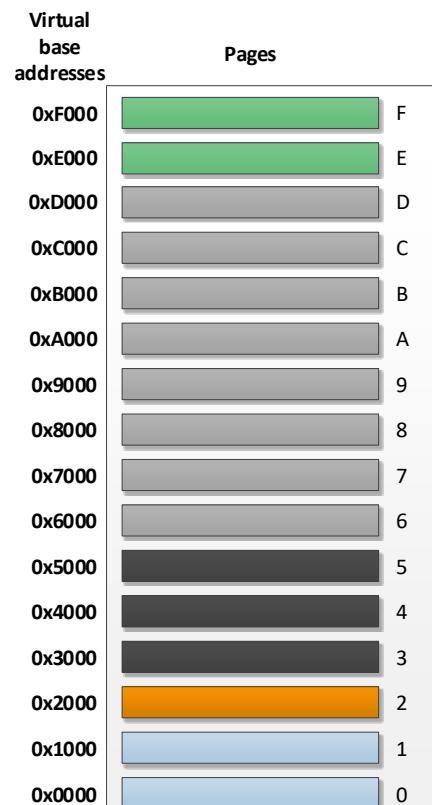
**Virtual memory**


Each process has its own virtual memory

**Main memory**


# Page table

Virtual memory

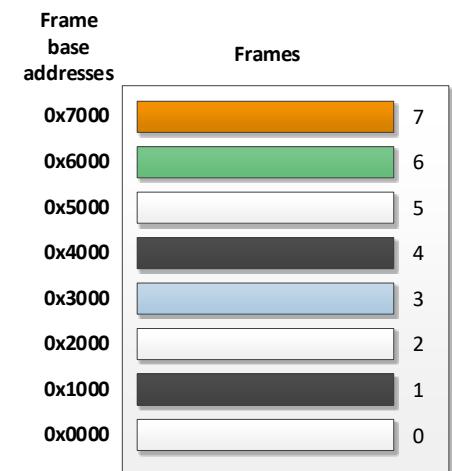


Each process has its own virtual memory

Each process has its own page table

CA 9 – MMU

Main memory



# Page table

Virtual memory

Virtual base addresses	Pages
0xF000	F
0xE000	E
0xD000	D
0xC000	C
0xB000	B
0xA000	A
0x9000	9
0x8000	8
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

Each process has its own virtual memory

Virtual addresses



Real addresses

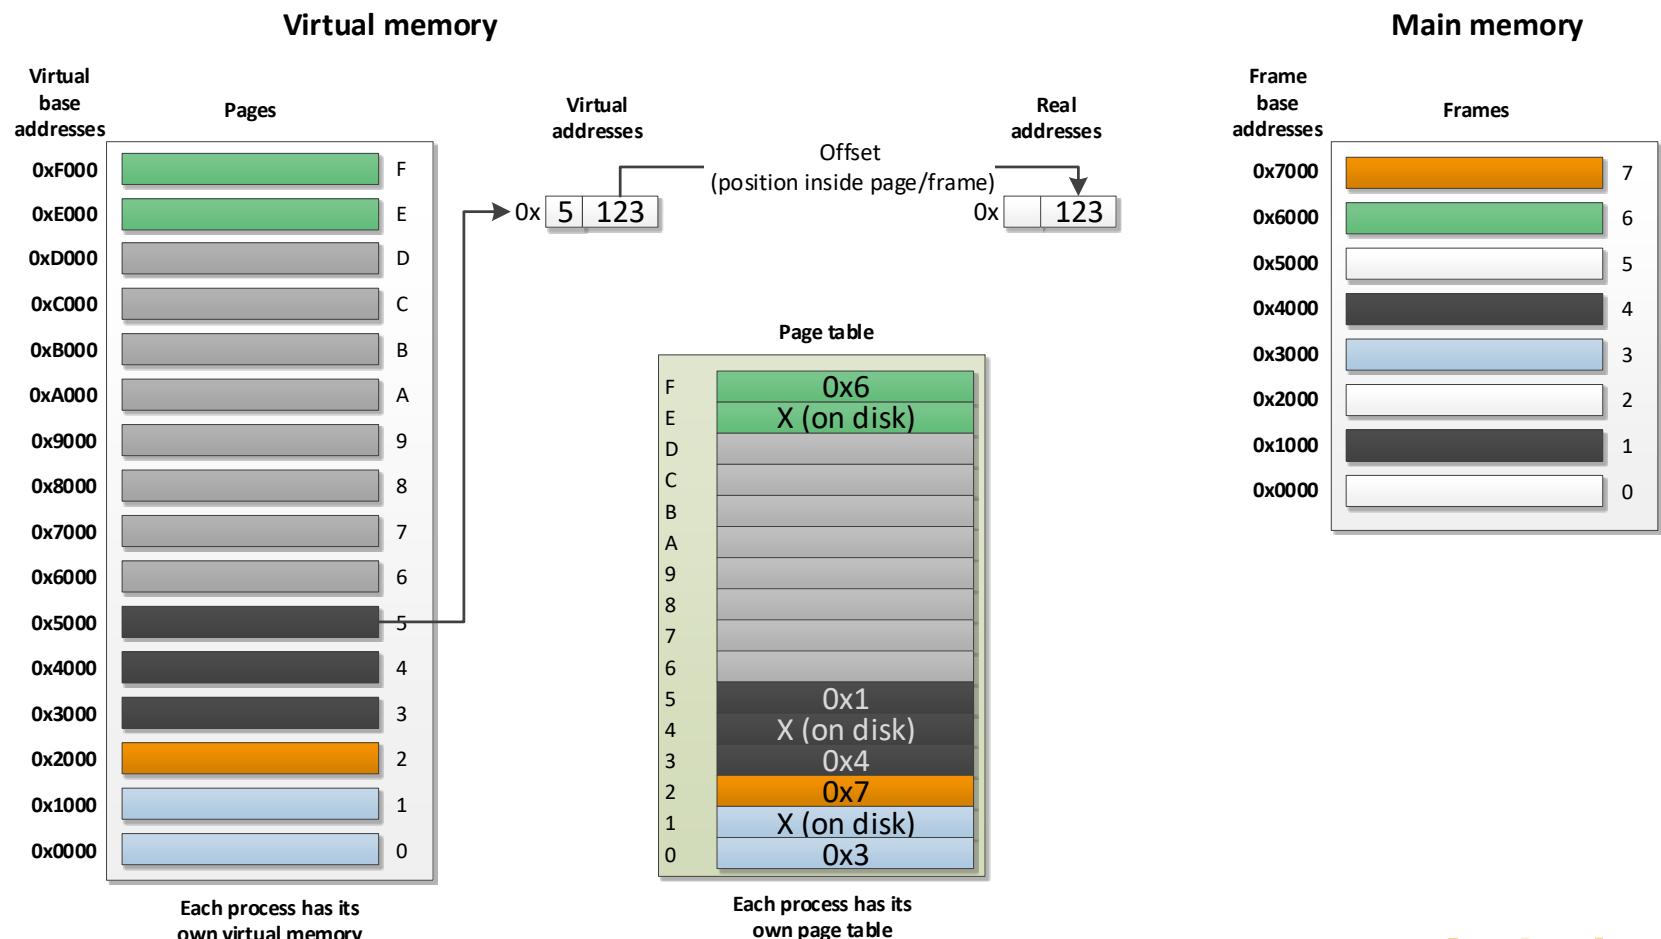
Page table	
F	0x6
E	X (on disk)
D	
C	
B	
A	
9	
8	
7	
6	
5	0x1
4	X (on disk)
3	0x4
2	0x7
1	X (on disk)
0	0x3

Each process has its own page table

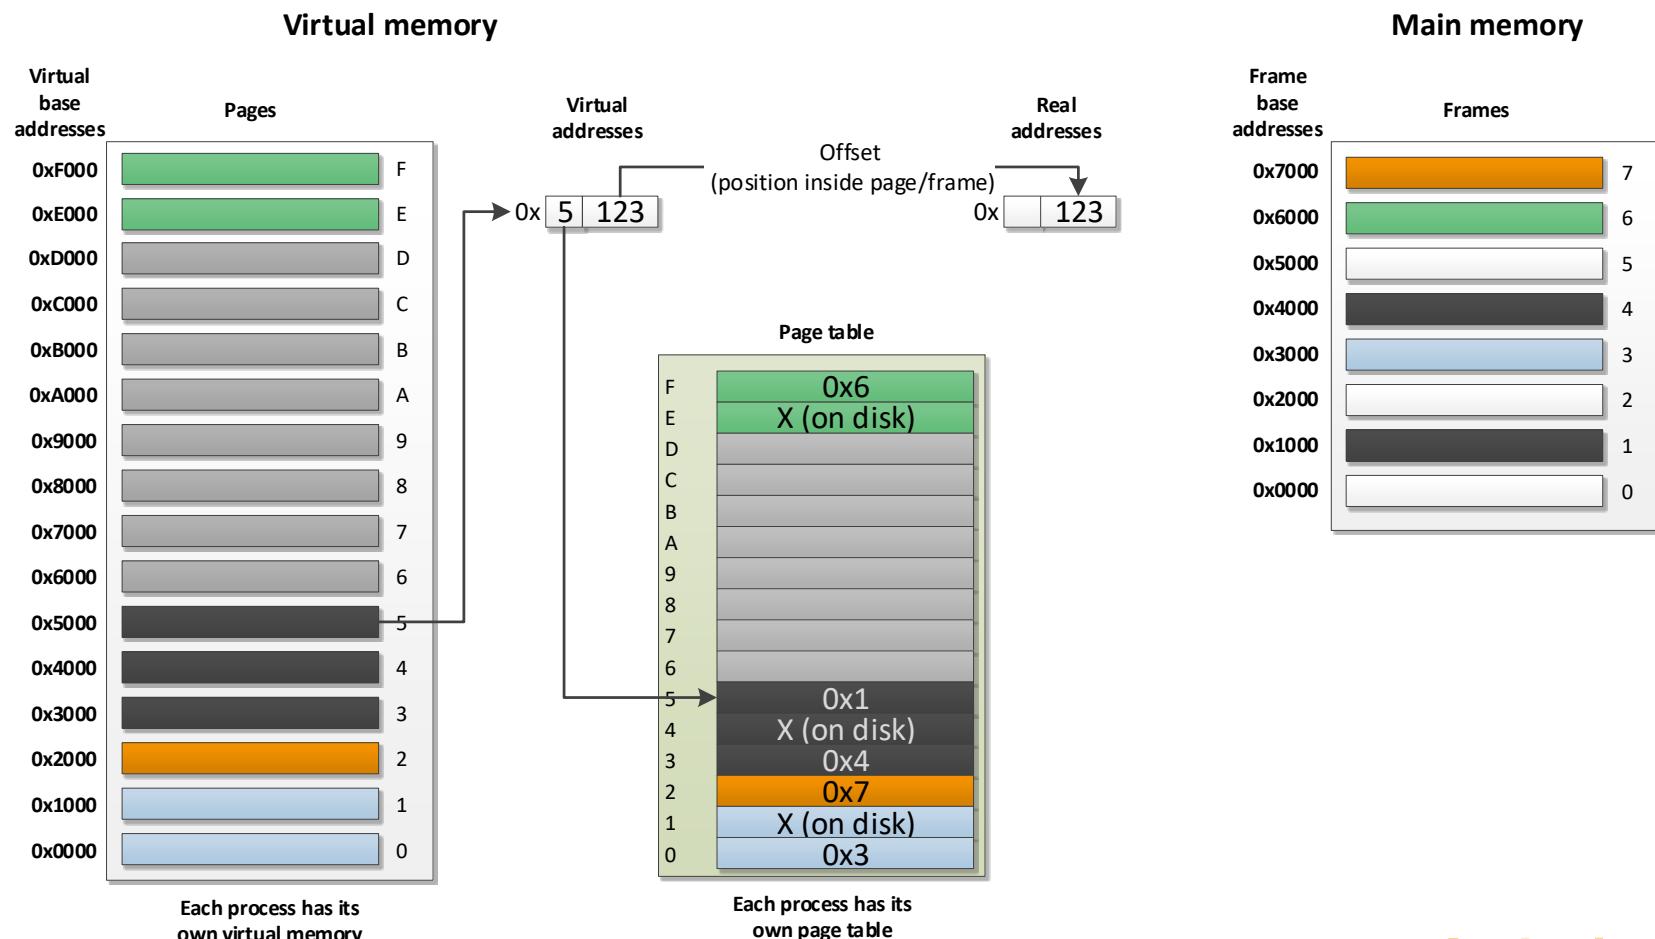
Main memory

Frame base addresses	Frames
0x7000	7
0x6000	6
0x5000	5
0x4000	4
0x3000	3
0x2000	2
0x1000	1
0x0000	0

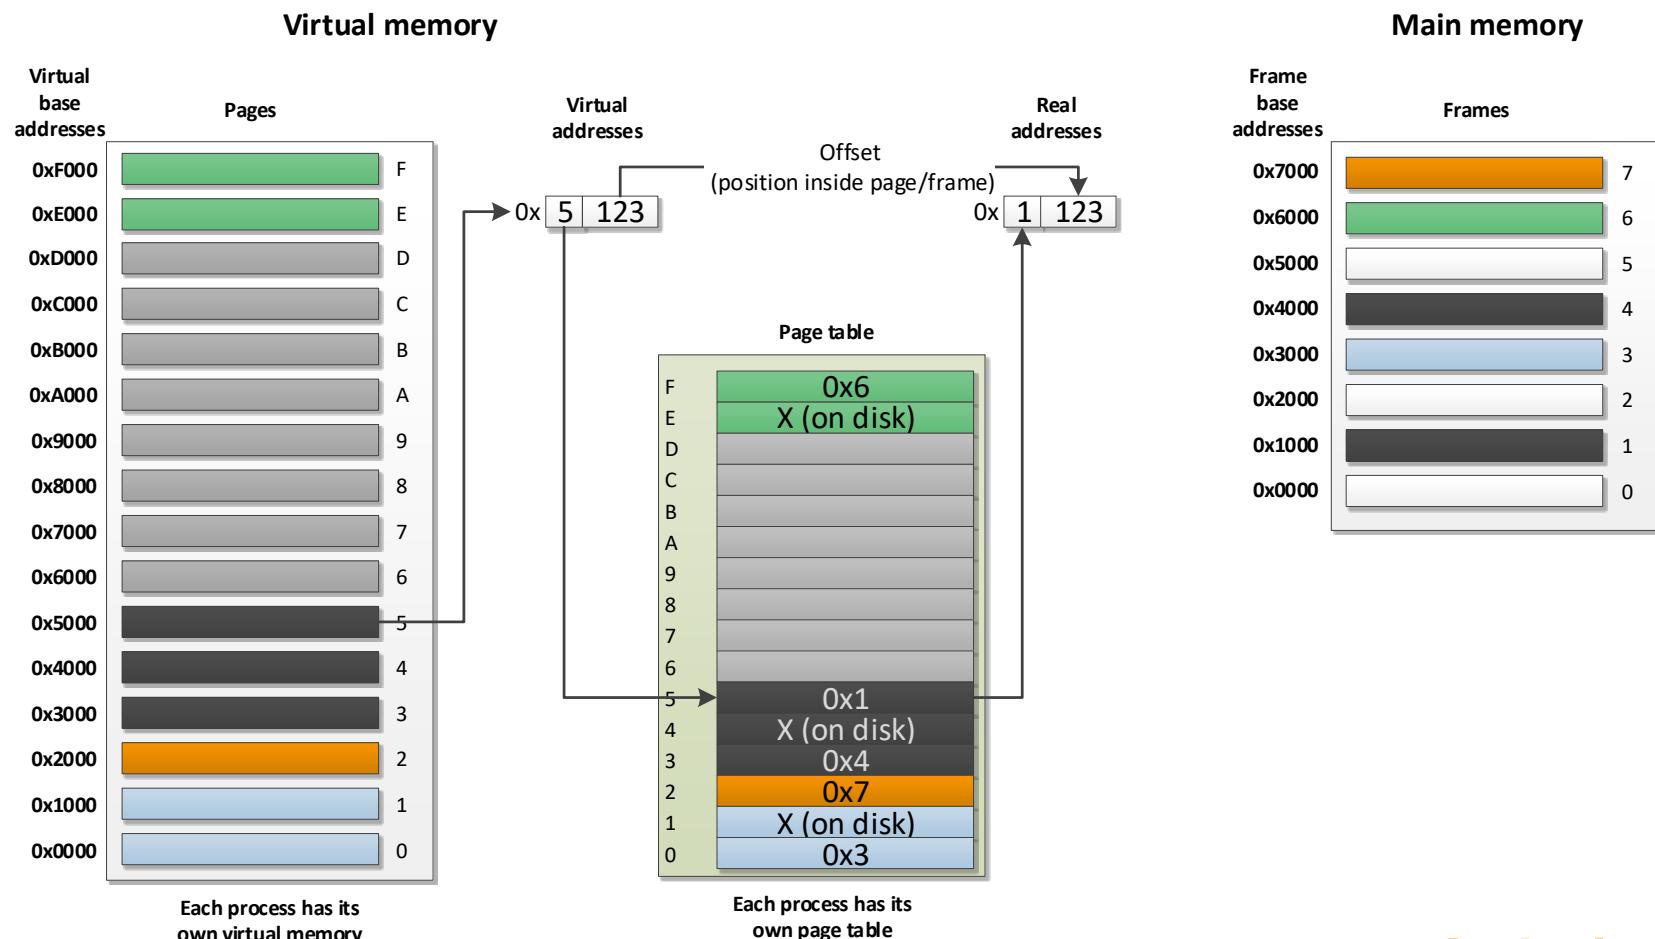
# Page table



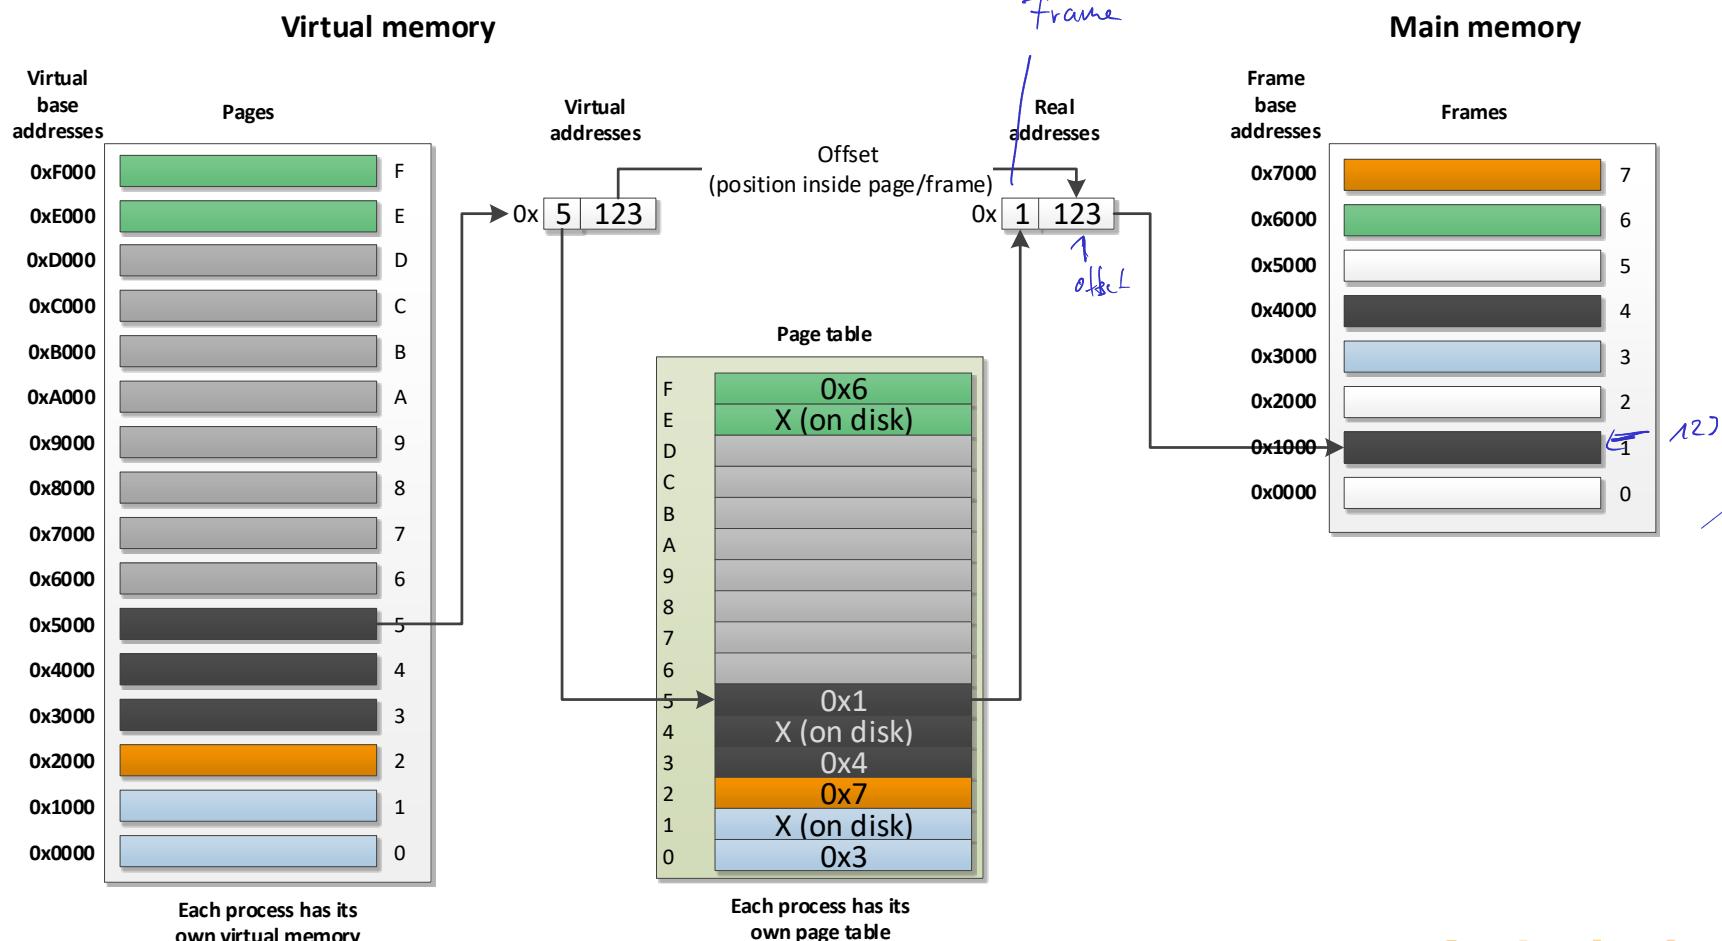
# Page table



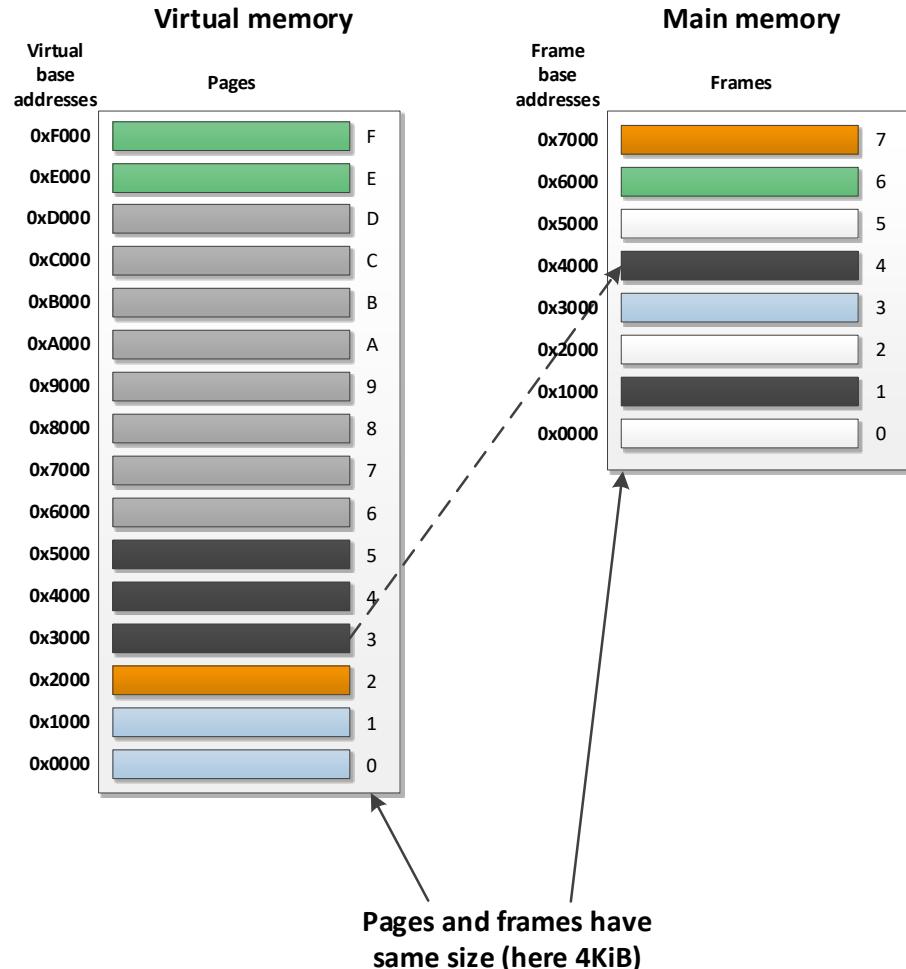
# Page table



# Page table



# Virtual address translation: example



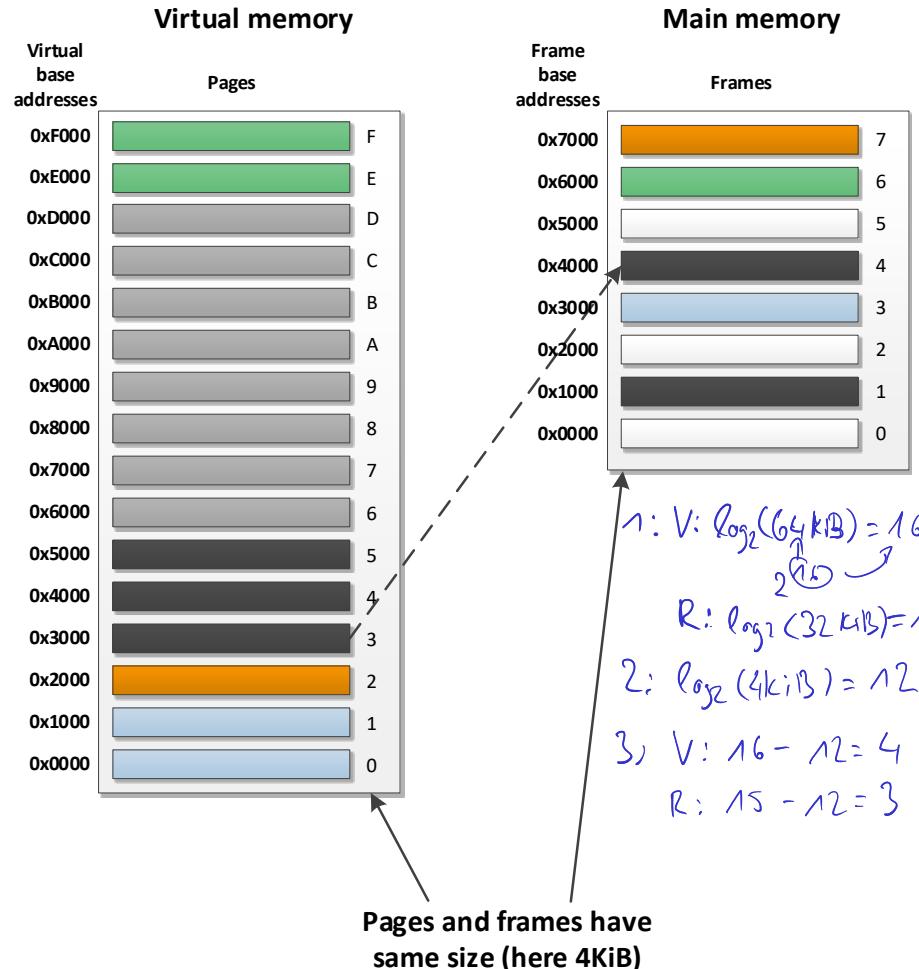
Given:

- Virtual address space: 64 KiB
- Real address space: 32 KiB
- Virtual address: 0x3D05

Questions:

- 1 Address width (virtual and real)?
- 2 Offset: number of bits?
- 3 Page/frame numbers: number of bits?
- 4 Number of pages and frames?
- 5 Real address?

# Virtual address translation: example


**Given:**

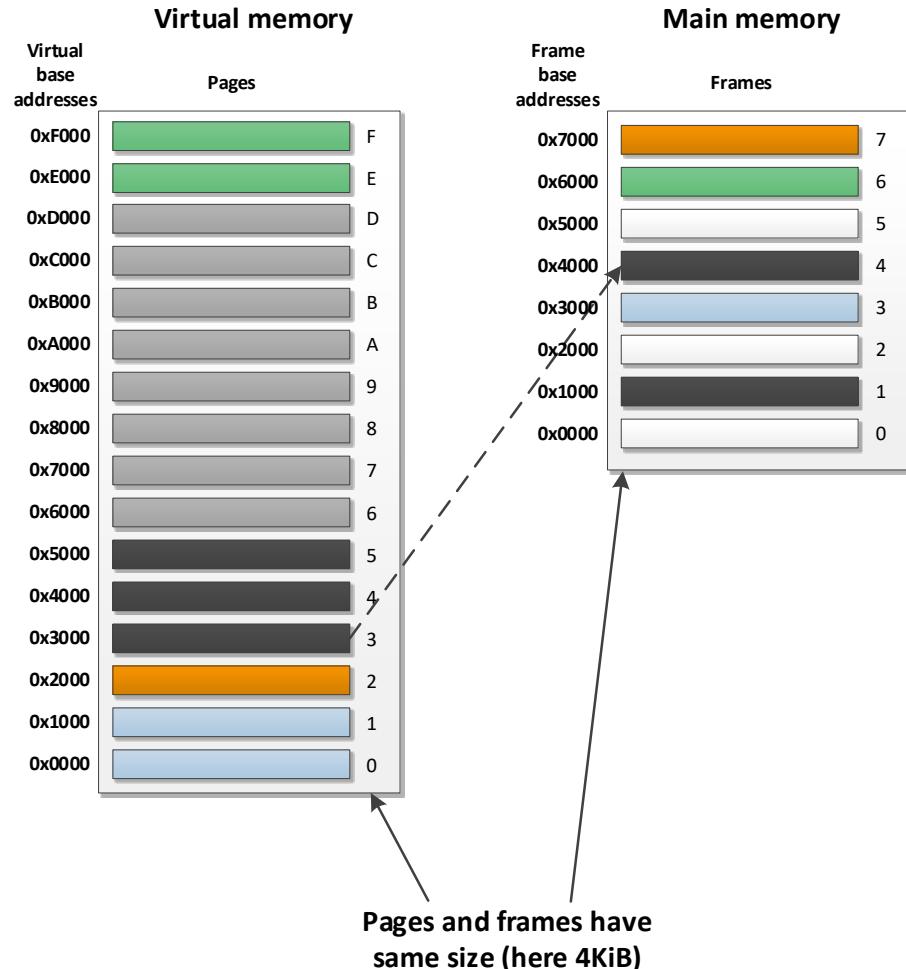
- Virtual address space: 64 KiB
- Real address space: 32 KiB
- Virtual address: 0x3D05

**Questions:**

- 1 Address width (virtual and real)?
- 2 Offset: number of bits?
- 3 Page/frame numbers: number of bits?
- 4 Number of pages and frames?
- 5 Real address?

$$\begin{aligned}
 1: & V: \log_2(64\text{KiB}) = 16 \\
 & R: \log_2(32\text{KiB}) = 15 \\
 2: & \log_2(4\text{KiB}) = 12 \\
 3: & V: 16 - 12 = 4 \\
 & R: 15 - 12 = 3 \\
 4: & V: 2^4 = 16 \text{ pages} \\
 & R: 2^3 = 8 \text{ frames} \\
 5: & V: 0x3\text{ }D05 \\
 & \downarrow \quad \downarrow \\
 & R: 0x41005
 \end{aligned}$$

# Virtual address translation: example



**Given:**

- Virtual address space: 64 KiB
- Real address space: 32 KiB
- Virtual address: 0x3D05

**Questions:**

- 1 Address width (virtual and real)?
- 2 Offset: number of bits?
- 3 Page/frame numbers: number of bits?
- 4 Number of pages and frames?
- 5 Real address?

**Results:**

- 1 virtual: 16 bit/real: 15 bit
- 2 offset: 12 bit
- 3 pages: 4 bit/frames: 3 bit
- 4 pages:  $2^4 = 16$ /frames:  $2^3 = 8$
- 5 real address: 0x4D05



# 1 level page table

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

# 1 level page table

$$\begin{aligned} \text{pages: } & 32 - 12 = 20 \text{ Bit} \\ & \text{offset} \\ & 2^{20} = 1 \text{ MiB} \end{aligned}$$

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

# 1 level page table

32 bit  $\geq$  4 byte

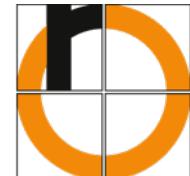
1 Mi<sup>2</sup> • 4 byte = 4MiB

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - $\Rightarrow$  less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk



# 1 level page table

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

# 1 level page table

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

# 1 level page table

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk

# 1 level page table

## Any problems with a 1 level page table?

- How many number of page table entries are required on an 32 bit system?
- How much space would the page table use?

## Possible solutions

- Increase page/frame size
  - i.e. 1 MiB, 2 MiB, ...
  - => less pages, less entries in page table, smaller page table
- Multilevel page table
  - 32 bit: 2 level page table with 4 KiB page size
  - 64 bit: 4 level page table with 4 KiB page size
- Swap page tables to disk



# Questions?

All right?  $\Rightarrow$



Question?  $\Rightarrow$



and use **chat**

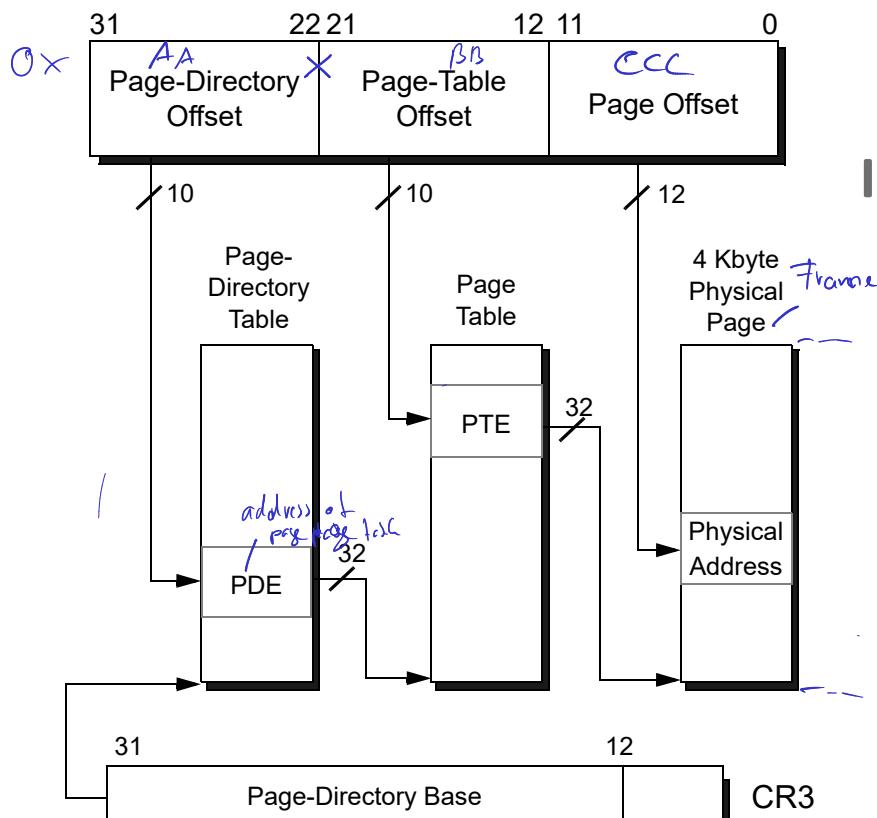
or

*speak after I  
ask you to*



Let's proceed with a  
**2 level page table.**

# 2 level page table

**Virtual Address**


## Intel x86/32 bit

- Usually uses 2 level page tables
- Offset: 12 bit
- $2^{12} = 4 \text{ KiB}$  page size
- 1 page directory per process
- $2^{10} = 1024$  page tables per process

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

# 2 level page table entries

PDE - page directory entry

- Page-Table Base Address (20 bit)  
= address of page table

PTE - page table entry

- Physical-Page Base Address (20 bit)  
= address of frame

Fields

- P = present (loaded into memory)
- R/W = read/write
- U/S = user/supervisor
- PWT = page level writethrough
- PCD = page level cache disable
- A = accessed
- D = dirty
- AVL = available to software (for OS)

31

	12	11	9	8	7	6	5	4	3	2	1	0
Page-Table Base Address	AVL		I G N	0	I G N	A	P C D	P W T	U /	R /	S W	P

Figure 5-5. 4-Kbyte PDE—Non-PAE Paging Legacy-Mode

31

	12	11	9	8	7	6	5	4	3	2	1	0
Physical-Page Base Address	AVL		G	P A T	D	A	P C D	P W T	U /	R /	S W	P

Figure 5-6. 4-Kbyte PTE—Non-PAE Paging Legacy-Mode

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

# 2 level page table entries

## PDE - page directory entry

- Page-Table Base Address (20 bit)  
= address of page table

## PTE - page table entry

- Physical-Page Base Address (20 bit)  
= address of frame

## Fields

- P = present (loaded into memory)
- R/W = read/write
- U/S = user/supervisor
- PWT = page level writethrough
- PCD = page level cache disable
- A = accessed
- D = dirty
- AVL = available to software (for OS)

31

	12	11	9	8	7	6	5	4	3	2	1	0
Page-Table Base Address	AVL		I G N	0	I G N	A	P C D	P W T	U S	R W	P	

Figure 5-5. 4-Kbyte PDE—Non-PAE Paging Legacy-Mode

31

	12	11	9	8	7	6	5	4	3	2	1	0
Physical-Page Base Address	AVL		G	P A T	D	A	P C D	P W T	U S	R W	P	

Figure 5-6. 4-Kbyte PTE—Non-PAE Paging Legacy-Mode

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

# 2 level page table entries

## PDE - page directory entry

- Page-Table Base Address (20 bit)  
= address of page table

## PTE - page table entry

- Physical-Page Base Address (20 bit)  
= address of frame

## Fields

- P = present (loaded into memory)
- R/W = read/write
- U/S = user/supervisor
- PWT = page level writethrough
- PCD = page level cache disable
- A = accessed
- D = dirty
- AVL = available to software (for OS)

31

Page-Table Base Address	AVL	I G N	G N	A	P C D	P W T	U S	R W	P

Figure 5-5. 4-Kbyte PDE—Non-PAE Paging Legacy-Mode

31

Physical-Page Base Address	AVL	G A T	P D	A	P C D	P W T	U S	R W	P
<i>Frame</i>									

Figure 5-6. 4-Kbyte PTE—Non-PAE Paging Legacy-Mode

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]

# 2 level page table entries

## PDE - page directory entry

- Page-Table Base Address (20 bit)  
= address of page table

## PTE - page table entry

- Physical-Page Base Address (20 bit)  
= address of frame

## Fields

- P = present (loaded into memory)
- R/W = read/write
- U/S = user/supervisor
- PWT = page level writethrough
- PCD = page level cache disable
- A = accessed
- D = dirty
- AVL = available to software (for OS)

31

Page-Table Base Address	AVL	I G N	0	I G N	A	P C D	P W T	U S	R W	P
-------------------------	-----	-------------	---	-------------	---	-------------	-------------	--------	--------	---

Figure 5-5. 4-Kbyte PDE—Non-PAE Paging Legacy-Mode

31

Physical-Page Base Address	AVL	G A T	P A D	D A	A	P C D	P W T	U S	R W	P
----------------------------	-----	-------------	-------------	--------	---	-------------	-------------	--------	--------	---

Figure 5-6. 4-Kbyte PTE—Non-PAE Paging Legacy-Mode

[source: AMD64 Architecture Programmer's Manual Volume 2: System Programming]



# Questions?

All right?  $\Rightarrow$



Question?  $\Rightarrow$



and use **chat**

or

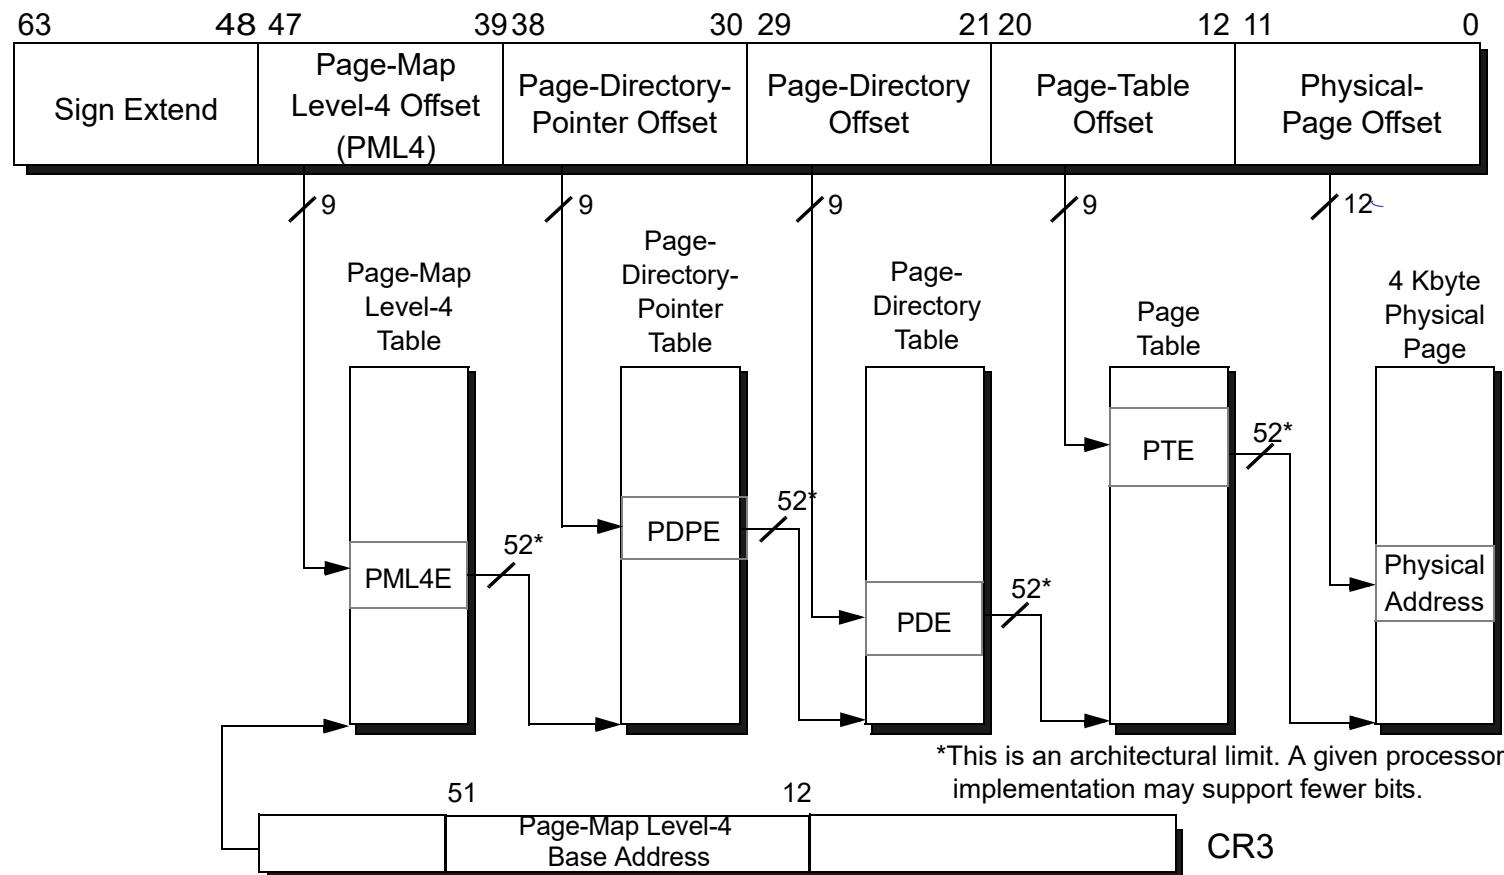
*speak after I  
ask you to*



And finally with  
**N level page tables.**

# 4 level page table

Virtual Address



# N level page table

**More details:**

**AMD64 Architecture Programmer's Manual Volume 2:  
System Programming** ([link](#))

Lets have a look on page 121, table 5-1.



# Questions?

All right?  $\Rightarrow$



Question?  $\Rightarrow$



and use **chat**

or

*speak after I  
ask you to*



# Summary and outlook

## Summary

- MMU
- Virtual addresses
- 1/2/N level page table

## Outlook

- Memory hierarchy
- Associative memory
- Translation lookaside buffer
- Cache
- Memory protection



# Summary and outlook

## Summary

- MMU
- Virtual addresses
- 1/2/N level page table

## Outlook

- Memory hierarchy
- Associative memory
- Translation lookaside buffer
- Cache
- Memory protection