



Start: 8:01

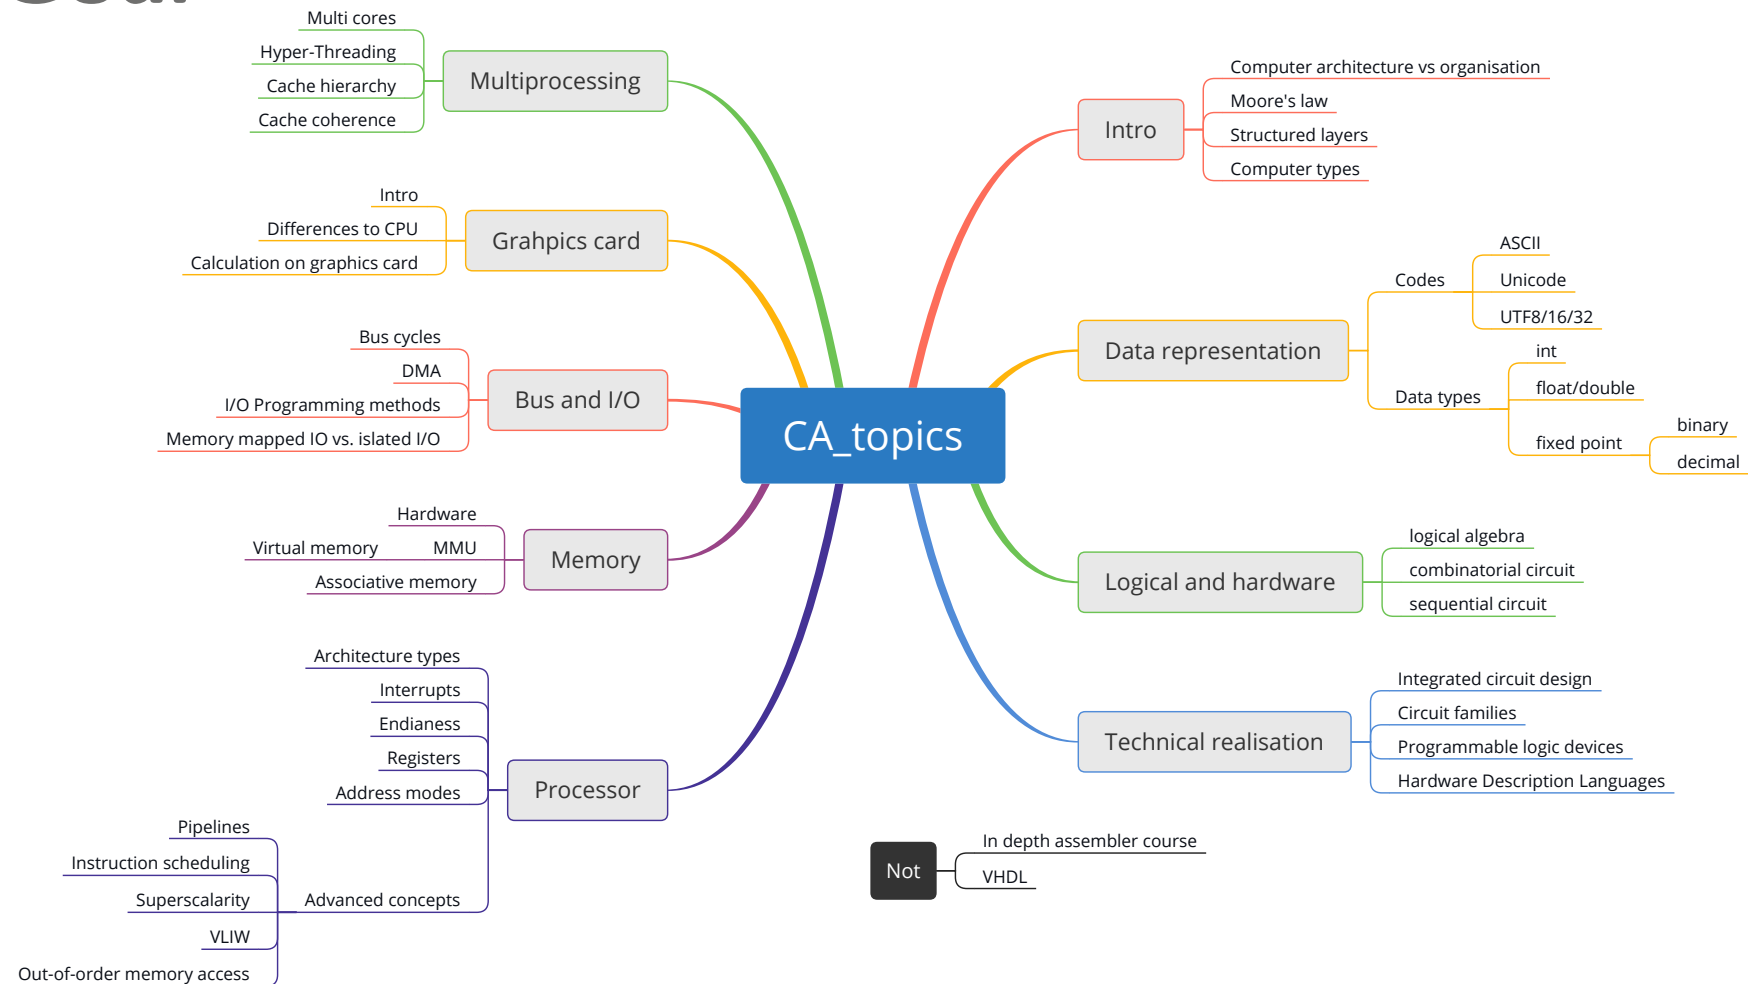
Prof. Dr. Florian Künzner

Technical University of Applied Sciences Rosenheim, Computer Science

CA 14 – Multiprocessing

The lecture is based on the work and the documents of Prof. Dr. Theodor Tempelmeier

Goal



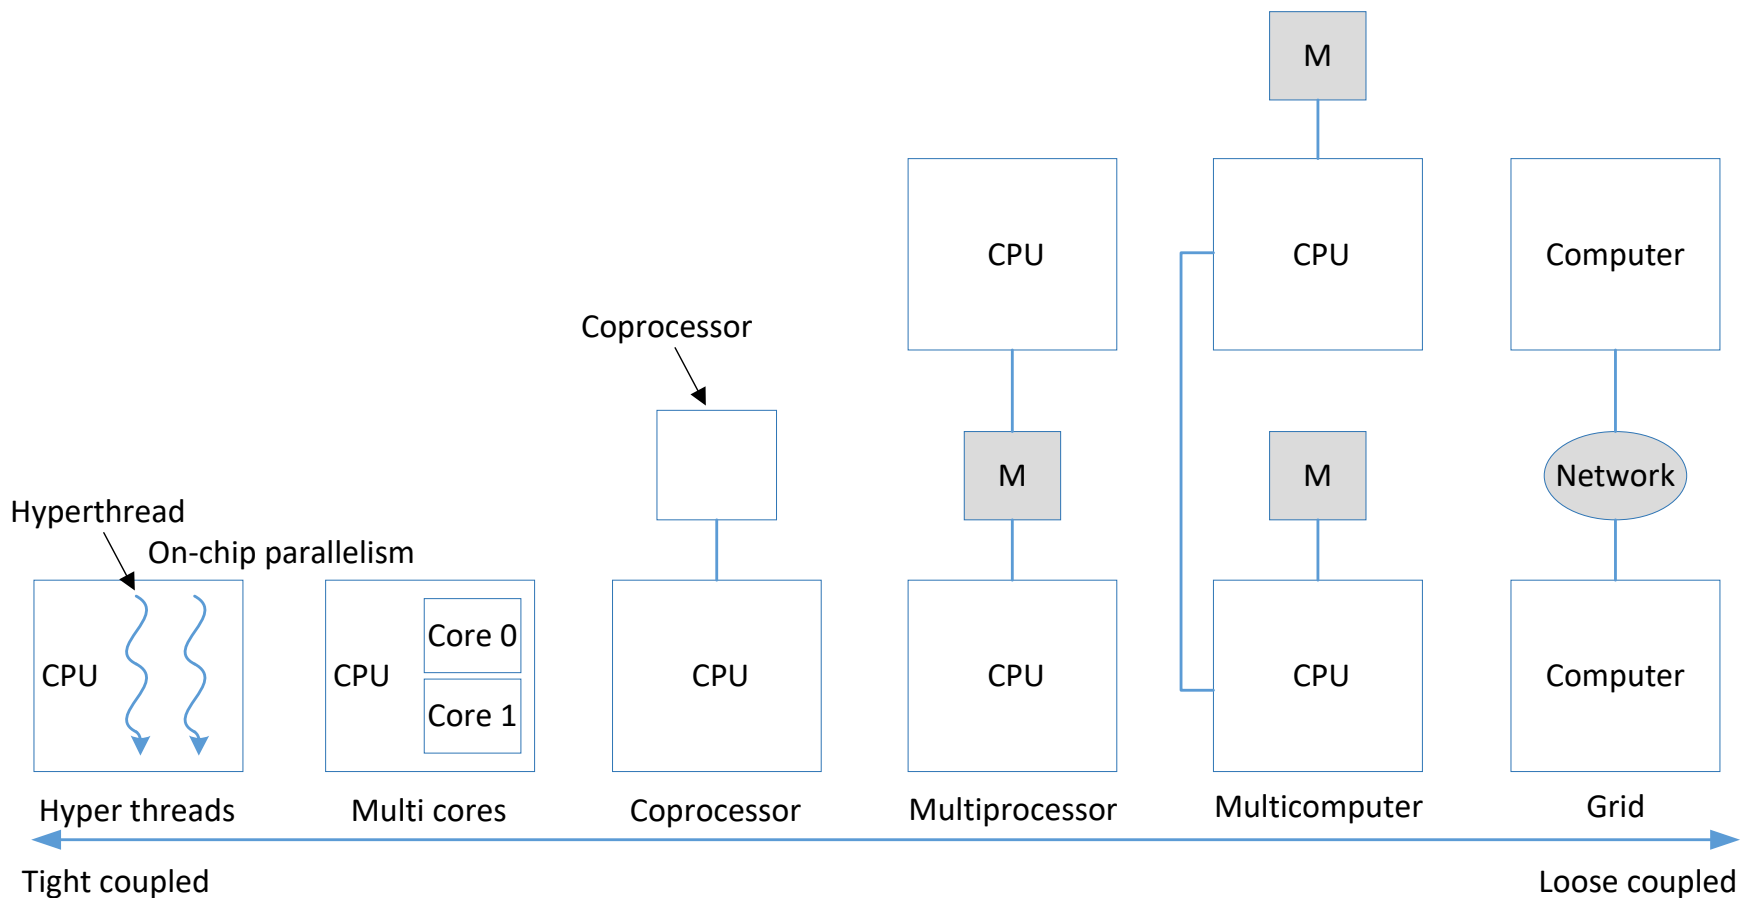
Types of multiprocessing

CPU (core)

Co-processor

Which types of multiprocessing do you know?

Types of multiprocessing



[cmp: [1, Fig. 8.1; P. 573]]

Hyper threading technology

Observation:

- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology

Observation:

- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology

Observation:

- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology

Observation:

- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology

Observation:

- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology

Observation:

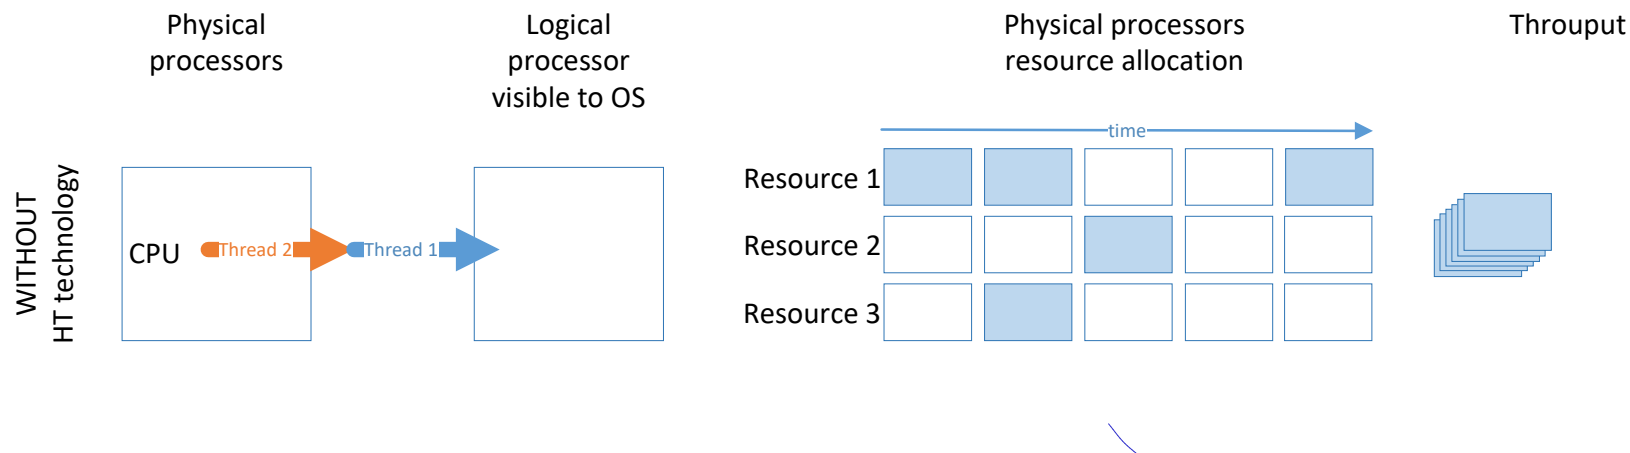
- Not all hardware parts are fully used inside one CPU (core)

Idea:

- Each physical processor is visible as two logical cores to the OS
- Spend some additional hardware to increase hardware utilisation inside one CPU (core)
- The main function of hyper-threading is to increase the number of independent instructions in the pipeline

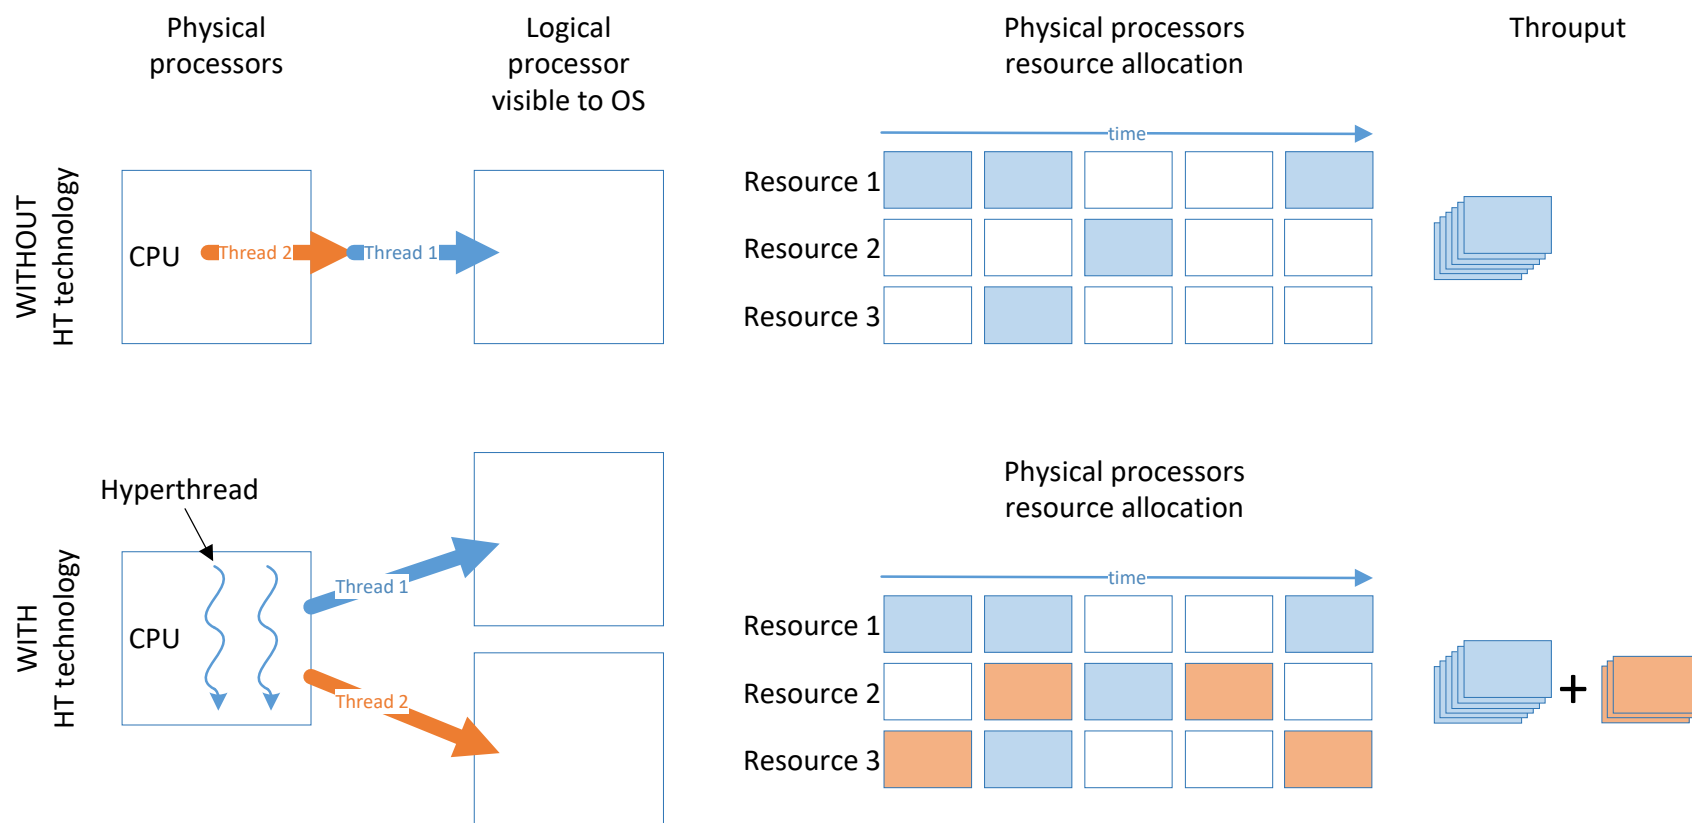
Only little additional hardware (5%) is required for hyper threading to speed-up execution up to 30%.

Hyper threading technology



[source (cmp): https://www.techpowerup.com/reviews/Intel/Core_i5_661/14.html]

Hyper threading technology



[source (cmp): https://www.techpowerup.com/reviews/Intel/Core_i5_661/14.html]

Multiple cores

Multicore CPU

Core 0

Core 1

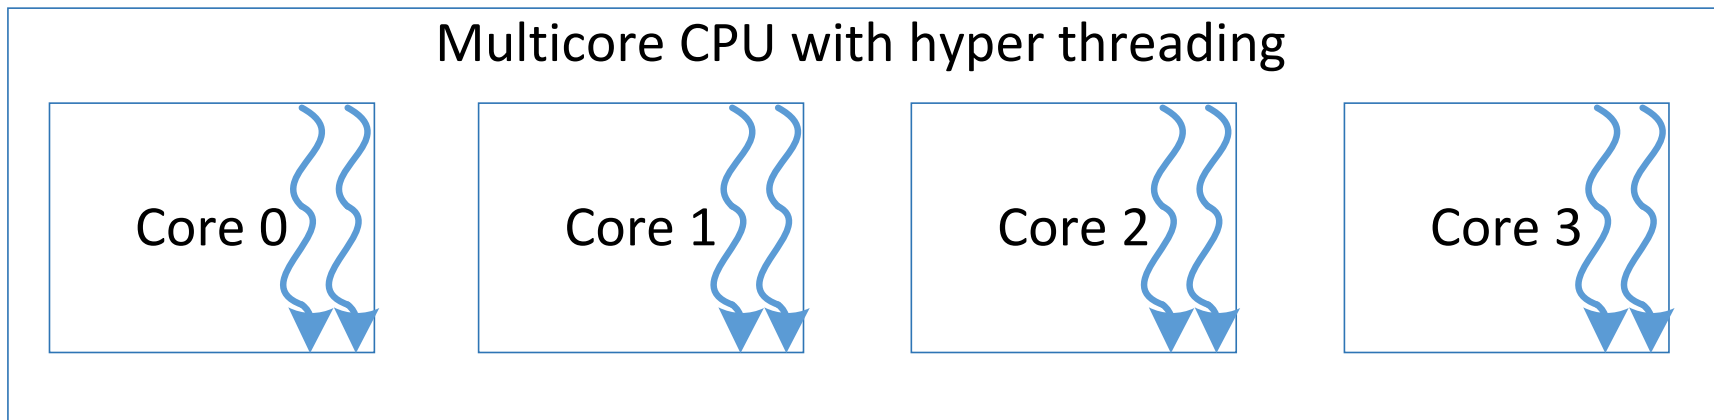
Core 2

Core 3



Multiple cores with hyper threading

Multicore CPU with hyper threading



Questions?

All right?



Question?



and use **chat**

or

speak *after* I
ask you to

CPU performance vs cost

	Single CPU (Power: 1)	Multiprocessor (Power: N)	Single CPU (Power: N)
Response time of a process (with unloaded machine)	X	X	$\frac{X}{N}$
Throughput of processes (with sufficient processes in ready state)	Y	$\approx N \cdot Y$	$N \cdot Y$
Response time of an application (consisting of several processes in ready state)	X	$\approx \frac{X}{N}$	$\frac{X}{N}$

Questions?

All right?



Question?

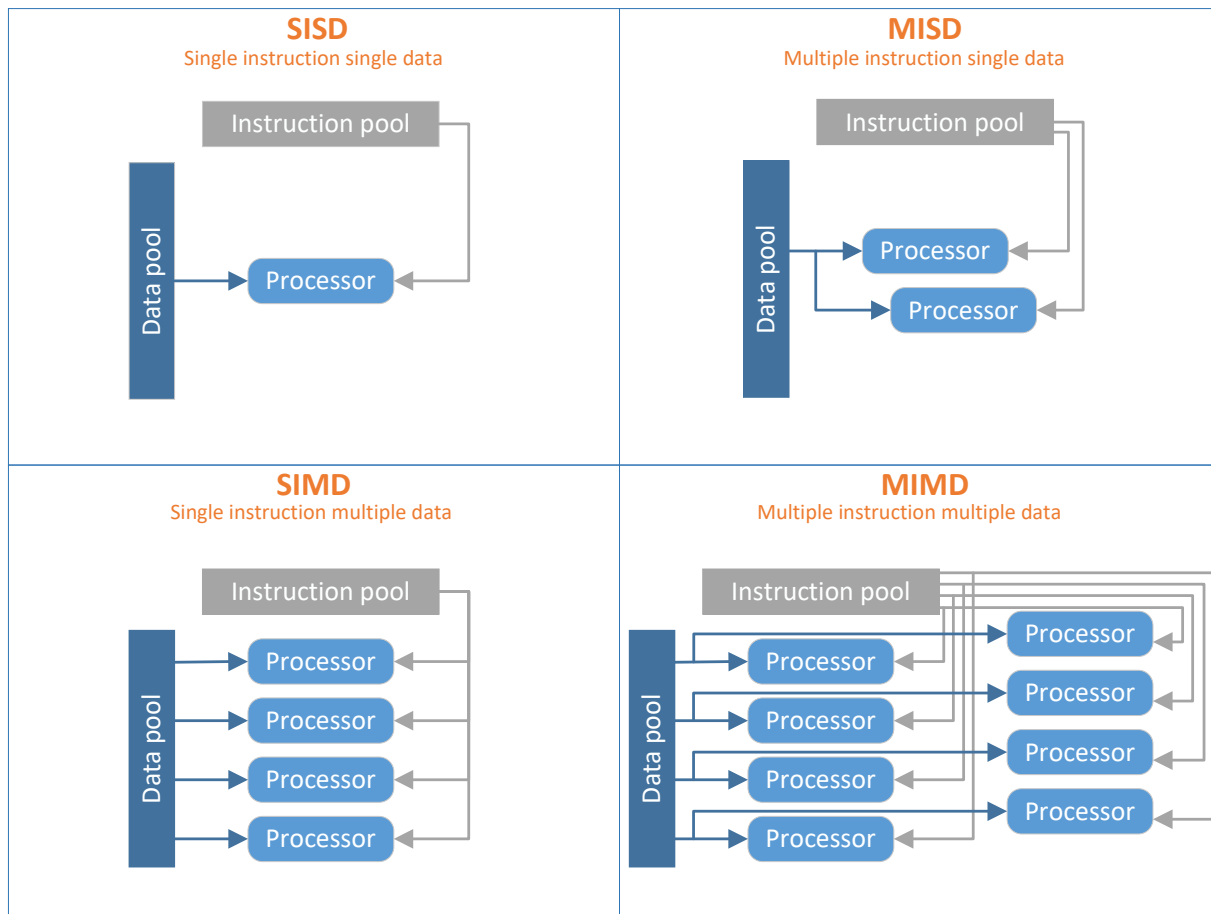


and use **chat**

or

speak *after* I
ask you to

Classification of Flynn



[source: <https://www.geeksforgeeks.org/computer-architecture-flynnns-taxonomy/>]

Questions?

All right?



Question?

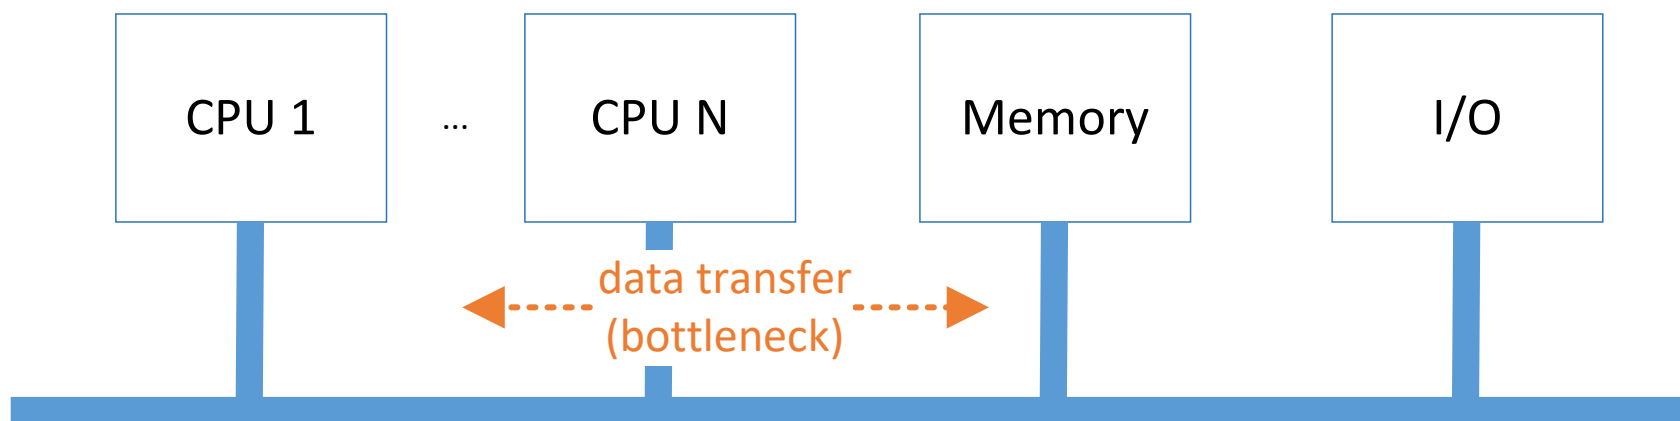


and use **chat**

or

speak *after* I
ask you to

CPU → memory bottleneck



Remember?

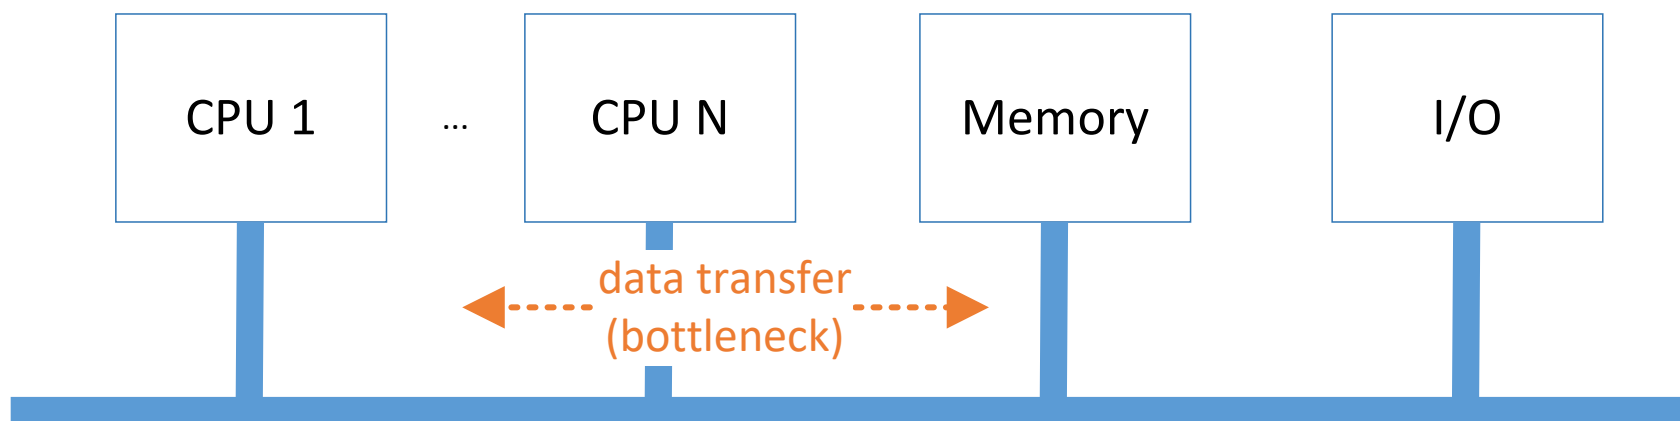
- CPU instructions are very fast ($< 1ns$)
- Memory access is slow ($< 30ns$)

With more CPUs, the problem gets even worse!

Solution:

- Local caches
- Large caches

CPU → memory bottleneck



Remember?

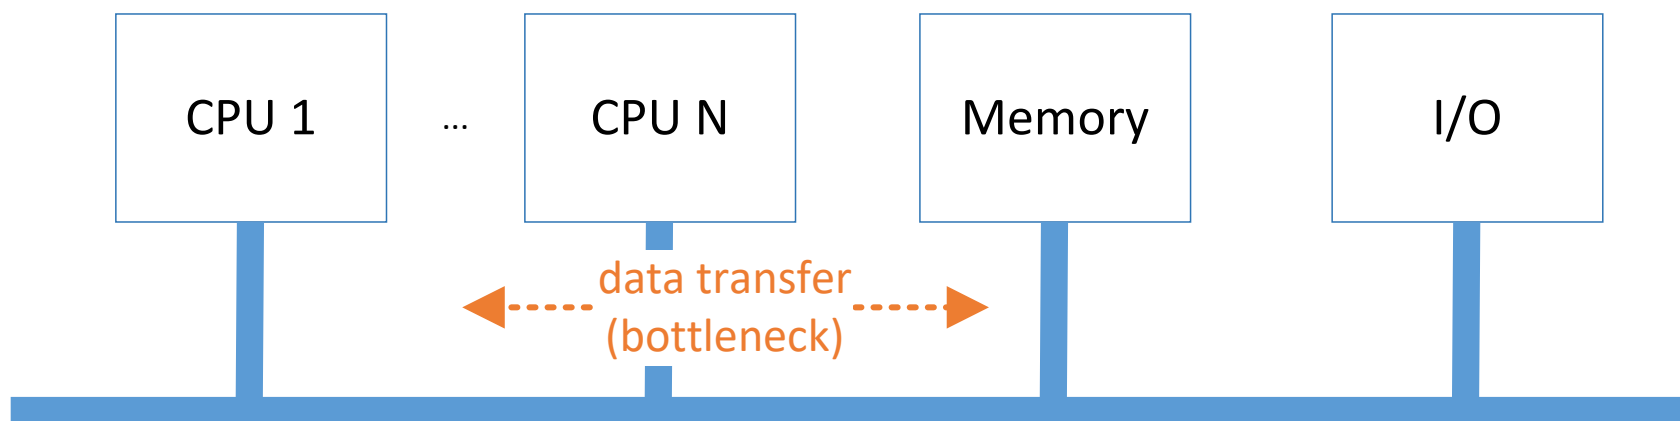
- CPU instructions are very fast ($< 1ns$)
- Memory access is slow ($< 30ns$)

With more CPUs, the problem gets even worse!

Solution:

- Local caches
- Large caches

CPU → memory bottleneck



Remember?

- CPU instructions are very fast ($< 1ns$)
- Memory access is slow ($< 30ns$)

With more CPUs, the problem gets even worse!

Solution:

- Local caches
- Large caches

Questions?

All right?



Question?



and use **chat**

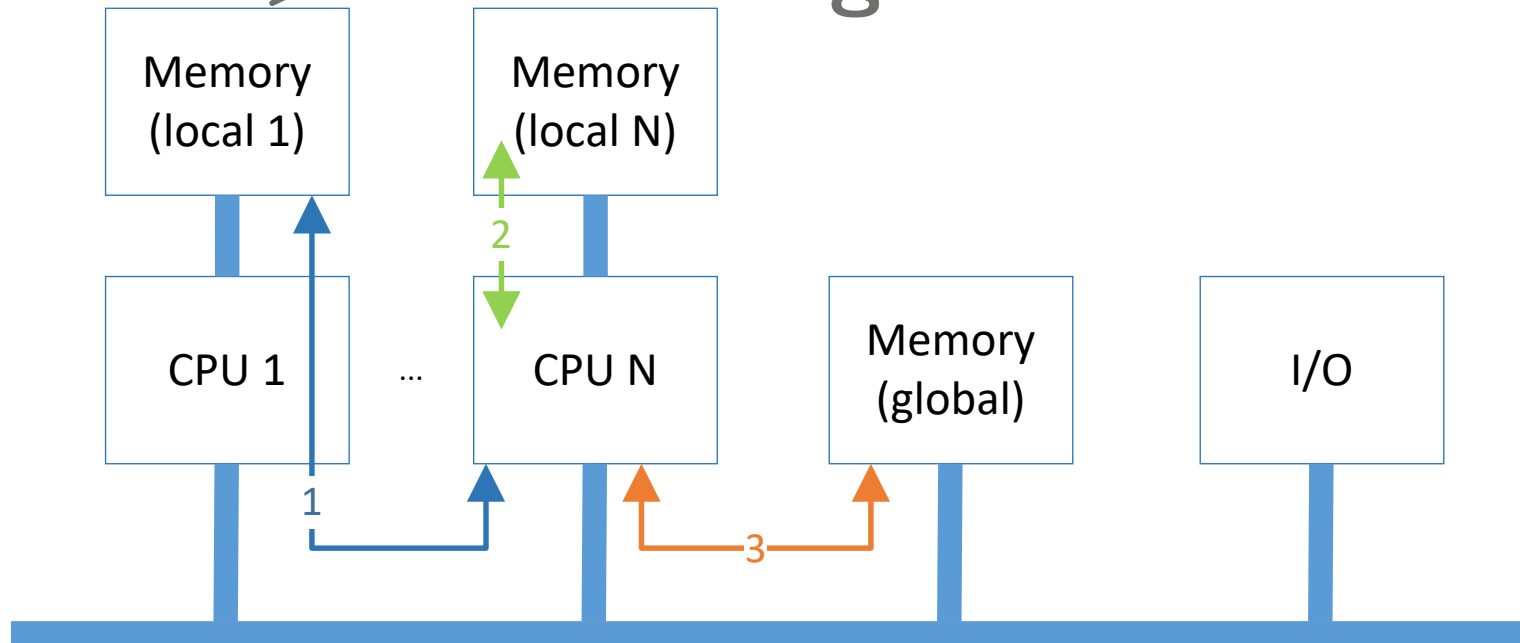
or

speak *after* I
ask you to

Structures of multi cores and multiprocessors

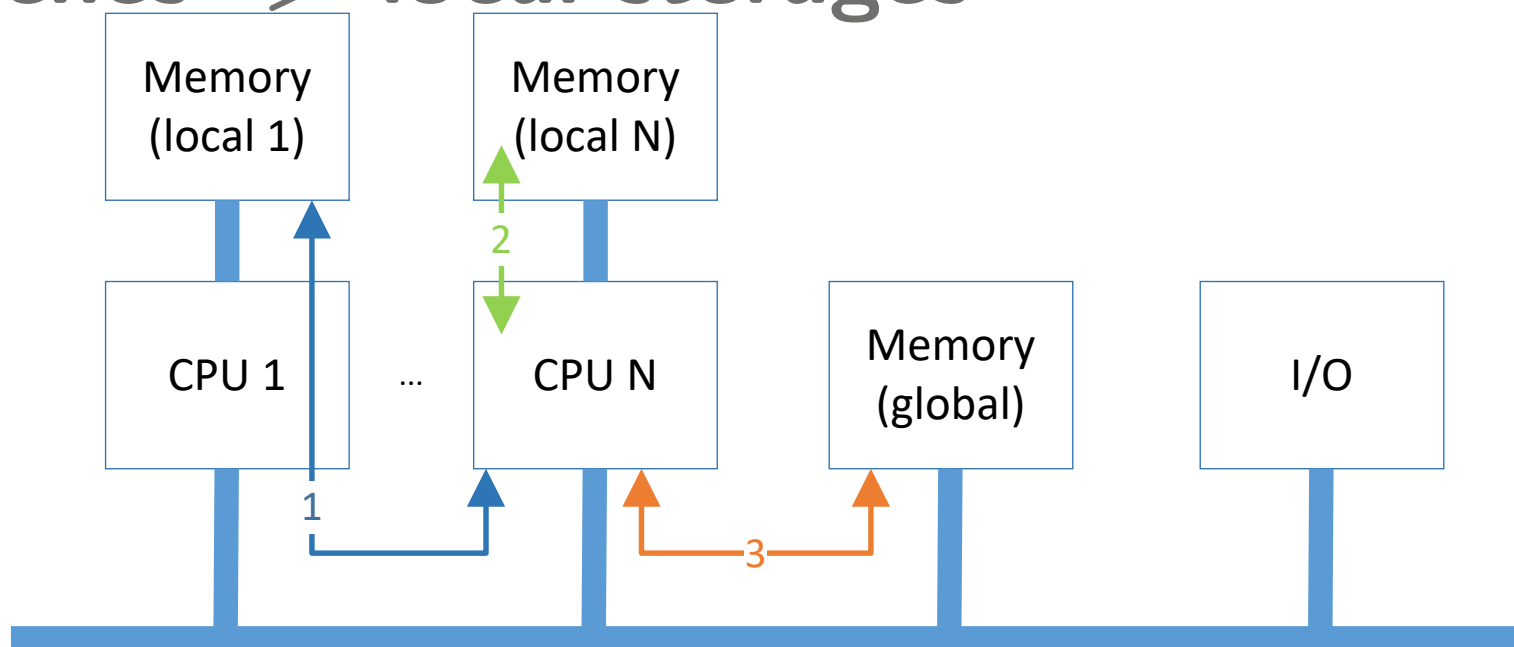
Structure of multi core CPUs with its caches.

Caches -> local storages



NUMA: Non-uniform memory access

Caches -> local storages

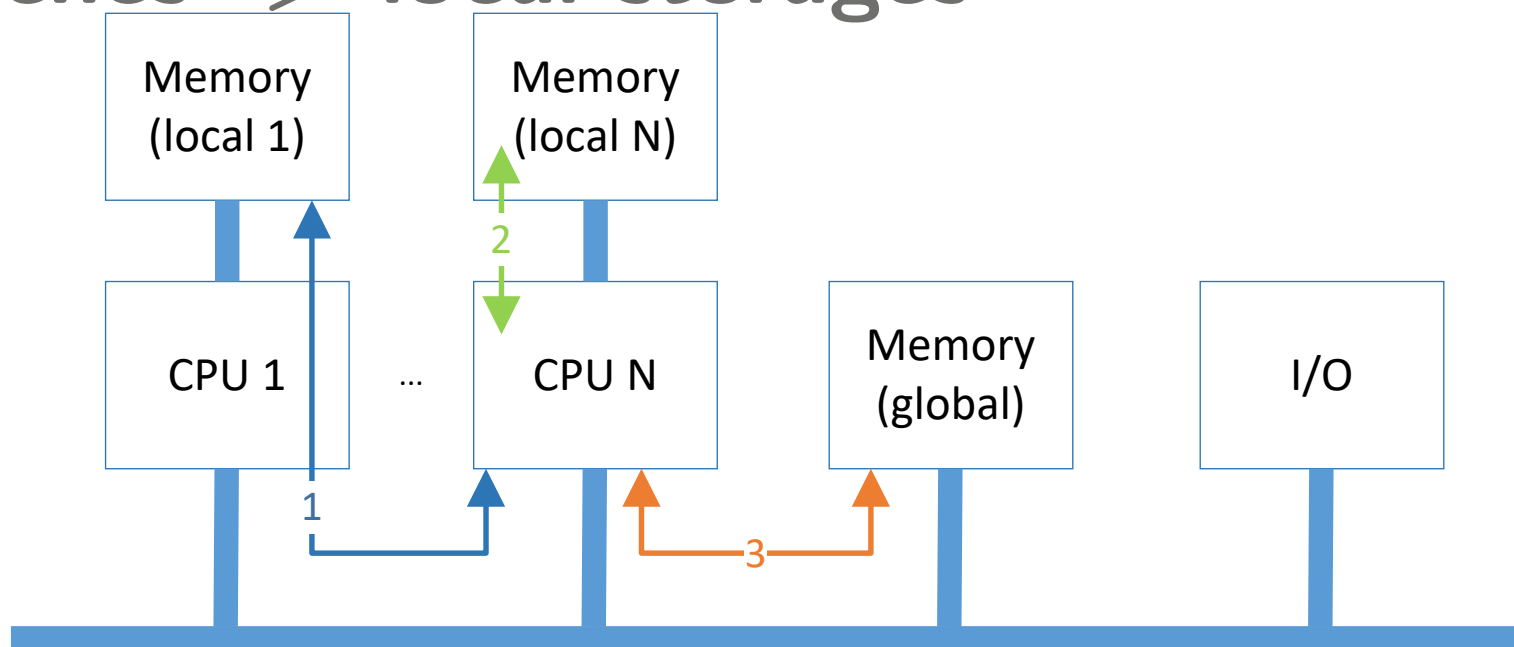


NUMA: Non-uniform memory access

- (1) CPU N access local memory 1
- (2) CPU N access local memory N
- (3) CPU N access global memory

Memory access time depends on the memory location.

Caches -> local storages

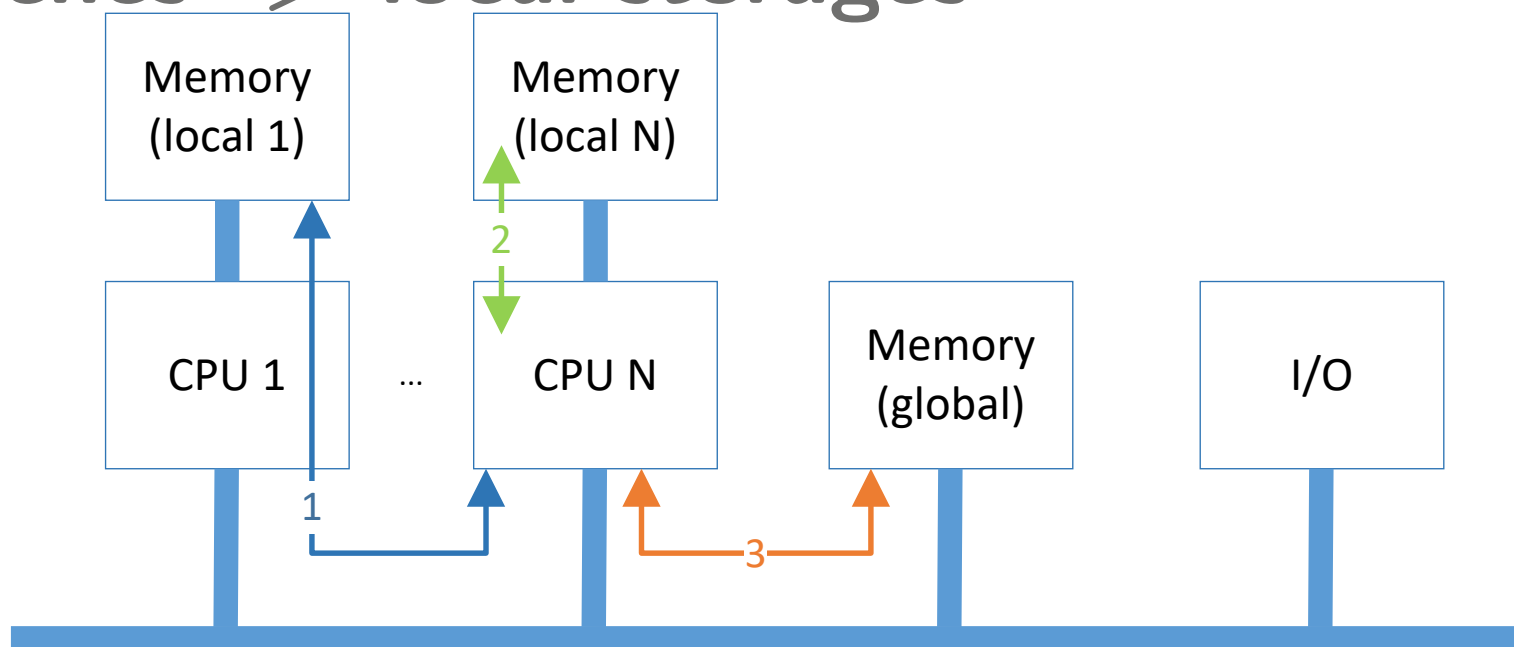


NUMA: Non-uniform memory access

- (1) CPU N access local memory 1
- (2) CPU N access local memory N
- (3) CPU N access global memory

Memory access time depends on the memory location.

Caches -> local storages

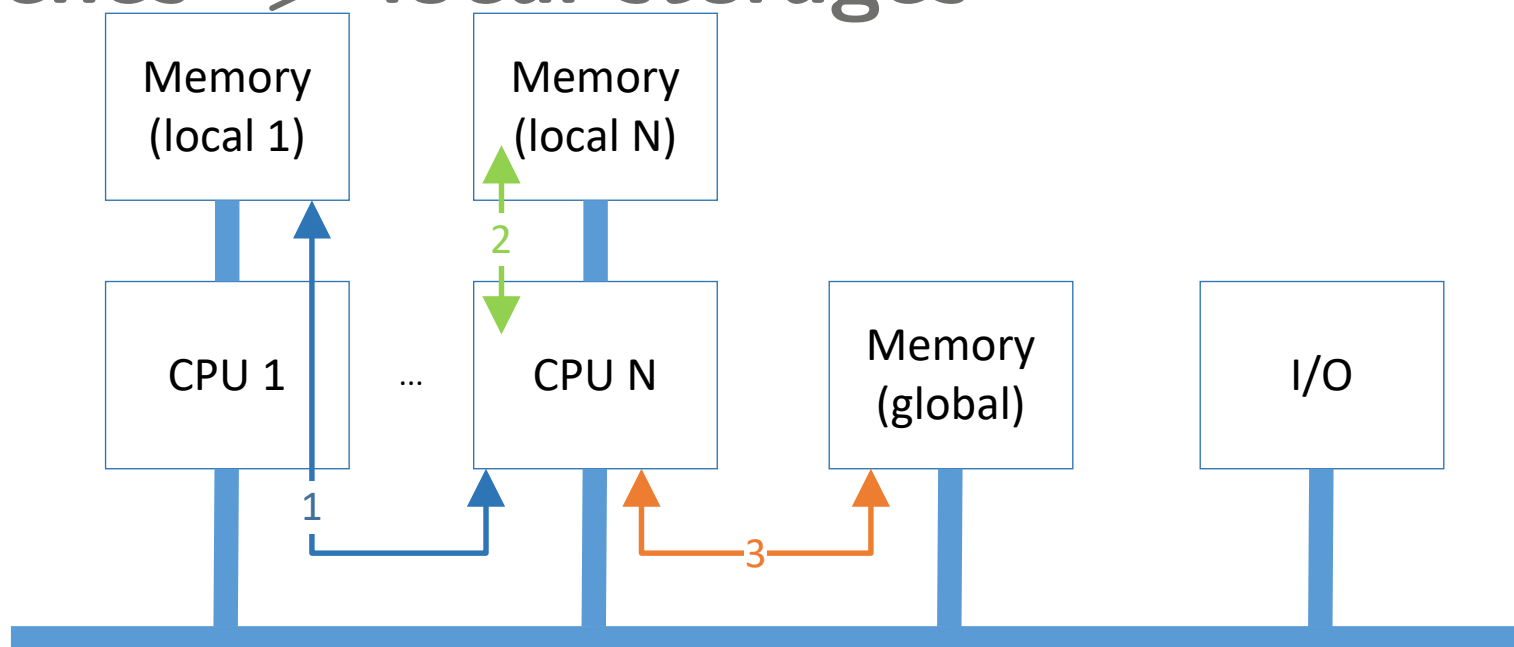


NUMA: Non-uniform memory access

- (1) CPU N access local memory 1
- (2) CPU N access local memory N
- (3) CPU N access global memory

Memory access time depends on the memory location.

Caches -> local storages



NUMA: Non-uniform memory access

- (1) CPU N access local memory 1
- (2) CPU N access local memory N
- (3) CPU N access global memory

Memory access time depends on the memory location.

Requirements

Important requirements for multiple cores and multiprocessing

Locking commands

- read-modify-write
- read/write cycles

Example:

- TAS: test and set
- CAS: compare and swap

Communication between processors must be possible

- Global memory
- Interrupt from processor to processor
- Identification of processors

Requirements

Important requirements for multiple cores and multiprocessing

Locking commands

- read-modify-write
- read/write cycles

Example:

- TAS: test and set
- CAS: compare and swap

Communication between processors must be possible

- Global memory
- Interrupt from processor to processor
- Identification of processors

Requirements

Important requirements for multiple cores and multiprocessing

Locking commands

- read-modify-write
- read/write cycles


Example:


- TAS: test and set
- CAS: compare and swap

Communication between processors must be possible

- Global memory
- Interrupt from processor to processor
- Identification of processors

Questions?

All right? \Rightarrow 

Question? \Rightarrow  and use **chat**

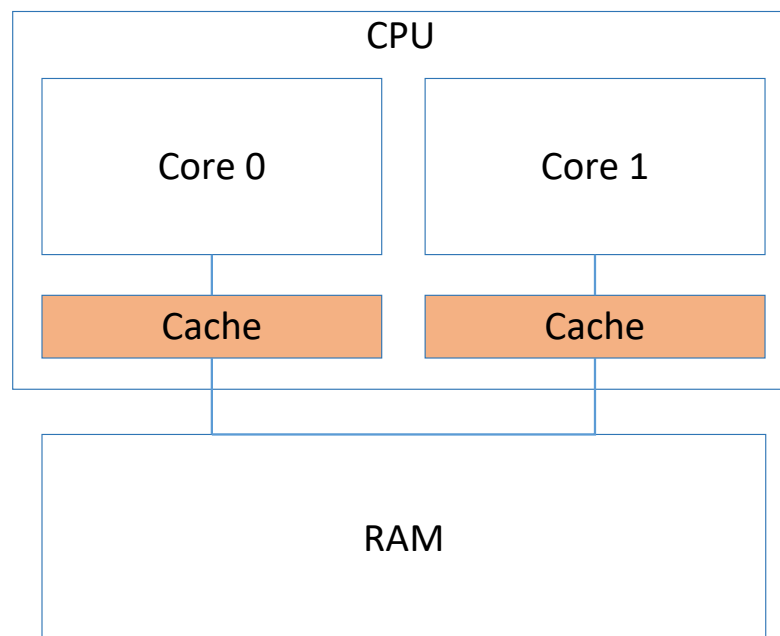
or

speak *after* I ask you to

Cache coherency

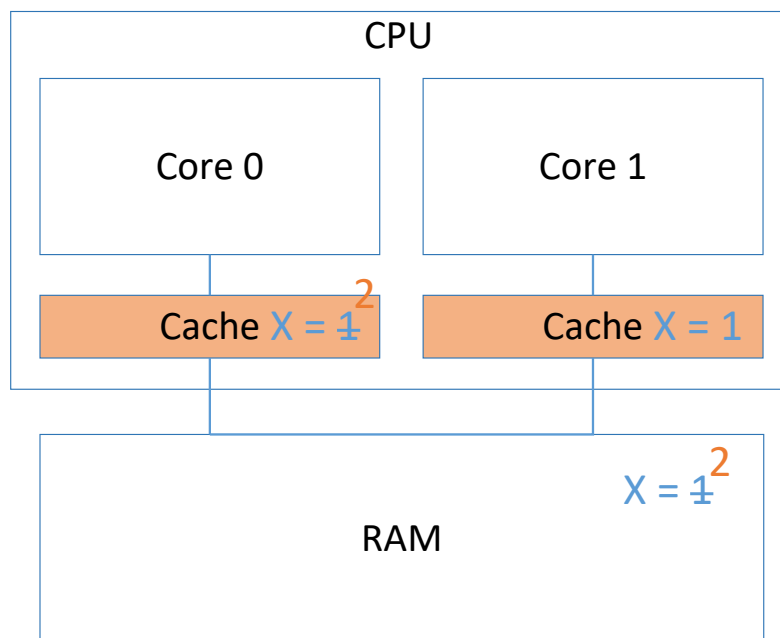
Is there a problem with local memory (caches)?

Cache coherency: Problem!



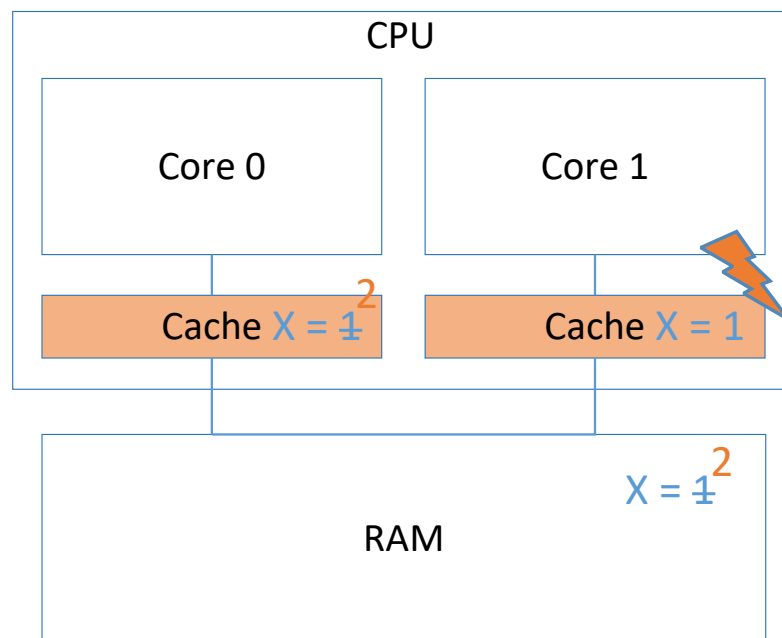
- A dual core CPU with local caches
- Variable X contains value 1
- Both caches contains a copy of X
- Core 0 sets $X=2$
(with **write through**)
- **Problem:** Core 1 still reads $X==1!!$

Cache coherency: Problem!



- A dual core CPU with local caches
- Variable X contains value 1
- Both caches contains a copy of X
- Core 0 sets $X=2$
(with **write through**)
- **Problem:** Core 1 still reads $X==1!!$

Cache coherency: Problem!



- A dual core CPU with local caches
- Variable X contains value 1
- Both caches contains a copy of X
- Core 0 sets $X=2$
(with **write through**)
- **Problem:** Core 1 still reads $X=1$!!

Towards a solution

How can we avoid such cache coherency problems?

- **Programming model:** Change the programming model, but then the consistency may not be ensured. Unpleasant thing!
- **Snoop protocols:** All caches operate with **write through** and **listen** to all write cycles on the bus and may set the corresponding entries to invalid. But this is **slow**!
- **MESI protocol:** Various protocols that (at least in some cases) avoid write through.

CC-NUMA: cache coherency with NUMA.

Cache coherency: The MESI protocol

The MESI protocol **ensures** the **cache coherency** in multi core and multiprocessor systems.

Finite state machine with 4 states:

- M: Modified exclusive
- E: Exclusive unmodified
- S: Shared unmodified
- I: Invalid

There is a finite state machine for

- every **cache line**
- in every **cache**

Cache coherency: The MESI protocol

The MESI protocol **ensures** the **cache coherency** in multi core and multiprocessor systems.

Finite state machine with 4 states:

- M: Modified exclusive
- E: Exclusive unmodified
- S: Shared unmodified
- I: Invalid

Cache coherency: The MESI protocol

The MESI protocol **ensures** the **cache coherency** in multi core and multiprocessor systems.

Finite state machine with 4 states:

- M: Modified exclusive
- E: Exclusive unmodified
- S: Shared unmodified
- I: Invalid

There is a finite state machine for

- every **cache line**
- in every **cache**

Cache coherency: The MESI protocol

The MESI protocol **ensures** the **cache coherency** in multi core and multiprocessor systems.

Finite state machine with 4 states:

- M: Modified exclusive
- E: Exclusive unmodified
- S: Shared unmodified
- I: Invalid

There is a finite state machine for

- every **cache line**
- in every **cache**

Cache coherency: The MESI protocol

The MESI protocol **ensures** the **cache coherency** in multi core and multiprocessor systems.

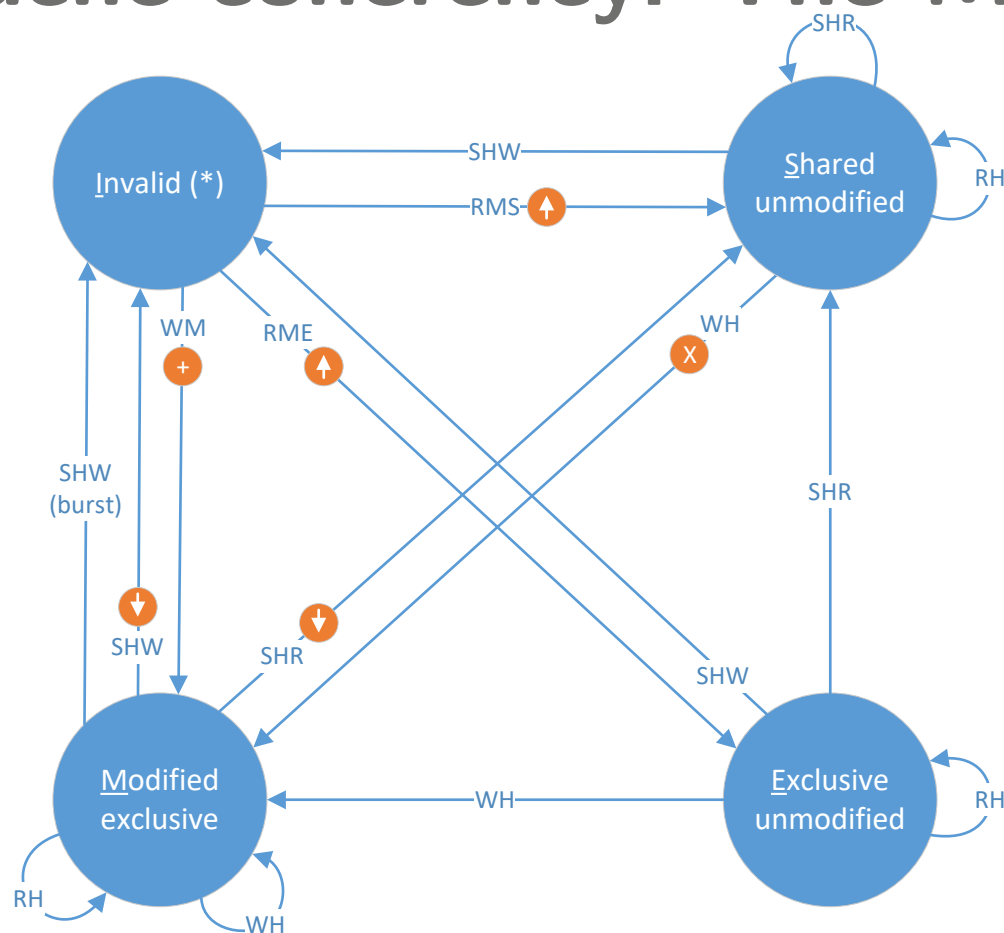
Finite state machine with 4 states:

- M: Modified exclusive
- E: Exclusive unmodified
- S: Shared unmodified
- I: Invalid

There is a finite state machine for

- every **cache line**
- in every **cache**

Cache coherency: The MESI protocol



RH Read Hit
RMS Read miss, shared
RME Read miss, exclusive
WH Write hit
WM Write miss
SHR Snoop hit on a read
SHW Snoop hit on a write or
Read-with-intent-to-modify or
invalidate

X Invalidate transaction
+ Read-with-intent-to-modify
↑ Cache line fill
↓ Dirty line copyback

(*) On a miss, the old line is first
invalidated and copied-back if
exclusive modified

Cache coherency: Modern MESI

The MESI protocol is with extensions still in use.

MOESI: AMD64

- MESI states as in the original MESI protocol
- O: Owned
- Allows **moving** a (**dirty**) cache line around caches without updating main memory

MESIF: Intel

- MESI states as in the original MESI protocol
- F: Forward
- Allows **copying** (**clean**) data from one cache to another without accessing main memory

[comparison: https://en.wikipedia.org/wiki/MESIF_protocol]

Cache coherency: Modern MESI

The MESI protocol is with extensions still in use.

MOESI: AMD64

- MESI states as in the original MESI protocol
- O: Owned
- Allows **moving** a (**dirty**) cache line around caches without updating main memory

MESIF: Intel

- MESI states as in the original MESI protocol
- F: Forward
- Allows **copying** (**clean**) data from one cache to another without accessing main memory

[comparison: https://en.wikipedia.org/wiki/MESIF_protocol]

Cache coherency: Modern MESI

The MESI protocol is with extensions still in use.

MOESI: AMD64

- MESI states as in the original MESI protocol
- O: Owned
- Allows **moving** a (**dirty**) cache line around caches without updating main memory

MESIF: Intel

- MESI states as in the original MESI protocol
- F: Forward
- Allows **copying** (**clean**) data from one cache to another without accessing main memory

[comparison: https://en.wikipedia.org/wiki/MESIF_protocol]

Questions?

All right?



Question?



and use **chat**

or

speak *after* I
ask you to

Summary

Summary

- Types of multiprocessing
- CPU performance vs cost
- Memory bottleneck
- Classification of Flynn
- Structures of multi cores and multiprocessors
- MESI protocol

Summary

Summary

- Types of multiprocessing
- CPU performance vs cost
- Memory bottleneck
- Classification of Flynn
- Structures of multi cores and multiprocessors
- MESI protocol