

Applications of & Introduction to Artificial Intelligence

Convolutional Neural Networks


Technische Hochschule Rosenheim

Sommer 2020

Prof. Dr. J. Schmidt


Example:

Typical MLP-Structure for Images



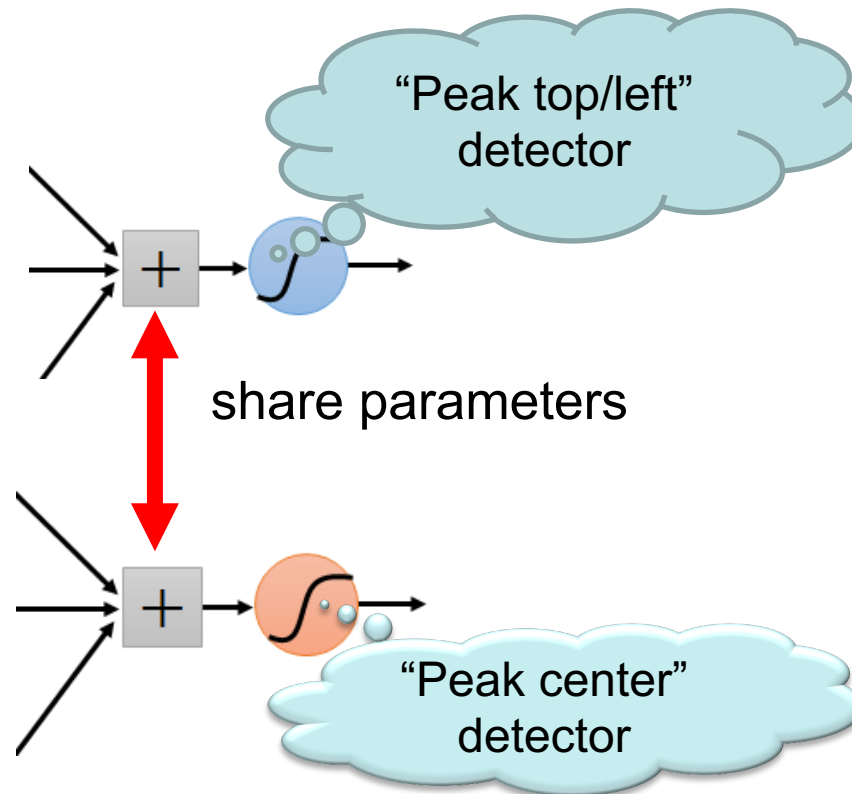

Recognition of Image Parts

- many patterns are smaller than the complete image
- for small regions: less parameters required




Similar Patterns

- similar patterns can be found in different image locations
- Idea: Train many small detectors that
 - ⊞ move over the image
 - ⊞ share parameters

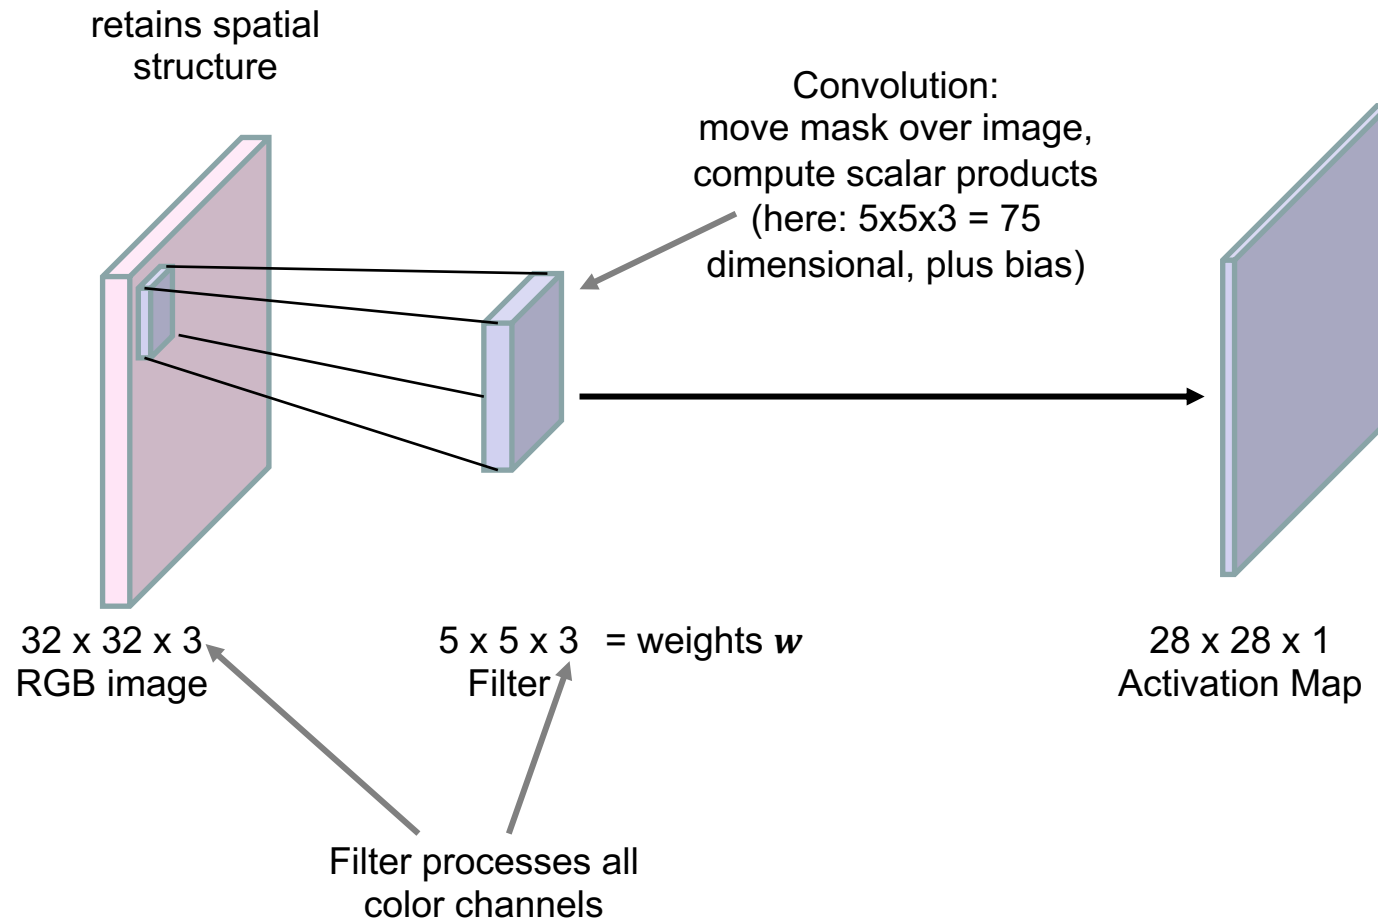


Convolutional Layer

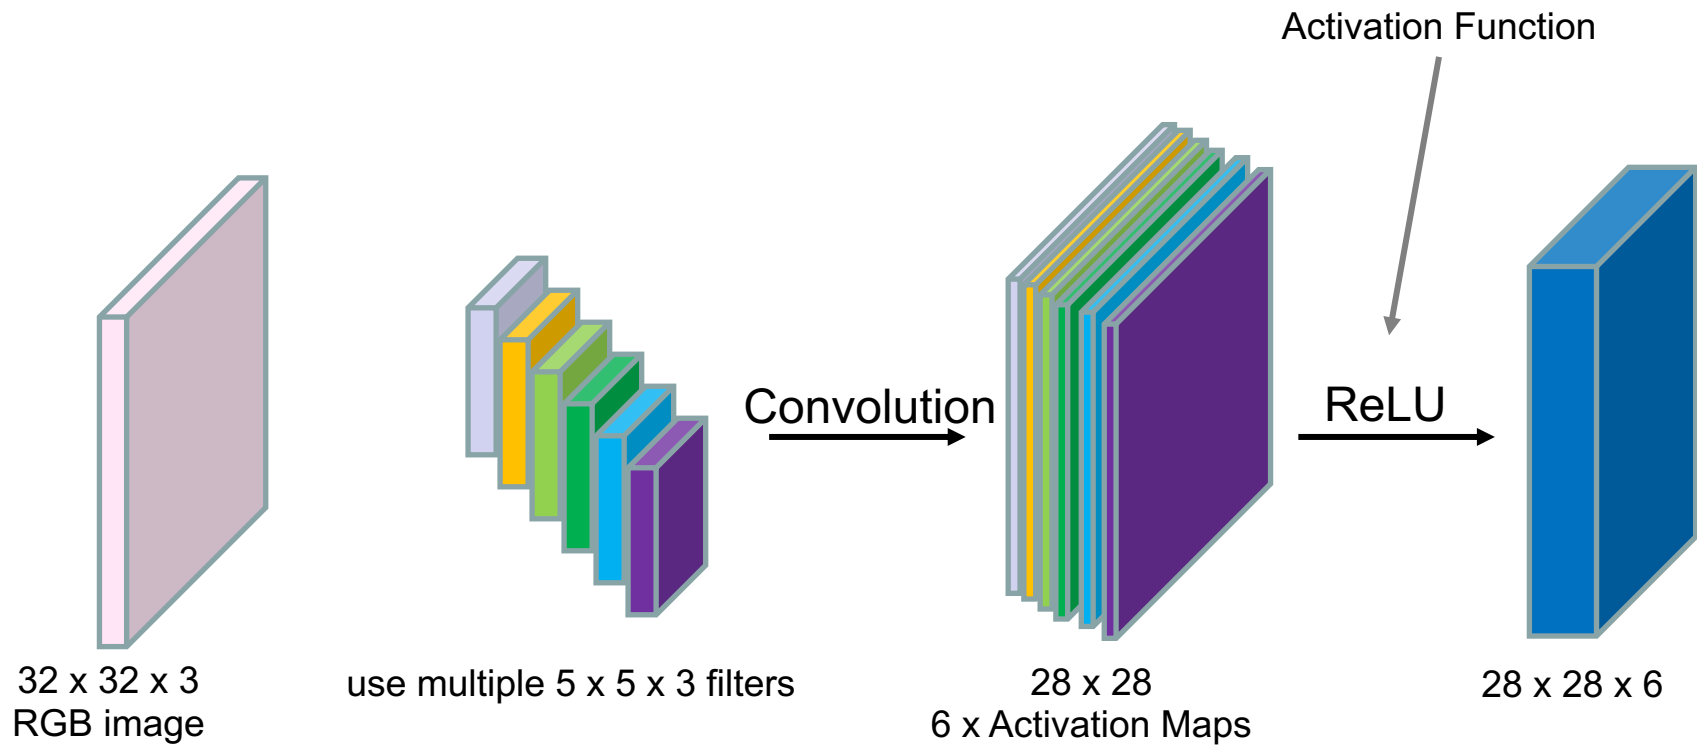
- hence: CNN – Convolutional Neuronal Network
- consists of (linear) convolution filters
- the filter masks are learned during training
- first used with backpropagation in LeNet (1989-1998):
LeCun, Bottou, Bengio, Haffner. Gradient-Based Learning Applied to Document Recognition. Proc. of the IEEE 86(11): 2278-2324, 1998.



Convolutional Layer




Convolutional Layer



Convolutional Layer

Convolution and activation are now repeated several times
Idea: Combine low-level features, combine again etc.



Hyperparameter – Stride

the filter mask can be moved by more than one pixel (stride)
this differs from the “normal” convolution operation

Example: 7x7 image with 3x3 filter

	X	X	X	X	X	
	X	X	X	X	X	
	X	X	X	X	X	
	X	X	X	X	X	
	X	X	X	X	X	

Stride 1
Output: 5x5


	X		X		X	
	X		X		X	
	X		X		X	

Stride 2
Output : 3x3

	X			X		
	X			X		

Stride 3
asymmetric border –
stride does not match

Hyperparameter – Stride



Stride S

Size of output: $\frac{N - F}{S} + 1$

If result is integer:
Stride and filter size match

Example $N = 7, F = 3$:

$$S = 1: \frac{7-3}{1} + 1 = 5$$

$$S = 2: \frac{7-3}{2} + 1 = 3$$

$$S = 3: \frac{7-3}{3} + 1 = 2,33$$

Hyperparameter – Pad

- Problem: Input size for a layer is getting smaller and smaller
- Solution: Padding of border
 - with zeros (Zero-Padding)
 - with copies of the border pixels

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

For filter size $F \times F$

$\frac{F-1}{2}$ values are lost at the border


Examples:

$F = 3$: Padding with 1

$F = 5$: Padding with 2


$F = 7$: Padding with 3

Example



Number of parameters for this layer:
 each filter has $5 \cdot 5 \cdot 3 + 1 = 76$ parameters (+1 because of bias)
 10 filters, total: $76 \cdot 10 = 760$ parameters

Example



Number of parameters for this layer:
each filter has $5 \cdot 5 \cdot 10 + 1 = 251$ parameters (+1 because of bias)
16 filters, total: $251 \cdot 16 = 4016$ parameters

Hyperparameters – Convolution

- Number K and size F of filters
- Stride S
- Size of padding P
- typical values:
 - ⊞ $K = \text{power of 2, e.g. } 32, 64, 128, 512$
 - ⊞ $F = 3, S = 1, P = 1$
 - ⊞ $F = 5, S = 1, P = 2$
 - ⊞ $F = 5, S = 2, P = \text{matching}$
 - ⊞ $F = 1, S = 1, P = 0$
- transforms a layer of size $W \times H \times D$ into a layer of size $W' \times H' \times D'$:

$$W' = \frac{W - F + 2P}{S} + 1, \quad H' = \frac{H - F + 2P}{S} + 1, \quad D' = K$$

- Number of weights: $(F \cdot F \cdot D) \cdot K + K$

Pooling

- scaling does not change the object
- objective: smaller-sized layers

Bird




Scaling
(subsampling)

Bird



Pooling

each activation map is processed separately




Pooling

MAX-Pooling: Use the largest element within a windows of size $F \times F$

Average-Pooling: Use the mean value of all elements within a windows of size

Example: MAX-Pooling using 2x2 windows and stride $S = 2$




Hyperparameters – Pooling


- Size F of windows
- Stride S
- Typical values:
 - ⊞ $F = 2, S = 2$
 - ⊞ $F = 3, S = 2$
- transforms a layer of size $W \times H \times D$ into a layer of size $W' \times H' \times D'$:
$$W' = \frac{W-F}{S} + 1, H' = \frac{H-F}{S} + 1, D' = D$$
- Number of weights: none

Fully Connected Layers / Flatten

- at the end: fully connected layers as before (MLP)
→ Flattening




Typical Architecture



- n ca. 3, up to ca. 5
- m large
- $0 \leq k \leq 2$
- General tendency:
 - ⊞ use smaller filter sizes and deeper architecture
 - ⊞ away from pooling/fully connected layers towards pure convolutional layers

Example: Alexnet



ImageNet Classification Challenge 2012

- 1000 classes
- 1.2 million training images
- 50,000 validation images
- 150,000 test images

Network:

- 650,000 neurons
- 60 million parameters
- used CNN with ReLU on GPU for the first time

Pre-Processing:

- Scale/Crop images to 256 x 256
(training uses random crops of size 224x224 from these)
- Subtract mean RGB image


Krizhevsky, Sutskever, Hinton: ImageNet Classification with Deep Convolutional Neural Networks. Commun. ACM 60(6):84-90, 2017.

Transfer-Learning

- Problem:
Huge data sets required for training so many weights
- What if you do not have 1,000,000 images?
 - ⊞ Get a CNN that was trained on similar data
 - ⊞ use transfer-learning on that CNN
- Trained models available from, e.g.,
<https://github.com/BVLC/caffe/wiki/Model-Zoo>
<https://github.com/tensorflow/models>
- Transfer-Learning is the rule, not the exception

Transfer-Learning

small data set:




instead of re-training the last layer of weights also: use CNN as feature extractor:

- remove output layer
- the output of the layer before produces features (here: 4096)
- use these features to train a classifier (e.g. SVM)

Advantage: overfitting is not an issue for small data sets

Transfer-Learning

large data set:



- the larger the data set, the more layers can be re-trained
- chose learning rate considerably smaller than that of the original network (e.g. 1/10)

Object Detection/-localization

- R-CNN – Region Proposal CNN (Girschick 2014)
<https://arxiv.org/abs/1311.2524>
- SSD – Single Shot MultiBox Detector (Liu, 2016)
<https://arxiv.org/abs/1512.02325>
- YOLO – You Only Look Once (Redmon 2016)
<https://pjreddie.com/darknet/yolo/>

Sources

- Goodfellow, Bengio, Courville: *Deep Learning*, MIT Press, 2017.
<http://www.deeplearningbook.org/>
- Li, Johnson, Yeung: *CS231n: Convolutional Neural Networks for Visual Recognition*. Vorlesung Stanford University, 2018.
<http://cs231n.stanford.edu/>
- Li: *Deep Learning and Its Applications*. Lecture University of Waterloo, 2017.
<https://cs.uwaterloo.ca/~mli/cs898-2017.html>
- Original research articles as stated on the slides